

03_Amazon_Fine_Food_Reviews_Analysis_KNN ass

January 4, 2019

1 Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454 Number of users: 256,059 Number of products: 74,258 Timespan: Oct 1999 - Oct 2012 Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective: Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative? [Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

2 [1]. Reading Data

2.1 [1.1] Loading the data

The dataset is available in two forms 1. .csv file 2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to “positive”. Otherwise, it will be set to “negative”.

```
In [0]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

In [0]: # using SQLite Table to read data.
con = sqlite3.connect('drive/ai/database.sqlite')

# filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
# SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 500000 data points.
# you can change the number to any other number based on your computing power

# filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000 """, con)
```

```

# for tsne assignment you can take 5k data points

filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 """, con)

# Give reviews with Score>3 a positive rating(1), and reviews with a score<3 a negative rating(0)
def partition(x):
    if x < 3:
        return 0
    return 1

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)

```

Number of data points in our data (525814, 10)

```

Out[0]:
   Id  ProductId  UserId  ProfileName \
0   1  B001E4KFG0  A3SGXH7AUHU8GW  delmartian
1   2  B00813GRG4  A1D87F6ZCVE5NK  dll pa
2   3  B000LQOCHO  ABXLMWJIXXAIN  Natalia Corres "Natalia Corres"

   HelpfulnessNumerator  HelpfulnessDenominator  Score  Time \
0                      1                      1      1  1303862400
1                      0                      0      0  1346976000
2                      1                      1      1  1219017600

   Summary  Text
0  Good Quality Dog Food  I have bought several of the Vitality canned d...
1  Not as Advertised  Product arrived labeled as Jumbo Salted Peanut...
2  "Delight" says it all  This is a confection that has been around a fe...

```

```

In [0]: display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", con)

```

```

In [0]: print(display.shape)
display.head()

```

(80668, 7)

```

Out[0]:
   UserId  ProductId  ProfileName  Time  Score \
0  #oc-R115TNMSPFT9I7  B007Y59HVM  Breyton  1331510400  2

```

1	#oc-R11D9D7SHXIJB9	B005HG9ETO	Louis E. Emory "hoppy"	1342396800	5
2	#oc-R11DNU2NBKQ23Z	B007Y59HVM	Kim Cieszykowski	1348531200	1
3	#oc-R1105J5ZVQE25C	B005HG9ETO	Penguin Chick	1346889600	5
4	#oc-R12KPB0DL2B5ZD	B0070SBE1U	Christopher P. Presta	1348617600	1

	Text	COUNT(*)
0	Overall its just OK when considering the price...	2
1	My wife has recurring extreme muscle spasms, u...	3
2	This coffee is horrible and unfortunately not ...	2
3	This will be the bottle that you grab from the...	3
4	I didnt like this coffee. Instead of telling y...	2

```
In [0]: display[display['UserId']=='AZY10LLTJ71NX']
```

```
Out[0]:
```

	UserId	ProductId	ProfileName	Time \
80638	AZY10LLTJ71NX	B006P7E5ZI	undertheshrine "undertheshrine"	1334707200

	Score	Text	COUNT(*)
80638	5	I was recommended to try green tea extract to ...	5

```
In [0]: display['COUNT(*)'].sum()
```

```
Out[0]: 393063
```

3 [2] Exploratory Data Analysis

3.1 [2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

```
In [0]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)
display.head()
```

```
Out[0]:
```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator \
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2
3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2
4	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2

	HelpfulnessDenominator	Score	Time \
0	2	5	1199577600

1	2	5	1199577600
2	2	5	1199577600
3	2	5	1199577600
4	2	5	1199577600

	Summary \
0	LOACKER QUADRATINI VANILLA WAFERS
1	LOACKER QUADRATINI VANILLA WAFERS
2	LOACKER QUADRATINI VANILLA WAFERS
3	LOACKER QUADRATINI VANILLA WAFERS
4	LOACKER QUADRATINI VANILLA WAFERS

	Text
0	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
1	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
2	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
3	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
4	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8) ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```
In [0]: #Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False)
```

```
In [0]: #Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first')
final.shape
```

```
Out[0]: (364173, 10)
```

```
In [0]: #Checking to see how much % of data still remains
(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

```
Out[0]: 69.25890143662969
```

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

```

In [0]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)

display.head()

Out[0]:
      Id  ProductId      UserId      ProfileName \
0  64422  B000MIDROQ  A161DK06JJMCYF  J. E. Stephens "Jeanne"
1  44737  B001EQ55RW  A2VOI904FH7ABY                      Ram

      HelpfulnessNumerator  HelpfulnessDenominator  Score      Time \
0                        3                        1      5  1224892800
1                        3                        2      4  1212883200

                        Summary \
0      Bought This for My Son at College
1  Pure cocoa taste with crunchy almonds inside

                        Text
0  My son loves spaghetti so I didn't hesitate or...
1  It was almost a 'love at first bite' - the per...

In [0]: final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]

In [0]: #Before starting the next phase of preprocessing lets see the number of entries left
print(final.shape)

#How many positive and negative reviews are present in our dataset?
final['Score'].value_counts()

(364171, 10)

Out[0]:
1      307061
0       57110
Name: Score, dtype: int64

```

4 [3] Preprocessing

4.1 [3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags

2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```
In [0]: # printing some random reviews
sent_0 = final['Text'].values[0]
print(sent_0)
print("="*50)
```

```
sent_1000 = final['Text'].values[1000]
print(sent_1000)
print("="*50)
```

```
sent_1500 = final['Text'].values[1500]
print(sent_1500)
print("="*50)
```

```
sent_4900 = final['Text'].values[4900]
print(sent_4900)
print("="*50)
```

```
this witty little book makes my son laugh at loud. i recite it in the car as we're driving along
=====
I was really looking forward to these pods based on the reviews. Starbucks is good, but I prefer
=====
Great ingredients although, chicken should have been 1st rather than chicken broth, the only thing
=====
Can't do sugar. Have tried scores of SF Syrups. NONE of them can touch the excellence of this
=====
```

```
In [0]: # https://codereview.stackexchange.com/questions/19663/http-url-validating
i, j = 0 , 0
for test in (final['Text'].values):
    i = i + 1
    if re.match('https?:/(?:www)?(?:[\w-]{2,255}(?:\.\w{2,6}){1,2})(?:/[\w%?#-]{1,300})?'):
        print(test)
        print("index = ",i)
        print("="*50)
        j = j + 1
    if j == 5:
        break;
```

```

http://www.amazon.com/gp/product/B0002XEI90/ref=cm_cr_rev_prod_title<br /><br />So far, the Sp
indes = 18044
=====
http://www.amazon.com/gp/product/B000ER3EFC/ref=cm_cr_rev_prod_img<br /><br />I don't have a b
indes = 53187
=====
http://www.amazon.com/gp/product/B000MVG6Y6/ref=cm_cr_rev_prod_img<br /><br />I have been in U
indes = 104802
=====
http://www.amazon.com/gp/product/B000UZIP92Y/ref=cm_cr_rev_prod_title<br /><br />This product is
indes = 123863
=====
http://www.consumeraffairs.com/news04/2009/06/fda_evangers02.html<br /><br />This company has
indes = 130020
=====

```

```

In [0]: # printing some urls test reviews
        sent_18044 = final['Text'].values[18044 - 1]
        print(sent_18044)
        print("="*50)

        sent_53187 = final['Text'].values[53187 - 1]
        print(sent_53187)
        print("="*50)

        sent_104802 = final['Text'].values[104802 - 1]
        print(sent_104802)
        print("="*50)

        sent_123863 = final['Text'].values[123863 - 1]
        print(sent_123863)
        print("="*50)

```

```

http://www.amazon.com/gp/product/B0002XEI90/ref=cm_cr_rev_prod_title<br /><br />So far, the Sp
=====
http://www.amazon.com/gp/product/B000ER3EFC/ref=cm_cr_rev_prod_img<br /><br />I don't have a b
=====
http://www.amazon.com/gp/product/B000MVG6Y6/ref=cm_cr_rev_prod_img<br /><br />I have been in U
=====
http://www.amazon.com/gp/product/B000UZIP92Y/ref=cm_cr_rev_prod_title<br /><br />This product is
=====

```

```

In [0]: # remove urls from text python: https://stackoverflow.com/a/40823105/4084039
        sent_18044 = re.sub(r"http\S+", "", sent_18044)
        sent_53187= re.sub(r"http\S+", "", sent_53187)
        sent_104802 = re.sub(r"http\S+", "", sent_104802)

```



```

sent_123863 = re.sub(r"http\S+", "", sent_123863)

print(sent_18044)
print("="*40)
print(sent_123863)

/><br />So far, the Sprout Bag is good for growing various legume / bean - Sprouts. I am also
=====
/><br />This product is supposed to be the "Herbal Revive" version but it does not exist in t

In [0]: # https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-remove-all-
        from bs4 import BeautifulSoup

        soup = BeautifulSoup(sent_18044, 'lxml')
        text = soup.get_text()
        print(text)
        print("="*50)

        soup = BeautifulSoup(sent_53187, 'lxml')
        text = soup.get_text()
        print(text)
        print("="*50)

        soup = BeautifulSoup(sent_104802, 'lxml')
        text = soup.get_text()
        print(text)
        print("="*50)

        soup = BeautifulSoup(sent_123863, 'lxml')
        text = soup.get_text()
        print(text)

/>So far, the Sprout Bag is good for growing various legume / bean - Sprouts. I am also using t
=====
/>I don't have a baby, but I am on soft foods for a while after stomach surgery. I ordered some
=====
/>I have been in Utah over 25 years, at an altitude of about 4200 feet. My internist says that
=====
/>This product is supposed to be the "Herbal Revive" version but it does not exist in the Ama

In [0]: # https://stackoverflow.com/a/47091490/4084039
        import re

        def decontracted(phrase):
            # specific
            phrase = re.sub(r"won't", "will not", phrase)
            phrase = re.sub(r"can't", "can not", phrase)

```

```

# general
phrase = re.sub(r"n't", " not", phrase)
phrase = re.sub(r"\ 're", " are", phrase)
phrase = re.sub(r"\ 's", " is", phrase)
phrase = re.sub(r"\ 'd", " would", phrase)
phrase = re.sub(r"\ 'll", " will", phrase)
phrase = re.sub(r"\ 't", " not", phrase)
phrase = re.sub(r"\ 've", " have", phrase)
phrase = re.sub(r"\ 'm", " am", phrase)
return phrase

```

```

In [0]: sent_18044 = decontracted(sent_18044)
print(sent_18044)
print("="*50)
sent_53187 = decontracted(sent_53187)
print(sent_53187)
print("="*50)
sent_123863 = decontracted(sent_123863)
print(sent_123863)
print("="*50)

```

```

/><br />So far, the Sprout Bag is good for growing various legume / bean - Sprouts. I am also
=====
/><br />I do not have a baby, but I am on soft foods for a while after stomach surgery. I ord
=====
/><br />This product is supposed to be the "Herbal Revive" version but it does not exist in t
=====

```

```

In [0]: # https://codereview.stackexchange.com/questions/19663/http-url-validating
i, j = 0 , 0
for test in (final['Text'].values):
    i = i + 1
    if re.match("\S*\d\S*",test):
        print(test)
        print("index = ",i)
        print("="*50)
        j = j + 1
        if j == 5:
            break;

```

```

[[ASIN:B001AGXEAG Beetlejuice (20th Anniversary Deluxe Edition)]<br /><br />Beetlejuice, the T
index = 391
=====
1st day, few stupid flies went in. 2nd day, it smell so so bad, had it relocated... 4th day, h
index = 752
=====
1st: my cats really don't eat this brand with much relish like they did when Wellness had the m

```

```
indes = 1699
```

```
=====
5 star. Our vet recommended this product for our 8 year old male Golden Retriever. He had son
indes = 2259
```

```
=====
24 packs for $19.94 comes to .83 per pack. That is a bargain and does not even consider the s
indes = 2805
=====
```

```
In [0]: #remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
sent = final['Text'].values[2805 - 1]
print(sent)
```

24 packs for \$19.94 comes to .83 per pack. That is a bargain and does not even consider the s

```
In [0]: #remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
sent = re.sub("\S*\d\S*", "", sent).strip()
print(sent)
```

packs for comes to per pack. That is a bargain and does not even consider the subscribe and

```
In [0]: #remove spacial character: https://stackoverflow.com/a/5843547/4084039
sent_18044 = re.sub('[^A-Za-z0-9]+', ' ', sent_18044)
print(sent_18044)
```

br So far the Sprout Bag is good for growing various legume bean Sprouts I am also using the

```
In [0]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
# instead of <br /> if we have <br/> these tags would have revmoved in the 1st step
```

```
stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves',
'you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him',
'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself', 't
'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', "th
'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'ha
'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as
'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through
'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'o
'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any
'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too
's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'n
've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't"
```

```

        "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mightn',
        "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't", 'won',
        "won't", 'wouldn', "wouldn't"])

```

```

In [0]: # Combining all the above stundents
from tqdm import tqdm
preprocessed_reviews = []
# tqdm is for printing the status bar
for sentence in tqdm(final['Text'].values):
    sentence = re.sub(r"http\S+", "", sentence)
    sentence = BeautifulSoup(sentence, 'lxml').get_text()
    sentence = decontracted(sentence)
    sentence = re.sub("\S*\d\S*", "", sentence).strip()
    sentence = re.sub('[^A-Za-z]+', ' ', sentence)
    # https://gist.github.com/sebleier/554280
    sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwords)
    preprocessed_reviews.append(sentence.strip())

```

100%|| 364171/364171 [02:38<00:00, 2298.24it/s]

```

In [0]: preprocessed_reviews[1500]

```

```

Out[0]: 'great ingredients although chicken rather chicken broth thing not think belongs canola

```

[3.2] Preprocessing Review Summary

```

In [0]: ## Similarly you can do preprocessing for review summary also.
        ## Similarly you can do preprocessing for review summary also.
preprocessed_reviews_summary = []
# tqdm is for printing the status bar
for sentence in tqdm(final['Summary'].values):
    sentence = re.sub(r"http\S+", "", sentence)
    sentence = BeautifulSoup(sentence, 'lxml').get_text()
    sentence = decontracted(sentence)
    sentence = re.sub("\S*\d\S*", "", sentence).strip()
    sentence = re.sub('[^A-Za-z]+', ' ', sentence)
    # https://gist.github.com/sebleier/554280
    sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwords)
    preprocessed_reviews_summary.append(sentence.strip())

```

100%|| 364171/364171 [01:46<00:00, 3419.59it/s]

```

In [0]: preprocessed_reviews_summary[1500]

```

```

Out[0]: 'would best canola oil left'

```

```

In [0]: # adding new columns ["CleanedText", "CleanedSummary"] to the data frame and save that
final["CleanedText"] = preprocessed_reviews
final["CleanedSummary"] = preprocessed_reviews_summary
final.tail()

```

```

Out[0]:
      Id  ProductId  UserId  ProfileName \
178145  193174  B009RSR8H0  A4P6AN2L435PV  romarc
173675  188389  B009SF0TN6  A1LOGWGRK4BYPT  Bety Robinson
204727  221795  B009SR40Q2  A32A6X5KCP7ARG  sicamar
5259    5703    B009WSNWC4  AMP7K1084DH1T  ESTY
302474  327601  B009WVB40S  A3ME78KVB31T21  K'la

      HelpfulnessNumerator  HelpfulnessDenominator  Score  Time \
178145                    0                    0      1  1350432000
173675                    0                    0      1  1350518400
204727                    1                    1      1  1350604800
5259                      0                    0      1  1351209600
302474                    0                    0      1  1351123200

      Summary \
178145      LOVE!! LOVE!!
173675  Amazing!! Great sauce for everything!
204727      Awesome Taste
5259      DELICIOUS
302474      Tasty!

      Text \
178145  LOVE, LOVE this sweetener!! I use it in all m...
173675  You have to try this sauce to believe it! It s...
204727  I bought this Hazelnut Paste (Nocciola Spread)...
5259    Purchased this product at a local store in NY ...
302474  I purchased this to send to my son who's away ...

      CleanedText \
178145  love love sweetener use baking unsweetened fla...
173675  try sauce believe starts little sweet honey ta...
204727  bought hazelnut paste nocciola spread local sh...
5259    purchased product local store ny kids love qui...
302474  purchased send son away college delivered righ...

      CleanedSummary
178145      love love
173675  amazing great sauce everything
204727      awesome taste
5259      delicious
302474      tasty

```

```

In [0]: # save the data frame as "preprocessed.pkl"
        final.to_pickle("drive/ai/preprocessed.pkl")

```

5 [4] Featurization

I am used to select 50k data points random and sort them using time based splitting

```

In [0]: # Split the data into train , test and crossvalidation datasets
# load "preprocessed.pkl" data frame
df = pd.read_pickle("drive/ai/preprocessed.pkl")
df.head(1)

Out[0]:
      Id  ProductId      UserId  ProfileName \
138706  150524  0006641040  ACITT7DI6IDDL  shari zychinski

      HelpfulnessNumerator  HelpfulnessDenominator  Score      Time \
138706                  0                      0      1  939340800

      Summary \
138706  EVERY book is educational

      Text \
138706  this witty little book makes my son laugh at l...

      CleanedText \
138706  witty little book makes son laugh loud recite ...

      CleanedSummary
138706  every book educational

In [0]: df.shape

Out[0]: (364171, 12)

In [0]: # take 50k sample data randomly
sample_data = df.sample(50000)
sample_data.shape

Out[0]: (50000, 12)

In [0]: # sorted the data using time based
sorted_data = sample_data.sort_values('Time', axis=0, inplace=False)
sorted_data.shape

Out[0]: (50000, 12)

In [0]: sorted_data['Score'].value_counts()

Out[0]: 1      42277
        0       7723
        Name: Score, dtype: int64

In [0]: from sklearn.model_selection import train_test_split

In [0]: X = np.array(sorted_data['CleanedText'])
        y = np.array(sorted_data['Score'])
        print(X.shape)
        print(y.shape)

```

```
(50000,)  
(50000,)
```

```
In [0]: # Simple cross validation  
# split the data sent into train and test  
train , test , train_y , test_y = train_test_split(X, y, test_size = 0.3, random_state=42)  
  
# split the train data set into cross validation train and cross validation test  
train, cv , train_y, cv_y = train_test_split(train, train_y, test_size=0.3, random_state=42)  
  
print("train data = ", train.shape)  
print("cross validation = ", cv.shape)  
print("test data = ", test.shape)
```

```
train data = (24500,)  
cross validation = (10500,)  
test data = (15000,)
```

```
In [0]:
```

5.1 [4.1] BAG OF WORDS

```
In [0]: #BoW  
count_vect = CountVectorizer() #in scikit-learn  
count_vect.fit(train)  
print("some feature names ", count_vect.get_feature_names()[:10])  
print('='*50)  
  
final_counts_train = count_vect.fit_transform(train)  
final_counts_cv = count_vect.transform(cv)  
final_counts_test = count_vect.transform(test)  
print("=====Train Data=====")  
print("the type of count vectorizer ",type(final_counts_train))  
print("the shape of out text BOW vectorizer ",final_counts_train.get_shape())  
print("the number of unique words ", final_counts_train.get_shape()[1])  
print("=====Cross validation Data=====")  
print("the type of count vectorizer ",type(final_counts_cv))  
print("the shape of out text BOW vectorizer ",final_counts_cv.get_shape())  
print("the number of unique words ", final_counts_cv.get_shape()[1])  
print("=====Test Data=====")  
print("the type of count vectorizer ",type(final_counts_test))  
print("the shape of out text BOW vectorizer ",final_counts_test.get_shape())  
print("the number of unique words ", final_counts_test.get_shape()[1])  
  
some feature names ['aa', 'aaa', 'aaaaaaarrrrrggghhh', 'aaah', 'aaahhs', 'aabsolutely', 'aaf  
=====
```

```

the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (24500, 30806)
the number of unique words 30806
=====Cross validation Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (10500, 30806)
the number of unique words 30806
=====Test Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (15000, 30806)
the number of unique words 30806

```

```

In [0]: # the time of bagofwords is sparse matrix now convert to lower dimensions
        # Dimensionality reduction using truncated SVD
        from sklearn.decomposition import TruncatedSVD

        bagofwords_svd = TruncatedSVD(n_components= 500) #Desired dimensionality of output data

        bow_train = bagofwords_svd.fit_transform(final_counts_train)
        print('Explained Variance Ratio is = %0.2f%%' % (bagofwords_svd.explained_variance_ratio_))
        bow_cv = bagofwords_svd.fit_transform(final_counts_cv)
        bow_test = bagofwords_svd.fit_transform(final_counts_test)
        print("trian data shape = ", bow_train.shape)
        print("CV data shape = ", bow_cv.shape)
        print("test data shape = ", bow_test.shape)

Explained Variance Ratio is = 67.43%
trian data shape = (24500, 500)
CV data shape = (10500, 500)
test data shape = (15000, 500)

```

```

In [0]: # save the datasets as numpy array
        bow_train = np.array(bow_train)
        bow_cv = np.array(bow_cv)
        bow_test = np.array(bow_test)

In [0]: np.save("drive/ai/bow_train.npy", bow_train)
        np.save("drive/ai/bow_train_y.npy", train_y)
        np.save("drive/ai/bow_cv.npy", bow_cv)
        np.save("drive/ai/bow_cv_y.npy", cv_y)
        np.save("drive/ai/bow_test.npy", bow_test)
        np.save("drive/ai/bow_test_y.npy", test_y)

```

```

In [0]: ## for kd_tree

```

5.2 for kd_tree

```

In [0]: # the time of bagofwords is sparse matrix now convert to lower dimensions
        # Dimensionality reduction using truncated SVD

```



```

from sklearn.decomposition import TruncatedSVD

bagofwords_svd = TruncatedSVD(n_components= 200) #Desired dimensionality of output data

bow_train = bagofwords_svd.fit_transform(final_counts_train)
print('Explained Variance Ratio is = %0.2f%%' % (bagofwords_svd.explained_variance_ratio_))
bow_cv = bagofwords_svd.fit_transform(final_counts_cv)
bow_test = bagofwords_svd.fit_transform(final_counts_test)
print("trian data shape = ", bow_train.shape)
print("CV data shape = ", bow_cv.shape)
print("test data shape = ", bow_test.shape)

```

```

Explained Variance Ratio is = 50.60%
trian data shape = (24500, 200)
CV data shape = (10500, 200)
test data shape = (15000, 200)

```

```

In [0]: # save the datasets as numpy array
        bow_train = np.array(bow_train)
        bow_cv = np.array(bow_cv)
        bow_test = np.array(bow_test)

In [0]: np.save("drive/ai/kd_bow_train.npy", bow_train)
        np.save("drive/ai/kd_bow_train_y.npy", train_y)
        np.save("drive/ai/kd_bow_cv.npy", bow_cv)
        np.save("drive/ai/kd_bow_cv_y.npy", cv_y)
        np.save("drive/ai/kd_bow_test.npy", bow_test)
        np.save("drive/ai/kd_bow_test_y.npy", test_y)

```

5.3 [4.2] Bi-Grams and n-Grams.

```

In [0]: #bi-gram, tri-gram and n-gram

#removing stop words like "not" should be avoided before building n-grams
# count_vect = CountVectorizer(ngram_range=(1,2))
# please do read the CountVectorizer documentation http://scikit-learn.org/stable/modu

# you can choose these numebrs min_df=10, max_features=5000, of your choice
count_vect = CountVectorizer(ngram_range=(1,2), min_df=10, max_features=5000)
final_bigram_counts = count_vect.fit_transform(preprocessed_reviews)
print("the type of count vectorizer ", type(final_bigram_counts))
print("the shape of out text BOW vectorizer ", final_bigram_counts.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_bigram_counts.get_shape()[0])

the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (364171, 5000)
the number of unique words including both unigrams and bigrams 5000

```

```
In [0]: tf_idf_vect = TfidfVectorizer(ngram_range=(1,2), min_df=10)
tf_idf_vect.fit(train)
print("some sample features(unique words in the corpus)",tf_idf_vect.get_feature_names)
print('='*50)
print("=====Train Data=====")
final_tf_idf_train = tf_idf_vect.fit_transform(train)
print("the type of count vectorizer ",type(final_tf_idf_train))
print("the shape of out text TFIDF vectorizer ",final_tf_idf_train.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_tf_idf_train.get_shape()[0])
print("=====CV Data=====")
final_tf_idf_cv = tf_idf_vect.transform(cv)
print("the type of count vectorizer ",type(final_tf_idf_cv))
print("the shape of out text TFIDF vectorizer ",final_tf_idf_cv.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_tf_idf_cv.get_shape()[0])
print("=====Test Data=====")
final_tf_idf_test = tf_idf_vect.transform(test)
print("the type of count vectorizer ",type(final_tf_idf_test))
print("the shape of out text TFIDF vectorizer ",final_tf_idf_test.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_tf_idf_test.get_shape()[0])
```

```
In [0]: # the time of bagofwords is sparse matrix now convert to lower dimensions
# Dimensionality reduction using truncated SVD

tfidf_svd = TruncatedSVD(n_components= 500) #Desired dimensionality of output data

tfidf_train = tfidf_svd.fit_transform(final_tf_idf_train)
print('Explained Variance Ratio is = %0.2f%%' % (tfidf_svd.explained_variance_ratio_.sum))
tfidf_cv = bagofwords_svd.fit_transform(final_tf_idf_cv)
tfidf_test = bagofwords_svd.fit_transform(final_tf_idf_test)
print("train data shape = ", tfidf_train.shape)
```

```
print("CV data shape = ", tfidf_cv.shape)
print("test data shape = ", tfidf_test.shape)
```

Explained Variance Ratio is = 34.12%
 trian data shape = (24500, 500)
 CV data shape = (10500, 500)
 test data shape = (15000, 500)

```
In [0]: # save the datasets as numpy array
        tfidf_train = np.array(tfidf_train)
        tfidf_cv = np.array(tfidf_cv)
        tfidf_test = np.array(tfidf_test)

In [0]: np.save("drive/ai/tfidf_train.npy", tfidf_train)
        np.save("drive/ai/tfidf_train_y.npy", train_y)
        np.save("drive/ai/tfidf_cv.npy", tfidf_cv)
        np.save("drive/ai/tfidf_cv_y.npy", cv_y)
        np.save("drive/ai/tfidf_test.npy", tfidf_test)
        np.save("drive/ai/tfidf_test_y.npy", test_y)
```

5.5 for kd_tree

```
In [0]: # the time of bagofwords is sparse matrix now convert to lower dimensions
        # Dimensionality reduction using truncated SVD
```

```
tfidf_svd = TruncatedSVD(n_components= 200) #Desired dimensionality of output data

tfidf_train = tfidf_svd.fit_transform(final_tf_idf_train)
print('Explained Variance Ratio is = %0.2f%%' % (tfidf_svd.explained_variance_ratio_.sum() * 100))
tfidf_cv = bagofwords_svd.fit_transform(final_tf_idf_cv)
tfidf_test = bagofwords_svd.fit_transform(final_tf_idf_test)
print("trian data shape = ", tfidf_train.shape)
print("CV data shape = ", tfidf_cv.shape)
print("test data shape = ", tfidf_test.shape)
```

Explained Variance Ratio is = 20.69%
 trian data shape = (24500, 200)
 CV data shape = (10500, 200)
 test data shape = (15000, 200)

```
In [0]: # save the datasets as numpy array
        tfidf_train = np.array(tfidf_train)
        tfidf_cv = np.array(tfidf_cv)
        tfidf_test = np.array(tfidf_test)

In [0]: np.save("drive/ai/kd_tfidf_train.npy", tfidf_train)
        np.save("drive/ai/kd_tfidf_train_y.npy", train_y)
```

```

np.save("drive/ai/kd_tfidf_cv.npy", tfidf_cv)
np.save("drive/ai/kd_tfidf_cv_y.npy", cv_y)
np.save("drive/ai/kd_tfidf_test.npy", tfidf_test)
np.save("drive/ai/kd_tfidf_test_y.npy", test_y)

```

5.6 [4.4] Word2Vec

```

In [0]: # Train your own Word2Vec model using your own text corpus
        ##### Train Set #####
        i=0
        list_of_train_sentence=[]
        for sentence in train:
            list_of_train_sentence.append(sentence.split())
        ##### CV Set #####
        i=0
        list_of_cv_sentence=[]
        for sentence in cv:
            list_of_cv_sentence.append(sentence.split())
        ##### Test Set #####
        i=0
        list_of_test_sentence=[]
        for sentence in test:
            list_of_test_sentence.append(sentence.split())
        print("Length of Train = ", len(list_of_train_sentence))
        print("Length of CV = ", len(list_of_cv_sentence))
        print("Length of Test = ", len(list_of_test_sentence))

```

```

Length of Train = 24500
Length of CV = 10500
Length of Test = 15000

```

```

In [0]: # Using Google News Word2Vectors

        # in this project we are using a pretrained model by google
        # its 3.3G file, once you load this into your memory
        # it occupies ~9Gb, so please do this step only if you have >12G of ram
        # we will provide a pickle file wich contains a dict ,
        # and it contains all our courpus words as keys and model[word] as values
        # To use this code-snippet, download "GoogleNews-vectors-negative300.bin"
        # from https://drive.google.com/file/d/0B7XkCwpI5KDYNlNUTTlSS21pQmM/edit
        # it's 1.9GB in size.

        # http://kavita-ganesan.com/gensim-word2vec-tutorial-starter-code/#.W17SRFAzZPY
        # you can comment this whole cell
        # or change these variable according to your need

```



```

In [0]: ##### Train data #####
        # average Word2Vec
        # compute average word2vec for each review.
sent_vectors_train = []; # the avg-w2v for each sentence/review is stored in this list
for sent in tqdm(list_of_train_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need to
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors_train.append(sent_vec)
print(len(sent_vectors_train))
print(len(sent_vectors_train[0]))

100%|| 24500/24500 [00:44<00:00, 551.45it/s]

24500
50

```

```

In [0]: ##### CV data #####
        # average Word2Vec
        # compute average word2vec for each review.
sent_vectors_cv = []; # the avg-w2v for each sentence/review is stored in this list
for sent in tqdm(list_of_cv_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need to
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors_cv.append(sent_vec)
print(len(sent_vectors_cv))
print(len(sent_vectors_cv[0]))

100%|| 10500/10500 [00:19<00:00, 536.26it/s]

10500
50

```

```

In [0]: ##### Test data #####
        # average Word2Vec
        # compute average word2vec for each review.
sent_vectors_test = []; # the avg-w2v for each sentence/review is stored in this list
for sent in tqdm(list_of_test_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need to
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors_test.append(sent_vec)
print(len(sent_vectors_test))
print(len(sent_vectors_test[0]))

```

100%|| 15000/15000 [00:27<00:00, 551.77it/s]

15000

50

```

In [0]: # save the datasets as numpy array
w2v_train = np.array(sent_vectors_train)
w2v_cv = np.array(sent_vectors_cv)
w2v_test = np.array(sent_vectors_test)

In [0]: np.save("drive/ai/w2v_train.npy", w2v_train)
np.save("drive/ai/w2v_train_y.npy", train_y)
np.save("drive/ai/w2v_cv.npy", w2v_cv)
np.save("drive/ai/w2v_cv_y.npy", cv_y)
np.save("drive/ai/w2v_test.npy", w2v_test)
np.save("drive/ai/w2v_test_y.npy", test_y)

```

[4.4.1.2] TFIDF weighted W2v

```

In [0]: model = TfidfVectorizer()
        #tf_idf_matrix = model.fit_transform(train)

        print("=====Train Data=====")
        final_tf_idf_train = model.fit_transform(train)

```

```

print("the type of count vectorizer ",type(final_tf_idf_train))
print("the shape of out text TFIDF vectorizer ",final_tf_idf_train.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_tf_idf_train.get_shape()[1])
print("=====CV Data=====")
final_tf_idf_cv = model.transform(cv)
print("the type of count vectorizer ",type(final_tf_idf_cv))
print("the shape of out text TFIDF vectorizer ",final_tf_idf_cv.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_tf_idf_cv.get_shape()[1])
print("=====Test Data=====")
final_tf_idf_test = model.transform(test)
print("the type of count vectorizer ",type(final_tf_idf_test))
print("the shape of out text TFIDF vectorizer ",final_tf_idf_test.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_tf_idf_test.get_shape()[1])

# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

```

```

=====Train Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (24500, 30961)
the number of unique words including both unigrams and bigrams 30961
=====CV Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (10500, 30961)
the number of unique words including both unigrams and bigrams 30961
=====Test Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (15000, 30961)
the number of unique words including both unigrams and bigrams 30961

```

```

In [0]: ##### Train #####
# TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

train_tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in train_tfidf_sent_vectors
row=0;
for sent in tqdm(list_of_train_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            #tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf valeus of word in this review

```



```

        tf_idf = dictionary[word]*(sent.count(word)/len(sent))
        sent_vec += (vec * tf_idf)
        weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    train_tfidf_sent_vectors.append(sent_vec)
    row += 1

```

100%|| 24500/24500 [08:09<00:00, 50.06it/s]

In [0]: ##### CV #####

```

# TF-IDF weighted Word2Vec
#tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

cv_tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this
row=0;
for sent in tqdm(list_of_cv_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            #tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    cv_tfidf_sent_vectors.append(sent_vec)
    row += 1

```

100%|| 10500/10500 [03:33<00:00, 49.12it/s]

In [0]: ##### Train #####

```

# TF-IDF weighted Word2Vec
#tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

test_tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in th
row=0;
for sent in tqdm(list_of_test_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/review

```

```

for word in sent: # for each word in a review/sentence
    if word in w2v_words and word in tfidf_feat:
        vec = w2v_model.wv[word]
        #tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
        # to reduce the computation we are
        # dictionary[word] = idf value of word in whole corpus
        # sent.count(word) = tf value of word in this review
        tf_idf = dictionary[word]*(sent.count(word)/len(sent))
        sent_vec += (vec * tf_idf)
        weight_sum += tf_idf
if weight_sum != 0:
    sent_vec /= weight_sum
test_tfidf_sent_vectors.append(sent_vec)
row += 1

```

100%|| 15000/15000 [04:52<00:00, 51.23it/s]

```

In [0]: # save the datasets as numpy array
tfidf_w2v_train = np.array(train_tfidf_sent_vectors)
tfidf_w2v_cv = np.array(cv_tfidf_sent_vectors)
tfidf_w2v_test = np.array(test_tfidf_sent_vectors)

In [0]: np.save("drive/ai/w2v_tfidf_train.npy", tfidf_w2v_train)
np.save("drive/ai/w2v_tfidf_train_y.npy", train_y)
np.save("drive/ai/w2v_tfidf_cv.npy", tfidf_w2v_cv)
np.save("drive/ai/w2v_tfidf_cv_y.npy", cv_y)
np.save("drive/ai/w2v_tfidf_test.npy", tfidf_w2v_test)
np.save("drive/ai/w2v_tfidf_test_y.npy", test_y)

```

6 [5] Assignment 3: KNN

Apply Knn(brute force version) on these feature sets

SET 1:Review text, preprocessed one converted into vectors using (BOW)

SET 2:Review text, preprocessed one converted into vectors using (TFIDF)

SET 3:Review text, preprocessed one converted into vectors using (AVG W2v)

SET 4:Review text, preprocessed one converted into vectors using (TFIDF W2v)

Apply Knn(kd tree version) on these feature sets NOTE: sklearn implementation of kd-tree accepts only dense matrices, you need to convert the sparse matrices of CountVectorizer/TfidfVectorizer into dense matrices. You can convert sparse matrices to dense using .toarray() attribute. For more information please visit this link

SET 5:Review text, preprocessed one converted into vectors using (BOW) but with restriction on maximum features generated.

SET 6:Review text, preprocessed one converted into vectors

<pre>

```
tf_idf_vect = TfidfVectorizer(min_df=10, max_features=500)
```

```

        tf_idf_vect.fit(preprocessed_reviews)
    </pre>
    </li>
    <li><font color='red'>SET 3:</font>Review text, preprocessed one converted into vectors</li>
    <li><font color='red'>SET 4:</font>Review text, preprocessed one converted into vectors</li>
</ul>
</li>
<br>
<li><strong>The hyper paramter tuning(find best K)</strong>
    <ul>
    <li>Find the best hyper parameter which will give the maximum <a href='https://www.appliedaicom'></a>
    <li>Find the best hyper paramter using k-fold cross validation or simple cross validation data
    <li>Use gridsearch cv or randomsearch cv or you can also write your own for loops to do this task
    </ul>
</li>
<br>
<li>
<strong>Representation of results</strong>
    <ul>
    <li>You need to plot the performance of model both on train data and cross validation data for
    <img src='train_cv_auc.JPG' width=300px></li>
    <li>Once after you found the best hyper parameter, you need to train your model with it, and find
    <img src='train_test_auc.JPG' width=300px></li>
    <li>Along with plotting ROC curve, you need to print the <a href='https://www.appliedaicom'></a>
    <img src='confusion_matrix.png' width=300px></li>
    </ul>
</li>
<br>
<li><strong>Conclusion</strong>
    <ul>
    <li>You need to summarize the results at the end of the notebook, summarize it in the table for
    <img src='summary.JPG' width=400px>
    </li>
    </ul>

```

Note: Data Leakage

1. There will be an issue of data-leakage if you vectorize the entire data and then split it into train/cv/test.
2. To avoid the issue of data-leakage, make sure to split your data first and then vectorize it.
3. While vectorizing your data, apply the method fit_transform() on your train data, and apply the method transform() on cv/test data.
4. For more details please go through this link.

6.1 [5.1] Applying KNN brute force

6.1.1 [5.1.1] Applying KNN brute force on BOW, SET 1

In [0]:

```
In [0]: from sklearn.neighbors import KNeighborsClassifier
        from sklearn.metrics import accuracy_score , f1_score , confusion_matrix
        from collections import Counter
        from sklearn.metrics import accuracy_score, roc_auc_score , roc_curve
        from sklearn.model_selection import train_test_split
```

```
In [0]: # Please write all the code with proper documentation
        # load Review text, preprocessed one converted into vectors using (BOW)
        bow_train = np.load("drive/ai/bow_train.npy")
        bow_train_y = np.load("drive/ai/bow_train_y.npy")
        bow_cv = np.load("drive/ai/bow_cv.npy")
        bow_cv_y = np.load("drive/ai/bow_cv_y.npy")
        bow_test = np.load("drive/ai/bow_test.npy")
        bow_test_y = np.load("drive/ai/bow_test_y.npy")
```

```
In [0]: print("Shape of Train = ", bow_train.shape)
        print("Shape of CV = ", bow_cv.shape)
        print("Shape of Test = ", bow_test.shape)
```

Shape of Train = (24500, 500)

Shape of CV = (10500, 500)

Shape of Test = (15000, 500)

```
In [0]: bow_train_auc = []
        bow_cv_auc = []
        for i in range(1,10,2):
            knn = KNeighborsClassifier(n_neighbors=i)
            knn.fit(bow_train, bow_train_y)
            # train data
            y_prob_train = knn.predict_proba(bow_train)[: ,1]
            y_pred = np.where(y_prob_train > 0.5, 1, 0)
            auc_roc_test = roc_auc_score(bow_train_y , y_prob_train)
            print('\nTrain AUC for k = %d is %0.2f%%' % (i, (auc_roc_test * float(100))))
            bow_train_auc.append(auc_roc_test)
            # CV
            y_prob_cv = knn.predict_proba(bow_cv)[: ,1]
            y_pred = np.where(y_prob_cv > 0.5, 1, 0)
            auc_roc_cv = roc_auc_score(bow_cv_y , y_prob_cv)
            print('\nCV AUC for k = %d is %0.2f%%' % (i, (auc_roc_cv * float(100))))
            bow_cv_auc.append(auc_roc_cv)
```

Train AUC for k = 1 is 99.96%

CV AUC for k = 1 is 53.47%

Train AUC for k = 3 is 92.31%

CV AUC for k = 3 is 55.03%

Train AUC for k = 5 is 88.99%

CV AUC for k = 5 is 55.29%

Train AUC for k = 7 is 87.13%

CV AUC for k = 7 is 55.61%

Train AUC for k = 9 is 86.04%

CV AUC for k = 9 is 56.72%

```
In [0]: k_values = [j for j in range(1,10,2)]
```

```
In [0]: # https://www.dataquest.io/blog/learning-curves-machine-learning/
```

```
import matplotlib.pyplot as plt
%matplotlib inline


plt.style.use('seaborn')

plt.plot(k_values, bow_train_auc, label = 'Training AUC')
plt.plot(k_values, bow_cv_auc, label = 'Validation AUC')

plt.ylabel('AUC', fontsize = 14)
plt.xlabel('k-values', fontsize = 14)
plt.title('Learning curves for a k-NN model', fontsize = 18, y = 1.03)
plt.legend()
```

```
Out[0]: <matplotlib.legend.Legend at 0x7f2bd83f5748>
```

Learning curves for a k-NN model



```
In [0]: knn = KNeighborsClassifier(n_neighbors=9)
knn.fit(bow_train, bow_train_y)
# train data
y_prob_train = knn.predict_proba(bow_train)[:,:1]
fpr, tpr, threshold = roc_curve(bow_train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(bow_train_y, y_prob_train)
print('\nTrain AUC for k = %d is %0.2f%%' % (9, (auc_roc_train * float(100))))
# CV
y_prob_cv = knn.predict_proba(bow_cv)[:,:1]
fpr, tpr, threshold = roc_curve(bow_cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(bow_cv_y, y_prob_cv)
print('\nCV AUC for k = %d is %0.2f%%' % (9, (auc_roc_cv * float(100))))
# Test
y_prob_test = knn.predict_proba(bow_test)[:,:1]
fpr, tpr, threshold = roc_curve(bow_test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
auc_roc_test = roc_auc_score(bow_test_y, y_prob_test)
print('\nTest AUC for k = %d is %0.2f%%' % (9, (auc_roc_test * float(100))))
print("="*50)
```

Train AUC for k = 9 is 86.04%


CV AUC for k = 9 is 56.72%

Test AUC for k = 9 is 52.45%

=====

In [0]: # <https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python>

```
import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b', label='AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y', label='AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label='AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```




```
In [0]: print("F1-Score on test set: %0.2f"%(f1_score(bow_test_y, y_pred_test)))
```

F1-Score on test set: 0.88

```
In [0]: df_cm = pd.DataFrame(confusion_matrix(bow_test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[0]: <matplotlib.axes._subplots.AxesSubplot at 0x7f2bd5b23198>
```



6.1.2 [5.1.2] Applying KNN brute force on TFIDE, SET 2

```
In [0]: # Please write all the code with proper documentation
# Please write all the code with proper documentation
# load Review text, preprocessed one converted into vectors using (BOW)
tfidf_train = np.load("drive/ai/tfidf_train.npy")
tfidf_train_y = np.load("drive/ai/tfidf_train_y.npy")
tfidf_cv = np.load("drive/ai/tfidf_cv.npy")
tfidf_cv_y = np.load("drive/ai/tfidf_cv_y.npy")
tfidf_test = np.load("drive/ai/tfidf_test.npy")
tfidf_test_y = np.load("drive/ai/tfidf_test_y.npy")
```



```
In [0]: print("Shape of Train = ", tfidf_train.shape)
        print("Shape of CV = ", tfidf_cv.shape)
        print("Shape of Test = ", tfidf_test.shape)
```

```
Shape of Train = (24500, 500)
Shape of CV = (10500, 500)
Shape of Test = (15000, 500)
```

```
In [0]: tfidf_train_auc = []
        tfidf_cv_auc = []
        for i in range(1,10,2):
            knn = KNeighborsClassifier(n_neighbors=i)
            knn.fit(tfidf_train, tfidf_train_y)
            # train data
            y_prob_train = knn.predict_proba(tfidf_train)[: ,1]
            y_pred = np.where(y_prob_train > 0.5, 1, 0)
            auc_roc_test = roc_auc_score(tfidf_train_y , y_prob_train)
            print('\nTrain AUC for k = %d is %0.2f%%' % (i, (auc_roc_test * float(100))))
            tfidf_train_auc.append(auc_roc_test)
            # CV
            y_prob_cv = knn.predict_proba(tfidf_cv)[: ,1]
            y_pred = np.where(y_prob_cv > 0.5, 1, 0)
            auc_roc_cv = roc_auc_score(tfidf_cv_y , y_prob_cv)
            print('\nCV AUC for k = %d is %0.2f%%' % (i, (auc_roc_cv * float(100))))
            tfidf_cv_auc.append(auc_roc_cv)
            print("="*50)
```

```
Train AUC for k = 1 is 99.96%
```

```
CV AUC for k = 1 is 49.98%
```

```
=====
```

```
Train AUC for k = 3 is 91.10%
```

```
CV AUC for k = 3 is 50.77%
```

```
=====
```

```
Train AUC for k = 5 is 83.34%
```

```
CV AUC for k = 5 is 50.96%
```

```
=====
```

```
Train AUC for k = 7 is 80.95%
```

```
CV AUC for k = 7 is 51.01%
```

```
=====
```

Train AUC for k = 9 is 79.26%

CV AUC for k = 9 is 51.53%

=====

```
In [0]: k_values = [j for j in range(1,10,2)]
```

```
In [0]: # https://www.dataquest.io/blog/learning-curves-machine-learning/
```

```
import matplotlib.pyplot as plt
```

```
%matplotlib inline
```

```
plt.style.use('seaborn')
```

```
plt.plot(k_values, tfidf_train_auc, label = 'Training AUC')
```

```
plt.plot(k_values, tfidf_cv_auc, label = 'Validation AUC')
```


```
plt.ylabel('AUC', fontsize = 14)
```

```
plt.xlabel('k-values', fontsize = 14)
```

```
plt.title('Learning curves for a k-NN model', fontsize = 18, y = 1.03)
```

```
plt.legend()
```

```
Out[0]: <matplotlib.legend.Legend at 0x7fcf996ce860>
```



```

In [0]: knn = KNeighborsClassifier(n_neighbors=9)
knn.fit(tfidf_train, tfidf_train_y)
# train data
y_prob_train = knn.predict_proba(tfidf_train)[: ,1]
fprt, tprt , thresholdt = roc_curve(tfidf_train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(tfidf_train_y , y_prob_train)
print('\nTrain AUC for k = %d is %0.2f%%' % (9, (auc_roc_train * float(100))))
# CV
y_prob_cv = knn.predict_proba(tfidf_cv)[: ,1]
fprc, tprc , thresholdc = roc_curve(tfidf_cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(tfidf_cv_y , y_prob_cv)
print('\nCV AUC for k = %d is %0.2f%%' % (9, (auc_roc_cv * float(100))))
# Test
y_prob_test = knn.predict_proba(tfidf_test)[: ,1]
fprts, tprts , thresholdts = roc_curve(tfidf_test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
auc_roc_test = roc_auc_score(tfidf_test_y , y_prob_test)
print('\nTest AUC for k = %d is %0.2f%%' % (9, (auc_roc_test * float(100))))
print("="*50)

```

Train AUC for k = 9 is 79.26%

CV AUC for k = 97 is 51.53%

Test AUC for k = 97 is 51.78%


=====

In [0]: # <https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python>

```

import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b' , label='AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y' , label='AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label='AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

```




```
In [0]: print("F1-Score on test set: %0.2f"%(f1_score(tfidf_test_y, y_pred_test)))
```

F1-Score on test set: 0.91

```
In [0]: df_cm = pd.DataFrame(confusion_matrix(tfidf_test_y, y_pred_test), range(2), range(2))  
sns.set(font_scale=1.4)  
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[0]: <matplotlib.axes._subplots.AxesSubplot at 0x7fcf96b31358>
```



6.1.3 [5.1.3] Applying KNN brute force on AVG W2V, SET 3

In [0]: *# Please write all the code with proper documentation*

```
In [0]: # Please write all the code with proper documentation
        # load Review text, preprocessed one converted into vectors using (w2v)
w2v_train = np.load("drive/ai/w2v_train.npy")
w2v_train_y = np.load("drive/ai/w2v_train_y.npy")
w2v_cv = np.load("drive/ai/w2v_cv.npy")
w2v_cv_y = np.load("drive/ai/w2v_cv_y.npy")
w2v_test = np.load("drive/ai/w2v_test.npy")
w2v_test_y = np.load("drive/ai/w2v_test_y.npy")
```

```
In [0]: print("Shape of Train = ", w2v_train.shape)
        print("Shape of CV = ", w2v_cv.shape)
        print("Shape of Test = ", w2v_test.shape)
```

Shape of Train = (24500, 50)

Shape of CV = (10500, 50)

Shape of Test = (15000, 50)

```

In [0]: w2v_train_auc = []
        w2v_cv_auc = []
        for i in range(1,40,2):
            knn = KNeighborsClassifier(n_neighbors=i)
            knn.fit(w2v_train, w2v_train_y)
            # train data
            y_prob_train = knn.predict_proba(w2v_train)[:,:1]
            y_pred = np.where(y_prob_train > 0.5, 1, 0)
            auc_roc_test = roc_auc_score(w2v_train_y , y_prob_train)
            print('\nTrain AUC for k = %d is %0.2f%%' % (i, (auc_roc_test * float(100))))
            w2v_train_auc.append(auc_roc_test)
            # CV
            y_prob_cv = knn.predict_proba(w2v_cv)[:,:1]
            y_pred = np.where(y_prob_cv > 0.5, 1, 0)
            auc_roc_cv = roc_auc_score(w2v_cv_y , y_prob_cv)
            print('\nCV AUC for k = %d is %0.2f%%' % (i, (auc_roc_cv * float(100))))
            w2v_cv_auc.append(auc_roc_cv)
        print("="*50)

```

Train AUC for k = 1 is 99.96%

CV AUC for k = 1 is 66.50%

=====

Train AUC for k = 3 is 95.18%

CV AUC for k = 3 is 75.20%

=====

Train AUC for k = 5 is 93.10%

CV AUC for k = 5 is 79.20%

=====

Train AUC for k = 7 is 92.02%

CV AUC for k = 7 is 81.47%

=====

Train AUC for k = 9 is 91.48%

CV AUC for k = 9 is 82.27%

=====

Train AUC for k = 11 is 91.01%

CV AUC for k = 11 is 82.93%

```

=====

Train AUC for k = 13 is 90.63%

CV AUC for k = 13 is 83.80%
=====

Train AUC for k = 15 is 90.24%

CV AUC for k = 15 is 84.16%
=====

Train AUC for k = 17 is 89.95%

CV AUC for k = 17 is 84.47%
=====

Train AUC for k = 19 is 89.69%

CV AUC for k = 19 is 84.94%
=====

Train AUC for k = 21 is 89.50%

CV AUC for k = 21 is 85.12%
=====

Train AUC for k = 23 is 89.32%

CV AUC for k = 23 is 85.38%
=====

Train AUC for k = 25 is 89.12%

CV AUC for k = 25 is 85.49%
=====

Train AUC for k = 27 is 89.02%

CV AUC for k = 27 is 85.62%
=====

Train AUC for k = 29 is 88.97%

CV AUC for k = 29 is 85.66%
=====

Train AUC for k = 31 is 88.94%

```

CV AUC for k = 31 is 85.79%
=====

Train AUC for k = 33 is 88.93%

CV AUC for k = 33 is 85.81%
=====

Train AUC for k = 35 is 88.90%

CV AUC for k = 35 is 85.92%
=====

Train AUC for k = 37 is 88.80%

CV AUC for k = 37 is 86.11%
=====

Train AUC for k = 39 is 88.77%

CV AUC for k = 39 is 86.24%
=====

```
In [0]: for i in range(41,50,2):
        knn = KNeighborsClassifier(n_neighbors=i)
        knn.fit(w2v_train, w2v_train_y)
        # train data
        y_prob_train = knn.predict_proba(w2v_train)[: ,1]
        y_pred = np.where(y_prob_train > 0.5, 1, 0)
        auc_roc_test = roc_auc_score(w2v_train_y , y_prob_train)
        print('\nTrain AUC for k = %d is %0.2f%%' % (i, (auc_roc_test * float(100))))
        w2v_train_auc.append(auc_roc_test)
        # CV
        y_prob_cv = knn.predict_proba(w2v_cv)[: ,1]
        y_pred = np.where(y_prob_cv > 0.5, 1, 0)
        auc_roc_cv = roc_auc_score(w2v_cv_y , y_prob_cv)
        print('\nCV AUC for k = %d is %0.2f%%' % (i, (auc_roc_cv * float(100))))
        w2v_cv_auc.append(auc_roc_cv)
        print("="*50)
```

Train AUC for k = 41 is 88.68%

CV AUC for k = 41 is 86.31%
=====

Train AUC for k = 43 is 88.57%

CV AUC for k = 43 is 86.35%

=====

Train AUC for k = 45 is 88.50%

CV AUC for k = 45 is 86.36%

=====

Train AUC for k = 47 is 88.44%

CV AUC for k = 47 is 86.45%

=====

Train AUC for k = 49 is 88.37%

CV AUC for k = 49 is 86.47%

=====

```
In [0]: for i in range(51,70,2):
        knn = KNeighborsClassifier(n_neighbors=i)
        knn.fit(w2v_train, w2v_train_y)
        # train data
        y_prob_train = knn.predict_proba(w2v_train)[: ,1]
        y_pred = np.where(y_prob_train > 0.5, 1, 0)
        auc_roc_test = roc_auc_score(w2v_train_y , y_prob_train)
        print('\nTrain AUC for k = %d is %0.2f%%' % (i, (auc_roc_test * float(100))))
        w2v_train_auc.append(auc_roc_test)
        # CV
        y_prob_cv = knn.predict_proba(w2v_cv)[: ,1]
        y_pred = np.where(y_prob_cv > 0.5, 1, 0)
        auc_roc_cv = roc_auc_score(w2v_cv_y , y_prob_cv)
        print('\nCV AUC for k = %d is %0.2f%%' % (i, (auc_roc_cv * float(100))))
        w2v_cv_auc.append(auc_roc_cv)
        print("="*50)
```

Train AUC for k = 51 is 88.36%

CV AUC for k = 51 is 86.48%

=====

Train AUC for k = 53 is 88.32%

CV AUC for k = 53 is 86.48%

=====

```

Train AUC for k = 55 is 88.30%

CV AUC for k = 55 is 86.56%
=====

Train AUC for k = 57 is 88.17%

CV AUC for k = 57 is 86.56%
=====

Train AUC for k = 59 is 88.13%

CV AUC for k = 59 is 86.60%
=====

Train AUC for k = 61 is 88.11%

CV AUC for k = 61 is 86.59%
=====

Train AUC for k = 63 is 88.10%

CV AUC for k = 63 is 86.61%
=====

Train AUC for k = 65 is 88.04%

CV AUC for k = 65 is 86.59%
=====

Train AUC for k = 67 is 88.01%

CV AUC for k = 67 is 86.64%
=====

Train AUC for k = 69 is 87.97%

CV AUC for k = 69 is 86.68%
=====

```

```

In [0]: for i in range(71,100,2):
        knn = KNeighborsClassifier(n_neighbors=i)
        knn.fit(w2v_train, w2v_train_y)
        # train data
        y_prob_train = knn.predict_proba(w2v_train)[: ,1]
        y_pred = np.where(y_prob_train > 0.5, 1, 0)

```

```

auc_roc_test = roc_auc_score(w2v_train_y , y_prob_train)
print('\nTrain AUC for k = %d is %0.2f%%' % (i, (auc_roc_test * float(100))))
w2v_train_auc.append(auc_roc_test)
# CV
y_prob_cv = knn.predict_proba(w2v_cv)[: ,1]
y_pred = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(w2v_cv_y , y_prob_cv)
print('\nCV AUC for k = %d is %0.2f%%' % (i, (auc_roc_cv * float(100))))
w2v_cv_auc.append(auc_roc_cv)
print("="*50)

```

Train AUC for k = 71 is 87.95%

CV AUC for k = 71 is 86.67%

=====

Train AUC for k = 73 is 87.90%

CV AUC for k = 73 is 86.67%

=====

Train AUC for k = 75 is 87.89%

CV AUC for k = 75 is 86.63%

=====

Train AUC for k = 77 is 87.87%

CV AUC for k = 77 is 86.65%

=====

Train AUC for k = 79 is 87.85%

CV AUC for k = 79 is 86.70%

=====

Train AUC for k = 81 is 87.81%

CV AUC for k = 81 is 86.69%

=====

Train AUC for k = 83 is 87.81%

CV AUC for k = 83 is 86.72%

=====

Train AUC for k = 85 is 87.77%

```

CV AUC for k = 85 is 86.75%
=====

Train AUC for k = 87 is 87.75%

CV AUC for k = 87 is 86.77%
=====

Train AUC for k = 89 is 87.74%

CV AUC for k = 89 is 86.75%
=====

Train AUC for k = 91 is 87.72%

CV AUC for k = 91 is 86.80%
=====

Train AUC for k = 93 is 87.70%

CV AUC for k = 93 is 86.82%
=====

Train AUC for k = 95 is 87.69%

CV AUC for k = 95 is 86.84%
=====

Train AUC for k = 97 is 87.70%

CV AUC for k = 97 is 86.87%
=====

Train AUC for k = 99 is 87.67%

CV AUC for k = 99 is 86.86%
=====

```

```

In [0]: k_values = [j for j in range(1,100,2)]

In [0]: # https://www.dataquest.io/blog/learning-curves-machine-learning/

import matplotlib.pyplot as plt
%matplotlib inline

plt.style.use('seaborn')

```

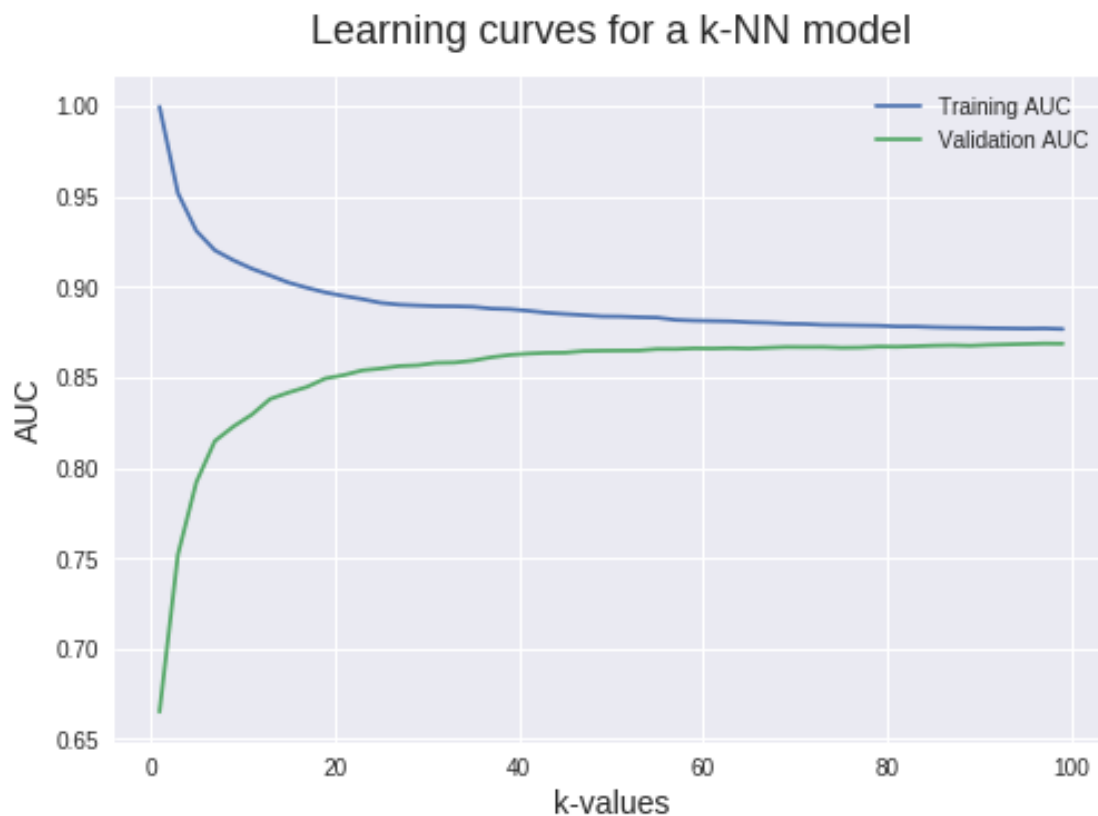
```

plt.plot(k_values, w2v_train_auc, label = 'Training AUC')
plt.plot(k_values, w2v_cv_auc, label = 'Validation AUC')

plt.ylabel('AUC', fontsize = 14)
plt.xlabel('k-values', fontsize = 14)
plt.title('Learning curves for a k-NN model', fontsize = 18, y = 1.03)
plt.legend()

```

Out[0]: <matplotlib.legend.Legend at 0x7fa51d8cbba8>



```

In [0]: knn = KNeighborsClassifier(n_neighbors=97)
knn.fit(w2v_train, w2v_train_y)
# train data
y_prob_train = knn.predict_proba(w2v_train)[: ,1]
fprt, tprt , thresholdt = roc_curve(w2v_train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(w2v_train_y , y_prob_train)
print('\nTrain AUC for k = %d is %0.2f%%' % (97, (auc_roc_test * float(100))))
# CV
y_prob_cv = knn.predict_proba(w2v_cv)[: ,1]

```

```

fprc, tprc , thresholdc = roc_curve(w2v_cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(w2v_cv_y , y_prob_cv)
print('\nCV AUC for k = %d is %0.2f%%' % (97, (auc_roc_cv * float(100))))
# Test
y_prob_test = knn.predict_proba(w2v_test)[: ,1]
fprts, tprts , thresholdts = roc_curve(w2v_test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
auc_roc_test = roc_auc_score(w2v_test_y , y_prob_test)
print('\nTest AUC for k = %d is %0.2f%%' % (97, (auc_roc_test * float(100))))
print("="*50)

```

Train AUC for k = 97 is 87.70%

CV AUC for k = 97 is 86.87%

Test AUC for k = 97 is 86.76%


=====

In [0]: # <https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python>

```

import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b' , label='AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y' , label='AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label='AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

```




```
In [0]: print("F1-Score on test set: %0.2f"%(f1_score(w2v_test_y, y_pred_test)))
```

F1-Score on test set: 0.92

```
In [0]: df_cm = pd.DataFrame(confusion_matrix(w2v_test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[0]: <matplotlib.axes._subplots.AxesSubplot at 0x7fa51d6676d8>
```



6.1.4 [5.1.4] Applying KNN brute force on TFIDF W2V, SET 4

```
In [0]: # Please write all the code with proper documentation

In [0]: # Please write all the code with proper documentation
# load Review text, preprocessed one converted into vectors using (w2v)
tfidf_w2v_train = np.load("drive/ai/w2v_tfidf_train.npy")
tfidf_w2v_train_y = np.load("drive/ai/w2v_tfidf_train_y.npy")
tfidf_w2v_cv = np.load("drive/ai/w2v_tfidf_cv.npy")
tfidf_w2v_cv_y = np.load("drive/ai/w2v_tfidf_cv_y.npy")
tfidf_w2v_test = np.load("drive/ai/w2v_tfidf_test.npy")
tfidf_w2v_test_y = np.load("drive/ai/w2v_tfidf_test_y.npy")

In [0]: print("Shape of Train = ", tfidf_w2v_train.shape)
        print("Shape of CV = ", tfidf_w2v_cv.shape)
        print("Shape of Test = ", tfidf_w2v_test.shape)

Shape of Train = (24500, 50)
Shape of CV = (10500, 50)
Shape of Test = (15000, 50)

In [0]: tfidf_w2v_train_auc = []
        tfidf_w2v_cv_auc = []
```



```

for i in range(1,50,2):
    knn = KNeighborsClassifier(n_neighbors=i)
    knn.fit(tfidf_w2v_train, tfidf_w2v_train_y)
    # train data
    y_prob_train = knn.predict_proba(tfidf_w2v_train)[: ,1]
    y_pred = np.where(y_prob_train > 0.5, 1, 0)
    auc_roc_test = roc_auc_score(tfidf_w2v_train_y , y_prob_train)
    print('\nTrain AUC for k = %d is %0.2f%%' % (i, (auc_roc_test * float(100))))
    tfidf_w2v_train_auc.append(auc_roc_test)
    # CV
    y_prob_cv = knn.predict_proba(tfidf_w2v_cv)[: ,1]
    y_pred = np.where(y_prob_cv > 0.5, 1, 0)
    auc_roc_cv = roc_auc_score(tfidf_w2v_cv_y , y_prob_cv)
    print('\nCV AUC for k = %d is %0.2f%%' % (i, (auc_roc_cv * float(100))))
    tfidf_w2v_cv_auc.append(auc_roc_cv)
print("="*50)

```

Train AUC for k = 1 is 99.96%

CV AUC for k = 1 is 62.63%

=====

Train AUC for k = 3 is 94.05%

CV AUC for k = 3 is 71.66%

=====

Train AUC for k = 5 is 91.67%

CV AUC for k = 5 is 75.75%

=====

Train AUC for k = 7 is 90.30%

CV AUC for k = 7 is 77.93%

=====

Train AUC for k = 9 is 89.37%

CV AUC for k = 9 is 79.13%

=====

Train AUC for k = 11 is 88.58%

CV AUC for k = 11 is 79.50%

=====

```

Train AUC for k = 13 is 88.04%
CV AUC for k = 13 is 80.21%
=====

Train AUC for k = 15 is 87.66%
CV AUC for k = 15 is 80.56%
=====

Train AUC for k = 17 is 87.48%
CV AUC for k = 17 is 80.81%
=====

Train AUC for k = 19 is 87.35%
CV AUC for k = 19 is 80.97%
=====

Train AUC for k = 21 is 87.16%
CV AUC for k = 21 is 81.16%
=====

Train AUC for k = 23 is 87.00%
CV AUC for k = 23 is 81.40%
=====

Train AUC for k = 25 is 86.82%
CV AUC for k = 25 is 81.36%
=====

Train AUC for k = 27 is 86.57%
CV AUC for k = 27 is 81.53%
=====

Train AUC for k = 29 is 86.44%
CV AUC for k = 29 is 81.59%
=====

Train AUC for k = 31 is 86.28%
CV AUC for k = 31 is 81.87%

```

```

=====

Train AUC for k = 33 is 86.17%

CV AUC for k = 33 is 82.15%
=====

Train AUC for k = 35 is 86.13%

CV AUC for k = 35 is 82.34%
=====

Train AUC for k = 37 is 86.06%

CV AUC for k = 37 is 82.47%
=====

Train AUC for k = 39 is 86.12%

CV AUC for k = 39 is 82.59%
=====

Train AUC for k = 41 is 86.01%

CV AUC for k = 41 is 82.61%
=====

Train AUC for k = 43 is 85.98%

CV AUC for k = 43 is 82.63%
=====

Train AUC for k = 45 is 85.91%

CV AUC for k = 45 is 82.70%
=====

Train AUC for k = 47 is 85.83%

CV AUC for k = 47 is 82.72%
=====

Train AUC for k = 49 is 85.79%

CV AUC for k = 49 is 82.79%
=====

```

```

In [0]: k_values = [j for j in range(1,50,2)]

In [0]: # https://www.dataquest.io/blog/learning-curves-machine-learning/

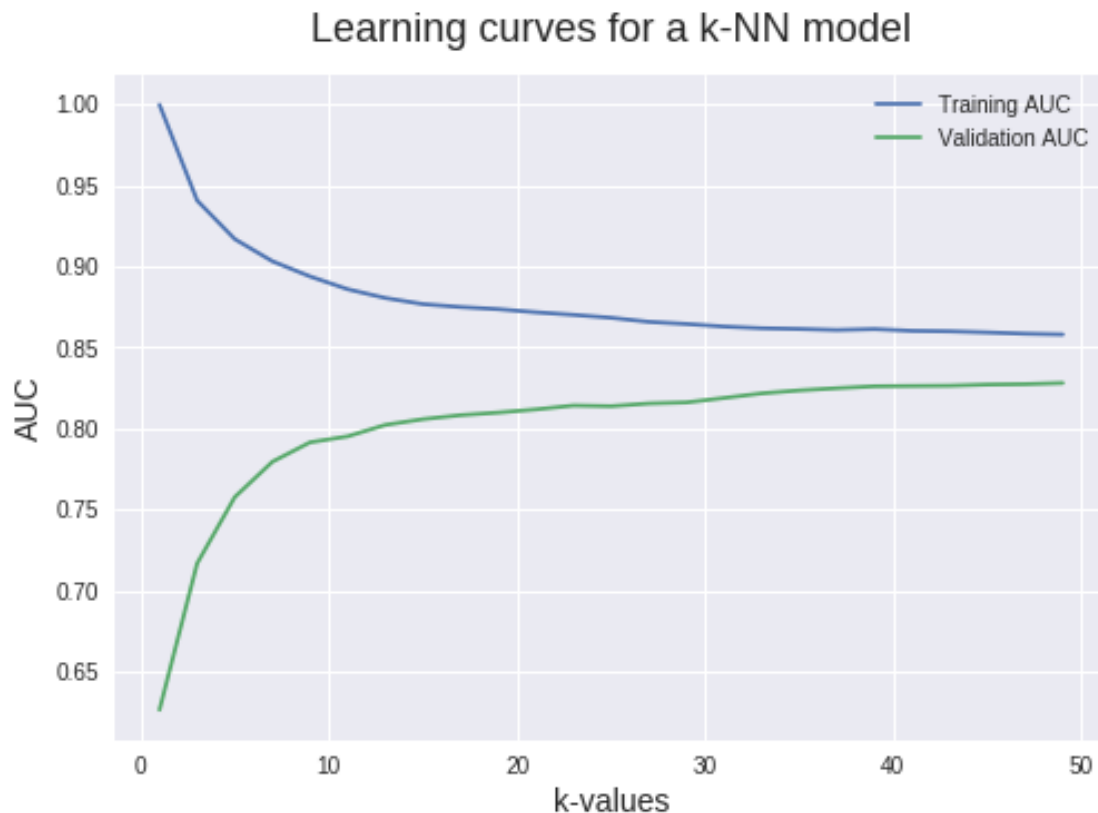
plt.style.use('seaborn')

plt.plot(k_values, tfidf_w2v_train_auc, label = 'Training AUC')
plt.plot(k_values, tfidf_w2v_cv_auc, label = 'Validation AUC')

plt.ylabel('AUC', fontsize = 14)
plt.xlabel('k-values', fontsize = 14)
plt.title('Learning curves for a k-NN model', fontsize = 18, y = 1.03)
plt.legend()

```

Out[0]: <matplotlib.legend.Legend at 0x7fa51d224b00>



```

In [0]: knn = KNeighborsClassifier(n_neighbors=49)
knn.fit(tfidf_w2v_train, tfidf_w2v_train_y)
# train data
y_prob_train = knn.predict_proba(tfidf_w2v_train)[:,:1]
fpr_t, tpr_t, threshold_t = roc_curve(tfidf_w2v_train_y, y_prob_train)

```

```

y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(tfidf_w2v_train_y , y_prob_train)
print('\nTrain AUC for k = %d is %0.2f%%' % (49, (auc_roc_train * float(100))))
# CV
y_prob_cv = knn.predict_proba(tfidf_w2v_cv)[: ,1]
fprc, tprc , thresholdc = roc_curve(tfidf_w2v_cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(tfidf_w2v_cv_y , y_prob_cv)
print('\nCV AUC for k = %d is %0.2f%%' % (49, (auc_roc_cv * float(100))))
# Test
y_prob_test = knn.predict_proba(tfidf_w2v_test)[: ,1]
fprts, tprts , thresholdts = roc_curve(tfidf_w2v_test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
auc_roc_test = roc_auc_score(tfidf_w2v_test_y , y_prob_test)
print('\nTest AUC for k = %d is %0.2f%%' % (97, (auc_roc_test * float(100))))
print("="*50)

```

Train AUC for k = 49 is 85.79%

CV AUC for k = 49 is 82.79%

Test AUC for k = 97 is 82.76%


=====

In [0]: # <https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python>

```

import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b' , label='AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y' , label='AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label='AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

```




```
In [0]: print("F1-Score on test set: %0.2f"%(f1_score(tfidf_w2v_test_y, y_pred_test)))
```

F1-Score on test set: 0.92

```
In [0]: df_cm = pd.DataFrame(confusion_matrix(tfidf_w2v_test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[0]: <matplotlib.axes._subplots.AxesSubplot at 0x7fa51d237898>
```



6.2 [5.2] Applying KNN kd-tree

6.2.1 [5.2.1] Applying KNN kd-tree on BOW, SET 5

In [0]: *# Please write all the code with proper documentation*

```
In [0]: # Please write all the code with proper documentation
# load Review text, preprocessed one converted into vectors using (BOW)
kd_bow_train = np.load("drive/ai/kd_bow_train.npy")
kd_bow_train_y = np.load("drive/ai/kd_bow_train_y.npy")
kd_bow_cv = np.load("drive/ai/kd_bow_cv.npy")
kd_bow_cv_y = np.load("drive/ai/kd_bow_cv_y.npy")
kd_bow_test = np.load("drive/ai/kd_bow_test.npy")
kd_bow_test_y = np.load("drive/ai/kd_bow_test_y.npy")
```

```
In [0]: print("Shape of Train = ", kd_bow_train.shape)
        print("Shape of CV = ", kd_bow_cv.shape)
        print("Shape of Test = ", kd_bow_test.shape)
```

```
Shape of Train = (24500, 200)
Shape of CV = (10500, 200)
Shape of Test = (15000, 200)
```

```

In [0]: bow_train_auc = []
        bow_cv_auc = []
        for i in range(1,20,2):
            knn = KNeighborsClassifier(n_neighbors=i, algorithm="kd_tree")
            knn.fit(kd_bow_train, kd_bow_train_y)
            # train data
            y_prob_train = knn.predict_proba(kd_bow_train)[: ,1]
            y_pred = np.where(y_prob_train > 0.5, 1, 0)
            auc_roc_test = roc_auc_score(kd_bow_train_y , y_prob_train)
            print('\nTrain AUC for k = %d is %0.2f%%' % (i, (auc_roc_test * float(100))))
            bow_train_auc.append(auc_roc_test)
            # CV
            y_prob_cv = knn.predict_proba(kd_bow_cv)[: ,1]
            y_pred = np.where(y_prob_cv > 0.5, 1, 0)
            auc_roc_cv = roc_auc_score(kd_bow_cv_y , y_prob_cv)
            print('\nCV AUC for k = %d is %0.2f%%' % (i, (auc_roc_cv * float(100))))
            bow_cv_auc.append(auc_roc_cv)

```

Train AUC for k = 1 is 99.89%

CV AUC for k = 1 is 55.49%

Train AUC for k = 3 is 92.33%

CV AUC for k = 3 is 59.16%

Train AUC for k = 5 is 89.09%

CV AUC for k = 5 is 60.28%

Train AUC for k = 7 is 87.52%

CV AUC for k = 7 is 61.62%

Train AUC for k = 9 is 86.61%

CV AUC for k = 9 is 62.11%

Train AUC for k = 11 is 85.79%

CV AUC for k = 11 is 62.28%

Train AUC for k = 13 is 84.96%

CV AUC for k = 13 is 62.50%

Train AUC for k = 15 is 84.59%

CV AUC for k = 15 is 62.82%

Train AUC for k = 17 is 84.22%

CV AUC for k = 17 is 62.82%

Train AUC for k = 19 is 83.84%

CV AUC for k = 19 is 62.86%

```
In [0]: k_values = [j for j in range(1,20,2)]
```

```
In [0]: # https://www.dataquest.io/blog/learning-curves-machine-learning/
```

```
import matplotlib.pyplot as plt
%matplotlib inline


plt.style.use('seaborn')

plt.plot(k_values, bow_train_auc, label = 'Training AUC')
plt.plot(k_values, bow_cv_auc, label = 'Validation AUC')

plt.ylabel('AUC', fontsize = 14)
plt.xlabel('k-values', fontsize = 14)
plt.title('Learning curves for a k-NN model', fontsize = 18, y = 1.03)
plt.legend()
```

```
Out[0]: <matplotlib.legend.Legend at 0x7f2bbb142860>
```

Learning curves for a k-NN model



```
In [0]: knn = KNeighborsClassifier(n_neighbors=19 , algorithm="kd_tree")
knn.fit(kd_bow_train, kd_bow_train_y)
# train data
y_prob_train = knn.predict_proba(kd_bow_train)[: ,1]
fprt, tprrt , thresholdt = roc_curve(kd_bow_train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(kd_bow_train_y , y_prob_train)
print('\nTrain AUC for k = %d is %0.2f%%' % (19, (auc_roc_train * float(100))))
# CV
y_prob_cv = knn.predict_proba(kd_bow_cv)[: ,1]
fprc, tprrc , thresholdc = roc_curve(kd_bow_cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(kd_bow_cv_y , y_prob_cv)
print('\nCV AUC for k = %d is %0.2f%%' % (19, (auc_roc_cv * float(100))))
# Test
y_prob_test = knn.predict_proba(kd_bow_test)[: ,1]
fprts, tprrts , thresholdts = roc_curve(kd_bow_test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
auc_roc_test = roc_auc_score(kd_bow_test_y , y_prob_test)
print('\nTest AUC for k = %d is %0.2f%%' % (19, (auc_roc_test * float(100))))
print("="*50)
```

Train AUC for k = 19 is 83.84%


CV AUC for k = 19 is 62.86%

Test AUC for k = 19 is 57.44%

=====

In [0]: # <https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python>

```
import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b', label='AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y', label='AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label='AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```




```
In [0]: print("F1-Score on test set: %0.2f"%(f1_score(kd_bow_test_y, y_pred_test)))
```

F1-Score on test set: 0.86

```
In [0]: df_cm = pd.DataFrame(confusion_matrix(kd_bow_test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[0]: <matplotlib.axes._subplots.AxesSubplot at 0x7f2bbd522f98>
```



6.2.2 [5.2.2] Applying KNN kd-tree on TFIDF, SET 6

```
In [0]: # Please write all the code with proper documentation
```

```
In [0]: # Please write all the code with proper documentation
# Please write all the code with proper documentation
# load Review text, preprocessed one converted into vectors using (BOW)
kd_tfidf_train = np.load("drive/ai/kd_tfidf_train.npy")
kd_tfidf_train_y = np.load("drive/ai/kd_tfidf_train_y.npy")
kd_tfidf_cv = np.load("drive/ai/kd_tfidf_cv.npy")
kd_tfidf_cv_y = np.load("drive/ai/kd_tfidf_cv_y.npy")
kd_tfidf_test = np.load("drive/ai/kd_tfidf_test.npy")
kd_tfidf_test_y = np.load("drive/ai/kd_tfidf_test_y.npy")
```

```
In [65]: print("Shape of Train = ", kd_tfidf_train.shape)
         print("Shape of CV = ", kd_tfidf_cv.shape)
         print("Shape of Test = ", kd_tfidf_test.shape)
```

```
Shape of Train = (24500, 200)
Shape of CV = (10500, 200)
Shape of Test = (15000, 200)
```

```
In [66]: tfidf_train_auc = []
         tfidf_cv_auc = []
         for i in range(1,10,2):
             knn = KNeighborsClassifier(n_neighbors=i, algorithm="kd_tree")
             knn.fit(kd_tfidf_train, kd_tfidf_train_y)
             # train data
             y_prob_train = knn.predict_proba(kd_tfidf_train)[:,:1]
             y_pred = np.where(y_prob_train > 0.5, 1, 0)
             auc_roc_test = roc_auc_score(kd_tfidf_train_y , y_prob_train)
             print('\nTrain AUC for k = %d is %0.2f%%' % (i, (auc_roc_test * float(100))))
             tfidf_train_auc.append(auc_roc_test)
             # CV
             y_prob_cv = knn.predict_proba(kd_tfidf_cv)[:,:1]
             y_pred = np.where(y_prob_cv > 0.5, 1, 0)
             auc_roc_cv = roc_auc_score(kd_tfidf_cv_y , y_prob_cv)
             print('\nCV AUC for k = %d is %0.2f%%' % (i, (auc_roc_cv * float(100))))
             tfidf_cv_auc.append(auc_roc_cv)
         print("="*50)
```

```
Train AUC for k = 1 is 99.89%
```

```
CV AUC for k = 1 is 56.31%
```

```
=====
```

```
Train AUC for k = 3 is 92.03%
```

```
CV AUC for k = 3 is 60.04%
```

```
=====
```

```
Train AUC for k = 5 is 89.02%
```

```
CV AUC for k = 5 is 61.82%
```

```
=====
```

```
Train AUC for k = 7 is 87.35%
```

```
CV AUC for k = 7 is 63.19%
```

```
=====
```

Train AUC for k = 9 is 86.26%

CV AUC for k = 9 is 63.70%

=====

```
In [0]: k_values = [j for j in range(1,10,2)]
```

```
In [68]: # https://www.dataquest.io/blog/learning-curves-machine-learning/
```


```
import matplotlib.pyplot as plt
%matplotlib inline

plt.style.use('seaborn')

plt.plot(k_values, tfidf_train_auc, label = 'Training AUC')
plt.plot(k_values, tfidf_cv_auc, label = 'Validation AUC')

plt.ylabel('AUC', fontsize = 14)
plt.xlabel('k-values', fontsize = 14)
plt.title('Learning curves for a k-NN model', fontsize = 18, y = 1.03)
plt.legend()
```

```
Out[68]: <matplotlib.legend.Legend at 0x7f2bbc27f940>
```



```

In [71]: knn = KNeighborsClassifier(n_neighbors=9, algorithm="kd_tree")
knn.fit(kd_tfidf_train, kd_tfidf_train_y)
# train data
y_prob_train = knn.predict_proba(kd_tfidf_train)[: ,1]
fpr_t, tpr_t , threshold_t = roc_curve(kd_tfidf_train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(kd_tfidf_train_y , y_prob_train)
print('\nTrain AUC for k = %d is %0.2f%%' % (9, (auc_roc_train * float(100))))
# CV
y_prob_cv = knn.predict_proba(kd_tfidf_cv)[: ,1]
fpr_c, tpr_c , threshold_c = roc_curve(kd_tfidf_cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(kd_tfidf_cv_y , y_prob_cv)
print('\nCV AUC for k = %d is %0.2f%%' % (9, (auc_roc_cv * float(100))))
# Test
y_prob_test = knn.predict_proba(kd_tfidf_test)[: ,1]
fpr_t, tpr_t , threshold_t = roc_curve(kd_tfidf_test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
auc_roc_test = roc_auc_score(kd_tfidf_test_y , y_prob_test)
print('\nTest AUC for k = %d is %0.2f%%' % (9, (auc_roc_test * float(100))))
print("="*50)

```

Train AUC for k = 9 is 86.26%

CV AUC for k = 9 is 63.70%

Test AUC for k = 9 is 60.71%

=====

```

In [73]: print("F1-Score on test set: %0.2f"%(f1_score(kd_tfidf_test_y, y_pred_test)))

```


F1-Score on test set: 0.89

```

In [74]: df_cm = pd.DataFrame(confusion_matrix(kd_tfidf_test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')

```

Out[74]: <matplotlib.axes._subplots.AxesSubplot at 0x7f2bbd29a7f0>



6.2.3 [5.2.3] Applying KNN kd-tree on AVG W2V, SET 3

In [0]: *# Please write all the code with proper documentation*

```
In [0]: # Please write all the code with proper documentation
# load Review text, preprocessed one converted into vectors using (w2v)
w2v_train = np.load("drive/ai/w2v_train.npy")
w2v_train_y = np.load("drive/ai/w2v_train_y.npy")
w2v_cv = np.load("drive/ai/w2v_cv.npy")
w2v_cv_y = np.load("drive/ai/w2v_cv_y.npy")
w2v_test = np.load("drive/ai/w2v_test.npy")
w2v_test_y = np.load("drive/ai/w2v_test_y.npy")
```

```
In [0]: print("Shape of Train = ", w2v_train.shape)
print("Shape of CV = ", w2v_cv.shape)
print("Shape of Test = ", w2v_test.shape)
```

Shape of Train = (24500, 50)

Shape of CV = (10500, 50)

Shape of Test = (15000, 50)

```
In [0]: w2v_train_auc = []
w2v_cv_auc = []
```



```

for i in range(1,100,2):
    knn = KNeighborsClassifier(n_neighbors=i, algorithm="kd_tree")
    knn.fit(w2v_train, w2v_train_y)
    # train data
    y_prob_train = knn.predict_proba(w2v_train)[:,-1]
    y_pred = np.where(y_prob_train > 0.5, 1, 0)
    auc_roc_test = roc_auc_score(w2v_train_y , y_prob_train)
    print('\nTrain AUC for k = %d is %0.2f%%' % (i, (auc_roc_test * float(100))))
    w2v_train_auc.append(auc_roc_test)
    # CV
    y_prob_cv = knn.predict_proba(w2v_cv)[:,-1]
    y_pred = np.where(y_prob_cv > 0.5, 1, 0)
    auc_roc_cv = roc_auc_score(w2v_cv_y , y_prob_cv)
    print('\nCV AUC for k = %d is %0.2f%%' % (i, (auc_roc_cv * float(100))))
    w2v_cv_auc.append(auc_roc_cv)
    print("="*50)

```

Train AUC for k = 1 is 99.96%

CV AUC for k = 1 is 66.50%

=====

Train AUC for k = 3 is 95.18%

CV AUC for k = 3 is 75.20%

=====

Train AUC for k = 5 is 93.10%

CV AUC for k = 5 is 79.20%

=====

Train AUC for k = 7 is 92.02%

CV AUC for k = 7 is 81.47%

=====

Train AUC for k = 9 is 91.48%

CV AUC for k = 9 is 82.27%

=====

Train AUC for k = 11 is 91.01%

CV AUC for k = 11 is 82.93%

=====

Train AUC for k = 13 is 90.63%

CV AUC for k = 13 is 83.80%

=====

Train AUC for k = 15 is 90.24%

CV AUC for k = 15 is 84.16%

=====

Train AUC for k = 17 is 89.95%

CV AUC for k = 17 is 84.47%

=====

Train AUC for k = 19 is 89.69%

CV AUC for k = 19 is 84.94%

=====

Train AUC for k = 21 is 89.50%

CV AUC for k = 21 is 85.12%

=====

Train AUC for k = 23 is 89.32%

CV AUC for k = 23 is 85.38%

=====

Train AUC for k = 25 is 89.12%

CV AUC for k = 25 is 85.49%

=====

Train AUC for k = 27 is 89.02%

CV AUC for k = 27 is 85.62%

=====

Train AUC for k = 29 is 88.97%

CV AUC for k = 29 is 85.66%

=====

Train AUC for k = 31 is 88.94%

CV AUC for k = 31 is 85.79%

```

=====

Train AUC for k = 33 is 88.93%

CV AUC for k = 33 is 85.81%
=====

Train AUC for k = 35 is 88.90%

CV AUC for k = 35 is 85.92%
=====

Train AUC for k = 37 is 88.80%

CV AUC for k = 37 is 86.11%
=====

Train AUC for k = 39 is 88.77%

CV AUC for k = 39 is 86.24%
=====

Train AUC for k = 41 is 88.68%

CV AUC for k = 41 is 86.31%
=====

Train AUC for k = 43 is 88.57%

CV AUC for k = 43 is 86.35%
=====

Train AUC for k = 45 is 88.50%

CV AUC for k = 45 is 86.36%
=====

Train AUC for k = 47 is 88.44%

CV AUC for k = 47 is 86.45%
=====

Train AUC for k = 49 is 88.37%

CV AUC for k = 49 is 86.47%
=====

Train AUC for k = 51 is 88.36%

```

```

CV AUC for k = 51 is 86.48%
=====

Train AUC for k = 53 is 88.32%

CV AUC for k = 53 is 86.48%
=====

Train AUC for k = 55 is 88.30%

CV AUC for k = 55 is 86.56%
=====

Train AUC for k = 57 is 88.17%

CV AUC for k = 57 is 86.56%
=====

Train AUC for k = 59 is 88.13%

CV AUC for k = 59 is 86.60%
=====

Train AUC for k = 61 is 88.11%

CV AUC for k = 61 is 86.59%
=====

Train AUC for k = 63 is 88.10%

CV AUC for k = 63 is 86.61%
=====

Train AUC for k = 65 is 88.04%

CV AUC for k = 65 is 86.59%
=====

Train AUC for k = 67 is 88.01%

CV AUC for k = 67 is 86.64%
=====

Train AUC for k = 69 is 87.97%

CV AUC for k = 69 is 86.68%
=====

```

Train AUC for k = 71 is 87.95%

CV AUC for k = 71 is 86.67%

=====

Train AUC for k = 73 is 87.90%

CV AUC for k = 73 is 86.67%

=====

Train AUC for k = 75 is 87.89%

CV AUC for k = 75 is 86.63%

=====

Train AUC for k = 77 is 87.87%

CV AUC for k = 77 is 86.65%

=====

Train AUC for k = 79 is 87.85%

CV AUC for k = 79 is 86.70%

=====

Train AUC for k = 81 is 87.81%

CV AUC for k = 81 is 86.69%

=====

Train AUC for k = 83 is 87.81%

CV AUC for k = 83 is 86.72%

=====

Train AUC for k = 85 is 87.77%

CV AUC for k = 85 is 86.75%

=====

Train AUC for k = 87 is 87.75%

CV AUC for k = 87 is 86.77%

=====

Train AUC for k = 89 is 87.74%

CV AUC for k = 89 is 86.75%
=====

Train AUC for k = 91 is 87.72%

CV AUC for k = 91 is 86.80%
=====

Train AUC for k = 93 is 87.70%

CV AUC for k = 93 is 86.82%
=====

Train AUC for k = 95 is 87.69%

CV AUC for k = 95 is 86.84%
=====

Train AUC for k = 97 is 87.70%

CV AUC for k = 97 is 86.87%
=====

Train AUC for k = 99 is 87.67%

CV AUC for k = 99 is 86.86%
=====

```
In [0]: k_values = [j for j in range(1,100,2)]
```

```
In [0]: # https://www.dataquest.io/blog/learning-curves-machine-learning/
```

```
import matplotlib.pyplot as plt
%matplotlib inline
```


```
plt.style.use('seaborn')
```

```
plt.plot(k_values, w2v_train_auc, label = 'Training AUC')
plt.plot(k_values, w2v_cv_auc, label = 'Validation AUC')
```

```
plt.ylabel('AUC', fontsize = 14)
plt.xlabel('k-values', fontsize = 14)
plt.title('Learning curves for a k-NN model', fontsize = 18, y = 1.03)
plt.legend()
```

```
Out[0]: <matplotlib.legend.Legend at 0x7f2bd4776c50>
```

Learning curves for a k-NN model



```
In [0]: knn = KNeighborsClassifier(n_neighbors=97, algorithm="kd_tree")
knn.fit(w2v_train, w2v_train_y)
# train data
y_prob_train = knn.predict_proba(w2v_train)[: ,1]
fpr_t, tpr_t , threshold_t = roc_curve(w2v_train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(w2v_train_y , y_prob_train)
print('\nTrain AUC for k = %d is %0.2f%%' % (97, (auc_roc_train * float(100))))
# CV
y_prob_cv = knn.predict_proba(w2v_cv)[: ,1]
fpr_cv, tpr_cv , threshold_cv = roc_curve(w2v_cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(w2v_cv_y , y_prob_cv)
print('\nCV AUC for k = %d is %0.2f%%' % (97, (auc_roc_cv * float(100))))
# Test
y_prob_test = knn.predict_proba(w2v_test)[: ,1]
fpr_test, tpr_test , threshold_test = roc_curve(w2v_test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
auc_roc_test = roc_auc_score(w2v_test_y , y_prob_test)
print('\nTest AUC for k = %d is %0.2f%%' % (97, (auc_roc_test * float(100))))
print("="*50)
```

Train AUC for k = 97 is 87.67%


CV AUC for k = 97 is 86.87%

Test AUC for k = 97 is 86.76%

=====

In [0]: # <https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python>

```
import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b', label='AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y', label='AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label='AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```





```
In [0]: print("F1-Score on test set: %0.2f"%(f1_score(w2v_test_y, y_pred_test)))
```

F1-Score on test set: 0.92

```
In [0]: df_cm = pd.DataFrame(confusion_matrix(w2v_test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[0]: <matplotlib.axes._subplots.AxesSubplot at 0x7f2bd4803908>
```



6.2.4 [5.2.4] Applying KNN kd-tree on TFIDF W2V, SET 4

```
In [0]: # Please write all the code with proper documentation
```

```
In [0]: # Please write all the code with proper documentation
# load Review text, preprocessed one converted into vectors using (w2v)
tfidf_w2v_train = np.load("drive/ai/w2v_tfidf_train.npy")
tfidf_w2v_train_y = np.load("drive/ai/w2v_tfidf_train_y.npy")
tfidf_w2v_cv = np.load("drive/ai/w2v_tfidf_cv.npy")
tfidf_w2v_cv_y = np.load("drive/ai/w2v_tfidf_cv_y.npy")
tfidf_w2v_test = np.load("drive/ai/w2v_tfidf_test.npy")
tfidf_w2v_test_y = np.load("drive/ai/w2v_tfidf_test_y.npy")
```

```
In [0]: print("Shape of Train = ", tfidf_w2v_train.shape)
        print("Shape of CV = ", tfidf_w2v_cv.shape)
        print("Shape of Test = ", tfidf_w2v_test.shape)
```

```
Shape of Train = (24500, 50)
Shape of CV = (10500, 50)
Shape of Test = (15000, 50)
```

```
In [0]: tfidf_w2v_train_auc = []
        tfidf_w2v_cv_auc = []
        for i in range(1,100,2):
            knn = KNeighborsClassifier(n_neighbors=i)
            knn.fit(tfidf_w2v_train, tfidf_w2v_train_y)
            # train data
            y_prob_train = knn.predict_proba(tfidf_w2v_train)[: ,1]
            y_pred = np.where(y_prob_train > 0.5, 1, 0)
            auc_roc_test = roc_auc_score(tfidf_w2v_train_y , y_prob_train)
            print('\nTrain AUC for k = %d is %0.2f%%' % (i, (auc_roc_test * float(100))))
            tfidf_w2v_train_auc.append(auc_roc_test)
            # CV
            y_prob_cv = knn.predict_proba(tfidf_w2v_cv)[: ,1]
            y_pred = np.where(y_prob_cv > 0.5, 1, 0)
            auc_roc_cv = roc_auc_score(tfidf_w2v_cv_y , y_prob_cv)
            print('\nCV AUC for k = %d is %0.2f%%' % (i, (auc_roc_cv * float(100))))
            tfidf_w2v_cv_auc.append(auc_roc_cv)
        print("="*50)
```

```
Train AUC for k = 1 is 99.96%
```

```
CV AUC for k = 1 is 62.63%
```

```
=====
```

```
Train AUC for k = 3 is 94.05%
```

```
CV AUC for k = 3 is 71.66%
```

```
=====
```

```
Train AUC for k = 5 is 91.67%
```

```
CV AUC for k = 5 is 75.75%
```

```
=====
```

```
Train AUC for k = 7 is 90.30%
```

```
CV AUC for k = 7 is 77.93%
```

```
=====
```

```

Train AUC for k = 9 is 89.37%

CV AUC for k = 9 is 79.13%
=====

Train AUC for k = 11 is 88.58%

CV AUC for k = 11 is 79.50%
=====

Train AUC for k = 13 is 88.04%

CV AUC for k = 13 is 80.21%
=====

Train AUC for k = 15 is 87.66%

CV AUC for k = 15 is 80.56%
=====

Train AUC for k = 17 is 87.48%

CV AUC for k = 17 is 80.81%
=====

Train AUC for k = 19 is 87.35%

CV AUC for k = 19 is 80.97%
=====

Train AUC for k = 21 is 87.16%

CV AUC for k = 21 is 81.16%
=====

Train AUC for k = 23 is 87.00%

CV AUC for k = 23 is 81.40%
=====

Train AUC for k = 25 is 86.82%

CV AUC for k = 25 is 81.36%
=====

Train AUC for k = 27 is 86.57%

```

```

CV AUC for k = 27 is 81.53%
=====

Train AUC for k = 29 is 86.44%

CV AUC for k = 29 is 81.59%
=====

Train AUC for k = 31 is 86.28%

CV AUC for k = 31 is 81.87%
=====

Train AUC for k = 33 is 86.17%

CV AUC for k = 33 is 82.15%
=====

Train AUC for k = 35 is 86.13%

CV AUC for k = 35 is 82.34%
=====

Train AUC for k = 37 is 86.06%

CV AUC for k = 37 is 82.47%
=====

Train AUC for k = 39 is 86.12%

CV AUC for k = 39 is 82.59%
=====

Train AUC for k = 41 is 86.01%

CV AUC for k = 41 is 82.61%
=====

Train AUC for k = 43 is 85.98%

CV AUC for k = 43 is 82.63%
=====

Train AUC for k = 45 is 85.91%

CV AUC for k = 45 is 82.70%
=====

```

Train AUC for k = 47 is 85.83%

CV AUC for k = 47 is 82.72%

=====

Train AUC for k = 49 is 85.79%

CV AUC for k = 49 is 82.79%

=====

Train AUC for k = 51 is 85.75%

CV AUC for k = 51 is 82.83%

=====

Train AUC for k = 53 is 85.63%

CV AUC for k = 53 is 82.93%

=====

Train AUC for k = 55 is 85.58%

CV AUC for k = 55 is 82.96%

=====

Train AUC for k = 57 is 85.49%

CV AUC for k = 57 is 82.96%

=====

Train AUC for k = 59 is 85.46%

CV AUC for k = 59 is 82.91%

=====

Train AUC for k = 61 is 85.42%

CV AUC for k = 61 is 83.02%

=====

Train AUC for k = 63 is 85.42%

CV AUC for k = 63 is 83.11%

=====

Train AUC for k = 65 is 85.36%

CV AUC for k = 65 is 83.06%

```

=====

Train AUC for k = 67 is 85.27%

CV AUC for k = 67 is 83.16%
=====

Train AUC for k = 69 is 85.20%

CV AUC for k = 69 is 83.21%
=====

Train AUC for k = 71 is 85.15%

CV AUC for k = 71 is 83.21%
=====

Train AUC for k = 73 is 85.14%

CV AUC for k = 73 is 83.19%
=====

Train AUC for k = 75 is 85.12%

CV AUC for k = 75 is 83.20%
=====

Train AUC for k = 77 is 85.06%

CV AUC for k = 77 is 83.21%
=====

Train AUC for k = 79 is 85.03%

CV AUC for k = 79 is 83.23%
=====

Train AUC for k = 81 is 85.02%

CV AUC for k = 81 is 83.18%
=====

Train AUC for k = 83 is 84.97%

CV AUC for k = 83 is 83.17%
=====

Train AUC for k = 85 is 84.90%

```

```

CV AUC for k = 85 is 83.20%
=====

Train AUC for k = 87 is 84.87%

CV AUC for k = 87 is 83.19%
=====

Train AUC for k = 89 is 84.84%

CV AUC for k = 89 is 83.23%
=====

Train AUC for k = 91 is 84.79%

CV AUC for k = 91 is 83.27%
=====

Train AUC for k = 93 is 84.75%

CV AUC for k = 93 is 83.24%
=====

Train AUC for k = 95 is 84.71%

CV AUC for k = 95 is 83.24%
=====

Train AUC for k = 97 is 84.64%

CV AUC for k = 97 is 83.26%
=====

Train AUC for k = 99 is 84.62%

CV AUC for k = 99 is 83.27%
=====

```

```
In [0]: k_values = [j for j in range(1,100,2)]
```


```
In [0]: # https://www.dataquest.io/blog/learning-curves-machine-learning/
```

```
plt.style.use('seaborn')
```

```
plt.plot(k_values, tfidf_w2v_train_auc, label = 'Training AUC')
plt.plot(k_values, tfidf_w2v_cv_auc, label = 'Validation AUC')
```

```
plt.ylabel('AUC', fontsize = 14)
plt.xlabel('k-values', fontsize = 14)
plt.title('Learning curves for a k-NN model', fontsize = 18, y = 1.03)
plt.legend()
```

Out[0]: <matplotlib.legend.Legend at 0x7f2bd4757780>



```
In [0]: knn = KNeighborsClassifier(n_neighbors=91, algorithm="kd_tree")
knn.fit(tfidf_w2v_train, tfidf_w2v_train_y)
# train data
y_prob_train = knn.predict_proba(tfidf_w2v_train)[: ,1]
fprt, tprt , thresholdt = roc_curve(tfidf_w2v_train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(tfidf_w2v_train_y , y_prob_train)
print('\nTrain AUC for k = %d is %0.2f%%' % (91, (auc_roc_test * float(100))))
# CV
y_prob_cv = knn.predict_proba(tfidf_w2v_cv)[: ,1]
fprc, tprc , thresholdc = roc_curve(tfidf_w2v_cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(tfidf_w2v_cv_y , y_prob_cv)
```



```

print('\nCV AUC for k = %d is %0.2f%%' % (91, (auc_roc_cv * float(100))))
# Test
y_prob_test = knn.predict_proba(tfidf_w2v_test)[: ,1]
fprts, tprts , thresholdts = roc_curve(tfidf_w2v_test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
auc_roc_test = roc_auc_score(tfidf_w2v_test_y , y_prob_test)
print('\nTest AUC for k = %d is %0.2f%%' % (91, (auc_roc_test * float(100))))
print("="*50)

```

Train AUC for k = 91 is 84.62%

CV AUC for k = 91 is 83.27%

Test AUC for k = 91 is 82.98%


=====

In [0]: # <https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python>

```


import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b' , label = 'AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y' , label = 'AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label = 'AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

```



```
In [0]: df_cm = pd.DataFrame(confusion_matrix(tfidf_w2v_test_y, y_pred_test), range(2), range(2))  
sns.set(font_scale=1.4)  
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[0]: <matplotlib.axes._subplots.AxesSubplot at 0x7f2bd470ec18>
```



7 [6] Conclusions

In [0]: *# Please compare all your models using Prettytable library*

In [78]: `from prettytable import PrettyTable`

```
x = PrettyTable(["Vectorizer", "Model", "Hyper parameter k", "AUC", "F1 Score"])

x.add_row(["BOW", "Brute", 9, 0.52, 0.88])
x.add_row(["BOW", "kd_tree", 19, 0.57, 0.86])
x.add_row(["TFIDF", "Brute", 9, 0.51, 0.91])
x.add_row(["TFIDF", "kd_tree", 9, 0.60, 0.89])
x.add_row(["W2V", "Brute", 97, 0.86, 0.92])
x.add_row(["W2V", "kd_tree", 97, 0.86, 0.92])
x.add_row(["TFIDFW2V", "Brute", 47, 0.82, 0.92])
x.add_row(["TFIDFW2V", "kd_tree", 91, 0.82, 0.92])

print(x.get_string(title="k-NN Model"))
```

```
+-----+-----+-----+-----+-----+
| Vectorizer | Model | Hyper parameter k | AUC | F1 Score |
```

| | | | | | | | | | | |
|--|----------|--|---------|--|----|--|------|--|------|--|
| | BOW | | Brute | | 9 | | 0.52 | | 0.88 | |
| | BOW | | kd_tree | | 19 | | 0.57 | | 0.86 | |
| | TFIDF | | Brute | | 9 | | 0.51 | | 0.91 | |
| | TFIDF | | kd_tree | | 9 | | 0.6 | | 0.89 | |
| | W2V | | Brute | | 97 | | 0.86 | | 0.92 | |
| | W2V | | kd_tree | | 97 | | 0.86 | | 0.92 | |
| | TFIDFW2V | | Brute | | 47 | | 0.82 | | 0.92 | |
| | TFIDFW2V | | kd_tree | | 91 | | 0.82 | | 0.92 | |

In [0]: