

06 Implement SGD

January 24, 2019

```
In [1]: import warnings
        warnings.filterwarnings("ignore")
        from sklearn.datasets import load_boston
        from random import seed
        from random import randrange
        from csv import reader
        from math import sqrt
        from sklearn import preprocessing
        import pandas as pd
        import numpy as np
        import matplotlib.pyplot as plt
        from prettytable import PrettyTable
        from sklearn.linear_model import SGDRegressor
        from sklearn import preprocessing
        from sklearn.metrics import mean_squared_error
```

```
In [2]: boston = load_boston()
```

```
In [3]: print(boston.data.shape)
```

```
(506, 13)
```

```
In [4]: print(boston.feature_names)
```

```
['CRIM' 'ZN' 'INDUS' 'CHAS' 'NOX' 'RM' 'AGE' 'DIS' 'RAD' 'TAX' 'PTRATIO'
 'B' 'LSTAT']
```

```
In [5]: X = load_boston().data
        Y = load_boston().target
```

```
In [6]: scaler = preprocessing.StandardScaler().fit(X)
        X = scaler.transform(X)
```

```
In [7]: bos = pd.DataFrame(boston.data)
        print(bos.head())
```

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | \ |
|---|---------|------|------|-----|-------|-------|------|--------|-----|-------|------|---|
| 0 | 0.00632 | 18.0 | 2.31 | 0.0 | 0.538 | 6.575 | 65.2 | 4.0900 | 1.0 | 296.0 | 15.3 | |
| 1 | 0.02731 | 0.0 | 7.07 | 0.0 | 0.469 | 6.421 | 78.9 | 4.9671 | 2.0 | 242.0 | 17.8 | |
| 2 | 0.02729 | 0.0 | 7.07 | 0.0 | 0.469 | 7.185 | 61.1 | 4.9671 | 2.0 | 242.0 | 17.8 | |
| 3 | 0.03237 | 0.0 | 2.18 | 0.0 | 0.458 | 6.998 | 45.8 | 6.0622 | 3.0 | 222.0 | 18.7 | |
| 4 | 0.06905 | 0.0 | 2.18 | 0.0 | 0.458 | 7.147 | 54.2 | 6.0622 | 3.0 | 222.0 | 18.7 | |

| | 11 | 12 |
|---|--------|------|
| 0 | 396.90 | 4.98 |
| 1 | 396.90 | 9.14 |
| 2 | 392.83 | 4.03 |
| 3 | 394.63 | 2.94 |
| 4 | 396.90 | 5.33 |

```
In [8]: bos['PRICE'] = boston.target
```

```

X = bos.drop('PRICE', axis = 1)
Y = bos['PRICE']

```

```
In [83]: scaler = preprocessing.StandardScaler().fit(X)
X = scaler.transform(X)
```

```
In [96]: from sklearn import cross_validation
X_train, X_test, Y_train, Y_test = cross_validation.train_test_split(X, Y, test_size = 0.2)
print(X_train.shape)
print(X_test.shape)
print(Y_train.shape)
print(Y_test.shape)
```

```

(339, 13)
(167, 13)
(339,)
(167,)

```

```
In [110]: clf = SGDRegressor(alpha=0.01)
clf.fit(X, Y)
print(mean_squared_error(Y_test, clf.predict(X_test)))
```

```
26.93189673099861
```

```
In [111]: sklearn_W = clf.coef_
sklearn_b = clf.intercept_
```

0.0.1 own sgd Implementation

```
In [127]: def SGD(X_train, Y_train, learning_rate=0.01, n_iter = 1000, batch_size = 200):
W = np.zeros(X_train.shape[1])
```

```

b = 0.0
r = learning_rate
rt_power = 0.25
for i in range(1,n_iter+1):
    idx = np.random.randint(0, len(X_train),batch_size)
    x_k = X_train[idx]
    y_k = Y_train[idx]
    N = float(batch_size)
    error = y_k - (np.dot(x_k, W) - b)
    W -= r * (-2/N) * x_k.T.dot(error)
    b -= r * np.sum(error)
    r = learning_rate / pow(i, rt_power)
return W, abs(b)

```

In [128]: W, b = SGD(X,Y)

In [129]: W

Out[129]: array([-0.67413219, 0.5806372 , -0.42776069, 0.77507576, -0.99597728,
3.17937914, -0.19316828, -2.10391887, 0.8655778 , -0.53894101,
-1.83358058, 0.91160332, -3.36927039])

In [130]: b

Out[130]: 22.595750477577788

In [131]: pred = X_test.dot(W) + b

In [132]: MSE = mean_squared_error(Y_test, pred)

In [133]: MSE

Out[133]: 27.05120620133145

0.0.2 predicted values v/s actual values


```

In [134]: # code source:https://medium.com/@haydar_ai/learning-data-science-day-9-linear-regre.
from sklearn.linear_model import LinearRegression
%matplotlib inline
import matplotlib.pyplot as plt

Y_pred = pred

plt.scatter(Y_test, Y_pred)
plt.xlabel("Prices: $Y_i$")
plt.ylabel("Predicted prices: $\hat{Y}_i$")
plt.title("Prices vs Predicted prices: $Y_i$ vs $\hat{Y}_i$")
plt.show()

```



0.0.3 chart of predicted values v/s actual values

```
In [135]: acutal = list(Y_test)
          predicted = list(Y_pred)

In [136]: print("actual values v/s predicted\n")
          for i in range(len(acutal)):
              print("%f \t %f"%(acutal[i] , predicted[i]))
```

actual values v/s predicted

| | |
|-----------|-----------|
| 37.600000 | 37.375569 |
| 27.900000 | 30.048159 |
| 22.600000 | 27.164871 |
| 13.800000 | 5.355200 |
| 35.200000 | 35.902943 |
| 10.400000 | 6.646771 |
| 23.900000 | 27.917514 |
| 29.000000 | 31.553187 |
| 22.800000 | 27.183857 |
| 23.200000 | 21.691713 |
| 33.200000 | 32.816749 |
| 19.000000 | 21.313788 |

| | |
|-----------|-----------|
| 20.300000 | 23.238460 |
| 36.100000 | 32.206616 |
| 24.400000 | 27.534659 |
| 17.200000 | 17.219575 |
| 17.900000 | 0.232393 |
| 19.600000 | 19.319231 |
| 19.700000 | 14.570873 |
| 15.000000 | 14.147063 |
| 8.100000 | 5.898771 |
| 23.000000 | 20.776961 |
| 44.800000 | 38.281860 |
| 23.100000 | 24.532152 |
| 32.200000 | 31.789490 |
| 10.800000 | 11.821311 |
| 23.100000 | 25.000656 |
| 21.200000 | 23.839150 |
| 22.200000 | 23.062710 |
| 24.100000 | 23.436042 |
| 17.300000 | 14.473725 |
| 7.000000 | 10.372603 |
| 12.700000 | 17.534814 |
| 17.800000 | 23.488554 |
| 26.400000 | 28.334708 |
| 19.600000 | 20.152030 |
| 25.100000 | 30.130981 |
| 8.300000 | 9.680069 |
| 48.800000 | 41.537782 |
| 34.900000 | 33.367630 |
| 13.800000 | 19.853368 |
| 14.400000 | 4.552592 |
| 30.100000 | 28.857315 |
| 12.700000 | 12.577188 |
| 27.100000 | 27.130407 |
| 24.800000 | 31.038018 |
| 7.000000 | -3.942816 |
| 20.500000 | 18.449914 |
| 21.500000 | 22.475579 |
| 14.000000 | 14.513333 |
| 20.400000 | 20.862346 |
| 22.200000 | 20.070116 |
| 21.400000 | 24.136552 |
| 13.500000 | 14.476340 |
| 19.400000 | 18.646675 |
| 24.700000 | 25.859367 |
| 43.800000 | 36.025187 |
| 14.100000 | 16.031596 |
| 28.600000 | 28.414920 |
| 19.700000 | 21.312331 |

| | |
|-----------|-----------|
| 16.800000 | 20.931281 |
| 23.200000 | 25.440515 |
| 16.200000 | 15.476305 |
| 41.300000 | 32.654584 |
| 22.700000 | 23.076221 |
| 8.300000 | 12.655471 |
| 18.400000 | 20.628324 |
| 24.700000 | 24.940677 |
| 21.700000 | 21.663592 |
| 20.600000 | 20.373578 |
| 16.700000 | 19.961516 |
| 22.100000 | 26.400737 |
| 19.400000 | 17.316550 |
| 27.500000 | 19.133106 |
| 27.900000 | 18.479923 |
| 30.100000 | 25.956618 |
| 17.400000 | 21.930549 |
| 15.400000 | 16.441173 |
| 31.000000 | 34.395415 |
| 14.200000 | 17.797422 |
| 19.600000 | 21.538310 |
| 50.000000 | 42.761446 |
| 21.700000 | 23.461283 |
| 11.700000 | 15.927102 |
| 19.400000 | 24.608908 |
| 13.000000 | 17.468966 |
| 17.500000 | 18.626358 |
| 9.700000 | 9.349649 |
| 20.300000 | 19.327596 |
| 18.600000 | 19.537240 |
| 50.000000 | 36.025385 |
| 19.600000 | 18.259078 |
| 21.400000 | 20.958557 |
| 18.400000 | 18.853227 |
| 22.600000 | 26.230720 |
| 25.000000 | 27.838200 |
| 15.600000 | 13.786135 |
| 26.600000 | 25.132454 |
| 22.400000 | 22.595148 |
| 13.100000 | 14.663969 |
| 23.000000 | 22.332422 |
| 24.500000 | 22.197823 |
| 13.100000 | 14.752905 |
| 50.000000 | 44.064727 |
| 8.800000 | 7.375388 |
| 20.600000 | 22.395009 |
| 12.100000 | 18.200549 |
| 50.000000 | 20.960225 |

| | |
|-----------|-----------|
| 24.100000 | 28.847790 |
| 16.100000 | 17.665832 |
| 23.900000 | 28.212302 |
| 24.300000 | 23.593537 |
| 13.100000 | 20.159721 |
| 30.300000 | 33.054078 |
| 15.200000 | 19.797778 |
| 13.800000 | 14.167284 |
| 26.400000 | 21.381511 |
| 16.600000 | 18.221857 |
| 18.900000 | 19.781838 |
| 17.600000 | 16.643520 |
| 18.700000 | 21.358814 |
| 33.400000 | 33.999282 |
| 20.700000 | 22.525399 |
| 17.100000 | 20.664044 |
| 23.400000 | 24.025095 |
| 26.500000 | 26.298582 |
| 21.400000 | 20.251396 |
| 21.500000 | 24.148385 |
| 19.200000 | 23.882737 |
| 50.000000 | 40.111622 |
| 50.000000 | 40.905903 |
| 23.000000 | 29.067877 |
| 10.500000 | 13.292929 |
| 17.800000 | 17.083084 |
| 10.900000 | 18.506063 |
| 21.000000 | 21.787066 |
| 13.800000 | 14.864155 |
| 10.500000 | 5.659188 |
| 22.200000 | 23.987468 |
| 30.500000 | 31.057239 |
| 19.400000 | 22.824881 |
| 15.600000 | 22.211114 |
| 20.200000 | 16.510354 |
| 19.300000 | 22.895657 |
| 34.600000 | 35.118097 |
| 50.000000 | 24.875253 |
| 24.000000 | 30.780451 |
| 18.700000 | 17.706931 |
| 19.800000 | 22.977074 |
| 22.500000 | 29.134828 |
| 13.300000 | 14.373256 |
| 50.000000 | 33.607708 |
| 11.800000 | 12.585493 |
| 11.000000 | 14.029992 |
| 23.700000 | 26.043526 |
| 35.400000 | 31.035287 |

| | |
|-----------|-----------|
| 15.200000 | 13.830639 |
| 24.400000 | 25.047405 |
| 33.400000 | 29.597219 |
| 31.600000 | 31.767734 |
| 13.400000 | 16.782462 |
| 34.900000 | 31.118824 |
| 14.400000 | 9.483307 |
| 35.400000 | 33.530729 |
| 25.300000 | 25.410127 |
| 18.300000 | 20.631338 |
| 16.600000 | 16.288110 |

0.0.4 compare the weights

```
In [137]: sklear_W = list(sklearn_W)
          own_W = list(W)
```

```
In [138]: print("Sklearn   v/s       own\n")
          print("="*50)
          for i in range(len(own_W)):
              print("%f \t %f"%(sklearn_W[i] , own_W[i]))
          print("\n\m intersept terms ")
          print("="*50)
          print("%f \t %f"%(sklearn_b, b))
```

| Sklearn | v/s | own |
|---------|-----|-----|
|---------|-----|-----|

| | |
|-----------|-----------|
| -0.652352 | -0.674132 |
| 0.547772 | 0.580637 |
| -0.318632 | -0.427761 |
| 0.748774 | 0.775076 |
| -1.195099 | -0.995977 |
| 2.993560 | 3.179379 |
| -0.085829 | -0.193168 |
| -2.042773 | -2.103919 |
| 0.959533 | 0.865578 |
| -0.463772 | -0.538941 |
| -1.927536 | -1.833581 |
| 0.812226 | 0.911603 |
| -3.551600 | -3.369270 |

\m intersept terms

| | |
|-----------|-----------|
| 22.363482 | 22.595750 |
|-----------|-----------|


```
In [140]: x = PrettyTable(["Implimentation", "learning_rate", "MSE"])

x.add_row(["sk_learn" , "0.01", 26.931])
x.add_row(["own", "0.01", 27.051])

print(x.get_string(title="SGD Model"))
```

```
+-----+-----+-----+
| Implimentation | learning_rate | MSE   |
+-----+-----+-----+
|      sk_learn  |         0.01   | 26.931 |
|         own    |         0.01   | 27.051 |
+-----+-----+-----+
```

```
In [ ]:
```