

07 Amazon Fine Food Reviews Analysis_Support Vector Machines

February 1, 2019

1 Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454 Number of users: 256,059 Number of products: 74,258 Timespan: Oct 1999 - Oct 2012 Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective: Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative? [Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

2 [1]. Reading Data

2.1 [1.1] Loading the data

The dataset is available in two forms 1. .csv file 2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to “positive”. Otherwise, it will be set to “negative”.

```
In [1]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

In [2]: # using SQLite Table to read data.
con = sqlite3.connect('database.sqlite')

# filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
# SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 500000 data points.
# you can change the number to any other number based on your computing power

# filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000 """, con)
```

```

# for tsne assignment you can take 5k data points

filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 5000

# Give reviews with Score>3 a positive rating(1), and reviews with a score<3 a negative rating(0)
def partition(x):
    if x < 3:
        return 0
    return 1

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)

```

Number of data points in our data (5000, 10)

```

Out[2]:
   Id  ProductId  UserId  ProfileName \
0   1  B001E4KFG0  A3SGXH7AUHU8GW  delmartian
1   2  B00813GRG4  A1D87F6ZCVE5NK  dll pa
2   3  B000LQOCHO  ABXLMWJIXXAIN  Natalia Corres "Natalia Corres"

   HelpfulnessNumerator  HelpfulnessDenominator  Score  Time \
0                      1                      1      1  1303862400
1                      0                      0      0  1346976000
2                      1                      1      1  1219017600

   Summary  Text
0  Good Quality Dog Food  I have bought several of the Vitality canned d...
1  Not as Advertised  Product arrived labeled as Jumbo Salted Peanut...
2  "Delight" says it all  This is a confection that has been around a fe...

```

```

In [3]: display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", con)

```

```

In [4]: print(display.shape)
display.head()

```

(80668, 7)

```

Out[4]:
   UserId  ProductId  ProfileName  Time  Score \
0  #oc-R115TNMSPFT9I7  B007Y59HVM  Breyton  1331510400  2

```

1	#oc-R11D9D7SHXIJB9	B005HG9ETO	Louis E. Emory "hoppy"	1342396800	5
2	#oc-R11DNU2NBKQ23Z	B007Y59HVM	Kim Cieszykowski	1348531200	1
3	#oc-R1105J5ZVQE25C	B005HG9ETO	Penguin Chick	1346889600	5
4	#oc-R12KPB0DL2B5ZD	B0070SBE1U	Christopher P. Presta	1348617600	1

	Text	COUNT(*)
0	Overall its just OK when considering the price...	2
1	My wife has recurring extreme muscle spasms, u...	3
2	This coffee is horrible and unfortunately not ...	2
3	This will be the bottle that you grab from the...	3
4	I didnt like this coffee. Instead of telling y...	2

```
In [5]: display[display['UserId']=='AZY10LLTJ71NX']
```

```
Out [5]:
```

	UserId	ProductId	ProfileName	Time \
80638	AZY10LLTJ71NX	B006P7E5ZI	undertheshrine "undertheshrine"	1334707200

	Score	Text	COUNT(*)
80638	5	I was recommended to try green tea extract to ...	5

```
In [6]: display['COUNT(*)'].sum()
```

```
Out [6]: 393063
```

3 [2] Exploratory Data Analysis

3.1 [2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

```
In [7]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)
display.head()
```

```
Out [7]:
```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator \
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2
3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2
4	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2

	HelpfulnessDenominator	Score	Time \
0	2	5	1199577600

1	2	5	1199577600
2	2	5	1199577600
3	2	5	1199577600
4	2	5	1199577600

	Summary \
0	LOACKER QUADRATINI VANILLA WAFERS
1	LOACKER QUADRATINI VANILLA WAFERS
2	LOACKER QUADRATINI VANILLA WAFERS
3	LOACKER QUADRATINI VANILLA WAFERS
4	LOACKER QUADRATINI VANILLA WAFERS

	Text
0	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
1	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
2	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
3	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
4	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8) ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```
In [8]: #Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False)
```

```
In [9]: #Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first')
final.shape
```

```
Out[9]: (4986, 10)
```

```
In [10]: #Checking to see how much % of data still remains
(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

```
Out[10]: 99.72
```

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

```

In [11]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)

display.head()

Out[11]:
```

	Id	ProductId	UserId	ProfileName	\
0	64422	B000MIDR0Q	A161DK06JJMCYF	J. E. Stephens	"Jeanne"
1	44737	B001EQ55RW	A2V0I904FH7ABY		Ram

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0	3	1	5	1224892800	
1	3	2	4	1212883200	

	Summary	\
0	Bought This for My Son at College	
1	Pure cocoa taste with crunchy almonds inside	

	Text
0	My son loves spaghetti so I didn't hesitate or...
1	It was almost a 'love at first bite' - the per...

```

In [12]: final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]

In [13]: #Before starting the next phase of preprocessing lets see the number of entries left
print(final.shape)

#How many positive and negative reviews are present in our dataset?
final['Score'].value_counts()

(4986, 10)

Out[13]: 1    4178
         0     808
         Name: Score, dtype: int64

```

4 [3] Preprocessing

4.1 [3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags

2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```
In [14]: # printing some random reviews
        sent_0 = final['Text'].values[0]
        print(sent_0)
        print("="*50)

        sent_1000 = final['Text'].values[1000]
        print(sent_1000)
        print("="*50)

        sent_1500 = final['Text'].values[1500]
        print(sent_1500)
        print("="*50)

        sent_4900 = final['Text'].values[4900]
        print(sent_4900)
        print("="*50)
```

```
Why is this $[...] when the same product is available for $[...] here?<br />http://www.amazon.
=====
I recently tried this flavor/brand and was surprised at how delicious these chips are.  The be
=====
Wow.  So far, two two-star reviews.  One obviously had no idea what they were ordering; the otl
=====
love to order my coffee on amazon.  easy and shows up quickly.<br />This k cup is great coffee
=====
```

```
In [15]: # remove urls from text python: https://stackoverflow.com/a/40823105/4084039
        sent_0 = re.sub(r"http\S+", "", sent_0)
        sent_1000 = re.sub(r"http\S+", "", sent_1000)
        sent_1500 = re.sub(r"http\S+", "", sent_1500)
        sent_4900 = re.sub(r"http\S+", "", sent_4900)

        print(sent_0)
```

```
Why is this $[...] when the same product is available for $[...] here?<br /> /><br />The Victor
```

```
In [16]: # https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-remove-all
        from bs4 import BeautifulSoup
```

```

soup = BeautifulSoup(sent_0, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1000, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1500, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_4900, 'lxml')
text = soup.get_text()
print(text)

```

Why is this \$[...] when the same product is available for \$[...] here? />The Victor M380 and M

=====

I recently tried this flavor/brand and was surprised at how delicious these chips are. The be

=====

Wow. So far, two two-star reviews. One obviously had no idea what they were ordering; the ot

=====

love to order my coffee on amazon. easy and shows up quickly.This k cup is great coffee. dca

```
In [17]: # https://stackoverflow.com/a/47091490/4084039
```

```

import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\'re", " are", phrase)
    phrase = re.sub(r"\'s", " is", phrase)
    phrase = re.sub(r"\'d", " would", phrase)
    phrase = re.sub(r"\'ll", " will", phrase)
    phrase = re.sub(r"\'t", " not", phrase)
    phrase = re.sub(r"\'ve", " have", phrase)
    phrase = re.sub(r"\'m", " am", phrase)
    return phrase

```

```
In [18]: sent_1500 = decontracted(sent_1500)
```



```
print(sent_1500)
print("="*50)
```

Wow. So far, two two-star reviews. One obviously had no idea what they were ordering; the other was
=====

```
In [19]: #remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
sent_0 = re.sub("\S*\d\S*", "", sent_0).strip()
print(sent_0)
```

Why is this \$[...] when the same product is available for \$[...] here?
 />
The Victor

```
In [20]: #remove spacial character: https://stackoverflow.com/a/5843547/4084039
sent_1500 = re.sub('[^A-Za-z0-9]+', ' ', sent_1500)
print(sent_1500)
```

Wow So far two two star reviews One obviously had no idea what they were ordering the other was

```
In [21]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
# instead of <br /> if we have <br/> these tags would have reumoved in the 1st step

stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves',
                'you'll', 'you'd', 'your', 'yours', 'yourself', 'yourselves', 'he', 'him',
                'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself',
                'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', 'that',
                'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had',
                'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as',
                'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through',
                'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'over',
                'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any',
                'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too',
                's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'n't',
                've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't",
                "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mightn',
                "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't",
                'won', "won't", 'wouldn', "wouldn't"])
```

```
In [22]: # Combining all the above stundents
from tqdm import tqdm
preprocessed_reviews = []
# tqdm is for printing the status bar
for sentence in tqdm(final['Text'].values):
    sentence = re.sub(r"http\S+", "", sentence)
```

```

sentence = BeautifulSoup(sentence, 'lxml').get_text()
sentence = decontracted(sentence)
sentence = re.sub("\S*\d\S*", "", sentence).strip()
sentence = re.sub('[^A-Za-z]+', ' ', sentence)
# https://gist.github.com/sebleier/554280
sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwords)
preprocessed_reviews.append(sentence.strip())

```

100%|| 4986/4986 [00:01<00:00, 3137.37it/s]

In [23]: preprocessed_reviews[1500]

Out[23]: 'wow far two two star reviews one obviously no idea ordering wants crispy cookies hey'

[3.2] Preprocessing Review Summary

In [6]: *## Similarly you can do preprocessing for review summary also.*

5 [4] Featurization

5.1 [4.1] BAG OF WORDS

In [25]: *#BoW*

```

count_vect = CountVectorizer() #in scikit-learn
count_vect.fit(preprocessed_reviews)
print("some feature names ", count_vect.get_feature_names()[:10])
print('='*50)

final_counts = count_vect.transform(preprocessed_reviews)
print("the type of count vectorizer ",type(final_counts))
print("the shape of out text BOW vectorizer ",final_counts.get_shape())
print("the number of unique words ", final_counts.get_shape()[1])

```

```

some feature names ['aa', 'aahhhs', 'aback', 'abandon', 'abates', 'abbott', 'abby', 'abdomina
=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (4986, 12997)
the number of unique words 12997

```

5.2 [4.2] Bi-Grams and n-Grams.

In [26]: *#bi-gram, tri-gram and n-gram*

```

#removing stop words like "not" should be avoided before building n-grams
# count_vect = CountVectorizer(ngram_range=(1,2))
# please do read the CountVectorizer documentation http://scikit-learn.org/stable/mod

```

```

# you can choose these numebrs min_df=10, max_features=5000, of your choice
count_vect = CountVectorizer(ngram_range=(1,2), min_df=10, max_features=5000)
final_bigram_counts = count_vect.fit_transform(preprocessed_reviews)
print("the type of count vectorizer ",type(final_bigram_counts))
print("the shape of out text BOW vectorizer ",final_bigram_counts.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_bigram_counts.get_shape()[0])

```

```

the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (4986, 3144)
the number of unique words including both unigrams and bigrams 3144

```

5.3 [4.3] TF-IDF

```

In [27]: tf_idf_vect = TfidfVectorizer(ngram_range=(1,2), min_df=10)
tf_idf_vect.fit(preprocessed_reviews)
print("some sample features(unique words in the corpus)",tf_idf_vect.get_feature_names())
print('='*50)

```

```

final_tf_idf = tf_idf_vect.transform(preprocessed_reviews)
print("the type of count vectorizer ",type(final_tf_idf))
print("the shape of out text TFIDF vectorizer ",final_tf_idf.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_tf_idf.get_shape()[0])

```

```

some sample features(unique words in the corpus) ['ability', 'able', 'able find', 'able get',
=====

```

```

the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (4986, 3144)
the number of unique words including both unigrams and bigrams 3144

```

5.4 [4.4] Word2Vec

```

In [28]: # Train your own Word2Vec model using your own text corpus
i=0
list_of_sentence=[]
for sentence in preprocessed_reviews:
    list_of_sentence.append(sentence.split())

```

```

In [42]: # Using Google News Word2Vectors

```

```

# in this project we are using a pretrained model by google
# its 3.3G file, once you load this into your memory
# it occupies ~9Gb, so please do this step only if you have >12G of ram
# we will provide a pickle file wich contains a dict ,
# and it contains all our courpus words as keys and model[word] as values
# To use this code-snippet, download "GoogleNews-vectors-negative300.bin"
# from https://drive.google.com/file/d/0B7XkCwpI5KDYNlNUTTlSS21pQmM/edit
# it's 1.9GB in size.

```

```

# http://kavita-ganesan.com/gensim-word2vec-tutorial-starter-code/#.W17SRFazZPY
# you can comment this whole cell
# or change these variable according to your need

is_your_ram_gt_16g=False
want_to_use_google_w2v = False
want_to_train_w2v = True

if want_to_train_w2v:
    # min_count = 5 considers only words that occurred atleast 5 times
    w2v_model=Word2Vec(list_of_sentence,min_count=5,size=50, workers=4)
    print(w2v_model.wv.most_similar('great'))
    print('='*50)
    print(w2v_model.wv.most_similar('worst'))

elif want_to_use_google_w2v and is_your_ram_gt_16g:
    if os.path.isfile('GoogleNews-vectors-negative300.bin'):
        w2v_model=KeyedVectors.load_word2vec_format('GoogleNews-vectors-negative300.b
        print(w2v_model.wv.most_similar('great'))
        print(w2v_model.wv.most_similar('worst'))
    else:
        print("you don't have gogole's word2vec file, keep want_to_train_w2v = True,

[('snack', 0.9951335191726685), ('calorie', 0.9946465492248535), ('wonderful', 0.9946032166481
=====
[('varieties', 0.9994194507598877), ('become', 0.9992934465408325), ('popcorn', 0.999275088310

In [36]: w2v_words = list(w2v_model.wv.vocab)
print("number of words that occurred minimum 5 times ",len(w2v_words))
print("sample words ", w2v_words[0:50])

number of words that occurred minimum 5 times 3817
sample words  ['product', 'available', 'course', 'total', 'pretty', 'stinky', 'right', 'nearby

```

5.5 [4.4.1] Converting text into vectors using Avg W2V, TFIDF-W2V

[4.4.1.1] Avg W2v

```

In [38]: # average Word2Vec
# compute average word2vec for each review.
sent_vectors = []; # the avg-w2v for each sentence/review is stored in this list
for sent in tqdm(list_of_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need t
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence

```

```

        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors.append(sent_vec)
print(len(sent_vectors))
print(len(sent_vectors[0]))

```

100%| 4986/4986 [00:03<00:00, 1330.47it/s]

4986

50

[4.4.1.2] TFIDF weighted W2v

```

In [39]: # S = ["abc def pqr", "def def def abc", "pqr pqr def"]
model = TfidfVectorizer()
tf_idf_matrix = model.fit_transform(preprocessed_reviews)
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

In [41]: # TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this l
row=0;
for sent in tqdm(list_of_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            # tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors.append(sent_vec)
    row += 1

```

6 [5] Assignment 7: SVM

Apply SVM on these feature sets

SET 1:Review text, preprocessed one converted into vectors using (BOW)

SET 2:Review text, preprocessed one converted into vectors using (TFIDF)

SET 3:Review text, preprocessed one converted into vectors using (AVG W2v)

SET 4:Review text, preprocessed one converted into vectors using (TFIDF W2v)

Procedure

You need to work with 2 versions of SVM

Linear kernel

RBF kernel

When you are working with linear kernel, use SGDClassifier with hinge loss because it is c

When you are working with SGDClassifier with hinge loss and trying to find the AUC

score, you would have to use <a href='https://scikit-learn.org/stable/modules/generated/sk

Similarly, like kdtree of knn, when you are working with RBF kernel it's better to reduce

the number of dimensions. You can put min_df = 10, max_features = 500 and consider a sample size of 40k points.

Hyper paramter tuning (find best alpha in range [10^{-4} to 10^4], and the best pena

Find the best hyper parameter which will give the maximum <a href='https://www.appliedaicom

Find the best hyper paramter using k-fold cross validation or simple cross validation data

Use gridsearch cv or randomsearch cv or you can also write your own for loops to do this t

Feature importance

When you are working on the linear kernel with BOW or TFIDF please print the top 10 best

features for each of the positive and negative classes.

Feature engineering

To increase the performance of your model, you can also experiment with with feature engin

Taking length of reviews as another feature.

```

        <li>Considering some features from review summary as well.</li>
    </ul>
</ul>
</li>
<br>
<li><strong>Representation of results</strong>
    <ul>
<li>You need to plot the performance of model both on train data and cross validation data for
<img src='train_cv_auc.JPG' width=300px></li>
<li>Once after you found the best hyper parameter, you need to train your model with it, and f
<img src='train_test_auc.JPG' width=300px></li>
<li>Along with plotting ROC curve, you need to print the <a href='https://www.appliedaicourse.
<img src='confusion_matrix.png' width=300px></li>
    </ul>
</li>
<br>
<li><strong>Conclusion</strong>
    <ul>
<li>You need to summarize the results at the end of the notebook, summarize it in the table for
    <img src='summary.JPG' width=400px>
</li>
    </ul>

```

Note: Data Leakage

1. There will be an issue of data-leakage if you vectorize the entire data and then split it into train/cv/test.
2. To avoid the issue of data-leakag, make sure to split your data first and then vectorize it.
3. While vectorizing your data, apply the method fit_transform() on you train data, and apply the method transform() on cv/test data.
4. For more details please go through this link.

7 Applying SVM

```

In [2]: # after preprocessing
df = pd.read_pickle("files/preprocessed.pkl")
df.shape

```

```
Out[2]: (364171, 12)
```

```
In [3]: df.head()
```

```

Out[3]:
      Id  ProductId  UserId  ProfileName \
138706  150524  0006641040  ACITT7DI6IDDL  shari zychinski
138688  150506  0006641040  A2IW4PEEK02R0U  Tracy
138689  150507  0006641040  A1S4A3IQ2MU7V4  sally sue "sally sue"
138690  150508  0006641040  AZGXZ2UUK6X  Catherine Hallberg "(Kate)"
138691  150509  0006641040  A3CMRKGEOP909G  Teresa

```

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time \
138706	0	0	1	939340800
138688	1	1	1	1194739200
138689	1	1	1	1191456000
138690	1	1	1	1076025600
138691	3	4	1	1018396800

	Summary \
138706	EVERY book is educational
138688	Love the book, miss the hard cover version
138689	chicken soup with rice months
138690	a good swingy rhythm for reading aloud
138691	A great way to learn the months

	Text \
138706	this witty little book makes my son laugh at l...
138688	I grew up reading these Sendak books, and watc...
138689	This is a fun way for children to learn their ...
138690	This is a great little book to read aloud- it ...
138691	This is a book of poetry about the months of t...

	CleanedText \
138706	witty little book makes son laugh loud recite ...
138688	grew reading sendak books watching really rosi...
138689	fun way children learn months year learn poems...
138690	great little book read aloud nice rhythm well ...
138691	book poetry months year goes month cute little...

	CleanedSummary
138706	every book educational
138688	love book miss hard cover version
138689	chicken soup rice months
138690	good swingy rhythm reading aloud
138691	great way learn months

```
In [48]: #code source: http://occam.olin.edu/sites/default/files/DataScienceMaterials/machine_
from sklearn.model_selection import train_test_split
from sklearn.grid_search import GridSearchCV
from sklearn.datasets import *
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score , f1_score , confusion_matrix
from collections import Counter
from sklearn.metrics import accuracy_score, roc_auc_score , roc_curve
from sklearn.model_selection import train_test_split
```

```
In [49]: # take 50k sample data randomly
sample_data = df.sample(100000)
sample_data.shape
```



```
Out[49]: (100000, 12)
```

```
In [50]: # sorted the data using time based
sorted_data = sample_data.sort_values('Time', axis=0, inplace=False)
sorted_data.shape
```

```
Out[50]: (100000, 12)
```

```
In [51]: sorted_data['Score'].value_counts()
```

```
Out[51]: 1      84429
         0      15571
         Name: Score, dtype: int64
```

```
In [52]: from sklearn.model_selection import train_test_split
```

```
In [53]: X = np.array(sorted_data['CleanedText'])
y = np.array(sorted_data['Score'])
print(X.shape)
print(y.shape)
```

```
(100000,)
(100000,)
```

```
In [54]: # Simple cross validation
# split the data sent into train and test
train , test , train_y , test_y = train_test_split(X, y, test_size = 0.3, random_state=42)

# split the train data set into cross validation train and cross validation test
train, cv , train_y, cv_y = train_test_split(train, train_y, test_size=0.3, random_state=42)

print("train data = ", train.shape)
print("cros validation = ", cv.shape)
print("test data = ", test.shape)
```

```
train data = (49000,)
cros validation = (21000,)
test data = (30000,)
```

7.1 [5.1] Linear SVM

7.1.1 [5.1.1] Applying Linear SVM on BOW, SET 1

```
In [3]: # Please write all the code with proper documentation
```

```
In [15]: # Please write all the code with proper documentation
# Please write all the code with proper documentation
#BoW
```

```

count_vect = CountVectorizer(min_df=10) #in scikit-learn
count_vect.fit(train)
print("some feature names ", count_vect.get_feature_names()[:10])
print('='*50)

bow_train = count_vect.fit_transform(train)
bow_cv = count_vect.transform(cv)
bow_test = count_vect.transform(test)
print("=====Train Data=====")
print("the type of count vectorizer ",type(bow_train))
print("the shape of out text BOW vectorizer ",bow_train.get_shape())
print("the number of unique words ", bow_train.get_shape()[1])
print("=====Cross validation Data=====")
print("the type of count vectorizer ",type(bow_cv))
print("the shape of out text BOW vectorizer ",bow_cv.get_shape())
print("the number of unique words ", bow_cv.get_shape()[1])
print("=====Test Data=====")
print("the type of count vectorizer ",type(bow_test))
print("the shape of out text BOW vectorizer ",bow_test.get_shape())
print("the number of unique words ", bow_test.get_shape()[1])

some feature names  ['aa', 'ability', 'able', 'abroad', 'absent', 'absolute', 'absolutely', 'al
=====
=====Train Data=====
the type of count vectorizer  <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer  (49000, 8771)
the number of unique words  8771
=====Cross validation Data=====
the type of count vectorizer  <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer  (21000, 8771)
the number of unique words  8771
=====Test Data=====
the type of count vectorizer  <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer  (30000, 8771)
the number of unique words  8771

```

```

In [61]: from sklearn.linear_model import SGDClassifier
         from sklearn.calibration import CalibratedClassifierCV

```

using L2 reglirization

```

In [37]: alpha = [pow(10,j) for j in range(-4,4,1)]

bow_train_auc = []
bow_cv_auc = []
for i in alpha:
    LSVM = SGDClassifier(loss='hinge',alpha=i, penalty='l2')

```

```

calibrated_clf = CalibratedClassifierCV(LSVM, cv=5, method='sigmoid')
calibrated_clf.fit(bow_train, train_y)
# train data
y_prob_train = calibrated_clf.predict_proba(bow_train)[: ,1]
y_pred = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y , y_prob_train)
print('\nTrain AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_train * float(100))))
bow_train_auc.append(auc_roc_train)
# CV
y_prob_cv = calibrated_clf.predict_proba(bow_cv)[: ,1]
y_pred = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
print('\nCV AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_cv * float(100))))
bow_cv_auc.append(auc_roc_cv)
print("="*50)

```

Train AUC for = 0.0001 is 97.49%

CV AUC for = 0.0001 is 93.87%

=====

Train AUC for = 0.001 is 96.37%

CV AUC for = 0.001 is 94.19%

=====

Train AUC for = 0.01 is 93.49%

CV AUC for = 0.01 is 92.52%

=====

Train AUC for = 0.1 is 78.23%

CV AUC for = 0.1 is 77.48%

=====

Train AUC for = 1 is 52.53%

CV AUC for = 1 is 52.90%

=====

Train AUC for = 10 is 55.79%

CV AUC for = 10 is 55.25%

=====

Train AUC for = 100 is 56.10%

CV AUC for = 100 is 55.56%

=====

Train AUC for = 1000 is 56.10%

CV AUC for = 1000 is 55.56%

=====

```
In [38]: hyper = [str(pow(10,j)) for j in range(-4,4)]
```

```
    # https://www.dataquest.io/blog/learning-curves-machine-learning/
```

```
import matplotlib.pyplot as plt
%matplotlib inline
```


```
plt.style.use('seaborn')
```

```
plt.plot(hyper,bow_train_auc,label = 'Training AUC')
plt.plot(hyper, bow_cv_auc, label = 'Validation AUC')
```

```
plt.ylabel('AUC', fontsize = 14)
plt.xlabel('\u03BB-values', fontsize = 14)
plt.title('Learning curves for a Logistic Regression model', fontsize = 18, y = 1.03)
plt.legend()
```

```
Out[38]: <matplotlib.legend.Legend at 0x7f70c50b87f0>
```

Learning curves for a Logistic Regression model



```
In [39]: i = 0.001
LSTM = SGDClassifier(loss='hinge',alpha=i, penalty='l2')
calibrated_clf = CalibratedClassifierCV(LSTM, cv=5, method='sigmoid')
calibrated_clf.fit(bow_train, train_y)
# train data
y_prob_train = calibrated_clf.predict_proba(bow_train)[:,-1]
fpr_t, tpr_t, threshold_t = roc_curve(train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y, y_prob_train)
print('\nTrain AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_train * float(100))))
# CV
y_prob_cv = calibrated_clf.predict_proba(bow_cv)[:,-1]
fpr_c, tpr_c, threshold_c = roc_curve(cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y, y_prob_cv)
print('\nCV AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_cv * float(100))))
# Test
y_prob_test = calibrated_clf.predict_proba(bow_test)[:,-1]
fpr_ts, tpr_ts, threshold_ts = roc_curve(test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
auc_roc_test = roc_auc_score(test_y, y_prob_test)
```

```
print('\nTest AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_test * float(100))))
```

Train AUC for $\lambda = 0.001$ is 96.40%

CV AUC for $\lambda = 0.001$ is 94.19%

Test AUC for $\lambda = 0.001$ is 93.92%

```
In [64]: LSVM = SGDClassifier(loss='hinge',alpha=0.001, penalty='l2')
LSVM.fit(bow_train, train_y)
```

```
# number of non-zero weights
```

```
w = LSVM.coef_
```

```
print("Number of non-zero weights : ",np.count_nonzero(w))
```

Number of non-zero weights : 8738

```
In [41]: # https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python
```

```
import matplotlib.pyplot as plt
```

```
plt.title('Receiver Operating Characteristic')
```

```
plt.plot(fprc, tprc, 'b' , label='AUC CV = %0.2f' % (auc_roc_cv * float(100)))
```

```
plt.plot(fprts, tprts, 'y' , label='AUC Test = %0.2f' % (auc_roc_test * float(100)))
```

```
plt.plot(fprt, tprt, 'g', label='AUC Train = %0.2f' % (auc_roc_train * float(100)))
```

```
plt.legend(loc = 'lower right')
```

```
plt.plot([0, 1], [0, 1], 'r')
```


```
plt.xlim([0, 1])
```

```
plt.ylim([0, 1])
```

```
plt.ylabel('True Positive Rate')
```

```
plt.xlabel('False Positive Rate')
```

```
plt.show()
```




```
In [42]: print("F1-Score on test set: %0.2f"%(f1_score(test_y, y_pred_test)))
```

F1-Score on test set: 0.95

```
In [43]: df_cm = pd.DataFrame(confusion_matrix(test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[43]: <matplotlib.axes._subplots.AxesSubplot at 0x7f70c51ba240>
```



Top 10 important features of positive class

```
In [67]: # Please write all the code with proper documentation
# Please write all the code with proper documentation
feature_weights = w.tolist()[0]
feature_names = count_vect.get_feature_names()
features = dict(zip(feature_weights, feature_names))

In [68]: features_df = pd.DataFrame.from_dict(features, orient='index')
sorted_features = features_df.sort_index(axis=0, ascending=False)

In [69]: sorted_features.head(10)

Out[69]:
```

	0
0.619961	excellent
0.611803	delicious
0.554702	highly
0.534308	perfect
0.526151	awesome
0.501679	smooth
0.497600	best
0.489443	amazing
0.485364	pleased
0.469049	great

7.1.2 Top 10 important features of negative class from

```
In [70]: # Please write all the code with proper documentation
sorted_features = features_df.sort_index(axis=0,ascending=True)
sorted_features.head(10)
```

```
Out [70]:
```

	0
-1.007436	worst
-0.880997	disappointing
-0.876918	terrible
-0.856525	disappointed
-0.758636	horrible
-0.754557	unfortunately
-0.738243	threw
-0.681141	return
-0.664826	disappointment
-0.624039	stale

using L1 regularization

```
In [32]: alpha = [pow(10,j) for j in range(-4,4,1)]

bow_train_auc = []
bow_cv_auc = []
for i in alpha:
    LSVM = SGDClassifier(loss='hinge',alpha=i, penalty='l1')
    calibrated_clf = CalibratedClassifierCV(LSVM, cv=5, method='sigmoid')
    calibrated_clf.fit(bow_train, train_y)
    # train data
    y_prob_train = calibrated_clf.predict_proba(bow_train)[: ,1]
    y_pred = np.where(y_prob_train > 0.5, 1, 0)
    auc_roc_train = roc_auc_score(train_y , y_prob_train)
    print('\nTrain AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_train * float(100))))
    bow_train_auc.append(auc_roc_train)
    # CV
    y_prob_cv = calibrated_clf.predict_proba(bow_cv)[: ,1]
    y_pred = np.where(y_prob_cv > 0.5, 1, 0)
    auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
    print('\nCV AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_cv * float(100))))
    bow_cv_auc.append(auc_roc_cv)
    print("="*50)
```

Train AUC for = 0.0001 is 95.73%

CV AUC for = 0.0001 is 93.49%

=====

Train AUC for = 0.001 is 91.17%

```

CV AUC for  = 0.001 is 90.66%
=====

Train AUC for  = 0.01 is 70.50%

CV AUC for  = 0.01 is 69.82%
=====

Train AUC for  = 0.1 is 54.77%

CV AUC for  = 0.1 is 54.42%
=====

Train AUC for  = 1 is 50.00%

CV AUC for  = 1 is 50.00%
=====

Train AUC for  = 10 is 50.00%

CV AUC for  = 10 is 50.00%
=====

Train AUC for  = 100 is 50.00%

CV AUC for  = 100 is 50.00%
=====

Train AUC for  = 1000 is 50.00%

CV AUC for  = 1000 is 50.00%
=====

```

```

In [33]: hyper = [str(pow(10,j)) for j in range(-4,4)]

          # https://www.dataquest.io/blog/learning-curves-machine-learning/

          import matplotlib.pyplot as plt
          %matplotlib inline

          plt.style.use('seaborn')


          plt.plot(hyper,bow_train_auc,label = 'Training AUC')
          plt.plot(hyper, bow_cv_auc, label = 'Validation AUC')

          plt.ylabel('AUC', fontsize = 14)

```

```
plt.xlabel('\u03BB-values', fontsize = 14)
plt.title('Learning curves for a Logistic Regression model', fontsize = 18, y = 1.03)
plt.legend()
```

Out [33]: <matplotlib.legend.Legend at 0x7f70c57b6c50>



```
In [34]: i = 0.0001
LSVM = SGDClassifier(loss='hinge',alpha=i, penalty='l1')
calibrated_clf = CalibratedClassifierCV(LSVM, cv=5, method='sigmoid')
calibrated_clf.fit(bow_train, train_y)
# train data
y_prob_train = calibrated_clf.predict_proba(bow_train)[:,-1]
fpr_t, tpr_t, threshold_t = roc_curve(train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y, y_prob_train)
print('\nTrain AUC for \u03BB = %s is %.2f%%' % (str(i), (auc_roc_train * float(100))))
# CV
y_prob_cv = calibrated_clf.predict_proba(bow_cv)[:,-1]
fpr_c, tpr_c, threshold_c = roc_curve(cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y, y_prob_cv)
```

```

print('\nCV AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_cv * float(100))))
# Test
y_prob_test = calibrated_clf.predict_proba(bow_test)[:,-1]
fprts, tprts, thresholdts = roc_curve(test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
auc_roc_test = roc_auc_score(test_y, y_prob_test)
print('\nTest AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_test * float(100))))

```

Train AUC for = 0.0001 is 95.77%

CV AUC for = 0.0001 is 93.50%

Test AUC for = 0.0001 is 93.07%

```

In [46]: LSVM = SGDClassifier(loss='hinge',alpha=0.0001, penalty='l1')
         LSVM.fit(bow_train, train_y)

```

```

# number of non-zero weights
w = LSVM.coef_
print("Number of weights : ", w.shape[1])
print("Number of non-zero weights : ",np.count_nonzero(w))

```

Number of weights : 8771

Number of non-zero weights : 2306

```


In [36]: # https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python

```

```

import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b', label='AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y', label='AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label='AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

```




```
In [44]: print("F1-Score on test set: %0.2f"%(f1_score(test_y, y_pred_test)))
```

F1-Score on test set: 0.95

```
In [45]: df_cm = pd.DataFrame(confusion_matrix(test_y, y_pred_test), range(2), range(2))  
sns.set(font_scale=1.4)  
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[45]: <matplotlib.axes._subplots.AxesSubplot at 0x7f70c5a679e8>
```



7.1.3 [5.1.2] Applying Linear SVM on TFIDF, SET 2

In [47]: *# Please write all the code with proper documentation*

```

model = TfidfVectorizer(min_df=20, ngram_range=(1,2))
#tf_idf_matrix = model.fit_transform(train)

print("=====Train Data=====")
tf_idf_train = model.fit_transform(train)
print("the type of count vectorizer ",type(tf_idf_train))
print("the shape of out text TFIDF vectorizer ",tf_idf_train.get_shape())
print("the number of unique words including both unigrams and bigrams ",tf_idf_train.get_shape()[0])
print("=====CV Data=====")
tf_idf_cv = model.transform(cv)
print("the type of count vectorizer ",type(tf_idf_cv))
print("the shape of out text TFIDF vectorizer ",tf_idf_cv.get_shape())
print("the number of unique words including both unigrams and bigrams ",tf_idf_cv.get_shape()[0])
print("=====Test Data=====")
tf_idf_test = model.transform(test)
print("the type of count vectorizer ",type(tf_idf_test))
print("the shape of out text TFIDF vectorizer ",tf_idf_test.get_shape())
print("the number of unique words including both unigrams and bigrams ", tf_idf_test.get_shape()[0])

```

```
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))
```

```
=====Train Data=====
```

```
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (49000, 13648)
the number of unique words including both unigrams and bigrams 13648
```

```
=====CV Data=====
```

```
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (21000, 13648)
the number of unique words including both unigrams and bigrams 13648
```

```
=====Test Data=====
```

```
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (30000, 13648)
the number of unique words including both unigrams and bigrams 13648
```

```
In [51]: alpha = [pow(10,j) for j in range(-4,4,1)]
```

```
tfidf_train_auc = []
```

```
tfidf_cv_auc = []
```

```
for i in alpha:
```

```
    LSVM = SGDClassifier(loss='hinge',alpha=i, penalty='l2')
```

```
    calibrated_clf = CalibratedClassifierCV(LSVM, cv=5, method='sigmoid')
```

```
    calibrated_clf.fit(tf_idf_train, train_y)
```

```
    # train data
```

```
    y_prob_train = calibrated_clf.predict_proba(tf_idf_train)[: ,1]
```

```
    y_pred = np.where(y_prob_train > 0.5, 1, 0)
```

```
    auc_roc_train = roc_auc_score(train_y , y_prob_train)
```

```
    print('\nTrain AUC for \u03BB = %s is %.2f%%' % (str(i), (auc_roc_train * float(100))))
```

```
    tfidf_train_auc.append(auc_roc_train)
```

```
    # CV
```

```
    y_prob_cv = calibrated_clf.predict_proba(tf_idf_cv)[: ,1]
```

```
    y_pred = np.where(y_prob_cv > 0.5, 1, 0)
```

```
    auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
```

```
    print('\nCV AUC for \u03BB = %s is %.2f%%' % (str(i), (auc_roc_cv * float(100))))
```

```
    tfidf_cv_auc.append(auc_roc_cv)
```

```
    print("="*50)
```

```
Train AUC for  = 0.0001 is 97.29%
```

```
CV AUC for  = 0.0001 is 95.73%
```

```
=====
```

```
Train AUC for  = 0.001 is 96.14%
```

```

CV AUC for  = 0.001 is 95.00%
=====

Train AUC for  = 0.01 is 96.17%

CV AUC for  = 0.01 is 95.04%
=====

Train AUC for  = 0.1 is 61.07%

CV AUC for  = 0.1 is 61.09%
=====

Train AUC for  = 1 is 61.07%

CV AUC for  = 1 is 61.09%
=====

Train AUC for  = 10 is 61.07%

CV AUC for  = 10 is 61.09%
=====

Train AUC for  = 100 is 61.07%

CV AUC for  = 100 is 61.09%
=====

Train AUC for  = 1000 is 61.07%

CV AUC for  = 1000 is 61.09%
=====

```

In [52]: # <https://www.dataquest.io/blog/learning-curves-machine-learning/>

```

import matplotlib.pyplot as plt
%matplotlib inline

plt.style.use('seaborn')

plt.plot(hyper,tfidf_train_auc,label = 'Training AUC')
plt.plot(hyper, tfidf_cv_auc, label = 'Validation AUC')

plt.ylabel('AUC', fontsize = 14)
plt.xlabel('\u03BB-values', fontsize = 14)
plt.title('Learning curves for a Logistic Regression model', fontsize = 18, y = 1.03)
plt.legend()

```


Out [52]: <matplotlib.legend.Legend at 0x7f70a948abe0>



```
In [53]: i = 0.0001
LSVM = SGDClassifier(loss='hinge',alpha=i, penalty='l2')
calibrated_clf = CalibratedClassifierCV(LSVM, cv=5, method='sigmoid')
calibrated_clf.fit(tf_idf_train, train_y)
# train data
y_prob_train = calibrated_clf.predict_proba(tf_idf_train)[:,-1]
fpr_t, tpr_t, threshold_t = roc_curve(train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y, y_prob_train)
print('\nTrain AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_train * float(100))))
# CV
y_prob_cv = calibrated_clf.predict_proba(tf_idf_cv)[:,-1]
fpr_c, tpr_c, threshold_c = roc_curve(cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y, y_prob_cv)
print('\nCV AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_cv * float(100))))
# Test
y_prob_test = calibrated_clf.predict_proba(tf_idf_test)[:,-1]
fpr_ts, tpr_ts, threshold_ts = roc_curve(test_y, y_prob_test)
```

```

y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
auc_roc_test = roc_auc_score(test_y , y_prob_test)
print('\nTest AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_test * float(100))))

```

Train AUC for = 0.0001 is 97.29%

CV AUC for = 0.0001 is 95.72%


Test AUC for = 0.0001 is 95.82%

In [54]: # <https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python>

```

import matplotlib.pyplot as plt
plt.clf()
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b' , label='AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y' , label='AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label='AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

```




```
In [55]: print("F1-Score on test set: %0.2f"%(f1_score(test_y, y_pred_test)))
```

F1-Score on test set: 0.96

```
In [56]: df_cm = pd.DataFrame(confusion_matrix(test_y, y_pred_test), range(2), range(2))  
sns.set(font_scale=1.4)  
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[56]: <matplotlib.axes._subplots.AxesSubplot at 0x7f70ade07f98>
```



```
In [71]: LSVM = SGDClassifier(loss='hinge',alpha=0.0001, penalty='l2')
LSVM.fit(tf_idf_train, train_y)
```

```
# number of non-zero weights
w = LSVM.coef_
print("Number of weights : ", w.shape[1])
print("Number of non-zero weights : ",np.count_nonzero(w))
```

Number of weights : 13648

Number of non-zero weights : 13626

Top 10 important features of positive class

```
In [72]: # Please write all the code with proper documentation
# Please write all the code with proper documentation
feature_weights = w.tolist()[0]
feature_names = model.get_feature_names()
features = dict(zip(feature_weights,feature_names))
```

```
In [73]: features_df = pd.DataFrame.from_dict(features, orient='index')
sorted_features = features_df.sort_index(axis=0,ascending=False )
```

```
In [74]: sorted_features.head(10)
```

```
Out[74]:
```

	0
3.520651	great
2.848719	best
2.816173	not disappointed
2.658403	delicious
2.384709	good
2.140364	perfect
2.136886	love
2.105295	excellent
1.771214	tasty
1.736261	loves

7.1.4 Top 10 important features of negative class

```
In [75]: # Please write all the code with proper documentation
sorted_features = features_df.sort_index(axis=0,ascending=True)
sorted_features.head(10)
```

```
Out[75]:
```

	0
-4.553429	disappointed
-3.342798	worst
-3.265935	terrible
-3.085888	awful
-2.990699	not buy
-2.934224	horrible
-2.857275	not worth
-2.834140	not good
-2.823619	disappointing
-2.783166	not

using L1 regularization

```
In [57]: alpha = [pow(10,j) for j in range(-4,4,1)]

tfidf_train_auc = []
tfidf_cv_auc = []
for i in alpha:
    LSVM = SGDClassifier(loss='hinge',alpha=i, penalty='l1')
    calibrated_clf = CalibratedClassifierCV(LSVM, cv=5, method='sigmoid')
    calibrated_clf.fit(tf_idf_train, train_y)
    # train data
    y_prob_train = calibrated_clf.predict_proba(tf_idf_train)[: ,1]
    y_pred = np.where(y_prob_train > 0.5, 1, 0)
    auc_roc_train = roc_auc_score(train_y , y_prob_train)
    print('\nTrain AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_train * float(100))))
    tfidf_train_auc.append(auc_roc_train)
# CV
```

```

y_prob_cv = calibrated_clf.predict_proba(tf_idf_cv)[: ,1]
y_pred = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
print('\nCV AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_cv * float(100))))
tfidf_cv_auc.append(auc_roc_cv)
print("="*50)

```

Train AUC for = 0.0001 is 94.34%

CV AUC for = 0.0001 is 93.81%

=====

Train AUC for = 0.001 is 76.74%

CV AUC for = 0.001 is 76.31%

=====

Train AUC for = 0.01 is 50.00%

CV AUC for = 0.01 is 50.00%

=====

Train AUC for = 0.1 is 50.00%

CV AUC for = 0.1 is 50.00%

=====

Train AUC for = 1 is 50.00%

CV AUC for = 1 is 50.00%

=====

Train AUC for = 10 is 50.00%

CV AUC for = 10 is 50.00%

=====

Train AUC for = 100 is 50.00%

CV AUC for = 100 is 50.00%

=====

Train AUC for = 1000 is 50.00%

CV AUC for = 1000 is 50.00%

=====

```
In [58]: # https://www.dataquest.io/blog/learning-curves-machine-learning/

import matplotlib.pyplot as plt
%matplotlib inline

plt.style.use('seaborn')

plt.plot(hyper, tfidf_train_auc, label = 'Training AUC')
plt.plot(hyper, tfidf_cv_auc, label = 'Validation AUC')

plt.ylabel('AUC', fontsize = 14)
plt.xlabel('\u03BB-values', fontsize = 14)
plt.title('Learning curves for a Logistic Regression model', fontsize = 18, y = 1.03)
plt.legend()

Out[58]: <matplotlib.legend.Legend at 0x7f70ad431668>
```



```
In [59]: i = 0.0001
LSVM = SGDClassifier(loss='hinge', alpha=i, penalty='l1')
calibrated_clf = CalibratedClassifierCV(LSVM, cv=5, method='sigmoid')
calibrated_clf.fit(tf_idf_train, train_y)
```

```

# train data
y_prob_train = calibrated_clf.predict_proba(tf_idf_train)[: ,1]
fpr_t, tpr_t, threshold_t = roc_curve(train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y , y_prob_train)
print('\nTrain AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_train * float(100))))
# CV
y_prob_cv = calibrated_clf.predict_proba(tf_idf_cv)[: ,1]
fpr_c, tpr_c, threshold_c = roc_curve(cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
print('\nCV AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_cv * float(100))))
# Test
y_prob_test = calibrated_clf.predict_proba(tf_idf_test)[: ,1]
fpr_t, tpr_t, threshold_t = roc_curve(test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
auc_roc_test = roc_auc_score(test_y , y_prob_test)
print('\nTest AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_test * float(100))))

```

Train AUC for = 0.0001 is 94.32%

CV AUC for = 0.0001 is 93.80%


Test AUC for = 0.0001 is 93.97%

In [60]: # <https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python>

```

import matplotlib.pyplot as plt
plt.clf()
plt.title('Receiver Operating Characteristic')
plt.plot(fpr_c, tpr_c, 'b' , label='AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fpr_t, tpr_t, 'y' , label='AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fpr_t, tpr_t, 'g', label='AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

```

```
In [61]: LSVM = SGDClassifier(loss='hinge',alpha=0.0001, penalty='l1')
LSVM.fit(tf_idf_train, train_y)
```

```
# number of non-zero weights
w = LSVM.coef_
print("Number of weights : ", w.shape[1])
print("Number of non-zero weights : ",np.count_nonzero(w))
```


```
Number of weights : 13648
Number of non-zero weights : 397
```

```
In [62]: print("F1-Score on test set: %0.2f"%(f1_score(test_y, y_pred_test)))
```

```
F1-Score on test set: 0.95
```

```
In [63]: df_cm = pd.DataFrame(confusion_matrix(test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[63]: <matplotlib.axes._subplots.AxesSubplot at 0x7f70ae92aac8>
```



7.1.5 [5.1.3] Applying Linear SVM on AVG W2V, SET 3

```
In [76]: # Train your own Word2Vec model using your own text corpus
##### Train Set #####
i=0
list_of_train_sentence=[]
for sentence in train:
    list_of_train_sentence.append(sentence.split())
##### CV Set #####
i=0
list_of_cv_sentence=[]
for sentence in cv:
    list_of_cv_sentence.append(sentence.split())
##### Test Set #####
i=0
list_of_test_sentence=[]
for sentence in test:
    list_of_test_sentence.append(sentence.split())
print("Length of Train = ", len(list_of_train_sentence))
print("Length of CV = ", len(list_of_cv_sentence))
print("Length of Test = ", len(list_of_test_sentence))
```

Length of Train = 49000

Length of CV = 21000
Length of Test = 30000

```
In [77]: w2v_model=Word2Vec(list_of_train_sentence,min_count=15,size=100, workers=4)
         print(w2v_model.wv.most_similar('great'))
         print('='*50)
         print(w2v_model.wv.most_similar('worst'))
```

```
[('terrific', 0.7772042751312256), ('excellent', 0.7740203142166138), ('fantastic', 0.76137471
=====
[('best', 0.7626129984855652), ('greatest', 0.7530427575111389), ('tastiest', 0.68608349561691
```

```
In [78]: w2v_words = list(w2v_model.wv.vocab)
         print("number of words that occurred minimum 5 times ",len(w2v_words))
         print("sample words ", w2v_words[0:50])
```

number of words that occurred minimum 5 times 7492
sample words ['fully', 'intended', 'take', 'work', 'share', 'friend', 'addictive', 'ended', 'e

```
In [79]: ##### Train data #####
         # average Word2Vec
         # compute average word2vec for each review.
         sent_vectors_train = []; # the avg-w2v for each sentence/review is stored in this list
         for sent in tqdm(list_of_train_sentence): # for each review/sentence
             sent_vec = np.zeros(100) # as word vectors are of zero length 50, you might need
             cnt_words = 0; # num of words with a valid vector in the sentence/review
             for word in sent: # for each word in a review/sentence
                 if word in w2v_words:
                     vec = w2v_model.wv[word]
                     sent_vec += vec
                     cnt_words += 1
             if cnt_words != 0:
                 sent_vec /= cnt_words
             sent_vectors_train.append(sent_vec)
         print(len(sent_vectors_train))
         print(len(sent_vectors_train[0]))
```

100%|| 49000/49000 [02:16<00:00, 407.72it/s]

49000
100

```
In [80]: ##### CV data #####
# average Word2Vec
# compute average word2vec for each review.
sent_vectors_cv = []; # the avg-w2v for each sentence/review is stored in this list
for sent in tqdm(list_of_cv_sentence): # for each review/sentence
    sent_vec = np.zeros(100) # as word vectors are of zero length 50, you might need
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors_cv.append(sent_vec)
print(len(sent_vectors_cv))
print(len(sent_vectors_cv[0]))

100%|| 21000/21000 [00:58<00:00, 356.99it/s]

21000
100
```

```
In [81]: ##### Test data #####
# average Word2Vec
# compute average word2vec for each review.
sent_vectors_test = []; # the avg-w2v for each sentence/review is stored in this list
for sent in tqdm(list_of_test_sentence): # for each review/sentence
    sent_vec = np.zeros(100) # as word vectors are of zero length 50, you might need
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors_test.append(sent_vec)
print(len(sent_vectors_test))
print(len(sent_vectors_test[0]))

100%|| 30000/30000 [01:20<00:00, 372.50it/s]

30000
100
```

```
In [82]: # save the datasets as numpy array
w2v_train = np.array(sent_vectors_train)
w2v_cv = np.array(sent_vectors_cv)
w2v_test = np.array(sent_vectors_test)
```

7.1.6 Using L2 regularization

```
In [87]: alpha = [pow(10,j) for j in range(-4,4,1)]

w2v_train_auc = []
w2v_cv_auc = []
for i in alpha:
    LSVM = SGDClassifier(loss='hinge',alpha=i, penalty='l2')
    calibrated_clf = CalibratedClassifierCV(LSVM, cv=5, method='sigmoid')
    calibrated_clf.fit(w2v_train, train_y)
    # train data
    y_prob_train = calibrated_clf.predict_proba(w2v_train)[:,-1]
    y_pred = np.where(y_prob_train > 0.5, 1, 0)
    auc_roc_train = roc_auc_score(train_y , y_prob_train)
    print('\nTrain AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_train * float(100))))
    w2v_train_auc.append(auc_roc_train)
    # CV
    y_prob_cv = calibrated_clf.predict_proba(w2v_cv)[:,-1]
    y_pred = np.where(y_prob_cv > 0.5, 1, 0)
    auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
    print('\nCV AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_cv * float(100))))
    w2v_cv_auc.append(auc_roc_cv)
    print("="*50)
```

Train AUC for = 0.0001 is 91.25%

CV AUC for = 0.0001 is 91.26%

=====

Train AUC for = 0.001 is 91.16%

CV AUC for = 0.001 is 91.19%

=====

Train AUC for = 0.01 is 90.69%

CV AUC for = 0.01 is 90.67%

=====

Train AUC for = 0.1 is 90.51%

CV AUC for = 0.1 is 90.48%

=====

Train AUC for = 1 is 90.49%

CV AUC for = 1 is 90.48%

=====

Train AUC for = 10 is 83.90%

CV AUC for = 10 is 83.83%

=====

Train AUC for = 100 is 65.58%

CV AUC for = 100 is 65.29%

=====

Train AUC for = 1000 is 65.58%

CV AUC for = 1000 is 65.29%

=====

```
In [88]: hyper = [str(pow(10,j)) for j in range(-4,4)]  
          # https://www.dataquest.io/blog/learning-curves-machine-learning/
```

```
import matplotlib.pyplot as plt  
%matplotlib inline
```

```
plt.style.use('seaborn')
```

```
plt.plot(hyper,w2v_train_auc,label = 'Training AUC')  
plt.plot(hyper, w2v_cv_auc, label = 'Validation AUC')
```

```
plt.ylabel('AUC', fontsize = 14)  
plt.xlabel('\u03BB-values', fontsize = 14)  
plt.title('Learning curves for a Logistic Regression model', fontsize = 18, y = 1.03)  
plt.legend()
```

```
Out[88]: <matplotlib.legend.Legend at 0x7f7088b65400>
```



```
In [90]: i = 0.0001
LSVM = SGDClassifier(loss='hinge',alpha=i, penalty='l2')
calibrated_clf = CalibratedClassifierCV(LSVM, cv=5, method='sigmoid')
calibrated_clf.fit(w2v_train, train_y)
# train data
y_prob_train = calibrated_clf.predict_proba(w2v_train)[:,-1]
fprt, tprt, thresholdt = roc_curve(train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y, y_prob_train)
print('\nTrain AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_train * float(100))))
# CV
y_prob_cv = calibrated_clf.predict_proba(w2v_cv)[:,-1]
fprc, tprc, thresholdc = roc_curve(cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y, y_prob_cv)
print('\nCV AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_cv * float(100))))
# Test
y_prob_test = calibrated_clf.predict_proba(w2v_test)[:,-1]
fprts, tprts, thresholdts = roc_curve(test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
auc_roc_test = roc_auc_score(test_y, y_prob_test)
```

```
print('\nTest AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_test * float(100))))
```


Train AUC for $\lambda = 0.0001$ is 91.20%

CV AUC for $\lambda = 0.0001$ is 91.26%

Test AUC for $\lambda = 0.0001$ is 91.17%

In [91]: # <https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python>

```
import matplotlib.pyplot as plt
plt.clf()
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b', label='AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y', label='AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label='AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```





```
In [92]: print("F1-Score on test set: %0.2f"%(f1_score(test_y, y_pred_test)))
```

F1-Score on test set: 0.94

```
In [93]: df_cm = pd.DataFrame(confusion_matrix(test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[93]: <matplotlib.axes._subplots.AxesSubplot at 0x7f7088b4deb8>
```



7.1.7 using L1

```
In [94]: alpha = [pow(10,j) for j in range(-4,4,1)]
```

```
w2v_train_auc = []
w2v_cv_auc = []
for i in alpha:
    LSVM = SGDClassifier(loss='hinge',alpha=i, penalty='l1')
    calibrated_clf = CalibratedClassifierCV(LSVM, cv=5, method='sigmoid')
```

```

calibrated_clf.fit(w2v_train, train_y)
# train data
y_prob_train = calibrated_clf.predict_proba(w2v_train)[:,1]
y_pred = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y , y_prob_train)
print('\nTrain AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_train * float(100))))
w2v_train_auc.append(auc_roc_train)
# CV
y_prob_cv = calibrated_clf.predict_proba(w2v_cv)[:,1]
y_pred = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
print('\nCV AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_cv * float(100))))
w2v_cv_auc.append(auc_roc_cv)
print("="*50)

```

Train AUC for = 0.0001 is 90.96%

CV AUC for = 0.0001 is 91.07%

=====

Train AUC for = 0.001 is 90.75%

CV AUC for = 0.001 is 90.77%

=====

Train AUC for = 0.01 is 86.41%

CV AUC for = 0.01 is 86.68%

=====

Train AUC for = 0.1 is 50.00%

CV AUC for = 0.1 is 50.00%

=====

Train AUC for = 1 is 50.00%

CV AUC for = 1 is 50.00%

=====

Train AUC for = 10 is 50.00%

CV AUC for = 10 is 50.00%

=====

Train AUC for = 100 is 50.00%

CV AUC for = 100 is 50.00%

=====

Train AUC for = 1000 is 50.00%

CV AUC for = 1000 is 50.00%

=====

```
In [95]: hyper = [str(pow(10,j)) for j in range(-4,4)]
          # https://www.dataquest.io/blog/learning-curves-machine-learning/


          import matplotlib.pyplot as plt
          %matplotlib inline

          plt.style.use('seaborn')

          plt.plot(hyper,w2v_train_auc,label = 'Training AUC')
          plt.plot(hyper, w2v_cv_auc, label = 'Validation AUC')

          plt.ylabel('AUC', fontsize = 14)
          plt.xlabel('\u03BB-values', fontsize = 14)
          plt.title('Learning curves for a Logistic Regression model', fontsize = 18, y = 1.03)
          plt.legend()

Out[95]: <matplotlib.legend.Legend at 0x7f7088a23908>
```



```
In [96]: i = 0.0001
LSVM = SGDClassifier(loss='hinge',alpha=i, penalty='l1')
calibrated_clf = CalibratedClassifierCV(LSVM, cv=5, method='sigmoid')
calibrated_clf.fit(w2v_train, train_y)
# train data
y_prob_train = calibrated_clf.predict_proba(w2v_train)[:,-1]
fpr_t, tpr_t, threshold_t = roc_curve(train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y, y_prob_train)
print('\nTrain AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_train * float(100))))
# CV
y_prob_cv = calibrated_clf.predict_proba(w2v_cv)[:,-1]
fpr_c, tpr_c, threshold_c = roc_curve(cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y, y_prob_cv)
print('\nCV AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_cv * float(100))))
# Test
y_prob_test = calibrated_clf.predict_proba(w2v_test)[:,-1]
fpr_ts, tpr_ts, threshold_ts = roc_curve(test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
auc_roc_test = roc_auc_score(test_y, y_prob_test)
```

```
print('\nTest AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_test * float(100))))
```


Train AUC for $\lambda = 0.0001$ is 91.18%

CV AUC for $\lambda = 0.0001$ is 91.10%

Test AUC for $\lambda = 0.0001$ is 91.15%

In [97]: # <https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python>

```
import matplotlib.pyplot as plt
plt.clf()
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b', label='AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y', label='AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label='AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```




```
In [98]: print("F1-Score on test set: %0.2f"%(f1_score(test_y, y_pred_test)))
```

F1-Score on test set: 0.94

```
In [99]: df_cm = pd.DataFrame(confusion_matrix(test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[99]: <matplotlib.axes._subplots.AxesSubplot at 0x7f7088a4c780>
```



7.1.8 [5.1.4] Applying Linear SVM on TFIDF W2V, SET 4

```
In [100]: # Please write all the code with proper documentation
```

```
model = TfidfVectorizer()
#tf_idf_matrix = model.fit_transform(train)

print("====Train Data====")
final_tf_idf_train = model.fit_transform(train)
```

```

print("the type of count vectorizer ",type(final_tf_idf_train))
print("the shape of out text TFIDF vectorizer ",final_tf_idf_train.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_tf_idf_train.get_shape()[0])
print("=====CV Data=====")
final_tf_idf_cv = model.transform(cv)
print("the type of count vectorizer ",type(final_tf_idf_cv))
print("the shape of out text TFIDF vectorizer ",final_tf_idf_cv.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_tf_idf_cv.get_shape()[0])
print("=====Test Data=====")
final_tf_idf_test = model.transform(test)
print("the type of count vectorizer ",type(final_tf_idf_test))
print("the shape of out text TFIDF vectorizer ",final_tf_idf_test.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_tf_idf_test.get_shape()[0])

# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

```

```

=====Train Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (49000, 43052)
the number of unique words including both unigrams and bigrams 43052
=====CV Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (21000, 43052)
the number of unique words including both unigrams and bigrams 43052
=====Test Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (30000, 43052)
the number of unique words including both unigrams and bigrams 43052

```

```

In [101]: ##### Train #####
# TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

train_tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in row=0;
for sent in tqdm(list_of_train_sentence): # for each review/sentence
    sent_vec = np.zeros(100) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            #tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf valeus of word in this review

```

```

        tf_idf = dictionary[word]*(sent.count(word)/len(sent))
        sent_vec += (vec * tf_idf)
        weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    train_tfidf_sent_vectors.append(sent_vec)
    row += 1

```

100%|| 49000/49000 [2:48:30<00:00, 4.85it/s]

In [102]: ##### CV #####

```

# TF-IDF weighted Word2Vec
#tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tfidf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

cv_tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in th
row=0;
for sent in tqdm(list_of_cv_sentence): # for each review/sentence
    sent_vec = np.zeros(100) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            #tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole courpus
            # sent.count(word) = tf valeus of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    cv_tfidf_sent_vectors.append(sent_vec)
    row += 1

```

100%|| 21000/21000 [16:02<00:00, 15.38it/s]

In [103]: ##### Train #####

```

# TF-IDF weighted Word2Vec
#tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tfidf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

test_tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in
row=0;
for sent in tqdm(list_of_test_sentence): # for each review/sentence
    sent_vec = np.zeros(100) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review

```



```

for word in sent: # for each word in a review/sentence
    if word in w2v_words and word in tfidf_feat:
        vec = w2v_model.wv[word]
        #tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
        # to reduce the computation we are
        # dictionary[word] = idf value of word in whole corpus
        # sent.count(word) = tf value of word in this review
        tf_idf = dictionary[word]*(sent.count(word)/len(sent))
        sent_vec += (vec * tf_idf)
        weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    test_tfidf_sent_vectors.append(sent_vec)
    row += 1

```

100%|| 30000/30000 [24:36<00:00, 20.31it/s]

```

In [104]: # save the datasets as numpy array
tfidf_w2v_train = np.array(train_tfidf_sent_vectors)
tfidf_w2v_cv = np.array(cv_tfidf_sent_vectors)
tfidf_w2v_test = np.array(test_tfidf_sent_vectors)

```

7.1.9 using L2 regularization

```

In [105]: alpha = [pow(10,j) for j in range(-4,4,1)]

```

```

tfidf_w2v_train_auc = []
tfidf_w2v_cv_auc = []
for i in alpha:
    LSVM = SGDClassifier(loss='hinge',alpha=i, penalty='l2')
    calibrated_clf = CalibratedClassifierCV(LSVM, cv=5, method='sigmoid')
    calibrated_clf.fit(tfidf_w2v_train, train_y)
    # train data
    y_prob_train = calibrated_clf.predict_proba(tfidf_w2v_train)[:,-1]
    y_pred = np.where(y_prob_train > 0.5, 1, 0)
    auc_roc_train = roc_auc_score(train_y , y_prob_train)
    print('\nTrain AUC for \u03BB = %s is %.2f%%' % (str(i), (auc_roc_train * float(100))))
    tfidf_w2v_train_auc.append(auc_roc_train)
    # CV
    y_prob_cv = calibrated_clf.predict_proba(tfidf_w2v_cv)[:,-1]
    y_pred = np.where(y_prob_cv > 0.5, 1, 0)
    auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
    print('\nCV AUC for \u03BB = %s is %.2f%%' % (str(i), (auc_roc_cv * float(100))))
    tfidf_w2v_cv_auc.append(auc_roc_cv)
    print("="*50)

```

Train AUC for = 0.0001 is 88.49%

```

CV AUC for  = 0.0001 is 88.65%
=====

Train AUC for  = 0.001 is 88.98%

CV AUC for  = 0.001 is 88.99%
=====

Train AUC for  = 0.01 is 88.45%

CV AUC for  = 0.01 is 88.45%
=====

Train AUC for  = 0.1 is 88.33%

CV AUC for  = 0.1 is 88.31%
=====

Train AUC for  = 1 is 88.39%

CV AUC for  = 1 is 88.41%
=====

Train AUC for  = 10 is 84.05%

CV AUC for  = 10 is 84.13%
=====

Train AUC for  = 100 is 64.66%

CV AUC for  = 100 is 64.65%
=====

Train AUC for  = 1000 is 64.66%

CV AUC for  = 1000 is 64.65%
=====

```

```

In [106]: hyper = [str(pow(10,j)) for j in range(-4,4)]
           # https://www.dataquest.io/blog/learning-curves-machine-learning/

           import matplotlib.pyplot as plt
           %matplotlib inline

           plt.style.use('seaborn')

```

```
plt.plot(hyper,tfidf_w2v_train_auc,label = 'Training AUC')
plt.plot(hyper, tfidf_w2v_cv_auc, label = 'Validation AUC')

plt.ylabel('AUC', fontsize = 14)
plt.xlabel('\u03BB-values', fontsize = 14)
plt.title('Learning curves for a Logistic Regression model', fontsize = 18, y = 1.03)
plt.legend()
```

Out[106]: <matplotlib.legend.Legend at 0x7f70ad2f8668>



```
In [107]: i = 0.001
LSVM = SGDClassifier(loss='hinge',alpha=i, penalty='l2')
calibrated_clf = CalibratedClassifierCV(LSVM, cv=5, method='sigmoid')
calibrated_clf.fit(tfidf_w2v_train, train_y)
# train data
y_prob_train = calibrated_clf.predict_proba(tfidf_w2v_train)[: ,1]
fprt, tprr, thresholdt = roc_curve(train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y , y_prob_train)
print('\nTrain AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_train * float(100)
# CV
```

```

y_prob_cv = calibrated_clf.predict_proba(tfidf_w2v_cv)[: ,1]
fprc, tprc, thresholdc = roc_curve(cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y, y_prob_cv)
print('\nCV AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_cv * float(100))))
# Test
y_prob_test = calibrated_clf.predict_proba(tfidf_w2v_test)[: ,1]
fprts, tprts, thresholdts = roc_curve(test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
auc_roc_test = roc_auc_score(test_y, y_prob_test)
print('\nTest AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_test * float(100))))

```

Train AUC for = 0.001 is 88.91%

CV AUC for = 0.001 is 88.92%


Test AUC for = 0.001 is 89.07%

In [108]: # <https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python>

```

import matplotlib.pyplot as plt
plt.clf()
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b', label='AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y', label='AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label='AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

```




```
In [109]: print("F1-Score on test set: %0.2f"%(f1_score(test_y, y_pred_test)))
```

F1-Score on test set: 0.93

```
In [110]: df_cm = pd.DataFrame(confusion_matrix(test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[110]: <matplotlib.axes._subplots.AxesSubplot at 0x7f70ad2faac8>
```



7.1.10 L1 regularization

```
In [111]: alpha = [pow(10,j) for j in range(-4,4,1)]
```

```
tfidf_w2v_train_auc = []
tfidf_w2v_cv_auc = []
for i in alpha:
    LSVM = SGDClassifier(loss='hinge',alpha=i, penalty='l1')
    calibrated_clf = CalibratedClassifierCV(LSVM, cv=5, method='sigmoid')
    calibrated_clf.fit(tfidf_w2v_train, train_y)
    # train data
    y_prob_train = calibrated_clf.predict_proba(tfidf_w2v_train)[:,:1]
    y_pred = np.where(y_prob_train > 0.5, 1, 0)
    auc_roc_train = roc_auc_score(train_y , y_prob_train)
    print('\nTrain AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_train * float(100))))
    tfidf_w2v_train_auc.append(auc_roc_train)
    # CV
    y_prob_cv = calibrated_clf.predict_proba(tfidf_w2v_cv)[:,:1]
    y_pred = np.where(y_prob_cv > 0.5, 1, 0)
    auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
    print('\nCV AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_cv * float(100))))
    tfidf_w2v_cv_auc.append(auc_roc_cv)
print("="*50)
```

```

Train AUC for  = 0.0001 is 88.57%

CV AUC for  = 0.0001 is 88.56%
=====

Train AUC for  = 0.001 is 88.57%

CV AUC for  = 0.001 is 88.62%
=====

Train AUC for  = 0.01 is 82.15%

CV AUC for  = 0.01 is 82.07%
=====

Train AUC for  = 0.1 is 50.00%

CV AUC for  = 0.1 is 50.00%
=====

Train AUC for  = 1 is 50.00%

CV AUC for  = 1 is 50.00%
=====

Train AUC for  = 10 is 50.00%

CV AUC for  = 10 is 50.00%
=====

Train AUC for  = 100 is 50.00%

CV AUC for  = 100 is 50.00%
=====

Train AUC for  = 1000 is 50.00%

CV AUC for  = 1000 is 50.00%
=====

```

```

In [112]: hyper = [str(pow(10,j)) for j in range(-4,4)]
           # https://www.dataquest.io/blog/learning-curves-machine-learning/

           import matplotlib.pyplot as plt
           %matplotlib inline

```

```

plt.style.use('seaborn')

plt.plot(hyper,tfidf_w2v_train_auc,label = 'Training AUC')
plt.plot(hyper, tfidf_w2v_cv_auc, label = 'Validation AUC')

plt.ylabel('AUC', fontsize = 14)
plt.xlabel('\u03BB-values', fontsize = 14)
plt.title('Learning curves for a Logistic Regression model', fontsize = 18, y = 1.03)
plt.legend()

```

Out[112]: <matplotlib.legend.Legend at 0x7f70afea3780>



```

In [113]: i = 0.001
LSVM = SGDClassifier(loss='hinge',alpha=i, penalty='l1')
calibrated_clf = CalibratedClassifierCV(LSVM, cv=5, method='sigmoid')
calibrated_clf.fit(w2v_train, train_y)
# train data
y_prob_train = calibrated_clf.predict_proba(w2v_train)[: ,1]
fprt, tprt, thresholdt = roc_curve(train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y , y_prob_train)

```



```

print('\nTrain AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_train * float(100))))
# CV
y_prob_cv = calibrated_clf.predict_proba(w2v_cv)[: ,1]
fprc, tprc, thresholdc = roc_curve(cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y, y_prob_cv)
print('\nCV AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_cv * float(100))))
# Test
y_prob_test = calibrated_clf.predict_proba(w2v_test)[: ,1]
fprts, tprts, thresholdts = roc_curve(test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
auc_roc_test = roc_auc_score(test_y, y_prob_test)
print('\nTest AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_test * float(100))))

```

Train AUC for = 0.001 is 90.75%

CV AUC for = 0.001 is 90.78%


Test AUC for = 0.001 is 90.90%

In [114]: # <https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python>

```

import matplotlib.pyplot as plt
plt.clf()
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b', label='AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y', label='AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label='AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

```




```
In [115]: print("F1-Score on test set: %0.2f"%(f1_score(test_y, y_pred_test)))
```

F1-Score on test set: 0.94

```
In [116]: df_cm = pd.DataFrame(confusion_matrix(test_y, y_pred_test), range(2), range(2))  
sns.set(font_scale=1.4)  
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[116]: <matplotlib.axes._subplots.AxesSubplot at 0x7f70afe91cf8>
```



7.2 [5.2] RBF SVM

7.2.1 I am used to select 20k data points randomly and sort them using time based splitting

```
In [3]: # Split the data into train , test and crossvalidation datasets
# load "preprocessed.pkl" data frame
df = pd.read_pickle("files/preprocessed.pkl")
df.head(1)
```

```
Out[3]:
```

	Id	ProductId	UserId	ProfileName	\
138706	150524	0006641040	ACITT7DI6IDDL	shari zychinski	

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
138706	0	0	1	939340800	

	Summary	\
138706	EVERY book is educational	

	Text	\
138706	this witty little book makes my son laugh at l...	

	CleanedText	\
138706	witty little book makes son laugh loud recite ...	

```
CleanedSummary
138706  every book educational
```

```
In [4]: # take 20k sample data randomly
sample_data = df.sample(20000)
sample_data.shape
```

```
Out[4]: (20000, 12)
```

```
In [5]: # sorted the data using time based
sorted_data = sample_data.sort_values('Time', axis=0, inplace=False)
sorted_data.shape
```

```
Out[5]: (20000, 12)
```

```
In [6]: X = np.array(sorted_data['CleanedText'])
y = np.array(sorted_data['Score'])
print(X.shape)
print(y.shape)
```

```
(20000,)
```

```
(20000,)
```

```
In [7]: # Simple cross validation
# split the data sent into train and test
train , test , train_y , test_y = train_test_split(X, y, test_size = 0.3, random_state=42)

# split the train data set into cross validation train and cross validation test
train, cv , train_y, cv_y = train_test_split(train, train_y, test_size=0.3, random_state=42)

print("train data = ", train.shape)
print("cross validation = ", cv.shape)
print("test data = ", test.shape)
```

```
train data = (9800,)
```

```
cross validation = (4200,)
```

```
test data = (6000,)
```

7.2.2 [5.2.1] Applying RBF SVM on BOW, SET 1

```
In [149]: # Please write all the code with proper documentation
#BoW
count_vect = CountVectorizer(min_df=10, max_features=500) #in scikit-learn
count_vect.fit(train)
print("some feature names ", count_vect.get_feature_names()[:10])
print('='*50)
```

```

bow_train = count_vect.fit_transform(train)
bow_cv = count_vect.transform(cv)
bow_test = count_vect.transform(test)
print("=====Train Data=====")
print("the type of count vectorizer ",type(bow_train))
print("the shape of out text BOW vectorizer ",bow_train.get_shape())
print("the number of unique words ", bow_train.get_shape()[1])
print("=====Cross validation Data=====")
print("the type of count vectorizer ",type(bow_cv))
print("the shape of out text BOW vectorizer ",bow_cv.get_shape())
print("the number of unique words ", bow_cv.get_shape()[1])
print("=====Test Data=====")
print("the type of count vectorizer ",type(bow_test))
print("the shape of out text BOW vectorizer ",bow_test.get_shape())
print("the number of unique words ", bow_test.get_shape()[1])

```

```

some feature names  ['able', 'absolutely', 'actually', 'add', 'added', 'adding', 'aftertaste',
=====
=====Train Data=====
the type of count vectorizer  <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer  (9800, 500)
the number of unique words  500
=====Cross validation Data=====
the type of count vectorizer  <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer  (4200, 500)
the number of unique words  500
=====Test Data=====
the type of count vectorizer  <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer  (6000, 500)
the number of unique words  500

```

```

In [8]: from sklearn.svm import SVC
        from sklearn.model_selection import StratifiedShuffleSplit , GridSearchCV

```

```

In [170]: C_range = list(np.logspace(-2, 10, 4))
          gamma_range = list(np.logspace(-9, 3, 4))
          bow_train_auc = []
          bow_cv_auc = []
          for i in C_range:
              for j in gamma_range:
                  svc = SVC(C = i , gamma = j , probability=True)
                  svc.fit(bow_train, train_y)
                  # train data
                  y_prob_train = svc.predict_proba(bow_train)[: ,1]
                  y_pred = np.where(y_prob_train > 0.5, 1, 0)
                  auc_roc_train = roc_auc_score(train_y , y_prob_train)

```

```

print('\nTrain AUC for \u03B3 = %s and C = %s is %0.2f%%' % (str(j),str(i), (auc_
bow_train_auc.append(auc_roc_train)
# CV
y_prob_cv = svc.predict_proba(bow_cv)[: ,1]
y_pred = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
print('\nCV AUC for \u03B3 = %s and C = %s is %0.2f%%' % (str(j),str(i), (auc_
bow_cv_auc.append(auc_roc_cv)
print("="*50)

```

Train AUC for $\gamma = 1e-09$ and $C = 0.01$ is 50.00%

CV AUC for $\gamma = 1e-09$ and $C = 0.01$ is 50.00%

=====

Train AUC for $\gamma = 1e-05$ and $C = 0.01$ is 75.32%

CV AUC for $\gamma = 1e-05$ and $C = 0.01$ is 74.61%

=====

Train AUC for $\gamma = 0.1$ and $C = 0.01$ is 99.24%

CV AUC for $\gamma = 0.1$ and $C = 0.01$ is 85.85%

=====

Train AUC for $\gamma = 1000.0$ and $C = 0.01$ is 0.21%

CV AUC for $\gamma = 1000.0$ and $C = 0.01$ is 49.89%

=====

Train AUC for $\gamma = 1e-09$ and $C = 100.0$ is 71.75%

CV AUC for $\gamma = 1e-09$ and $C = 100.0$ is 72.45%

=====

Train AUC for $\gamma = 1e-05$ and $C = 100.0$ is 90.41%

CV AUC for $\gamma = 1e-05$ and $C = 100.0$ is 89.38%

=====

Train AUC for $\gamma = 0.1$ and $C = 100.0$ is 99.96%

CV AUC for $\gamma = 0.1$ and $C = 100.0$ is 85.36%

=====

Train AUC for $\gamma = 1000.0$ and $C = 100.0$ is 0.17%

```

CV AUC for  = 1000.0 and C = 100.0 is 49.89%
=====

Train AUC for  = 1e-09 and C = 1000000.0 is 71.36%

CV AUC for  = 1e-09 and C = 1000000.0 is 72.01%
=====

Train AUC for  = 1e-05 and C = 1000000.0 is 92.03%

CV AUC for  = 1e-05 and C = 1000000.0 is 88.38%
=====

Train AUC for  = 0.1 and C = 1000000.0 is 99.98%

CV AUC for  = 0.1 and C = 1000000.0 is 85.36%
=====

Train AUC for  = 1000.0 and C = 1000000.0 is 0.17%

CV AUC for  = 1000.0 and C = 1000000.0 is 49.89%
=====

Train AUC for  = 1e-09 and C = 10000000000.0 is 71.16%

CV AUC for  = 1e-09 and C = 10000000000.0 is 71.77%
=====

Train AUC for  = 1e-05 and C = 10000000000.0 is 91.70%

CV AUC for  = 1e-05 and C = 10000000000.0 is 88.01%
=====

Train AUC for  = 0.1 and C = 10000000000.0 is 99.92%

CV AUC for  = 0.1 and C = 10000000000.0 is 85.36%
=====

Train AUC for  = 1000.0 and C = 10000000000.0 is 0.17%

CV AUC for  = 1000.0 and C = 10000000000.0 is 49.89%
=====


```

```

In [174]: scores = np.array(bow_cv_auc).reshape(len(C_range),len(gamma_range))
          # plot the mean cross-validation scores
          mglearn.tools.heatmap(scores, xlabel='gamma', xticklabels=gamma_range,
                                ylabel='C', yticklabels=C_range, cmap="viridis")

```

Out[174]: <matplotlib.collections.PolyCollection at 0x7f70c56dff98>



```
In [176]: i = 100
          j = 1e-05
          svc = SVC(C = i , gamma = j , probability=True)
          svc.fit(bow_train, train_y)
          # train data
          y_prob_train = svc.predict_proba(bow_train)[: ,1]
          fprt, tprt, thresholdt = roc_curve(train_y, y_prob_train)
          y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
          auc_roc_train = roc_auc_score(train_y , y_prob_train)
          print('\nTrain AUC for \u03B3 = %s and C = %s is %0.2f%%' % (str(j),str(i),(auc_roc_train)))
          # CV
          y_prob_cv = svc.predict_proba(bow_cv)[: ,1]
          fprc, tprc, thresholdc = roc_curve(cv_y, y_prob_cv)
          y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
          auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
          print('\nCV AUC for \u03B3 = %s and C = %s is %0.2f%%' % (str(j),str(i),(auc_roc_cv)))
          # Test
          y_prob_test = svc.predict_proba(bow_test)[: ,1]
```



```

fprts, tprts, thresholdts = roc_curve(test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
auc_roc_test = roc_auc_score(test_y, y_prob_test)
print('\nTest AUC for \u03B2 = %s and C = %s is %0.2f%%' % (str(j),str(i),auc_roc_t

```

Train AUC for $\beta = 1e-05$ and $C = 100$ is 90.41%

CV AUC for $\beta = 1e-05$ and $C = 100$ is 89.38%


Test AUC for $\beta = 1e-05$ and $C = 100$ is 86.75%

In [177]: # <https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python>

```

import matplotlib.pyplot as plt
plt.clf()
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b', label='AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y', label='AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label='AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

```




```
In [178]: print("F1-Score on test set: %0.2f"%(f1_score(test_y, y_pred_test)))
```

F1-Score on test set: 0.93

```
In [179]: df_cm = pd.DataFrame(confusion_matrix(test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[179]: <matplotlib.axes._subplots.AxesSubplot at 0x7f70c5644f28>
```



7.2.3 [5.2.2] Applying RBF SVM on TFIDF, SET 2

In [186]: *# Please write all the code with proper documentation*
Please write all the code with proper documentation

```

model = TfidfVectorizer(min_df=10, ngram_range=(1,2), max_features=500)
#tf_idf_matrix = model.fit_transform(train)

print("=====Train Data=====")
tf_idf_train = model.fit_transform(train)
print("the type of count vectorizer ",type(tf_idf_train))
print("the shape of out text TFIDF vectorizer ",tf_idf_train.get_shape())
print("the number of unique words including both unigrams and bigrams ",tf_idf_train.get_shape()[1])
print("=====CV Data=====")
tf_idf_cv = model.transform(cv)
print("the type of count vectorizer ",type(tf_idf_cv))
print("the shape of out text TFIDF vectorizer ",tf_idf_cv.get_shape())
print("the number of unique words including both unigrams and bigrams ",tf_idf_cv.get_shape()[1])
print("=====Test Data=====")
tf_idf_test = model.transform(test)
print("the type of count vectorizer ",type(tf_idf_test))
print("the shape of out text TFIDF vectorizer ",tf_idf_test.get_shape())

```

```

print("the number of unique words including both unigrams and bigrams ", tf_idf_test

# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

=====Train Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (9800, 500)
the number of unique words including both unigrams and bigrams 500
=====CV Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (4200, 500)
the number of unique words including both unigrams and bigrams 500
=====Test Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (6000, 500)
the number of unique words including both unigrams and bigrams 500

In [187]: C_range = list(np.logspace(-2, 10, 4))
gamma_range = list(np.logspace(-9, 3, 4))
tfidf_train_auc = []
tfidf_cv_auc = []
for i in C_range:
    for j in gamma_range:
        svc = SVC(C = i , gamma = j , probability=True)
        svc.fit(tf_idf_train, train_y)
        # train data
        y_prob_train = svc.predict_proba(tf_idf_train)[: ,1]
        y_pred = np.where(y_prob_train > 0.5, 1, 0)
        auc_roc_train = roc_auc_score(train_y , y_prob_train)
        print('\nTrain AUC for \u03B3 = %s and C = %s is %0.2f%%' % (str(j),str(i), (auc_roc_train)))
        tfidf_train_auc.append(auc_roc_train)
        # CV
        y_prob_cv = svc.predict_proba(tf_idf_cv)[: ,1]
        y_pred = np.where(y_prob_cv > 0.5, 1, 0)
        auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
        print('\nCV AUC for \u03B3 = %s and C = %s is %0.2f%%' % (str(j),str(i), (auc_roc_cv)))
        tfidf_cv_auc.append(auc_roc_cv)
    print("="*50)

```

Train AUC for = 1e-09 and C = 0.01 is 50.31%

CV AUC for = 1e-09 and C = 0.01 is 16.63%

=====

Train AUC for = 1e-05 and C = 0.01 is 49.79%

```

CV AUC for  = 1e-05 and C = 0.01 is 82.97%
=====

Train AUC for  = 0.1 and C = 0.01 is 49.58%

CV AUC for  = 0.1 and C = 0.01 is 88.05%
=====

Train AUC for  = 1000.0 and C = 0.01 is 50.03%

CV AUC for  = 1000.0 and C = 0.01 is 49.91%
=====

Train AUC for  = 1e-09 and C = 100.0 is 49.70%

CV AUC for  = 1e-09 and C = 100.0 is 83.37%
=====

Train AUC for  = 1e-05 and C = 100.0 is 50.12%

CV AUC for  = 1e-05 and C = 100.0 is 88.05%
=====

Train AUC for  = 0.1 and C = 100.0 is 49.89%

CV AUC for  = 0.1 and C = 100.0 is 89.10%
=====

Train AUC for  = 1000.0 and C = 100.0 is 50.03%

CV AUC for  = 1000.0 and C = 100.0 is 49.91%
=====

Train AUC for  = 1e-09 and C = 1000000.0 is 49.69%

CV AUC for  = 1e-09 and C = 1000000.0 is 83.37%
=====

Train AUC for  = 1e-05 and C = 1000000.0 is 49.85%

CV AUC for  = 1e-05 and C = 1000000.0 is 87.81%
=====

Train AUC for  = 0.1 and C = 1000000.0 is 49.91%

CV AUC for  = 0.1 and C = 1000000.0 is 88.70%
=====

```

Train AUC for $\gamma = 1000.0$ and $C = 1000000.0$ is 50.03%

CV AUC for $\gamma = 1000.0$ and $C = 1000000.0$ is 49.91%
=====

Train AUC for $\gamma = 1e-09$ and $C = 10000000000.0$ is 49.70%

CV AUC for $\gamma = 1e-09$ and $C = 10000000000.0$ is 83.37%
=====

Train AUC for $\gamma = 1e-05$ and $C = 10000000000.0$ is 50.34%

CV AUC for $\gamma = 1e-05$ and $C = 10000000000.0$ is 86.31%
=====

Train AUC for $\gamma = 0.1$ and $C = 10000000000.0$ is 49.91%

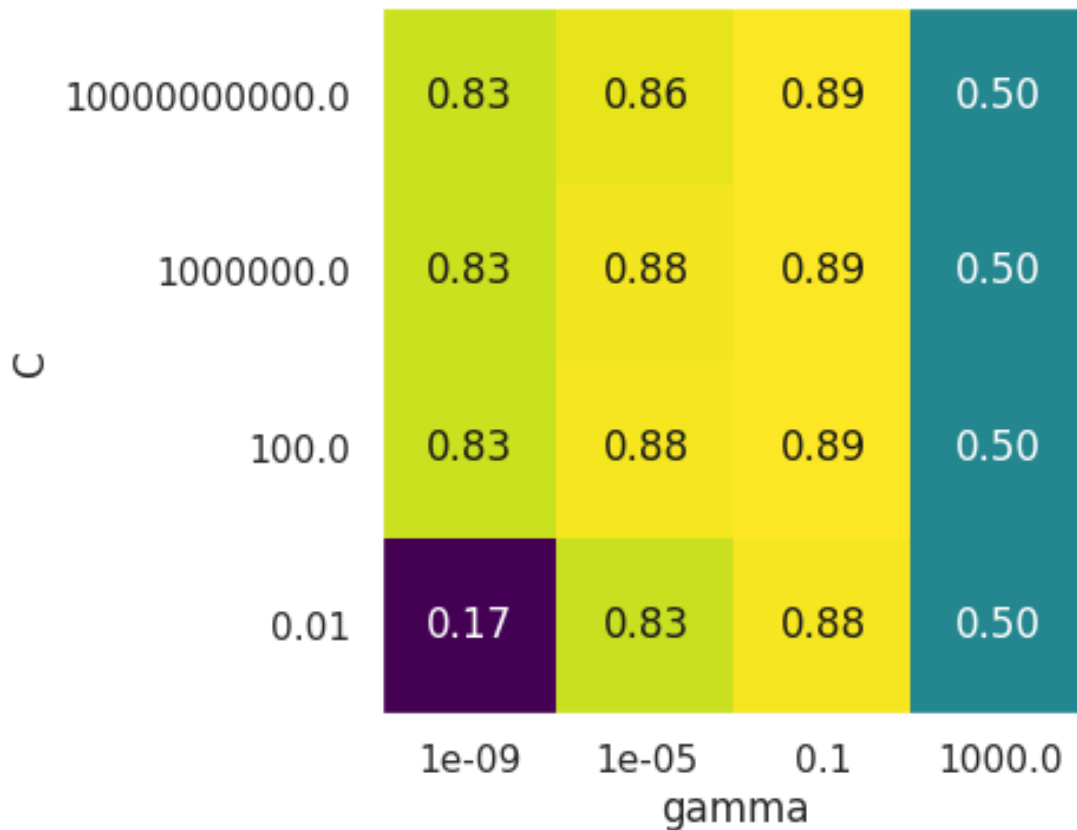
CV AUC for $\gamma = 0.1$ and $C = 10000000000.0$ is 88.70%
=====

Train AUC for $\gamma = 1000.0$ and $C = 10000000000.0$ is 50.03%

CV AUC for $\gamma = 1000.0$ and $C = 10000000000.0$ is 49.91%
=====

```
In [189]: scores = np.array(tfidf_cv_auc).reshape(len(C_range),len(gamma_range))
          # plot the mean cross-validation scores
          mglearn.tools.heatmap(scores, xlabel='gamma', xticklabels=gamma_range,
                                ylabel='C', yticklabels=C_range, cmap="viridis")
```

```
Out[189]: <matplotlib.collections.PolyCollection at 0x7f70af2ba080>
```



```
In [191]: i = 100
          j = 0.1
          svc = SVC(C = i , gamma = j , probability=True)
          svc.fit(tf_idf_train, train_y)
          # train data
          y_prob_train = svc.predict_proba(tf_idf_train)[: ,1]
          fprt, tprt, thresholdt = roc_curve(train_y, y_prob_train)
          y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
          auc_roc_train = roc_auc_score(train_y , y_prob_train)
          print('\nTrain AUC for \u03B3 = %s and C = %s is %0.2f%%' % (str(j),str(i),(auc_roc_train)))
          # CV
          y_prob_cv = svc.predict_proba(tf_idf_cv)[: ,1]
          fprc, tprc, thresholdc = roc_curve(cv_y, y_prob_cv)
          y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
          auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
          print('\nCV AUC for \u03B3 = %s and C = %s is %0.2f%%' % (str(j),str(i),(auc_roc_cv)))
          # Test
          y_prob_test = svc.predict_proba(tf_idf_test)[: ,1]
          fprts, tprts, thresholdts = roc_curve(test_y, y_prob_test)
          y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
```

```

auc_roc_test = roc_auc_score(test_y , y_prob_test)
print('\nTest AUC for \u03B2 = %s and C = %s is %0.2f%%' % (str(j),str(i),(auc_roc_test * 100)))

```

Train AUC for $\beta = 0.1$ and $C = 100$ is 99.89%

CV AUC for $\beta = 0.1$ and $C = 100$ is 89.10%


Test AUC for $\beta = 0.1$ and $C = 100$ is 89.04%

In [193]: # <https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python>

```

import matplotlib.pyplot as plt
plt.clf()
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b' , label='AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y' , label='AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label='AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

```





```
In [194]: print("F1-Score on test set: %.2f"%(f1_score(test_y, y_pred_test)))
```

F1-Score on test set: 0.94

```
In [195]: df_cm = pd.DataFrame(confusion_matrix(test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[195]: <matplotlib.axes._subplots.AxesSubplot at 0x7f70aba36128>
```



7.2.4 [5.2.3] Applying RBF SVM on AVG W2V, SET 3

```
In [9]: # Please write all the code with proper documentation
# Train your own Word2Vec model using your own text corpus
##### Train Set #####
i=0
list_of_train_sentence=[]
for sentence in train:
    list_of_train_sentence.append(sentence.split())
```

```

##### CV Set #####
i=0
list_of_cv_sentence=[]
for sentence in cv:
    list_of_cv_sentence.append(sentence.split())
##### Test Set #####
i=0
list_of_test_sentence=[]
for sentence in test:
    list_of_test_sentence.append(sentence.split())
print("Length of Train = ", len(list_of_train_sentence))
print("Length of CV = ", len(list_of_cv_sentence))
print("Length of Test = ", len(list_of_test_sentence))

```

```

Length of Train = 9800
Length of CV = 4200
Length of Test = 6000

```

```

In [10]: w2v_model=Word2Vec(list_of_train_sentence,min_count=15,size=50, workers=4)
         print(w2v_model.wv.most_similar('great'))
         print('='*50)
         print(w2v_model.wv.most_similar('worst'))

```

```

[('excellent', 0.8764510750770569), ('good', 0.8501964807510376), ('especially', 0.84445118904
=====
[('american', 0.9778188467025757), ('none', 0.9739011526107788), ('favorites', 0.9726344347000

```

```

In [11]: w2v_words = list(w2v_model.wv.vocab)
         print("number of words that occurred minimum 5 times ",len(w2v_words))
         print("sample words ", w2v_words[0:50])

```

```

number of words that occurred minimum 5 times 2955
sample words ['smell', 'coming', 'cans', 'bought', 'three', 'separate', 'bad', 'thought', 'wo

```

```

In [12]: ##### Train data #####
         # average Word2Vec
         # compute average word2vec for each review.
         sent_vectors_train = []; # the avg-w2v for each sentence/review is stored in this lis
         for sent in tqdm(list_of_train_sentence): # for each review/sentence
             sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need t
             cnt_words = 0; # num of words with a valid vector in the sentence/review
             for word in sent: # for each word in a review/sentence
                 if word in w2v_words:
                     vec = w2v_model.wv[word]
                     sent_vec += vec
                     cnt_words += 1

```

```

        if cnt_words != 0:
            sent_vec /= cnt_words
        sent_vectors_train.append(sent_vec)
    print(len(sent_vectors_train))
    print(len(sent_vectors_train[0]))

```

100%|| 9800/9800 [00:19<00:00, 502.95it/s]

9800

50

In [13]: ##### CV data #####

```

# average Word2Vec
# compute average word2vec for each review.
sent_vectors_cv = []; # the avg-w2v for each sentence/review is stored in this list
for sent in tqdm(list_of_cv_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need t
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors_cv.append(sent_vec)
print(len(sent_vectors_cv))
print(len(sent_vectors_cv[0]))

```

100%|| 4200/4200 [00:08<00:00, 472.86it/s]

4200

50

In [14]: ##### Test data #####

```

# average Word2Vec
# compute average word2vec for each review.
sent_vectors_test = []; # the avg-w2v for each sentence/review is stored in this list
for sent in tqdm(list_of_test_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need t
    cnt_words = 0; # num of words with a valid vector in the sentence/review

```

```

    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors_test.append(sent_vec)
print(len(sent_vectors_test))
print(len(sent_vectors_test[0]))

```

100%| 6000/6000 [00:12<00:00, 482.55it/s]

6000

50

```

In [15]: # save the datasets as numpy array
w2v_train = np.array(sent_vectors_train)
w2v_cv = np.array(sent_vectors_cv)
w2v_test = np.array(sent_vectors_test)

```

```

In [16]: C_range = [pow(10,j) for j in range(-2,2,1)]
gamma_range = [pow(10,j) for j in range(-2,2,1)]
w2v_train_auc = []
w2v_cv_auc = []
for i in C_range:
    for j in gamma_range:
        svc = SVC(C = i , gamma = j , probability=True)
        svc.fit(w2v_train, train_y)
        # train data
        y_prob_train = svc.predict_proba(w2v_train)[: ,1]
        y_pred = np.where(y_prob_train > 0.5, 1, 0)
        auc_roc_train = roc_auc_score(train_y , y_prob_train)
        print('\nTrain AUC for \u03B3 = %s and C = %s is %0.2f%%' % (str(j),str(i),(auc_roc_train)))
        w2v_train_auc.append(auc_roc_train)
        # CV
        y_prob_cv = svc.predict_proba(w2v_cv)[: ,1]
        y_pred = np.where(y_prob_cv > 0.5, 1, 0)
        auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
        print('\nCV AUC for \u03B3 = %s and C = %s is %0.2f%%' % (str(j),str(i),(auc_roc_cv)))
        w2v_cv_auc.append(auc_roc_cv)
    print("="*50)

```

Train AUC for $\gamma = 0.01$ and $C = 0.01$ is 82.13%

```

CV AUC for  = 0.01 and C = 0.01 is 81.95%
=====

Train AUC for  = 0.1 and C = 0.01 is 86.04%

CV AUC for  = 0.1 and C = 0.01 is 84.35%
=====

Train AUC for  = 1 and C = 0.01 is 87.31%

CV AUC for  = 1 and C = 0.01 is 83.03%
=====

Train AUC for  = 10 and C = 0.01 is 99.18%

CV AUC for  = 10 and C = 0.01 is 74.43%
=====

Train AUC for  = 0.01 and C = 0.1 is 85.53%

CV AUC for  = 0.01 and C = 0.1 is 84.10%
=====

Train AUC for  = 0.1 and C = 0.1 is 87.01%

CV AUC for  = 0.1 and C = 0.1 is 85.06%
=====

Train AUC for  = 1 and C = 0.1 is 87.39%

CV AUC for  = 1 and C = 0.1 is 83.07%
=====

Train AUC for  = 10 and C = 0.1 is 99.17%

CV AUC for  = 10 and C = 0.1 is 74.44%
=====

Train AUC for  = 0.01 and C = 1 is 86.89%

CV AUC for  = 0.01 and C = 1 is 85.02%
=====

Train AUC for  = 0.1 and C = 1 is 87.07%

CV AUC for  = 0.1 and C = 1 is 85.10%
=====

```

```

Train AUC for  = 1 and C = 1 is 88.10%

CV AUC for  = 1 and C = 1 is 83.34%
=====

Train AUC for  = 10 and C = 1 is 99.60%

CV AUC for  = 10 and C = 1 is 74.31%
=====

Train AUC for  = 0.01 and C = 10 is 86.97%

CV AUC for  = 0.01 and C = 10 is 85.08%
=====

Train AUC for  = 0.1 and C = 10 is 87.24%

CV AUC for  = 0.1 and C = 10 is 85.27%
=====

Train AUC for  = 1 and C = 10 is 92.85%

CV AUC for  = 1 and C = 10 is 84.24%
=====

Train AUC for  = 10 and C = 10 is 99.93%


CV AUC for  = 10 and C = 10 is 72.54%
=====

```

```
In [19]: import mglearn
```

```
In [20]: scores = np.array(w2v_cv_auc).reshape(len(C_range),len(gamma_range))
         # plot the mean cross-validation scores
         mglearn.tools.heatmap(scores, xlabel='gamma', xticklabels=gamma_range,
                                ylabel='C', yticklabels=C_range, cmap="viridis")
```

```
Out[20]: <matplotlib.collections.PolyCollection at 0x7f3018032ef0>
```



```
In [21]: i = 10
         j = 0.1
         svc = SVC(C = i , gamma = j , probability=True)
         svc.fit(w2v_train, train_y)
         # train data
         y_prob_train = svc.predict_proba(w2v_train)[: ,1]
         fprt, tprt, thresholdt = roc_curve(train_y, y_prob_train)
         y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
         auc_roc_train = roc_auc_score(train_y , y_prob_train)
         print('\nTrain AUC for \u03B3 = %s and C = %s is %0.2f%%' % (str(j),str(i),(auc_roc_train * 100)))
         # CV
         y_prob_cv = svc.predict_proba(w2v_cv)[: ,1]
         fprc, tprc, thresholdc = roc_curve(cv_y, y_prob_cv)
         y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
         auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
         print('\nCV AUC for \u03B3 = %s and C = %s is %0.2f%%' % (str(j),str(i),(auc_roc_cv * 100)))
         # Test
         y_prob_test = svc.predict_proba(w2v_test)[: ,1]
         fprts, tprts, thresholdts = roc_curve(test_y, y_prob_test)
         y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
         auc_roc_test = roc_auc_score(test_y , y_prob_test)
         print('\nTest AUC for \u03B3 = %s and C = %s is %0.2f%%' % (str(j),str(i),(auc_roc_test * 100)))
```


Train AUC for $\gamma = 0.1$ and $C = 10$ is 87.24%

CV AUC for $\lambda = 0.1$ and $C = 10$ is 85.27%

Test AUC for $\lambda = 0.1$ and $C = 10$ is 85.45%

In [22]: # <https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python>

```
import matplotlib.pyplot as plt
plt.clf()
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b', label='AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y', label='AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label='AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```




In [23]: print("F1-Score on test set: %0.2f"%(f1_score(test_y, y_pred_test)))

F1-Score on test set: 0.93

```
In [24]: df_cm = pd.DataFrame(confusion_matrix(test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

Out[24]: <matplotlib.axes._subplots.AxesSubplot at 0x7f30010d65f8>



7.2.5 [5.2.4] Applying RBF SVM on TFIDF W2V, SET 4

In [26]: *# Please write all the code with proper documentation*

```
model = TfidfVectorizer(min_df=10 , max_features=500, ngram_range=(1,2))
#tf_idf_matrix = model.fit_transform(train)

print("=====Train Data=====")
final_tf_idf_train = model.fit_transform(train)
print("the type of count vectorizer ",type(final_tf_idf_train))
print("the shape of out text TFIDF vectorizer ",final_tf_idf_train.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_tf_idf.get_shape()[0])
print("=====CV Data=====")
final_tf_idf_cv = model.transform(cv)
print("the type of count vectorizer ",type(final_tf_idf_cv))
print("the shape of out text TFIDF vectorizer ",final_tf_idf_cv.get_shape())
```

```

print("the number of unique words including both unigrams and bigrams ", final_tf_idf)
print("=====Test Data=====")
final_tf_idf_test = model.transform(test)
print("the type of count vectorizer ",type(final_tf_idf_test))
print("the shape of out text TFIDF vectorizer ",final_tf_idf_test.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_tf_idf)

# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

```

```

=====Train Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (9800, 500)
the number of unique words including both unigrams and bigrams 500
=====CV Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (4200, 500)
the number of unique words including both unigrams and bigrams 500
=====Test Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (6000, 500)
the number of unique words including both unigrams and bigrams 500

```

```

In [28]: ##### Train #####
# TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

train_tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in
row=0;
for sent in tqdm(list_of_train_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            #tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole courpus
            # sent.count(word) = tf valeus of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    train_tfidf_sent_vectors.append(sent_vec)
    row += 1

```

100%|| 9800/9800 [00:26<00:00, 373.16it/s]

```
In [29]: ##### CV #####
# TF-IDF weighted Word2Vec
#tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

cv_tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in tqdm(list_of_cv_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            #tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    cv_tfidf_sent_vectors.append(sent_vec)
    row += 1
```

100%|| 4200/4200 [00:10<00:00, 397.58it/s]

```
In [31]: ##### Train #####
# TF-IDF weighted Word2Vec
#tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

test_tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in tqdm(list_of_test_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            #tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf value of word in this review
```

```

        tf_idf = dictionary[word]*(sent.count(word)/len(sent))
        sent_vec += (vec * tf_idf)
        weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    test_tfidf_sent_vectors.append(sent_vec)
    row += 1

```

100%|| 6000/6000 [00:16<00:00, 358.53it/s]

In [32]: *# save the datasets as numpy array*

```

tfidf_w2v_train = np.array(train_tfidf_sent_vectors)
tfidf_w2v_cv = np.array(cv_tfidf_sent_vectors)
tfidf_w2v_test = np.array(test_tfidf_sent_vectors)

```

In [34]: C_range = [pow(10,j) for j in range(-4,3,1)]

gamma_range = [pow(10,j) for j in range(-4,3,1)]

tfidf_w2v_train_auc = []

tfidf_w2v_cv_auc = []

for i in C_range:

for j in gamma_range:

svc = SVC(C = i , gamma = j , probability=True)

svc.fit(tfidf_w2v_train, train_y)

train data

y_prob_train = svc.predict_proba(tfidf_w2v_train)[:,-1]

y_pred = np.where(y_prob_train > 0.5, 1, 0)

auc_roc_train = roc_auc_score(train_y , y_prob_train)

print('\nTrain AUC for \u03B3 = %s and C = %s is %0.2f%%' % (str(j),str(i),auc_roc_train))

tfidf_w2v_train_auc.append(auc_roc_train)

CV

y_prob_cv = svc.predict_proba(tfidf_w2v_cv)[:,-1]

y_pred = np.where(y_prob_cv > 0.5, 1, 0)

auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)

print('\nCV AUC for \u03B3 = %s and C = %s is %0.2f%%' % (str(j),str(i),auc_roc_cv))

tfidf_w2v_cv_auc.append(auc_roc_cv)

print("="*50)

Train AUC for = 0.0001 and C = 0.0001 is 50.00%

CV AUC for = 0.0001 and C = 0.0001 is 50.00%

Train AUC for = 0.001 and C = 0.0001 is 75.05%

CV AUC for = 0.001 and C = 0.0001 is 72.48%

Train AUC for $\lambda = 0.01$ and $C = 0.0001$ is 75.93%

CV AUC for $\lambda = 0.01$ and $C = 0.0001$ is 73.40%

=====

Train AUC for $\lambda = 0.1$ and $C = 0.0001$ is 71.06%

CV AUC for $\lambda = 0.1$ and $C = 0.0001$ is 69.93%

=====

Train AUC for $\lambda = 1$ and $C = 0.0001$ is 67.47%

CV AUC for $\lambda = 1$ and $C = 0.0001$ is 64.72%

=====

Train AUC for $\lambda = 10$ and $C = 0.0001$ is 98.53%

CV AUC for $\lambda = 10$ and $C = 0.0001$ is 56.96%

=====

Train AUC for $\lambda = 100$ and $C = 0.0001$ is 99.94%

CV AUC for $\lambda = 100$ and $C = 0.0001$ is 52.95%

=====

Train AUC for $\lambda = 0.0001$ and $C = 0.001$ is 75.25%

CV AUC for $\lambda = 0.0001$ and $C = 0.001$ is 72.68%

=====

Train AUC for $\lambda = 0.001$ and $C = 0.001$ is 75.05%

CV AUC for $\lambda = 0.001$ and $C = 0.001$ is 72.48%

=====

Train AUC for $\lambda = 0.01$ and $C = 0.001$ is 77.59%

CV AUC for $\lambda = 0.01$ and $C = 0.001$ is 76.06%

=====

Train AUC for $\lambda = 0.1$ and $C = 0.001$ is 81.17%

CV AUC for $\lambda = 0.1$ and $C = 0.001$ is 78.83%

=====

Train AUC for $\lambda = 1$ and $C = 0.001$ is 85.89%

CV AUC for $\lambda = 1$ and $C = 0.001$ is 74.94%

```

=====

Train AUC for  = 10 and C = 0.001 is 99.90%

CV AUC for  = 10 and C = 0.001 is 60.23%
=====

Train AUC for  = 100 and C = 0.001 is 0.05%

CV AUC for  = 100 and C = 0.001 is 46.55%
=====

Train AUC for  = 0.0001 and C = 0.01 is 75.25%

CV AUC for  = 0.0001 and C = 0.01 is 72.69%
=====

Train AUC for  = 0.001 and C = 0.01 is 78.74%

CV AUC for  = 0.001 and C = 0.01 is 77.44%
=====

Train AUC for  = 0.01 and C = 0.01 is 77.97%

CV AUC for  = 0.01 and C = 0.01 is 75.87%
=====

Train AUC for  = 0.1 and C = 0.01 is 83.66%

CV AUC for  = 0.1 and C = 0.01 is 79.57%
=====

Train AUC for  = 1 and C = 0.01 is 90.35%

CV AUC for  = 1 and C = 0.01 is 76.95%
=====

Train AUC for  = 10 and C = 0.01 is 99.84%

CV AUC for  = 10 and C = 0.01 is 67.91%
=====

Train AUC for  = 100 and C = 0.01 is 0.17%

CV AUC for  = 100 and C = 0.01 is 44.22%
=====

Train AUC for  = 0.0001 and C = 0.1 is 79.53%

```

```

CV AUC for  = 0.0001 and C = 0.1 is 78.11%
=====

Train AUC for  = 0.001 and C = 0.1 is 77.20%

CV AUC for  = 0.001 and C = 0.1 is 74.85%
=====

Train AUC for  = 0.01 and C = 0.1 is 81.75%

CV AUC for  = 0.01 and C = 0.1 is 79.03%
=====

Train AUC for  = 0.1 and C = 0.1 is 85.16%

CV AUC for  = 0.1 and C = 0.1 is 80.61%
=====

Train AUC for  = 1 and C = 0.1 is 90.56%

CV AUC for  = 1 and C = 0.1 is 77.05%
=====

Train AUC for  = 10 and C = 0.1 is 99.83%

CV AUC for  = 10 and C = 0.1 is 68.17%
=====

Train AUC for  = 100 and C = 0.1 is 0.09%

CV AUC for  = 100 and C = 0.1 is 44.14%
=====

Train AUC for  = 0.0001 and C = 1 is 79.20%

CV AUC for  = 0.0001 and C = 1 is 77.08%
=====

Train AUC for  = 0.001 and C = 1 is 80.82%

CV AUC for  = 0.001 and C = 1 is 77.90%
=====

Train AUC for  = 0.01 and C = 1 is 85.00%

CV AUC for  = 0.01 and C = 1 is 81.34%
=====

```

Train AUC for $\lambda = 0.1$ and $C = 1$ is 85.22%

CV AUC for $\lambda = 0.1$ and $C = 1$ is 80.70%

=====

Train AUC for $\lambda = 1$ and $C = 1$ is 91.48%

CV AUC for $\lambda = 1$ and $C = 1$ is 77.39%

=====

Train AUC for $\lambda = 10$ and $C = 1$ is 99.82%

CV AUC for $\lambda = 10$ and $C = 1$ is 68.18%

=====

Train AUC for $\lambda = 100$ and $C = 1$ is 99.98%

CV AUC for $\lambda = 100$ and $C = 1$ is 55.81%

=====

Train AUC for $\lambda = 0.0001$ and $C = 10$ is 78.14%

CV AUC for $\lambda = 0.0001$ and $C = 10$ is 75.21%

=====

Train AUC for $\lambda = 0.001$ and $C = 10$ is 84.16%

CV AUC for $\lambda = 0.001$ and $C = 10$ is 80.78%

=====

Train AUC for $\lambda = 0.01$ and $C = 10$ is 85.21%

CV AUC for $\lambda = 0.01$ and $C = 10$ is 81.50%

=====

Train AUC for $\lambda = 0.1$ and $C = 10$ is 86.04%

CV AUC for $\lambda = 0.1$ and $C = 10$ is 80.94%

=====

Train AUC for $\lambda = 1$ and $C = 10$ is 96.83%

CV AUC for $\lambda = 1$ and $C = 10$ is 76.74%

=====

Train AUC for $\lambda = 10$ and $C = 10$ is 99.92%

CV AUC for $\gamma = 10$ and $C = 10$ is 67.59%
=====

Train AUC for $\gamma = 100$ and $C = 10$ is 99.86%

CV AUC for $\gamma = 100$ and $C = 10$ is 55.52%
=====

Train AUC for $\gamma = 0.0001$ and $C = 100$ is 82.99%

CV AUC for $\gamma = 0.0001$ and $C = 100$ is 79.89%
=====

Train AUC for $\gamma = 0.001$ and $C = 100$ is 84.96%

CV AUC for $\gamma = 0.001$ and $C = 100$ is 81.48%
=====

Train AUC for $\gamma = 0.01$ and $C = 100$ is 85.22%

CV AUC for $\gamma = 0.01$ and $C = 100$ is 81.51%
=====

Train AUC for $\gamma = 0.1$ and $C = 100$ is 89.12%

CV AUC for $\gamma = 0.1$ and $C = 100$ is 81.18%
=====

Train AUC for $\gamma = 1$ and $C = 100$ is 99.80%

CV AUC for $\gamma = 1$ and $C = 100$ is 72.67%
=====

Train AUC for $\gamma = 10$ and $C = 100$ is 99.92%


CV AUC for $\gamma = 10$ and $C = 100$ is 67.59%
=====

Train AUC for $\gamma = 100$ and $C = 100$ is 99.85%

CV AUC for $\gamma = 100$ and $C = 100$ is 55.64%
=====

```
In [42]: scores = np.array(tfidf_w2v_cv_auc).reshape(len(C_range),len(gamma_range))
         # plot the mean cross-validation scores
         mglearn.tools.heatmap(scores, xlabel='gamma', xticklabels=gamma_range,
                                ylabel='C', yticklabels=C_range, cmap="viridis")
```

Out[42]: <matplotlib.collections.PolyCollection at 0x7f2ff6c8af28>



```
In [43]: i = 10
j = 0.1
svc = SVC(C = i , gamma = j , probability=True)
svc.fit(tfidf_w2v_train, train_y)
# train data
y_prob_train = svc.predict_proba(tfidf_w2v_train)[:,-1]
fprt, tprt, thresholdt = roc_curve(train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y , y_prob_train)
print('\nTrain AUC for \u03B3 = %s and C = %s is %0.2f%%' % (str(j),str(i),(auc_roc_train * 100)))
# CV
y_prob_cv = svc.predict_proba(tfidf_w2v_cv)[:,-1]
fprc, tprc, thresholdc = roc_curve(cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
print('\nCV AUC for \u03B3 = %s and C = %s is %0.2f%%' % (str(j),str(i),(auc_roc_cv * 100)))
# Test
y_prob_test = svc.predict_proba(tfidf_w2v_test)[:,-1]
fprts, tprts, thresholdts = roc_curve(test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
auc_roc_test = roc_auc_score(test_y , y_prob_test)
print('\nTest AUC for \u03B3 = %s and C = %s is %0.2f%%' % (str(j),str(i),(auc_roc_test * 100)))
```


Train AUC for $\lambda = 0.1$ and $C = 10$ is 86.04%

CV AUC for $\lambda = 0.1$ and $C = 10$ is 80.94%

Test AUC for $\lambda = 0.1$ and $C = 10$ is 81.09%

In [44]: [# https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python](https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python)

```
import matplotlib.pyplot as plt
plt.clf()
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b', label='AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y', label='AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label='AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```




```
In [45]: print("F1-Score on test set: %0.2f"%(f1_score(test_y, y_pred_test)))
```

F1-Score on test set: 0.92

```
In [46]: df_cm = pd.DataFrame(confusion_matrix(test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[46]: <matplotlib.axes._subplots.AxesSubplot at 0x7f2ff74efa58>
```



```
In [55]: # Please compare all your models using Prettytable library
X_review = np.array(sorted_data['CleanedText'])
X_summary = np.array(sorted_data['Summary'])
y = np.array(sorted_data['Score'])
print(X_review.shape)
print(X_summary.shape)
print(y.shape)
```

```
(100000,)
(100000,)
(100000,)
```

```
In [56]: # Simple cross validation
# split the data sent into train and test
```

```

train_review , test_review, train_summary, test_summary , train_y , test_y = train_test_split(
    train_data, train_labels, test_size=0.2, random_state=42)

# split the train data set into cross validation train and cross validation test
train_review, cv_review, train_summary, cv_summary , train_y, cv_y = train_test_split(
    train_data, train_labels, test_size=0.2, random_state=42)

print("train review", train_review.shape)
print("cv review", cv_review.shape)
print("test review", test_review.shape)
print("train summary", train_summary.shape)
print("cv summary", cv_summary.shape)
print("test summary", test_summary.shape)

train review (49000,)
cv review (21000,)
test review (30000,)
train summary (49000,)
cv summary (21000,)
test summary (30000,)

In [59]: # bow for summary
# Please write all the code with proper documentation
#BoW
count_vect = CountVectorizer(min_df=15, ngram_range=(1,2)) #in scikit-learn
count_vect.fit(train_summary)
print("some feature names ", count_vect.get_feature_names()[:10])
print('='*50)

bow_train_summary = count_vect.fit_transform(train_summary)
bow_cv_summary = count_vect.transform(cv_summary)
bow_test_summary = count_vect.transform(test_summary)
print("=====Train Data=====")
print("the type of count vectorizer ",type(bow_train_summary))
print("the shape of out text BOW vectorizer ",bow_train_summary.get_shape())
print("the number of unique words ", bow_train_summary.get_shape()[1])
print("=====Cross validation Data=====")
print("the type of count vectorizer ",type(bow_cv_summary))
print("the shape of out text BOW vectorizer ",bow_cv_summary.get_shape())
print("the number of unique words ", bow_cv_summary.get_shape()[1])
print("=====Test Data=====")
print("the type of count vectorizer ",type(bow_test_summary))
print("the shape of out text BOW vectorizer ",bow_test_summary.get_shape())
print("the number of unique words ", bow_test_summary.get_shape()[1])

some feature names  ['10', '100', '12', '15', '20', '24', '50', 'about', 'about it', 'about the
=====
=====Train Data=====
the type of count vectorizer  <class 'scipy.sparse.csr.csr_matrix'>

```

```

the shape of out text BOW vectorizer (49000, 2299)
the number of unique words 2299
=====Cross validation Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (21000, 2299)
the number of unique words 2299
=====Test Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (30000, 2299)
the number of unique words 2299

```

```
In [62]: alpha = [pow(10,j) for j in range(-4,4,1)]
```

```

bow_train_auc = []
bow_cv_auc = []
for i in alpha:
    LSMV = SGDClassifier(loss='hinge',alpha=i, penalty='l2')
    calibrated_clf = CalibratedClassifierCV(LSMV, cv=5, method='sigmoid')
    calibrated_clf.fit(bow_train_summary, train_y)
    # train data
    y_prob_train = calibrated_clf.predict_proba(bow_train_summary)[: ,1]
    y_pred = np.where(y_prob_train > 0.5, 1, 0)
    auc_roc_train = roc_auc_score(train_y , y_prob_train)
    print('\nTrain AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_train * float(100))))
    bow_train_auc.append(auc_roc_train)
    # CV
    y_prob_cv = calibrated_clf.predict_proba(bow_cv_summary)[: ,1]
    y_pred = np.where(y_prob_cv > 0.5, 1, 0)
    auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
    print('\nCV AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_cv * float(100))))
    bow_cv_auc.append(auc_roc_cv)
    print("="*50)

```

```
Train AUC for  = 0.0001 is 94.04%
```

```
CV AUC for  = 0.0001 is 92.10%
```

```
=====
```

```
Train AUC for  = 0.001 is 92.38%
```

```
CV AUC for  = 0.001 is 91.07%
```

```
=====
```

```
Train AUC for  = 0.01 is 91.13%
```

```
CV AUC for  = 0.01 is 89.95%
```

=====

Train AUC for $\alpha = 0.1$ is 75.82%

CV AUC for $\alpha = 0.1$ is 75.30%

=====

Train AUC for $\alpha = 1$ is 69.94%

CV AUC for $\alpha = 1$ is 69.41%

=====

Train AUC for $\alpha = 10$ is 69.94%

CV AUC for $\alpha = 10$ is 69.41%

=====

Train AUC for $\alpha = 100$ is 69.94%

CV AUC for $\alpha = 100$ is 69.41%

=====

Train AUC for $\alpha = 1000$ is 69.94%

CV AUC for $\alpha = 1000$ is 69.41%

=====

```
In [64]: hyper = [str(pow(10,j)) for j in range(-4,4)]
import matplotlib.pyplot as plt
%matplotlib inline


plt.style.use('seaborn')

plt.plot(hyper,bow_train_auc,label = 'Training AUC')
plt.plot(hyper, bow_cv_auc, label = 'Validation AUC')

plt.ylabel('AUC')
plt.xlabel('alpha-values', fontsize = 10)
plt.title('Learning curves for a MultinomialNB model', fontsize = 18, y = 1.03)
plt.legend()
```

Out[64]: <matplotlib.legend.Legend at 0x7f2ff82f14a8>

Learning curves for a MultinomialNB model



In [65]: *# summary text*

i = 0.0001

LSVM = SGDClassifier(loss='hinge',alpha=i, penalty='l2')

calibrated_clf = CalibratedClassifierCV(LSVM, cv=5, method='sigmoid')

calibrated_clf.fit(bow_train_summary, train_y)

train data

y_prob_train = calibrated_clf.predict_proba(bow_train_summary)[: ,1]

fprt, tprt, thresholdt = roc_curve(train_y, y_prob_train)

y_pred_train = np.where(y_prob_train > 0.5, 1, 0)

auc_roc_train = roc_auc_score(train_y , y_prob_train)

*print('\nTrain AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_train * float(100))))*

CV

y_prob_cv = calibrated_clf.predict_proba(bow_cv_summary)[: ,1]

fprc, tprc, thresholdc = roc_curve(cv_y, y_prob_cv)

y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)

auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)

*print('\nCV AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_cv * float(100))))*

Test

y_prob_test = calibrated_clf.predict_proba(bow_test_summary)[: ,1]

fprts, tprts, thresholdts = roc_curve(test_y, y_prob_test)

y_pred_test = np.where(y_prob_test > 0.5, 1, 0)


```

auc_roc_test = roc_auc_score(test_y , y_prob_test)
print('\nTest AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_test * float(100))))

```

Train AUC for = 0.0001 is 94.21%

CV AUC for = 0.0001 is 92.34%

Test AUC for = 0.0001 is 92.48%

```

In [66]: # bow for review
# Please write all the code with proper documentation
#BoW
count_vect = CountVectorizer(min_df=15, ngram_range=(1,2)) #in scikit-learn
count_vect.fit(train_review)
print("some feature names ", count_vect.get_feature_names()[:10])
print('='*50)

bow_train_review = count_vect.fit_transform(train_review)
bow_cv_review = count_vect.transform(cv_review)
bow_test_review = count_vect.transform(test_review)
print("=====Train Data=====")
print("the type of count vectorizer ",type(bow_train_review))
print("the shape of out text BOW vectorizer ",bow_train_review.get_shape())
print("the number of unique words ", bow_train_review.get_shape()[1])
print("=====Cross validation Data=====")
print("the type of count vectorizer ",type(bow_cv_review))
print("the shape of out text BOW vectorizer ",bow_cv_review.get_shape())
print("the number of unique words ", bow_cv_review.get_shape()[1])
print("=====Test Data=====")
print("the type of count vectorizer ",type(bow_test_review))
print("the shape of out text BOW vectorizer ",bow_test_review.get_shape())
print("the number of unique words ", bow_test_review.get_shape()[1])

```

```

some feature names  ['ability', 'able', 'able buy', 'able drink', 'able eat', 'able enjoy', 'al
=====
=====Train Data=====
the type of count vectorizer  <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer  (49000, 18421)
the number of unique words  18421
=====Cross validation Data=====
the type of count vectorizer  <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer  (21000, 18421)
the number of unique words  18421
=====Test Data=====
the type of count vectorizer  <class 'scipy.sparse.csr.csr_matrix'>

```

the shape of out text BOW vectorizer (30000, 18421)
the number of unique words 18421

In [67]: *# bow for review*

```
i = 0.0001
LSVM = SGDClassifier(loss='hinge',alpha=i, penalty='l1')
calibrated_clf = CalibratedClassifierCV(LSVM, cv=5, method='sigmoid')
calibrated_clf.fit(bow_train_review, train_y)
# train data
y_prob_train = calibrated_clf.predict_proba(bow_train_review)[:,-1]
fprt, tprt, thresholdt = roc_curve(train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y , y_prob_train)
print('\nTrain AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_train * float(100))))
# CV
y_prob_cv = calibrated_clf.predict_proba(bow_cv_review)[:,-1]
fprc, tprc, thresholdc = roc_curve(cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
print('\nCV AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_cv * float(100))))
# Test
y_prob_test_review = calibrated_clf.predict_proba(bow_test_review)[:,-1]
fprts, tprts, thresholdts = roc_curve(test_y, y_prob_test_review)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
auc_roc_test = roc_auc_score(test_y , y_prob_test)
print('\nTest AUC for \u03BB = %s is %0.2f%%' % (str(i), (auc_roc_test * float(100))))
```

Train AUC for = 0.0001 is 96.88%

CV AUC for = 0.0001 is 94.20%

Test AUC for = 0.0001 is 92.48%

In [68]: *# adding both summary and review test probabilitys and average*


```
new_proba = (y_prob_test_review + y_prob_test) / 2
```

In [69]: fprts, tprts, thresholdts = roc_curve(test_y, new_proba)

```
y_pred_test = np.where(new_proba > 0.5, 1, 0)
auc_roc_test = roc_auc_score(test_y , new_proba)
print('\nTest AUC for alpha = %s is %0.2f%%' % (str(i), (auc_roc_test * float(100))))
```

Test AUC for alpha = 0.0001 is 97.16%

```
In [70]: import matplotlib.pyplot as plt
plt.clf()
plt.title('Receiver Operating Characteristic')
plt.plot(fprts, tprts, 'y', label='AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```




```
In [71]: print("F1-Score on test set: %0.2f"%(f1_score(test_y, y_pred_test)))
```

F1-Score on test set: 0.96

```
In [72]: df_cm = pd.DataFrame(confusion_matrix(test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

Out[72]: <matplotlib.axes._subplots.AxesSubplot at 0x7f2ff79204e0>



Observation : adding two models we get better results. we see the 97.16% AUC and F1 score is 0.96

8 [6] Conclusions

8.1 Linear SVM

In [73]: *# Please compare all your models using Prettytable library*
`from prettytable import PrettyTable`

```
x = PrettyTable(["Vectorizer" , "Regularization","Hyper parameter \u03BB", "AUC", "F1"])

x.add_row(["BOW", "L2" ,0.001, "93.92%", 0.95])
x.add_row(["BOW","L1" , 0.0001,"93.07%", 0.95])
x.add_row(["TFIDF", "L2" ,0.0001, "95.82%", 0.96])
x.add_row(["TFIDF","L1" , 0.0001,"93.97%", 0.95])
x.add_row(["AVG-W2V", "L2" ,0.0001, "91.97%", 0.94])
x.add_row(["AVG-W2V","L1" , 0.0001,"91.15%", 0.94])
x.add_row(["TFIDF-W2V", "L2" ,0.001, "89.07%", 0.94])
x.add_row(["TFIDF-W2V","L1" , 0.001,"91.06%", 0.94])

print(x.get_string(title="Linear SVM"))
```

Vectorizer	Regularization	Hyper parameter	AUC	F1 Score
BOW	L2	0.001	93.92%	0.95
BOW	L1	0.0001	93.07%	0.95
TFIDF	L2	0.0001	95.82%	0.96
TFIDF	L1	0.0001	93.97%	0.95
AVG-W2V	L2	0.0001	91.97%	0.94
AVG-W2V	L1	0.0001	91.15%	0.94
TFIDF-W2V	L2	0.001	89.07%	0.94
TFIDF-W2V	L1	0.001	91.06%	0.94

8.2 RBF

```
In [74]: x = PrettyTable(["Vectorizer", "Hyper parameter C", "gamma", "AUC", "F1 Score"])
```

```
x.add_row(["BOW" ,100,"1e-05" ,"86.72%", 0.93])
x.add_row(["TFIDF" ,100,0.1 ,"89.04%", 0.94])
x.add_row(["AVG-W2V",10 ,0.1, "85.45%", 0.93])
x.add_row(["TFIDF-W2V",10 ,0.1, "81.09%", 0.92])
```

```
print(x.get_string(title="RBF SVM"))
```

Vectorizer	Hyper parameter C	gamma	AUC	F1 Score
BOW	100	1e-05	86.72%	0.93
TFIDF	100	0.1	89.04%	0.94
AVG-W2V	10	0.1	85.45%	0.93
TFIDF-W2V	10	0.1	81.09%	0.92

8.3 Feature engineering

```
In [75]: x = PrettyTable(["Feature" , "Vectorizer" , "Regularization", "Hyper parameter \u03BB"])
```

```
x.add_row(["Test length", "BOW", "L2", 0.0001, "92.48%", 0.95])
x.add_row(["Summary + Review", "BOW", "L2", 0.0001 , "97.16%", 0.97])
```

```
print(x.get_string(title="Linear SVM Model"))
```

Feature	Vectorizer	Regularization	Hyper parameter	AUC	F1 Score
Test length	BOW	L2	0.0001	92.48%	0.95

Summary + Review	BOW		L2		0.0001		97.16%		0.97	
+-----+	-----+	-----+	-----+	-----+	-----+	-----+	-----+	-----+	-----+	-----+