

# Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454

Number of users: 256,059

Number of products: 74,258

Timespan: Oct 1999 - Oct 2012

Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

## Objective:

Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative?

[Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

## [1]. Reading Data

### [1.1] Loading the data

The dataset is available in two forms

1. .csv file
2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

```
In [1]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
```

```

import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

```

```

In [2]: # using SQLite Table to read data.
con = sqlite3.connect('database.sqlite')

# filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
# SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 50
0000 data points
# you can change the number to any other number based on your computing
power

# filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Sco
re != 3 LIMIT 500000""", con)
# for tsne assignment you can take 5k data points

```

```

filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score
!= 3 LIMIT 5000""", con)

# Give reviews with Score>3 a positive rating(1), and reviews with a sc
ore<3 a negative rating(0).
def partition(x):
    if x < 3:
        return 0
    return 1

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)

```

Number of data points in our data (5000, 10)

Out[2]:

|   | Id | ProductId  | UserId         | ProfileName                              | HelpfulnessNumerator | HelpfulnessDenomin |
|---|----|------------|----------------|--|----------------------|--------------------|
| 0 | 1  | B001E4KFG0 | A3SGXH7AUHU8GW | delmartian                               | 1                    |                    |
| 1 | 2  | B00813GRG4 | A1D87F6ZCVE5NK | dll pa                                   | 0                    |                    |
| 2 | 3  | B000LQOCH0 | ABXLMWJIXXAIN  | Natalia<br>Corres<br>"Natalia<br>Corres" | 1                    |                    |

```
In [3]: display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", con)
```

```
In [4]: print(display.shape)
display.head()
```

```
(80668, 7)
```

Out[4]:

|   | UserId             | ProductId  | ProfileName            | Time       | Score | Text  | COUNT(*) |
|---|--------------------|------------|------------------------|------------|-------|---|----------|
| 0 | #oc-R115TNMSPFT9I7 | B007Y59HVM | Breyton                | 1331510400 | 2     | Overall its just OK when considering the price... | 2        |
| 1 | #oc-R11D9D7SHXIJB9 | B005HG9ET0 | Louis E. Emory "hoppy" | 1342396800 | 5     | My wife has recurring extreme muscle spasms, u... | 3        |
| 2 | #oc-R11DNU2NBKQ23Z | B007Y59HVM | Kim Cieszykowski       | 1348531200 | 1     | This coffee is horrible and unfortunately not ... | 2        |
| 3 | #oc-R11O5J5ZVQE25C | B005HG9ET0 | Penguin Chick          | 1346889600 | 5     | This will be the bottle that you grab from the... | 3        |
| 4 | #oc-R12KPBODL2B5ZD | B007OSBE1U | Christopher P. Presta  | 1348617600 | 1     | I didnt like this coffee. Instead of telling y... | 2        |

```
In [5]: display[display['UserId']=='AZY10LLTJ71NX']
```

Out[5]:

|  | UserId | ProductId | ProfileName | Time | Score | Text | COUNT(*) |
|--|--------|-----------|-------------|------|-------|------|----------|
|--|--------|-----------|-------------|------|-------|------|----------|

|       | UserId        | ProductId  | ProfileName                        | Time       | Score | Text  | COUNT(*) |
|-------|---------------|------------|------------------------------------|------------|-------|---|----------|
| 80638 | AZY10LLTJ71NX | B006P7E5ZI | undertheshrine<br>"undertheshrine" | 1334707200 | 5     | I was<br>recommended<br>to try green<br>tea extract to<br>... | 5        |

In [6]: `display['COUNT(*)'].sum()`

Out[6]: 393063

## [2] Exploratory Data Analysis

### [2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

In [7]: `display= pd.read_sql_query("""  
SELECT *  
FROM Reviews  
WHERE Score != 3 AND UserId="AR5J8UI46CURR"  
ORDER BY ProductID  
""", con)  
display.head()`

Out[7]:

|   | Id    | ProductId  | UserId        | ProfileName        | HelpfulnessNumerator | HelpfulnessDenon |
|---|-------|------------|---------------|--------------------|----------------------|------------------|
| 0 | 78445 | B000HDL1RQ | AR5J8UI46CURR | Geetha<br>Krishnan | 2                    |                  |

|   | Id     | ProductId  | UserId        | ProfileName     | HelpfulnessNumerator | HelpfulnessDenon |
|---|--------|------------|---------------|-----------------|----------------------|------------------|
| 1 | 138317 | B000HDOPYC | AR5J8UI46CURR | Geetha Krishnan | 2                    |                  |
| 2 | 138277 | B000HDOPYM | AR5J8UI46CURR | Geetha Krishnan | 2                    |                  |
| 3 | 73791  | B000HDOPZG | AR5J8UI46CURR | Geetha Krishnan | 2                    |                  |
| 4 | 155049 | B000PAQ75C | AR5J8UI46CURR | Geetha Krishnan | 2                    |                  |

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that

ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8)

ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```
In [8]: #Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False, kind='quicksort', na_position='last')
```

```
In [9]: #Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first', inplace=False)
final.shape
```

```
Out[9]: (4986, 10)
```

```
In [10]: #Checking to see how much % of data still remains
(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

```
Out[10]: 99.72
```

**Observation:-** It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

```
In [11]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)

display.head()
```

```
Out[11]:
```



|   | Id    | ProductId  | UserId         | ProfileName                | HelpfulnessNumerator | HelpfulnessDenom |
|---|-------|------------|----------------|----------------------------|----------------------|------------------|
| 0 | 64422 | B000MIDROQ | A161DK06JJMCYF | J. E. Stephens<br>"Jeanne" | 3                    |                  |
| 1 | 44737 | B001EQ55RW | A2V0I904FH7ABY | Ram                        | 3                    |                  |

```
In [12]: final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

```
In [13]: #Before starting the next phase of preprocessing lets see the number of
         entries left
         print(final.shape)

         #How many positive and negative reviews are present in our dataset?
         final['Score'].value_counts()
```

```
(4986, 10)
```

```
Out[13]: 1    4178
         0     808
         Name: Score, dtype: int64
```

## [3] Preprocessing

### [3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```
In [14]: # printing some random reviews
sent_0 = final['Text'].values[0]
print(sent_0)
print("="*50)

sent_1000 = final['Text'].values[1000]
print(sent_1000)
print("="*50)

sent_1500 = final['Text'].values[1500]
print(sent_1500)
print("="*50)

sent_4900 = final['Text'].values[4900]
print(sent_4900)
print("="*50)
```

```
Why is this $[...] when the same product is available for $[...] here?<
br />http://www.amazon.com/VICTOR-FLY-MAGNET-BAIT-REFILL/dp/B00004RBDY<
br /><br />The Victor M380 and M502 traps are unreal, of course -- tota
l fly genocide. Pretty stinky, but only right nearby.
```

```
=====
```

I recently tried this flavor/brand and was surprised at how delicious these chips are. The best thing was that there were a lot of "brown" chips in the bsg (my favorite), so I bought some more through amazon and shared with family and friends. I am a little disappointed that there are not, so far, very many brown chips in these bags, but the flavor is still very good. I like them better than the yogurt and green onion flavor because they do not seem to be as salty, and the onion flavor is better. If you haven't eaten Kettle chips before, I recommend that you try a bag before buying bulk. They are thicker and crunchier than Lays but just as fresh out of the bag.

=====  
Wow. So far, two two-star reviews. One obviously had no idea what they were ordering; the other wants crispy cookies. Hey, I'm sorry; but these reviews do nobody any good beyond reminding us to look before ordering.<br /><br />These are chocolate-oatmeal cookies. If you don't like that combination, don't order this type of cookie. I find the combo quite nice, really. The oatmeal sort of "calms" the rich chocolate flavor and gives the cookie sort of a coconut-type consistency. Now let's also remember that tastes differ; so, I've given my opinion.<br /><br />Then, these are soft, chewy cookies -- as advertised. They are not "crispy" cookies, or the blurb would say "crispy," rather than "chewy." I happen to like raw cookie dough; however, I don't see where these taste like raw cookie dough. Both are soft, however, so is this the confusion? And, yes, they stick together. Soft cookies tend to do that. They aren't individually wrapped, which would add to the cost. Oh yeah, chocolate chip cookies tend to be somewhat sweet.<br /><br />So, if you want something hard and crisp, I suggest Nabiso's Ginger Snaps. If you want a cookie that's soft, chewy and tastes like a combination of chocolate and oatmeal, give these a try. I'm here to place my second order.

=====  
love to order my coffee on amazon. easy and shows up quickly.<br />This k cup is great coffee. dcaf is very good as well  
=====

```
In [15]: # remove urls from text python: https://stackoverflow.com/a/40823105/4084039
sent_0 = re.sub(r"http\S+", "", sent_0)
sent_1000 = re.sub(r"http\S+", "", sent_1000)
sent_150 = re.sub(r"http\S+", "", sent_1500)
```

```
sent_4900 = re.sub(r"http\S+", "", sent_4900)

print(sent_0)
```

Why is this \$[...] when the same product is available for \$[...] here?<br /> /><br />The Victor M380 and M502 traps are unreal, of course -- total fly genocide. Pretty stinky, but only right nearby.

In [16]: *# https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-remove-all-tags-from-an-element*

```
from bs4 import BeautifulSoup

soup = BeautifulSoup(sent_0, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1000, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1500, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_4900, 'lxml')
text = soup.get_text()
print(text)
```

Why is this \$[...] when the same product is available for \$[...] here? />The Victor M380 and M502 traps are unreal, of course -- total fly genocide. Pretty stinky, but only right nearby.

=====

I recently tried this flavor/brand and was surprised at how delicious these chips are. The best thing was that there were a lot of "brown" chips in the bsg (my favorite), so I bought some more through amazon and shared with family and friends. I am a little disappointed that there

are not, so far, very many brown chips in these bags, but the flavor is still very good. I like them better than the yogurt and green onion flavor because they do not seem to be as salty, and the onion flavor is better. If you haven't eaten Kettle chips before, I recommend that you try a bag before buying bulk. They are thicker and crunchier than Lays but just as fresh out of the bag.

Wow. So far, two two-star reviews. One obviously had no idea what they were ordering; the other wants crispy cookies. Hey, I'm sorry; but these reviews do nobody any good beyond reminding us to look before ordering. These are chocolate-oatmeal cookies. If you don't like that combination, don't order this type of cookie. I find the combo quite nice, really. The oatmeal sort of "calms" the rich chocolate flavor and gives the cookie sort of a coconut-type consistency. Now let's also remember that tastes differ; so, I've given my opinion. Then, these are soft, chewy cookies -- as advertised. They are not "crispy" cookies, or the blurb would say "crispy," rather than "chewy." I happen to like raw cookie dough; however, I don't see where these taste like raw cookie dough. Both are soft, however, so is this the confusion? And, yes, they stick together. Soft cookies tend to do that. They aren't individually wrapped, which would add to the cost. Oh yeah, chocolate chip cookies tend to be somewhat sweet. So, if you want something hard and crisp, I suggest Nabisco's Ginger Snaps. If you want a cookie that's soft, chewy and tastes like a combination of chocolate and oatmeal, give these a try. I'm here to place my second order.

I love to order my coffee on amazon. easy and shows up quickly. This k cup is great coffee. dcaf is very good as well

```
In [17]: # https://stackoverflow.com/a/47091490/4084039
import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
```

```

phrase = re.sub(r"\ 're", " are", phrase)
phrase = re.sub(r"\ 's", " is", phrase)
phrase = re.sub(r"\ 'd", " would", phrase)
phrase = re.sub(r"\ 'll", " will", phrase)
phrase = re.sub(r"\ 't", " not", phrase)
phrase = re.sub(r"\ 've", " have", phrase)
phrase = re.sub(r"\ 'm", " am", phrase)
return phrase

```

```

In [18]: sent_1500 = decontracted(sent_1500)
print(sent_1500)
print("="*50)

```

Wow. So far, two two-star reviews. One obviously had no idea what they were ordering; the other wants crispy cookies. Hey, I am sorry; but these reviews do nobody any good beyond reminding us to look before ordering.<br /><br />These are chocolate-oatmeal cookies. If you do not like that combination, do not order this type of cookie. I find the combo quite nice, really. The oatmeal sort of "calms" the rich chocolate flavor and gives the cookie sort of a coconut-type consistency. Now let it also remember that tastes differ; so, I have given my opinion.<br /><br />Then, these are soft, chewy cookies -- as advertised. They are not "crispy" cookies, or the blurb would say "crispy," rather than "chewy." I happen to like raw cookie dough; however, I do not see where these taste like raw cookie dough. Both are soft, however, so is this the confusion? And, yes, they stick together. Soft cookies tend to do that. They are not individually wrapped, which would add to the cost. Oh yeah, chocolate chip cookies tend to be somewhat sweet.<br /><br />So, if you want something hard and crisp, I suggest Nabisco's Ginger Snaps. If you want a cookie that is soft, chewy and tastes like a combination of chocolate and oatmeal, give these a try. I am here to place my second order.

=====

```

In [19]: #remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
sent_0 = re.sub("\S*\d\S*", "", sent_0).strip()
print(sent_0)

```

Why is this \$[...] when the same product is available for \$[...] here?<br /> /><br />The Victor and traps are unreal, of course -- total fly genocide. Pretty stinky, but only right nearby.

```
In [20]: #remove spacial character: https://stackoverflow.com/a/5843547/4084039
sent_1500 = re.sub('[^A-Za-z0-9]+', ' ', sent_1500)
print(sent_1500)
```

Wow So far two two star reviews One obviously had no idea what they were ordering the other wants crispy cookies Hey I am sorry but these reviews do nobody any good beyond reminding us to look before ordering br br These are chocolate oatmeal cookies If you do not like that combination do not order this type of cookie I find the combo quite nice really The oatmeal sort of calms the rich chocolate flavor and gives the cookie sort of a coconut type consistency Now let is also remember that tastes differ so I have given my opinion br br Then these are soft chewy cookies as advertised They are not crispy cookies or the blurb would say crispy rather than chewy I happen to like raw cookie dough however I do not see where these taste like raw cookie dough Both are soft however so is this the confusion And yes they stick together Soft cookies tend to do that They are not individually wrapped which would add to the cost Oh yeah chocolate chip cookies tend to be somewhat sweet br br So if you want something hard and crisp I suggest Nabiso is Ginger Snaps If you want a cookie that is soft chewy and tastes like a combination of chocolate and oatmeal give these a try I am here to place my second order

```
In [21]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
# instead of <br /> if we have <br/> these tags would have revmoved in the 1st step

stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're", "you've", \
               "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him', 'his', 'himself', \
               'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'it
```

```
s', 'itself', 'they', 'them', 'their', \
    'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'th
is', 'that', "that'll", 'these', 'those', \
    'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'h
ave', 'has', 'had', 'having', 'do', 'does', \
    'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or',
    'because', 'as', 'until', 'while', 'of', \
    'at', 'by', 'for', 'with', 'about', 'against', 'between',
    'into', 'through', 'during', 'before', 'after', \
    'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out',
    'on', 'off', 'over', 'under', 'again', 'further', \
    'then', 'once', 'here', 'there', 'when', 'where', 'why', 'h
ow', 'all', 'any', 'both', 'each', 'few', 'more', \
    'most', 'other', 'some', 'such', 'only', 'own', 'same', 's
o', 'than', 'too', 'very', \
    's', 't', 'can', 'will', 'just', 'don', "don't", 'should',
    "should've", 'now', 'd', 'll', 'm', 'o', 're', \
    've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't",
    'didn', "didn't", 'doesn', "doesn't", 'hadn', \
    "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "is
n't", 'ma', 'mightn', "mightn't", 'mustn', \
    "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn',
    "shouldn't", 'wasn', "wasn't", 'weren', "weren't", \
    'won', "won't", 'wouldn', "wouldn't"])
```

```
In [22]: # Combining all the above students
from tqdm import tqdm
preprocessed_reviews = []
# tqdm is for printing the status bar
for sentence in tqdm(final['Text'].values):
    sentence = re.sub(r"http\S+", "", sentence)
    sentence = BeautifulSoup(sentence, 'lxml').get_text()
    sentence = decontracted(sentence)
    sentence = re.sub("\S*\d\S*", "", sentence).strip()
    sentence = re.sub('[^A-Za-z]+', ' ', sentence)
    # https://gist.github.com/sebleier/554280
    sentence = ' '.join(e.lower() for e in sentence.split() if e.lower
() not in stopwords)
    preprocessed_reviews.append(sentence.strip())
```



[illegible]

```
preprocessed_reviews[1500]
```

'wow far two two star reviews one obviously no idea ordering wants crispy cookies hey sorry reviews nobody good beyond reminding us look ordering chocolate oatmeal cookies not like combination not order type cookie find combo quite nice really oatmeal sort calms rich chocolate flavor gives cookie sort coconut type consistency let also remember tastes differ given opinion soft chewy cookies advertised not crispy cookies blurry would say crispy rather chewy happen like raw cookie dough however not see taste like raw cookie dough soft however confusion yes stick together soft cookies tend not individually wrapped would add cost oh yeah chocolate chip cookies tend somewhat sweet want something hard crisp suggest nabisco ginger snaps want cookie soft chewy tastes like combination chocolate oatmeal give try place second order'

### [3.2] Preprocessing Review Summary

```
## Similarly you can do preprocessing for review summary also.
```

## [4] Featurization

## [4.1] BAG OF WORDS

```
#Bow
count_vect = CountVectorizer() #in scikit-learn
count_vect.fit(preprocessed_reviews)
print("some feature names ", count_vect.get_feature_names()[:10])
print('='*50)

final_counts = count_vect.transform(preprocessed_reviews)
```

```
print("the type of count vectorizer ",type(final_counts))
print("the shape of out text BOW vectorizer ",final_counts.get_shape())
print("the number of unique words ", final_counts.get_shape()[1])
```

some feature names ['aa', 'aahhs', 'aback', 'abandon', 'abates', 'abbott', 'abby', 'abdominal', 'abiding', 'ability']

=====

```
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (4986, 12997)
the number of unique words 12997
```

## [4.2] Bi-Grams and n-Grams.

In [26]: *#bi-gram, tri-gram and n-gram*

```
#removing stop words like "not" should be avoided before building n-grams
# count_vect = CountVectorizer(ngram_range=(1,2))
# please do read the CountVectorizer documentation http://scikit-learn.org/stable/modules/generated/sklearn.feature_extraction.text.CountVectorizer.html
```

```
# you can choose these numebrs min_df=10, max_features=5000, of your choice
count_vect = CountVectorizer(ngram_range=(1,2), min_df=10, max_features=5000)
final_bigram_counts = count_vect.fit_transform(preprocessed_reviews)
print("the type of count vectorizer ",type(final_bigram_counts))
print("the shape of out text BOW vectorizer ",final_bigram_counts.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_bigram_counts.get_shape()[1])
```

```
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (4986, 3144)
the number of unique words including both unigrams and bigrams 3144
```

## [4.3] TF-IDF

```
In [27]: tf_idf_vect = TfidfVectorizer(ngram_range=(1,2), min_df=10)
tf_idf_vect.fit(preprocessed_reviews)
print("some sample features(unique words in the corpus)",tf_idf_vect.get_feature_names()[0:10])
print('='*50)

final_tf_idf = tf_idf_vect.transform(preprocessed_reviews)
print("the type of count vectorizer ",type(final_tf_idf))
print("the shape of out text TFIDF vectorizer ",final_tf_idf.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_tf_idf.get_shape()[1])

some sample features(unique words in the corpus) ['ability', 'able', 'able find', 'able get', 'absolute', 'absolutely', 'absolutely delicious', 'absolutely love', 'absolutely no', 'according']
=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (4986, 3144)
the number of unique words including both unigrams and bigrams 3144
```

## [4.4] Word2Vec

```
In [28]: # Train your own Word2Vec model using your own text corpus
i=0
list_of_sentence=[]
for sentence in preprocessed_reviews:
    list_of_sentence.append(sentence.split())
```

```
In [42]: # Using Google News Word2Vectors

# in this project we are using a pretrained model by google
# its 3.3G file, once you load this into your memory
# it occupies ~9Gb, so please do this step only if you have >12G of ram
```

```
# we will provide a pickle file wich contains a dict ,
# and it contains all our courpus words as keys and model[word] as values
# To use this code-snippet, download "GoogleNews-vectors-negative300.bin"
# from https://drive.google.com/file/d/0B7XkCwpI5KDYNlNUTTlSS21pQmM/edit
# it's 1.9GB in size.
```

```
# http://kavita-ganesan.com/gensim-word2vec-tutorial-starter-code/#.W17SRFAzZPY
# you can comment this whole cell
# or change these variable according to your need
```

```
is_your_ram_gt_16g=False
want_to_use_google_w2v = False
want_to_train_w2v = True
```

```
if want_to_train_w2v:
    # min_count = 5 considers only words that occurred at least 5 times
    w2v_model=Word2Vec(list_of_sentence,min_count=5,size=50, workers=4)
    print(w2v_model.wv.most_similar('great'))
    print('='*50)
    print(w2v_model.wv.most_similar('worst'))
```

```
elif want_to_use_google_w2v and is_your_ram_gt_16g:
    if os.path.isfile('GoogleNews-vectors-negative300.bin'):
        w2v_model=KeyedVectors.load_word2vec_format('GoogleNews-vectors-negative300.bin', binary=True)
        print(w2v_model.wv.most_similar('great'))
        print(w2v_model.wv.most_similar('worst'))
    else:
        print("you don't have gogole's word2vec file, keep want_to_train_w2v = True, to train your own w2v ")
```

```
[('snack', 0.9951335191726685), ('calorie', 0.9946465492248535), ('wonderful', 0.9946032166481018), ('excellent', 0.9944332838058472), ('especially', 0.9941144585609436), ('baked', 0.9940600395202637), ('salted', 0.994047224521637), ('alternative', 0.9937226176261902), ('tasty', 0.99
```

```
36816692352295), ('healthy', 0.9936649799346924)]
=====
[('varieties', 0.9994194507598877), ('become', 0.9992934465408325), ('p
opcorn', 0.9992750883102417), ('de', 0.9992610216140747), ('miss', 0.99
92451071739197), ('melitta', 0.999218761920929), ('choice', 0.999210238
4567261), ('american', 0.9991837739944458), ('beef', 0.999178051948547
4), ('finish', 0.9991567134857178)]
```

```
In [36]: w2v_words = list(w2v_model.wv.vocab)
print("number of words that occurred minimum 5 times ", len(w2v_words))
print("sample words ", w2v_words[0:50])
```

```
number of words that occurred minimum 5 times 3817
sample words ['product', 'available', 'course', 'total', 'pretty', 'st
inky', 'right', 'nearby', 'used', 'ca', 'not', 'beat', 'great', 'receiv
ed', 'shipment', 'could', 'hardly', 'wait', 'try', 'love', 'call', 'ins
tead', 'removed', 'easily', 'daughter', 'designed', 'printed', 'use',
'car', 'windows', 'beautifully', 'shop', 'program', 'going', 'lot', 'fu
n', 'everywhere', 'like', 'tv', 'computer', 'really', 'good', 'idea',
'final', 'outstanding', 'window', 'everybody', 'asks', 'bought', 'mad
e']
```

## [4.4.1] Converting text into vectors using Avg W2V, TFIDF-W2V

### [4.4.1.1] Avg W2v

```
In [38]: # average Word2Vec
# compute average word2vec for each review.
sent_vectors = []; # the avg-w2v for each sentence/review is stored in
this list
for sent in tqdm(list_of_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, yo
u might need to change this to 300 if you use google's w2v
    cnt_words = 0; # num of words with a valid vector in the sentence/re
view
```

```
100%|███████████████████████████████████████████████████████████████████████████  
██████████ | 4986/4986 [00:03<00:00, 1330.47it/s]
```



### 3. Graphviz

- Visualize your decision tree with Graphviz. It helps you to understand how a decision is being made, given a new vector.
- Since feature names are not obtained from word2vec related models, visualize only BOW & TFIDF decision trees using Graphviz
- Make sure to print the words in each node of the decision tree instead of printing its index.
- Just for visualization purpose, limit max\_depth to 2 or 3 and either embed the generated images of graphviz in your notebook, or directly upload them as .png files.



### 4. Feature importance

- Find the top 10 features of positive class and top 10 features of negative class for both feature sets **Set 1** and **Set 2** using `feature\_importances\_` method of [Decision Tree Classifier](#) and print their corresponding feature names

### 5. Feature engineering

- To increase the performance of your model, you can also experiment with with feature engineering like :
  - Taking length of reviews as another feature.
  - Considering some features from review summary as well.

### 6. Representation of results

- You need to plot the performance of model both on train data and cross validation data for each hyper parameter, like shown in the figure.  
 Once after you found the best hyper parameter, you need to train your model with it, and find the AUC on test data and plot the ROC curve on both train and test.  
 Along with plotting ROC curve, you need to print the [confusion matrix](#) with predicted and original labels of test data points. Please visualize your confusion matrices using [seaborn heatmaps](#).





## 7. Conclusion

- [You need to summarize the results at the end of the notebook, summarize it in the table format. To print out a table please refer to this prettytable library link](#)



### Note: Data Leakage

1. There will be an issue of data-leakage if you vectorize the entire data and then split it into train/cv/test.
2. To avoid the issue of data-leakag, make sure to split your data first and then vectorize it.
3. While vectorizing your data, apply the method fit\_transform() on you train data, and apply the method transform() on cv/test data.
4. For more details please go through this [link](#).

## Applying Decision Trees

```
In [2]: # after preprocessing
df = pd.read_pickle("files/preprocessed.pkl")
df.shape
```

Out[2]: (364171, 12)

```
In [3]: df.head()
```

Out[3]:

|  | <b>Id</b> | <b>ProductId</b> | <b>UserId</b> | <b>ProfileName</b> | <b>HelpfulnessNumerator</b> | <b>Helpfulness</b> |
|--|-----------|------------------|---------------|--------------------|-----------------------------|--------------------|
|--|-----------|------------------|---------------|--------------------|-----------------------------|--------------------|

|        | Id     | ProductId  | UserId         | ProfileName                        | HelpfulnessNumerator | HelpfulnessDenominator |
|--------|--------|------------|----------------|------------------------------------|----------------------|------------------------|
| 138706 | 150524 | 0006641040 | ACITT7DI6IDDL  | shari<br>zychinski                 | 0                    | 0                      |
| 138688 | 150506 | 0006641040 | A2IW4PEEKO2R0U | Tracy                              | 1                    | 1                      |
| 138689 | 150507 | 0006641040 | A1S4A3IQ2MU7V4 | sally sue<br>"sally sue"           | 1                    | 1                      |
| 138690 | 150508 | 0006641040 | AZGXZ2UUK6X    | Catherine<br>Hallberg "<br>(Kate)" | 1                    | 1                      |
| 138691 | 150509 | 0006641040 | A3CMRKGE0P909G | Teresa                             | 3                    | 3                      |

```
In [4]: from sklearn.model_selection import train_test_split
        from sklearn.grid_search import GridSearchCV
```

```
from sklearn.datasets import *
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score , f1_score , confusion_matri
X
from collections import Counter
from sklearn.metrics import accuracy_score, roc_auc_score , roc_curve
from sklearn.model_selection import train_test_split
```

/home/prasad/anaconda3/lib/python3.6/site-packages/sklearn/cross\_validation.py:41: DeprecationWarning: This module was deprecated in version 0.18 in favor of the model\_selection module into which all the refactored classes and functions are moved. Also note that the interface of the new CV iterators are different from that of this module. This module will be removed in 0.20.

"This module will be removed in 0.20.", DeprecationWarning)

/home/prasad/anaconda3/lib/python3.6/site-packages/sklearn/grid\_search.py:42: DeprecationWarning: This module was deprecated in version 0.18 in favor of the model\_selection module into which all the refactored classes and functions are moved. This module will be removed in 0.20.

DeprecationWarning)

```
In [5]: # take 50k sample data randomly
sample_data = df.sample(100000)
sample_data.shape
```

Out[5]: (100000, 12)

```
In [6]: # sorted the data using time based
sorted_data = sample_data.sort_values('Time', axis=0, inplace=False)
sorted_data.shape
```

Out[6]: (100000, 12)

```
In [7]: sorted_data['Score'].value_counts()
```

```
Out[7]: 1    84419
        0    15581
        Name: Score, dtype: int64
```

```
In [8]: X = np.array(sorted_data['CleanedText'])
y = np.array(sorted_data['Score'])
print(X.shape)
print(y.shape)

(100000,)
(100000,)
```

```
In [9]: # Simple cross validation
# split the data sent into train and test
train , test , train_y , test_y = train_test_split(X, y, test_size = 0.3, random_state=None)

# split the train data set into cross validation train and cross validation test
train, cv , train_y, cv_y = train_test_split(train, train_y, test_size=0.3, random_state=None)

print("train data = ", train.shape)
print("cross validation = ", cv.shape)
print("test data = ", test.shape)

train data = (49000,)
cross validation = (21000,)
test data = (30000,)
```

## [5.1] Applying Decision Trees on BOW, SET 1

```
In [10]: # Please write all the code with proper documentation
#BOW
count_vect = CountVectorizer(min_df=20) #in scikit-learn
count_vect.fit(train)
print("some feature names ", count_vect.get_feature_names()[:10])
print('='*50)

bow_train = count_vect.fit_transform(train)
bow_cv = count_vect.transform(cv)
```

```

bow_test = count_vect.transform(test)
print("=====Train Data=====")
print("the type of count vectorizer ",type(bow_train))
print("the shape of out text BOW vectorizer ",bow_train.get_shape())
print("the number of unique words ", bow_train.get_shape()[1])
print("=====Cross validation Data=====")
print("the type of count vectorizer ",type(bow_cv))
print("the shape of out text BOW vectorizer ",bow_cv.get_shape())
print("the number of unique words ", bow_cv.get_shape()[1])
print("=====Test Data=====")
print("the type of count vectorizer ",type(bow_test))
print("the shape of out text BOW vectorizer ",bow_test.get_shape())
print("the number of unique words ", bow_test.get_shape()[1])

```

```

some feature names ['ability', 'able', 'absolute', 'absolutely', 'abso
rb', 'absorbed', 'absorbs', 'abundance', 'acai', 'accent']

```

```

=====
=====Train Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (49000, 5975)
the number of unique words 5975
=====Cross validation Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (21000, 5975)
the number of unique words 5975
=====Test Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (30000, 5975)
the number of unique words 5975

```

```

In [11]: from sklearn.tree import DecisionTreeClassifier

```

```

In [12]: depth = [1, 5, 10, 50, 100, 500, 1000]
samples = [5, 10, 100, 500]
bow_train_auc = []
bow_cv_auc = []
for i in depth:
    for j in samples:
        DTC = DecisionTreeClassifier(criterion='gini',max_depth=i, min_

```

```

samples_split=j )
    DTC.fit(bow_train, train_y)
    # train data
    y_prob_train = DTC.predict_proba(bow_train)[: ,1]
    y_pred = np.where(y_prob_train > 0.5, 1, 0)
    auc_roc_train = roc_auc_score(train_y , y_prob_train)
    print('\nTrain AUC for min_sample = %s and max_depth = %s is %
0.2f%%' % (str(j),str(i),(auc_roc_train * float(100))))
    bow_train_auc.append(auc_roc_train)
    # CV
    y_prob_cv = DTC.predict_proba(bow_cv)[: ,1]
    y_pred = np.where(y_prob_cv > 0.5, 1, 0)
    auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
    print('\nCV AUC for min_sample = %s and max_depth = %s is %0.2
f%%' % (str(j),str(i),(auc_roc_cv * float(100))))
    bow_cv_auc.append(auc_roc_cv)
    print("="*50)

```

Train AUC for min\_sample = 5 and max\_depth = 1 is 60.69%

CV AUC for min\_sample = 5 and max\_depth = 1 is 60.20%

=====

Train AUC for min\_sample = 10 and max\_depth = 1 is 60.69%

CV AUC for min\_sample = 10 and max\_depth = 1 is 60.20%

=====

Train AUC for min\_sample = 100 and max\_depth = 1 is 60.69%

CV AUC for min\_sample = 100 and max\_depth = 1 is 60.20%

=====

Train AUC for min\_sample = 500 and max\_depth = 1 is 60.69%

CV AUC for min\_sample = 500 and max\_depth = 1 is 60.20%

=====

Train AUC for min\_sample = 5 and max\_depth = 5 is 71.30%

```
CV AUC for min_sample = 5 and max_depth = 5 is 71.06%
=====

Train AUC for min_sample = 10 and max_depth = 5 is 71.30%

CV AUC for min_sample = 10 and max_depth = 5 is 71.06%
=====

Train AUC for min_sample = 100 and max_depth = 5 is 71.25%

CV AUC for min_sample = 100 and max_depth = 5 is 71.08%
=====

Train AUC for min_sample = 500 and max_depth = 5 is 71.17%

CV AUC for min_sample = 500 and max_depth = 5 is 71.00%
=====

Train AUC for min_sample = 5 and max_depth = 10 is 79.30%

CV AUC for min_sample = 5 and max_depth = 10 is 76.22%
=====

Train AUC for min_sample = 10 and max_depth = 10 is 79.23%

CV AUC for min_sample = 10 and max_depth = 10 is 76.61%
=====

Train AUC for min_sample = 100 and max_depth = 10 is 78.58%

CV AUC for min_sample = 100 and max_depth = 10 is 77.23%
=====

Train AUC for min_sample = 500 and max_depth = 10 is 78.10%

CV AUC for min_sample = 500 and max_depth = 10 is 77.51%
=====

Train AUC for min_sample = 5 and max_depth = 50 is 97.20%
```

```
CV AUC for min_sample = 5 and max_depth = 50 is 70.66%
=====

Train AUC for min_sample = 10 and max_depth = 50 is 96.79%

CV AUC for min_sample = 10 and max_depth = 50 is 71.77%
=====

Train AUC for min_sample = 100 and max_depth = 50 is 94.54%

CV AUC for min_sample = 100 and max_depth = 50 is 78.02%
=====

Train AUC for min_sample = 500 and max_depth = 50 is 90.04%

CV AUC for min_sample = 500 and max_depth = 50 is 82.73%
=====

Train AUC for min_sample = 5 and max_depth = 100 is 99.32%

CV AUC for min_sample = 5 and max_depth = 100 is 66.80%
=====

Train AUC for min_sample = 10 and max_depth = 100 is 99.06%

CV AUC for min_sample = 10 and max_depth = 100 is 67.62%
=====

Train AUC for min_sample = 100 and max_depth = 100 is 97.04%

CV AUC for min_sample = 100 and max_depth = 100 is 76.43%
=====

Train AUC for min_sample = 500 and max_depth = 100 is 92.80%

CV AUC for min_sample = 500 and max_depth = 100 is 81.75%
=====
```



```
Train AUC for min_sample = 5 and max_depth = 500 is 99.93%
CV AUC for min_sample = 5 and max_depth = 500 is 69.20%
=====

Train AUC for min_sample = 10 and max_depth = 500 is 99.69%
CV AUC for min_sample = 10 and max_depth = 500 is 70.69%
=====

Train AUC for min_sample = 100 and max_depth = 500 is 97.92%
CV AUC for min_sample = 100 and max_depth = 500 is 76.69%
=====

Train AUC for min_sample = 500 and max_depth = 500 is 93.93%
CV AUC for min_sample = 500 and max_depth = 500 is 81.10%
=====

Train AUC for min_sample = 5 and max_depth = 1000 is 99.93%
CV AUC for min_sample = 5 and max_depth = 1000 is 69.13%
=====

Train AUC for min_sample = 10 and max_depth = 1000 is 99.71%
CV AUC for min_sample = 10 and max_depth = 1000 is 70.51%
=====


Train AUC for min_sample = 100 and max_depth = 1000 is 97.92%
CV AUC for min_sample = 100 and max_depth = 1000 is 77.00%
=====

Train AUC for min_sample = 500 and max_depth = 1000 is 93.96%
CV AUC for min_sample = 500 and max_depth = 1000 is 81.11%
=====
```

```
In [13]: import mglearn
```


```
In [38]: # auc on cv
cv_scores = np.array(bow_cv_auc).reshape(len(depth),len(samples))
df_cm = pd.DataFrame(cv_scores, depth, samples)
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 10}, fmt='.3g', linewidths=.5, cbar_kws={"orientation": "horizontal"})
```

```
Out[38]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3a956792b0>
```



```
In [51]: # train auc
train_scores = np.array(bow_train_auc).reshape(len(depth),len(samples))
df_cm = pd.DataFrame(train_scores, depth, samples)
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 10}, fmt='.3g', linewidths=.5, cbar_kws={"orientation": "horizontal"})
```

```
Out[51]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3aa6387438>
```



```
In [55]: depth = [1, 5, 10, 50, 100, 500, 1000]
# https://www.dataquest.io/blog/learning-curves-machine-learning/

import matplotlib.pyplot as plt
%matplotlib inline


plt.style.use('seaborn')

plt.plot(depth, train_scores[:, -1], label = 'Training AUC')
plt.plot(depth, cv_scores[:, -1], label = 'Validation AUC')

plt.ylabel('AUC', fontsize = 14)
plt.xlabel('depth', fontsize = 14)
plt.title('Learning curves for a Desision trees model', fontsize = 18,
y = 1.03)
plt.legend()
```

```
Out[55]: <matplotlib.legend.Legend at 0x7f3aa60a55f8>
```

### Learning curves for a Decision trees model



```
In [15]: i = 50
j = 500

DTC = DecisionTreeClassifier(criterion='gini' , max_depth=i, min_sample
s_split=j)
DTC.fit(bow_train, train_y)
# train data
y_prob_train = DTC.predict_proba(bow_train)[: ,1]
fprt, tprt, thresholdt = roc_curve(train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y , y_prob_train)
print('\nTrain AUC for min_sample = %s and depth = %s is %0.2f%%' % (st
r(j),str(i),(auc_roc_train * float(100))))
# CV
y_prob_cv = DTC.predict_proba(bow_cv)[: ,1]
fprc, tprc, thresholdc = roc_curve(cv_y, y_prob_cv)
```

```

y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
print('\nCV AUC for min_sample = %s and depth = %s is %0.2f%%' % (str(j)
),str(i),(auc_roc_cv * float(100)))
# Test
y_prob_test = DTC.predict_proba(bow_test)[: ,1]
fprts, tprts, thresholdts = roc_curve(test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
auc_roc_test = roc_auc_score(test_y , y_prob_test)
print('\nTest AUC for min_sample = %s and depth = %s is %0.2f%%' % (str
(j),str(i),(auc_roc_test * float(100)))

```

Train AUC for min\_sample = 500 and depth = 50 is 90.02%

CV AUC for min\_sample = 500 and depth = 50 is 82.81%


Test AUC for min\_sample = 500 and depth = 50 is 82.67%

In [16]: `# https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python`

```

import matplotlib.pyplot as plt
plt.clf()
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b' , label='AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y' , label='AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label='AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

```




```
In [17]: print("F1-Score on test set: %0.2f"%(f1_score(test_y, y_pred_test)))
```

F1-Score on test set: 0.92

```
In [18]: df_cm = pd.DataFrame(confusion_matrix(test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[18]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3aafe7c080>
```



### [5.1.1] Top 10 important features of from SET 1

```
In [19]: # Please write all the code with proper documentation
feature_imp = DTC.feature_importances_
feature_names = count_vect.get_feature_names()
features = dict(zip(feature_imp, feature_names))
```

```
In [20]: features_df = pd.DataFrame.from_dict(features, orient='index')
```

```
In [21]: sorted_features = features_df.sort_index(axis=0, ascending=False)
```

```
In [75]: # top 10 important features
sorted_features.head(10)
```

Out[75]:

|          | 0     |
|----------|-------|
| 0.095068 | not   |
| 0.060579 | great |

|          | 0            |
|----------|--------------|
| 0.047388 | worst        |
| 0.039606 | disappointed |
| 0.029607 | horrible     |
| 0.028482 | best         |
| 0.026107 | waste        |
| 0.025539 | delicious    |
| 0.024585 | money        |
| 0.020553 | love         |

### [5.1.2] Top 10 unimportant features from SET 1

In [76]: *# Please write all the code with proper documentation*  
*# top 10 unimportant features*  
 sorted\_features.tail(10)

Out[76]:

|          | 0          |
|----------|------------|
| 0.000195 | rise       |
| 0.000195 | justify    |
| 0.000194 | edible     |
| 0.000192 | skinny     |
| 0.000132 | flavor     |
| 0.000132 | italy      |
| 0.000131 | outrageous |
| 0.000130 | thirst     |
| 0.000050 | results    |



### [5.1.3] Graphviz visualization of Decision Tree on BOW, SET 1

```
In [77]: # Please write all the code with proper documentation
i = 3
j = 2

DTC = DecisionTreeClassifier(criterion='gini' , max_depth=i, min_sample
s_split=j)
DTC.fit(bow_train, train_y)
# train data
y_prob_train = DTC.predict_proba(bow_train)[: ,1]
fprt, tprt, thresholdt = roc_curve(train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y , y_prob_train)
print('\nTrain AUC for min_sample = %s and depth = %s is %0.2f%%' % (str(j),str(i),(auc_roc_train * float(100))))
# CV
y_prob_cv = DTC.predict_proba(bow_cv)[: ,1]
fprc, tprc, thresholdc = roc_curve(cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
print('\nCV AUC for min_sample = %s and depth = %s is %0.2f%%' % (str(j),str(i),(auc_roc_cv * float(100))))
# Test
y_prob_test = DTC.predict_proba(bow_test)[: ,1]
fprts, tprts, thresholdts = roc_curve(test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
auc_roc_test = roc_auc_score(test_y , y_prob_test)
print('\nTest AUC for min_sample = %s and depth = %s is %0.2f%%' % (str(j),str(i),(auc_roc_test * float(100))))
```

Train AUC for min\_sample = 2 and depth = 3 is 65.07%

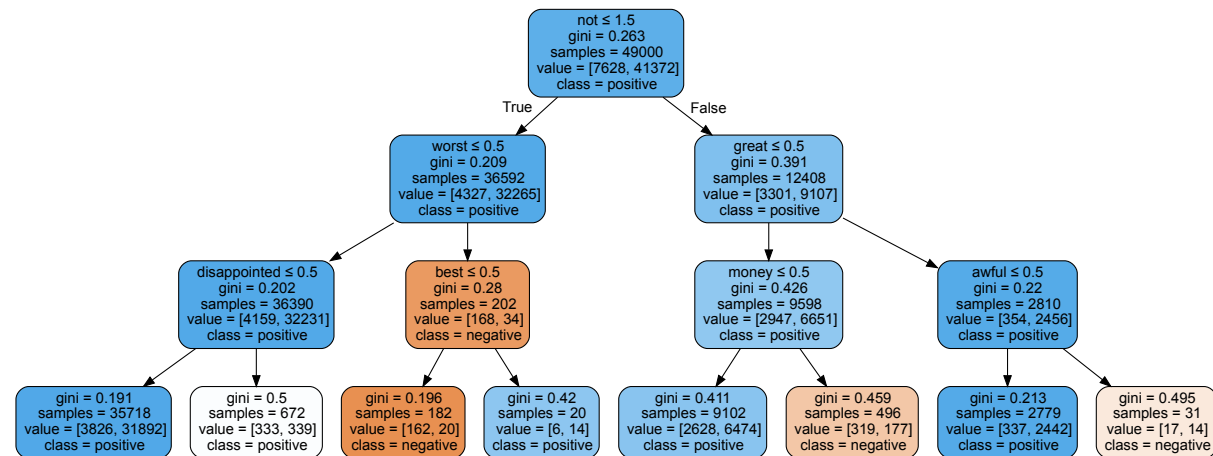
CV AUC for min\_sample = 2 and depth = 3 is 65.44%

Test AUC for min\_sample = 2 and depth = 3 is 64.88%

```
In [80]: import graphviz
        from sklearn import tree
```

```
In [87]: dot_data = tree.export_graphviz(DTC , out_file=None, feature_names =count_vect.get_feature_names() , class_names=['negative', 'positive'], filled=True, rounded=True, special_characters=True)
        graph = graphviz.Source(dot_data)
        graph
```

Out[87]:



```
In [94]: graph.view()
```

Out[94]: 'bow\_graph.gv.pdf'

## [5.2] Applying Decision Trees on TFIDF, SET 2

```
In [56]: # Please write all the code with proper documentation

        model = TfidfVectorizer(min_df=20, ngram_range=(1,2))
        #tf_idf_matrix = model.fit_transform(train)
```

```

print("=====Train Data=====")
tf_idf_train = model.fit_transform(train)
print("the type of count vectorizer ",type(tf_idf_train))
print("the shape of out text TFIDF vectorizer ",tf_idf_train.get_shape
())
print("the number of unique words including both unigrams and bigrams "
,tf_idf_train.get_shape()[1])
print("=====CV Data=====")
tf_idf_cv = model.transform(cv)
print("the type of count vectorizer ",type(tf_idf_cv))
print("the shape of out text TFIDF vectorizer ",tf_idf_cv.get_shape())
print("the number of unique words including both unigrams and bigrams "
,tf_idf_cv.get_shape()[1])
print("=====Test Data=====")
tf_idf_test = model.transform(test)
print("the type of count vectorizer ",type(tf_idf_test))
print("the shape of out text TFIDF vectorizer ",tf_idf_test.get_shape
())
print("the number of unique words including both unigrams and bigrams "
, tf_idf_test.get_shape()[1])

# we are converting a dictionary with word as a key, and the idf as a v
alue
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

```

```

=====Train Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (49000, 13730)
the number of unique words including both unigrams and bigrams 13730
=====CV Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (21000, 13730)
the number of unique words including both unigrams and bigrams 13730
=====Test Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (30000, 13730)
the number of unique words including both unigrams and bigrams 13730

```

```

In [57]: depth = [1, 5, 10, 50, 100, 500, 1000]
samples = [5, 10, 100, 500]
tfidf_train_auc = []
tfidf_cv_auc = []
for i in depth:
    for j in samples:
        DTC = DecisionTreeClassifier(criterion='gini',max_depth=i, min_
samples_split=j )
        DTC.fit(tf_idf_train, train_y)
        # train data
        y_prob_train = DTC.predict_proba(tf_idf_train)[:,-1]
        y_pred = np.where(y_prob_train > 0.5, 1, 0)
        auc_roc_train = roc_auc_score(train_y , y_prob_train)
        print('\nTrain AUC for min_sample = %s and max_depth = %s is %
0.2f%%' % (str(j),str(i),(auc_roc_train * float(100))))
        tfidf_train_auc.append(auc_roc_train)
        # CV
        y_prob_cv = DTC.predict_proba(tf_idf_cv)[:,-1]
        y_pred = np.where(y_prob_cv > 0.5, 1, 0)
        auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
        print('\nCV AUC for min_sample = %s and max_depth = %s is %0.2
f%%' % (str(j),str(i),(auc_roc_cv * float(100))))
        tfidf_cv_auc.append(auc_roc_cv)
        print("="*50)

```

Train AUC for min\_sample = 5 and max\_depth = 1 is 60.97%

CV AUC for min\_sample = 5 and max\_depth = 1 is 60.23%

=====

Train AUC for min\_sample = 10 and max\_depth = 1 is 60.97%

CV AUC for min\_sample = 10 and max\_depth = 1 is 60.23%

=====

Train AUC for min\_sample = 100 and max\_depth = 1 is 60.97%

CV AUC for min\_sample = 100 and max\_depth = 1 is 60.23%

=====

```
Train AUC for min_sample = 500 and max_depth = 1 is 60.97%
CV AUC for min_sample = 500 and max_depth = 1 is 60.23%
=====

Train AUC for min_sample = 5 and max_depth = 5 is 67.71%
CV AUC for min_sample = 5 and max_depth = 5 is 66.22%
=====

Train AUC for min_sample = 10 and max_depth = 5 is 67.71%
CV AUC for min_sample = 10 and max_depth = 5 is 66.22%
=====

Train AUC for min_sample = 100 and max_depth = 5 is 67.67%
CV AUC for min_sample = 100 and max_depth = 5 is 66.33%
=====

Train AUC for min_sample = 500 and max_depth = 5 is 67.62%
CV AUC for min_sample = 500 and max_depth = 5 is 66.32%
=====

Train AUC for min_sample = 5 and max_depth = 10 is 76.33%
CV AUC for min_sample = 5 and max_depth = 10 is 72.44%
=====

Train AUC for min_sample = 10 and max_depth = 10 is 76.22%
CV AUC for min_sample = 10 and max_depth = 10 is 72.75%
=====

Train AUC for min_sample = 100 and max_depth = 10 is 75.81%
CV AUC for min_sample = 100 and max_depth = 10 is 73.73%
=====
```

```
Train AUC for min_sample = 500 and max_depth = 10 is 75.43%
CV AUC for min_sample = 500 and max_depth = 10 is 73.89%
=====

Train AUC for min_sample = 5 and max_depth = 50 is 95.67%
CV AUC for min_sample = 5 and max_depth = 50 is 72.05%
=====

Train AUC for min_sample = 10 and max_depth = 50 is 95.30%
CV AUC for min_sample = 10 and max_depth = 50 is 74.09%
=====

Train AUC for min_sample = 100 and max_depth = 50 is 93.54%
CV AUC for min_sample = 100 and max_depth = 50 is 80.28%
=====

Train AUC for min_sample = 500 and max_depth = 50 is 91.06%
CV AUC for min_sample = 500 and max_depth = 50 is 82.81%
=====

Train AUC for min_sample = 5 and max_depth = 100 is 98.64%
CV AUC for min_sample = 5 and max_depth = 100 is 69.09%
=====

Train AUC for min_sample = 10 and max_depth = 100 is 98.52%
CV AUC for min_sample = 10 and max_depth = 100 is 70.28%
=====

Train AUC for min_sample = 100 and max_depth = 100 is 96.92%
CV AUC for min_sample = 100 and max_depth = 100 is 77.28%
=====
```

Train AUC for min\_sample = 500 and max\_depth = 100 is 94.98%

CV AUC for min\_sample = 500 and max\_depth = 100 is 81.04%

=====

Train AUC for min\_sample = 5 and max\_depth = 500 is 99.97%

CV AUC for min\_sample = 5 and max\_depth = 500 is 70.68%

=====

Train AUC for min\_sample = 10 and max\_depth = 500 is 99.88%

CV AUC for min\_sample = 10 and max\_depth = 500 is 71.88%

=====

Train AUC for min\_sample = 100 and max\_depth = 500 is 98.79%

CV AUC for min\_sample = 100 and max\_depth = 500 is 76.20%

=====

Train AUC for min\_sample = 500 and max\_depth = 500 is 97.28%

CV AUC for min\_sample = 500 and max\_depth = 500 is 78.31%

=====

Train AUC for min\_sample = 5 and max\_depth = 1000 is 99.97%

CV AUC for min\_sample = 5 and max\_depth = 1000 is 71.10%

=====

Train AUC for min\_sample = 10 and max\_depth = 1000 is 99.88%

CV AUC for min\_sample = 10 and max\_depth = 1000 is 72.42%

=====

Train AUC for min\_sample = 100 and max\_depth = 1000 is 98.82%

CV AUC for min\_sample = 100 and max\_depth = 1000 is 76.43%

=====  
Train AUC for min\_sample = 500 and max\_depth = 1000 is 97.23%

CV AUC for min\_sample = 500 and max\_depth = 1000 is 78.34%  
=====

```
In [60]: # cross validation auc
cv_scores = np.array(tfidf_cv_auc).reshape(len(depth), len(samples))
df_cm = pd.DataFrame(cv_scores, depth, samples)
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 10}, fmt='.3g', linewidths=.5, cbar_kws={"orientation": "horizontal"})
```

Out[60]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7f3ab4360160>



```
In [61]: # train auc
train_scores = np.array(tfidf_train_auc).reshape(len(depth), len(samples))
```

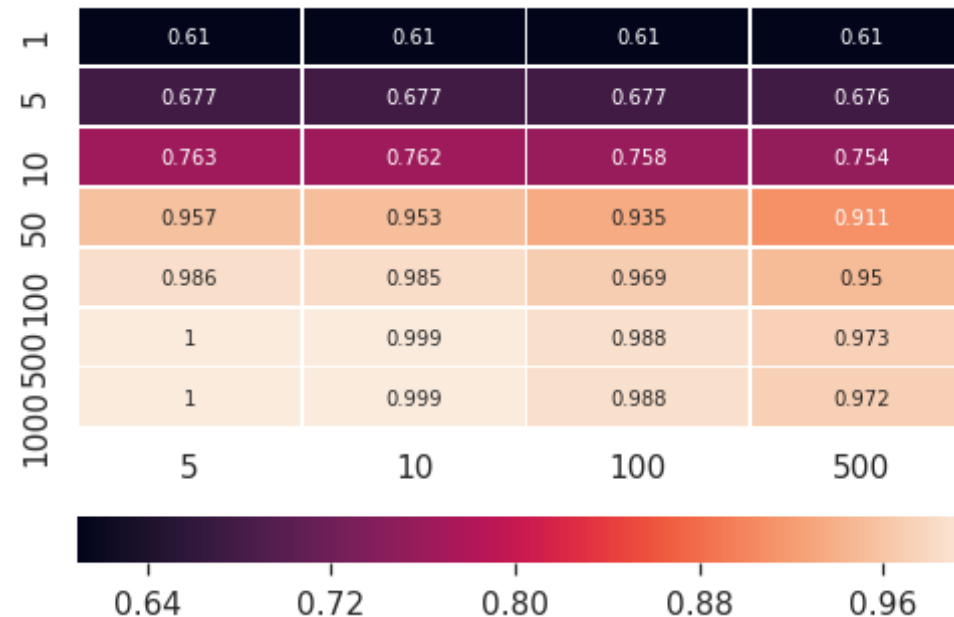


```

))
df_cm = pd.DataFrame(train_scores, depth, samples)
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 10}, fmt='.3g', linewidths=.5, cbar_kws={"orientation": "horizontal"})

```

Out[61]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7f3ab45b3e80>



```

In [62]: depth = [1, 5, 10, 50, 100, 500, 1000]
# https://www.dataquest.io/blog/learning-curves-machine-learning/

import matplotlib.pyplot as plt
%matplotlib inline

plt.style.use('seaborn')

plt.plot(depth, train_scores[:, -1], label = 'Training AUC')
plt.plot(depth, cv_scores[:, -1], label = 'Validation AUC')

```

```
plt.ylabel('AUC', fontsize = 14)
plt.xlabel('depth', fontsize = 14)
plt.title('Learning curves for a Desision trees model', fontsize = 18,
y = 1.03)
plt.legend()
```

Out[62]: <matplotlib.legend.Legend at 0x7f3ab49074a8>



```
In [158]: # Please write all the code with proper documentation
i = 50
j = 500

DTC = DecisionTreeClassifier(criterion='gini' , max_depth=i, min_sample
s_split=j)
DTC.fit(tf_idf_train, train_y)
# train data
```

```

y_prob_train = DTC.predict_proba(tf_idf_train)[: ,1]
fprt, tprt, thresholdt = roc_curve(train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y , y_prob_train)
print('\nTrain AUC for min_sample = %s and depth = %s is %0.2f%%' % (str(j),str(i),(auc_roc_train * float(100))))
# CV
y_prob_cv = DTC.predict_proba(tf_idf_cv)[: ,1]
fprc, tprc, thresholdc = roc_curve(cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
print('\nCV AUC for min_sample = %s and depth = %s is %0.2f%%' % (str(j),str(i),(auc_roc_cv * float(100))))
# Test
y_prob_test = DTC.predict_proba(tf_idf_test)[: ,1]
fprts, tprts, thresholdts = roc_curve(test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
auc_roc_test = roc_auc_score(test_y , y_prob_test)
print('\nTest AUC for min_sample = %s and depth = %s is %0.2f%%' % (str(j),str(i),(auc_roc_test * float(100))))

```

Train AUC for min\_sample = 500 and depth = 50 is 90.03%

CV AUC for min\_sample = 500 and depth = 50 is 82.36%

Test AUC for min\_sample = 500 and depth = 50 is 82.06%


In [159]: `# https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python`

```

import matplotlib.pyplot as plt
plt.clf()
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b' , label='AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y' , label='AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label='AUC Train = %0.2f' % (auc_roc_train * float(100)))

```

```
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```




```
In [161]: print("F1-Score on test set: %0.2f"%(f1_score(test_y, y_pred_test)))
```

F1-Score on test set: 0.92

```
In [162]: df_cm = pd.DataFrame(confusion_matrix(test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

Out[162]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7fce87c36240>



### [5.2.1] Top 10 important features from SET 2

```
In [164]: # Please write all the code with proper documentation
feature_imp = DTC.feature_importances_
feature_names = model.get_feature_names()
features = dict(zip(feature_imp, feature_names))
```

```
In [165]: features_df = pd.DataFrame.from_dict(features, orient='index')
```

```
In [166]: sorted_features = features_df.sort_index(axis=0, ascending=False)
```

```
In [167]: # top 10 important features
sorted_features.head(10)
```

Out[167]:

|          | 0     |
|----------|-------|
| 0.089370 | not   |
| 0.045175 | great |

|          | 0            |
|----------|--------------|
| 0.041375 | disappointed |
| 0.038242 | worst        |
| 0.034612 | not buy      |
| 0.030820 | waste money  |
| 0.025641 | money        |
| 0.025234 | horrible     |
| 0.019566 | refund       |
| 0.018403 | awful        |

### [5.2.2] Top 10 unimportant features from SET 2

In [168]: *# Please write all the code with proper documentation*  
*# top 10 unimportant features*  
 sorted\_features.tail(10)

Out[168]:

|          | 0             |
|----------|---------------|
| 0.000346 | big           |
| 0.000345 | one first     |
| 0.000345 | taste food    |
| 0.000345 | expensive not |
| 0.000344 | quantity      |
| 0.000343 | taste flavor  |
| 0.000343 | say not       |
| 0.000342 | dehydrated    |
| 0.000341 | requirements  |

### [5.2.3] Graphviz visualization of Decision Tree on TFIDF, SET 2

```
In [169]: # Please write all the code with proper documentation
i = 3
j = 2

DTC = DecisionTreeClassifier(criterion='gini' , max_depth=i, min_sample
s_split=j)
DTC.fit(tf_idf_train, train_y)
# train data
y_prob_train = DTC.predict_proba(tf_idf_train)[: ,1]
fprt, tprt, thresholdt = roc_curve(train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y , y_prob_train)
print('\nTrain AUC for min_sample = %s and depth = %s is %0.2f%%' % (str(j),str(i),(auc_roc_train * float(100))))
# CV
y_prob_cv = DTC.predict_proba(tf_idf_cv)[: ,1]
fprc, tprc, thresholdc = roc_curve(cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
print('\nCV AUC for min_sample = %s and depth = %s is %0.2f%%' % (str(j),str(i),(auc_roc_cv * float(100))))
# Test
y_prob_test = DTC.predict_proba(tf_idf_test)[: ,1]
fprts, tprts, thresholdts = roc_curve(test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
auc_roc_test = roc_auc_score(test_y , y_prob_test)
print('\nTest AUC for min_sample = %s and depth = %s is %0.2f%%' % (str(j),str(i),(auc_roc_test * float(100))))
```

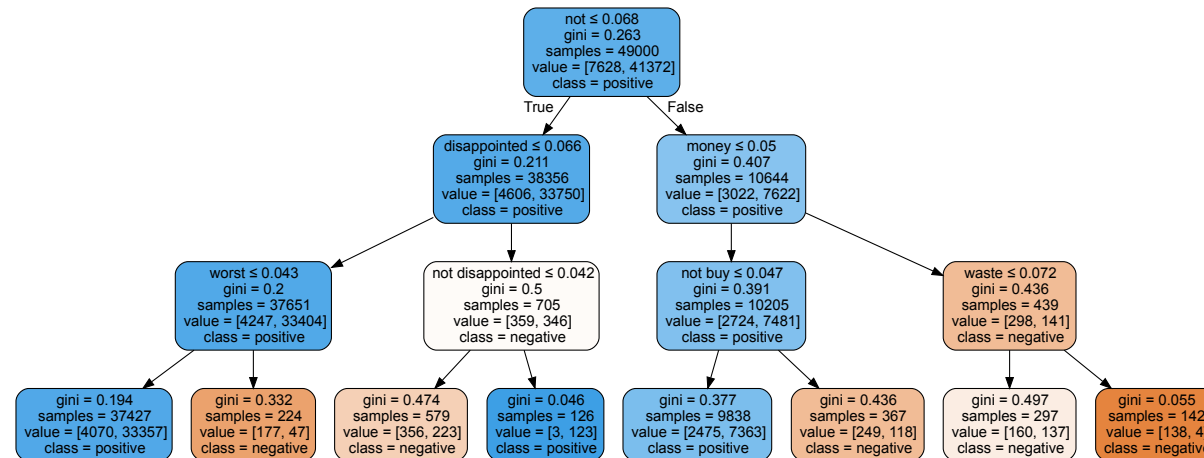
Train AUC for min\_sample = 2 and depth = 3 is 64.91%

CV AUC for min\_sample = 2 and depth = 3 is 65.01%

Test AUC for min\_sample = 2 and depth = 3 is 65.08%

```
In [171]: dot_data = tree.export_graphviz(DTC , out_file=None, feature_names=model.get_feature_names() , class_names=['negative', 'positive'], filled=True, rounded=True, special_characters=True)
graph = graphviz.Source(dot_data)
graph
```

Out[171]:



```
In [172]: graph.view()
```

Out[172]: 'Source.gv.pdf'

## [5.3] Applying Decision Trees on AVG W2V, SET 3

```
In [63]: # Train your own Word2Vec model using your own text corpus
##### Train Set #####
i=0
list_of_train_sentence=[]
for sentence in train:
    list_of_train_sentence.append(sentence.split())
##### CV Set #####
```



```

i=0
list_of_cv_sentence=[]
for sentence in cv:
    list_of_cv_sentence.append(sentence.split())
##### Test Set #####
i=0
list_of_test_sentence=[]
for sentence in test:
    list_of_test_sentence.append(sentence.split())
print("Length of Train = ", len(list_of_train_sentence))
print("Length of CV = ", len(list_of_cv_sentence))
print("Length of Test = ", len(list_of_test_sentence))

```

```

Length of Train = 49000
Length of CV = 21000
Length of Test = 30000

```

```

In [64]: w2v_model=Word2Vec(list_of_train_sentence,min_count=15,size=100, worker
s=4)
print(w2v_model.wv.most_similar('great'))
print('='*50)
print(w2v_model.wv.most_similar('worst'))

```

```

[('awesome', 0.7750544548034668), ('excellent', 0.7643860578536987),
('fantastic', 0.762516975402832), ('good', 0.7511047720909119), ('wonderful', 0.7507315874099731), ('fabulous', 0.70582115650177), ('terrific', 0.6960405707359314), ('amazing', 0.6841104626655579), ('perfect', 0.6794580817222595), ('decent', 0.5999100208282471)]
=====

```

```

[('greatest', 0.8142168521881104), ('best', 0.7570872902870178), ('tastiest', 0.728318452835083), ('smoothest', 0.681471586227417), ('disgusting', 0.5995469093322754), ('honestly', 0.5933093428611755), ('closest', 0.5928040742874146), ('superior', 0.5815048217773438), ('horrible', 0.5754640698432922), ('freshest', 0.5535082817077637)]

```

```

In [65]: w2v_words = list(w2v_model.wv.vocab)
print("number of words that occurred minimum 5 times ",len(w2v_words))
print("sample words ", w2v_words[0:50])

```

```
number of words that occurred minimum 5 times 7519
sample words ['fresh', 'nuts', 'excellent', 'service', 'shipped', 'fast', 'made', 'hazelnut', 'cake', 'hit', 'party', 'definitely', 'use', 'fruits', 'vegetables', 'skin', 'edible', 'negative', 'washed', 'produce', 'much', 'shorter', 'shelf', 'life', 'hate', 'bananas', 'not', 'big', 'fan', 'banana', 'nut', 'bread', 'either', 'expect', 'like', 'wrong', 'loved', 'good', 'cold', 'warmed', 'seconds', 'prefer', 'warm', 'favorite', 'part', 'fact', 'microwaved', 'plastic', 'bag', 'came']
```

```
In [66]: ##### Train data #####
# average Word2Vec
# compute average word2vec for each review.
sent_vectors_train = []; # the avg-w2v for each sentence/review is stored in this list
for sent in tqdm(list_of_train_sentence): # for each review/sentence
    sent_vec = np.zeros(100) # as word vectors are of zero length 50, you might need to change this to 300 if you use google's w2v
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors_train.append(sent_vec)
print(len(sent_vectors_train))
print(len(sent_vectors_train[0]))
```

```
100%|██████████| 49000/49000 [02:08<00:00, 380.99it/s]
```

```
49000
100
```

```
In [67]: ##### CV data #####
# average Word2Vec
# compute average word2vec for each review.
```

```

sent_vectors_cv = []; # the avg-w2v for each sentence/review is stored
                        in this list
for sent in tqdm(list_of_cv_sentence): # for each review/sentence
    sent_vec = np.zeros(100) # as word vectors are of zero length 50, y
    ou might need to change this to 300 if you use google's w2v
    cnt_words = 0; # num of words with a valid vector in the sentence/re
    view
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors_cv.append(sent_vec)
print(len(sent_vectors_cv))
print(len(sent_vectors_cv[0]))

```

```

100%|██████████| 21000/21000 [00:56<00:00, 374.22it/s]

```

```

21000

```

```

100

```

In [68]:

```

##### Test data #####
# average Word2Vec
# compute average word2vec for each review.
sent_vectors_test = []; # the avg-w2v for each sentence/review is store
d in this list
for sent in tqdm(list_of_test_sentence): # for each review/sentence
    sent_vec = np.zeros(100) # as word vectors are of zero length 50, y
    ou might need to change this to 300 if you use google's w2v
    cnt_words = 0; # num of words with a valid vector in the sentence/re
    view
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:

```

```

        sent_vec /= cnt_words
        sent_vectors_test.append(sent_vec)
print(len(sent_vectors_test))
print(len(sent_vectors_test[0]))

```

```

100%|██████████| 30000/30000 [01:26<00:00, 347.61it/s]

```

```

30000
100

```

```

In [69]: # save the datasets as numpy array
w2v_train = np.array(sent_vectors_train)
w2v_cv = np.array(sent_vectors_cv)
w2v_test = np.array(sent_vectors_test)

```

```

In [70]: depth = [1, 5, 10, 50, 100, 500, 1000]
samples = [5, 10, 100, 500]
w2v_train_auc = []
w2v_cv_auc = []
for i in depth:
    for j in samples:
        DTC = DecisionTreeClassifier(criterion='gini',max_depth=i, min_
samples_split=j )
        DTC.fit(w2v_train, train_y)
        # train data
        y_prob_train = DTC.predict_proba(w2v_train)[: ,1]
        y_pred = np.where(y_prob_train > 0.5, 1, 0)
        auc_roc_train = roc_auc_score(train_y , y_prob_train)
        print('\nTrain AUC for min_sample = %s and max_depth = %s is %
0.2f%%' % (str(j),str(i),(auc_roc_train * float(100))))
        w2v_train_auc.append(auc_roc_train)
        # CV
        y_prob_cv = DTC.predict_proba(w2v_cv)[: ,1]
        y_pred = np.where(y_prob_cv > 0.5, 1, 0)
        auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
        print('\nCV AUC for min_sample = %s and max_depth = %s is %0.2
f%%' % (str(j),str(i),(auc_roc_cv * float(100))))

```

```
w2v_cv_auc.append(auc_roc_cv)
print("="*50)
```

Train AUC for min\_sample = 5 and max\_depth = 1 is 65.07%

CV AUC for min\_sample = 5 and max\_depth = 1 is 64.50%

=====

Train AUC for min\_sample = 10 and max\_depth = 1 is 65.07%

CV AUC for min\_sample = 10 and max\_depth = 1 is 64.50%

=====

Train AUC for min\_sample = 100 and max\_depth = 1 is 65.07%

CV AUC for min\_sample = 100 and max\_depth = 1 is 64.50%

=====

Train AUC for min\_sample = 500 and max\_depth = 1 is 65.07%

CV AUC for min\_sample = 500 and max\_depth = 1 is 64.50%

=====

Train AUC for min\_sample = 5 and max\_depth = 5 is 82.38%

CV AUC for min\_sample = 5 and max\_depth = 5 is 80.20%

=====

Train AUC for min\_sample = 10 and max\_depth = 5 is 82.38%

CV AUC for min\_sample = 10 and max\_depth = 5 is 80.20%

=====

Train AUC for min\_sample = 100 and max\_depth = 5 is 82.38%

CV AUC for min\_sample = 100 and max\_depth = 5 is 80.20%

=====

Train AUC for min\_sample = 500 and max\_depth = 5 is 82.38%

```
CV AUC for min_sample = 500 and max_depth = 5 is 80.20%
=====

Train AUC for min_sample = 5 and max_depth = 10 is 91.74%

CV AUC for min_sample = 5 and max_depth = 10 is 78.43%
=====

Train AUC for min_sample = 10 and max_depth = 10 is 91.61%

CV AUC for min_sample = 10 and max_depth = 10 is 79.12%
=====

Train AUC for min_sample = 100 and max_depth = 10 is 89.95%

CV AUC for min_sample = 100 and max_depth = 10 is 81.56%
=====

Train AUC for min_sample = 500 and max_depth = 10 is 86.87%

CV AUC for min_sample = 500 and max_depth = 10 is 81.69%
=====

Train AUC for min_sample = 5 and max_depth = 50 is 99.97%

CV AUC for min_sample = 5 and max_depth = 50 is 67.36%
=====

Train AUC for min_sample = 10 and max_depth = 50 is 99.73%

CV AUC for min_sample = 10 and max_depth = 50 is 69.27%
=====

Train AUC for min_sample = 100 and max_depth = 50 is 94.58%

CV AUC for min_sample = 100 and max_depth = 50 is 78.89%
=====

Train AUC for min_sample = 500 and max_depth = 50 is 88.23%
```

```
CV AUC for min_sample = 500 and max_depth = 50 is 81.53%
=====

Train AUC for min_sample = 5 and max_depth = 100 is 99.97%

CV AUC for min_sample = 5 and max_depth = 100 is 67.30%
=====

Train AUC for min_sample = 10 and max_depth = 100 is 99.73%

CV AUC for min_sample = 10 and max_depth = 100 is 69.19%
=====

Train AUC for min_sample = 100 and max_depth = 100 is 94.60%

CV AUC for min_sample = 100 and max_depth = 100 is 78.81%
=====

Train AUC for min_sample = 500 and max_depth = 100 is 88.24%

CV AUC for min_sample = 500 and max_depth = 100 is 81.56%
=====

Train AUC for min_sample = 5 and max_depth = 500 is 99.97%

CV AUC for min_sample = 5 and max_depth = 500 is 67.81%
=====

Train AUC for min_sample = 10 and max_depth = 500 is 99.73%

CV AUC for min_sample = 10 and max_depth = 500 is 68.99%
=====

Train AUC for min_sample = 100 and max_depth = 500 is 94.60%

CV AUC for min_sample = 100 and max_depth = 500 is 78.84%
=====
```

Train AUC for min\_sample = 500 and max\_depth = 500 is 88.25%

CV AUC for min\_sample = 500 and max\_depth = 500 is 81.48%

=====

Train AUC for min\_sample = 5 and max\_depth = 1000 is 99.97%

CV AUC for min\_sample = 5 and max\_depth = 1000 is 67.11%

=====

Train AUC for min\_sample = 10 and max\_depth = 1000 is 99.73%

CV AUC for min\_sample = 10 and max\_depth = 1000 is 69.11%

=====

Train AUC for min\_sample = 100 and max\_depth = 1000 is 94.61%

CV AUC for min\_sample = 100 and max\_depth = 1000 is 78.84%

=====

Train AUC for min\_sample = 500 and max\_depth = 1000 is 88.24%


CV AUC for min\_sample = 500 and max\_depth = 1000 is 81.51%

=====

```
In [77]: # cv auc
cv_scores = np.array(w2v_cv_auc).reshape(len(depth),len(samples))
df_cm = pd.DataFrame(cv_scores, depth, samples)
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 10}, fmt='.3g', linewidths=.5, cbar_kws={"orientation": "horizontal"})
```


```
Out[77]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3a7d9acb38>
```





```
In [80]: # train auc
train_scores = np.array(w2v_train_auc).reshape(len(depth), len(samples))
df_cm = pd.DataFrame(train_scores, depth, samples)
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 10}, fmt='.3g', linewidths=.5, cbar_kws={"orientation": "horizontal"})
```

Out[80]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7f3a7d8215f8>



```
In [81]: depth = [1, 5, 10, 50, 100, 500, 1000]
# https://www.dataquest.io/blog/learning-curves-machine-learning/

import matplotlib.pyplot as plt
%matplotlib inline


plt.style.use('seaborn')

plt.plot(depth, train_scores[:, -1], label = 'Training AUC')
plt.plot(depth, cv_scores[:, -1], label = 'Validation AUC')

plt.ylabel('AUC', fontsize = 14)
plt.xlabel('depth', fontsize = 14)
plt.title('Learning curves for a Decision trees model', fontsize = 18,
y = 1.03)
plt.legend()
```

Out[81]: <matplotlib.legend.Legend at 0x7f3a7d70e978>

## Learning curves for a Decision trees model



```
In [183]: # Please write all the code with proper documentation
i = 10
j = 500

DTC = DecisionTreeClassifier(criterion='gini' , max_depth=i, min_sample
s_split=j)
DTC.fit(w2v_train, train_y)
# train data
y_prob_train = DTC.predict_proba(w2v_train)[: ,1]
fprt, tprt, thresholdt = roc_curve(train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y , y_prob_train)
print('\nTrain AUC for min_sample = %s and depth = %s is %0.2f%%' % (st
r(j),str(i),(auc_roc_train * float(100))))
# CV
y_prob_cv = DTC.predict_proba(w2v_cv)[: ,1]
```

```
fprc, tprc, thresholdc = roc_curve(cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y, y_prob_cv)
print('\nCV AUC for min_sample = %s and depth = %s is %0.2f%%' % (str(j), str(i), (auc_roc_cv * float(100))))
# Test
y_prob_test = DTC.predict_proba(w2v_test)[: ,1]
fprts, tprts, thresholdts = roc_curve(test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
auc_roc_test = roc_auc_score(test_y, y_prob_test)
print('\nTest AUC for min_sample = %s and depth = %s is %0.2f%%' % (str(j), str(i), (auc_roc_test * float(100))))
```


Train AUC for min\_sample = 500 and depth = 10 is 87.13%

CV AUC for min\_sample = 500 and depth = 10 is 83.00%

Test AUC for min\_sample = 500 and depth = 10 is 83.11%

In [184]: *# <https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python>*

```
import matplotlib.pyplot as plt
plt.clf()
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b', label='AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y', label='AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label='AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```




```
In [185]: print("F1-Score on test set: %0.2f"%(f1_score(test_y, y_pred_test)))
```

F1-Score on test set: 0.92

```
In [186]: df_cm = pd.DataFrame(confusion_matrix(test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[186]: <matplotlib.axes._subplots.AxesSubplot at 0x7fce6c0f2160>
```



## [5.4] Applying Decision Trees on TFIDF W2V, SET 4

```
In [187]: # Please write all the code with proper documentation

model = TfidfVectorizer()
#tf_idf_matrix = model.fit_transform(train)

print("=====Train Data=====")
final_tf_idf_train = model.fit_transform(train)
print("the type of count vectorizer ",type(final_tf_idf_train))
print("the shape of out text TFIDF vectorizer ",final_tf_idf_train.get_shape())
print("the number of unique words including both unigrams and bigrams "
, final_tf_idf_train.get_shape()[1])
print("=====CV Data=====")
final_tf_idf_cv = model.transform(cv)
print("the type of count vectorizer ",type(final_tf_idf_cv))
print("the shape of out text TFIDF vectorizer ",final_tf_idf_cv.get_shape())
print("the number of unique words including both unigrams and bigrams "
, final_tf_idf_cv.get_shape()[1])
```

```

print("=====Test Data=====")
final_tf_idf_test = model.transform(test)
print("the type of count vectorizer ",type(final_tf_idf_test))
print("the shape of out text TFIDF vectorizer ",final_tf_idf_test.get_shape())
print("the number of unique words including both unigrams and bigrams "
, final_tf_idf_test.get_shape()[1])

# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

=====Train Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (49000, 42866)
the number of unique words including both unigrams and bigrams 42866
=====CV Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (21000, 42866)
the number of unique words including both unigrams and bigrams 42866
=====Test Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (30000, 42866)
the number of unique words including both unigrams and bigrams 42866

```

```

In [188]: ##### Train #####
# TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

train_tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in tqdm(list_of_train_sentence): # for each review/sentence
    sent_vec = np.zeros(100) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence

```

```

        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            #tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    train_tfidf_sent_vectors.append(sent_vec)
    row += 1

```

100%|██████████| 49000/49000 [50:42<00:00, 20.65it/s]

```

In [189]: ##### CV #####
# TF-IDF weighted Word2Vec
#tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

cv_tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in tqdm(list_of_cv_sentence): # for each review/sentence
    sent_vec = np.zeros(100) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            #tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf

```



```

    if weight_sum != 0:
        sent_vec /= weight_sum
    cv_tfidf_sent_vectors.append(sent_vec)
    row += 1

```

100%|██████████| 21000/21000 [15:39<00:00, 22.34it/s]

```

In [190]: ##### Train #####
# TF-IDF weighted Word2Vec
#tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and ce
ll_val = tfidf

test_tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review
is stored in this list
row=0;
for sent in tqdm(list_of_test_sentence): # for each review/sentence
    sent_vec = np.zeros(100) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/r
review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            #tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole courpus
            # sent.count(word) = tf valeus of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    test_tfidf_sent_vectors.append(sent_vec)
    row += 1

```

100%|██████████| 30000/30000 [24:17<00:00, 20.58it/s]

```

In [191]: # save the datasets as numpy array
tfidf_w2v_train = np.array(train_tfidf_sent_vectors)

```

```
tfidf_w2v_cv = np.array(cv_tfidf_sent_vectors)
tfidf_w2v_test = np.array(test_tfidf_sent_vectors)
```

```
In [193]: depth = [1, 5, 10, 50, 100, 500, 1000]
samples = [5, 10, 100, 500]
tfidf_w2v_train_auc = []
tfidf_w2v_cv_auc = []
for i in depth:
    for j in samples:
        DTC = DecisionTreeClassifier(criterion='gini',max_depth=i, min_
samples_split=j )
        DTC.fit(tfidf_w2v_train, train_y)
        # train data
        y_prob_train = DTC.predict_proba(tfidf_w2v_train)[: ,1]
        y_pred = np.where(y_prob_train > 0.5, 1, 0)
        auc_roc_train = roc_auc_score(train_y , y_prob_train)
        print('\nTrain AUC for min_sample = %s and max_depth = %s is %
0.2f%%' % (str(j),str(i),(auc_roc_train * float(100))))
        tfidf_w2v_train_auc.append(auc_roc_train)
        # CV
        y_prob_cv = DTC.predict_proba(tfidf_w2v_cv)[: ,1]
        y_pred = np.where(y_prob_cv > 0.5, 1, 0)
        auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
        print('\nCV AUC for min_sample = %s and max_depth = %s is %0.2
f%%' % (str(j),str(i),(auc_roc_cv * float(100))))
        tfidf_w2v_cv_auc.append(auc_roc_cv)
    print("="*50)
```

Train AUC for min\_sample = 5 and max\_depth = 1 is 66.20%

CV AUC for min\_sample = 5 and max\_depth = 1 is 66.25%

=====

Train AUC for min\_sample = 10 and max\_depth = 1 is 66.20%

CV AUC for min\_sample = 10 and max\_depth = 1 is 66.25%

=====

Train AUC for min\_sample = 100 and max\_depth = 1 is 66.20%

```
CV AUC for min_sample = 100 and max_depth = 1 is 66.25%
=====

Train AUC for min_sample = 500 and max_depth = 1 is 66.20%

CV AUC for min_sample = 500 and max_depth = 1 is 66.25%
=====

Train AUC for min_sample = 5 and max_depth = 5 is 80.38%

CV AUC for min_sample = 5 and max_depth = 5 is 79.14%
=====

Train AUC for min_sample = 10 and max_depth = 5 is 80.38%

CV AUC for min_sample = 10 and max_depth = 5 is 79.14%
=====

Train AUC for min_sample = 100 and max_depth = 5 is 80.38%

CV AUC for min_sample = 100 and max_depth = 5 is 79.14%
=====

Train AUC for min_sample = 500 and max_depth = 5 is 80.38%

CV AUC for min_sample = 500 and max_depth = 5 is 79.14%
=====

Train AUC for min_sample = 5 and max_depth = 10 is 89.89%

CV AUC for min_sample = 5 and max_depth = 10 is 77.01%
=====

Train AUC for min_sample = 10 and max_depth = 10 is 89.73%

CV AUC for min_sample = 10 and max_depth = 10 is 77.45%
=====
```

```
Train AUC for min_sample = 100 and max_depth = 10 is 87.93%
CV AUC for min_sample = 100 and max_depth = 10 is 79.73%
=====

Train AUC for min_sample = 500 and max_depth = 10 is 84.57%
CV AUC for min_sample = 500 and max_depth = 10 is 80.43%
=====

Train AUC for min_sample = 5 and max_depth = 50 is 99.96%
CV AUC for min_sample = 5 and max_depth = 50 is 65.32%
=====

Train AUC for min_sample = 10 and max_depth = 50 is 99.69%
CV AUC for min_sample = 10 and max_depth = 50 is 67.01%
=====

Train AUC for min_sample = 100 and max_depth = 50 is 93.85%
CV AUC for min_sample = 100 and max_depth = 50 is 76.35%
=====

Train AUC for min_sample = 500 and max_depth = 50 is 86.34%
CV AUC for min_sample = 500 and max_depth = 50 is 79.89%
=====

Train AUC for min_sample = 5 and max_depth = 100 is 99.96%
CV AUC for min_sample = 5 and max_depth = 100 is 65.52%
=====

Train AUC for min_sample = 10 and max_depth = 100 is 99.70%
CV AUC for min_sample = 10 and max_depth = 100 is 67.20%
=====
```

Train AUC for min\_sample = 100 and max\_depth = 100 is 93.76%  
CV AUC for min\_sample = 100 and max\_depth = 100 is 76.50%  
=====

Train AUC for min\_sample = 500 and max\_depth = 100 is 86.33%  
CV AUC for min\_sample = 500 and max\_depth = 100 is 79.82%  
=====

Train AUC for min\_sample = 5 and max\_depth = 500 is 99.96%  
CV AUC for min\_sample = 5 and max\_depth = 500 is 65.88%  
=====

Train AUC for min\_sample = 10 and max\_depth = 500 is 99.69%  
CV AUC for min\_sample = 10 and max\_depth = 500 is 66.95%  
=====

Train AUC for min\_sample = 100 and max\_depth = 500 is 93.78%  
CV AUC for min\_sample = 100 and max\_depth = 500 is 76.58%  
=====

Train AUC for min\_sample = 500 and max\_depth = 500 is 86.35%  
CV AUC for min\_sample = 500 and max\_depth = 500 is 79.86%  
=====


Train AUC for min\_sample = 5 and max\_depth = 1000 is 99.96%  
CV AUC for min\_sample = 5 and max\_depth = 1000 is 65.64%  
=====

Train AUC for min\_sample = 10 and max\_depth = 1000 is 99.70%  
CV AUC for min\_sample = 10 and max\_depth = 1000 is 66.86%

```
=====
Train AUC for min_sample = 100 and max_depth = 1000 is 93.81%
CV AUC for min_sample = 100 and max_depth = 1000 is 76.50%
=====
Train AUC for min_sample = 500 and max_depth = 1000 is 86.32%
CV AUC for min_sample = 500 and max_depth = 1000 is 79.87%
=====
```


```
In [82]: # cv auc
cv_scores = np.array(tfidf_w2v_cv_auc).reshape(len(depth), len(samples))
df_cm = pd.DataFrame(cv_scores, depth, samples)
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 10}, fmt='.3g', linewidths=.5, cbar_kws={"orientation": "horizontal"})
```

```
Out[82]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3a7d6bdc18>
```



```
In [83]: # train auc
train_scores = np.array(tfidf_w2v_cv_auc).reshape(len(depth), len(samples))
df_cm = pd.DataFrame(train_scores, depth, samples)
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 10}, fmt='.3g', linewidths=.5, cbar_kws={"orientation": "horizontal"})
```

Out[83]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7f3a7d636390>



```
In [84]: depth = [1, 5, 10, 50, 100, 500, 1000]
# https://www.dataquest.io/blog/learning-curves-machine-learning/

import matplotlib.pyplot as plt
%matplotlib inline

plt.style.use('seaborn')


plt.plot(depth, train_scores[:, -1], label = 'Training AUC')
plt.plot(depth, cv_scores[:, -1], label = 'Validation AUC')

plt.ylabel('AUC', fontsize = 14)
plt.xlabel('depth', fontsize = 14)
plt.title('Learning curves for a Decision trees model', fontsize = 18,
y = 1.03)
plt.legend()
```

Out[84]: <matplotlib.legend.Legend at 0x7f3a7d53c3c8>



## Learning curves for a Decision trees model



```
In [195]: # Please write all the code with proper documentation
i = 10
j = 500

DTC = DecisionTreeClassifier(criterion='gini' , max_depth=i, min_sample
s_split=j)
DTC.fit(tfidf_w2v_train, train_y)
# train data
y_prob_train = DTC.predict_proba(tfidf_w2v_train)[: ,1]
fprt, tprt, thresholdt = roc_curve(train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y , y_prob_train)
print('\nTrain AUC for min_sample = %s and depth = %s is %0.2f%%' % (str
(j),str(i),(auc_roc_train * float(100))))
# CV
y_prob_cv = DTC.predict_proba(tfidf_w2v_cv)[: ,1]
```

```
fprc, tprc, thresholdc = roc_curve(cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y, y_prob_cv)
print('\nCV AUC for min_sample = %s and depth = %s is %0.2f%%' % (str(j), str(i), (auc_roc_cv * float(100))))
# Test
y_prob_test = DTC.predict_proba(tfidf_w2v_test)[: ,1]
fprts, tprts, thresholdts = roc_curve(test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
auc_roc_test = roc_auc_score(test_y, y_prob_test)
print('\nTest AUC for min_sample = %s and depth = %s is %0.2f%%' % (str(j), str(i), (auc_roc_test * float(100))))
```


Train AUC for min\_sample = 500 and depth = 10 is 84.57%

CV AUC for min\_sample = 500 and depth = 10 is 80.42%

Test AUC for min\_sample = 500 and depth = 10 is 79.82%

In [196]: `# https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python`

```
import matplotlib.pyplot as plt
plt.clf()
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b', label='AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y', label='AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label='AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```




```
In [197]: print("F1-Score on test set: %0.2f"%(f1_score(test_y, y_pred_test)))
```

F1-Score on test set: 0.92

```
In [198]: df_cm = pd.DataFrame(confusion_matrix(test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[198]: <matplotlib.axes._subplots.AxesSubplot at 0x7fce88ee08d0>
```



## [6] Conclusions

In [199]:

```
# Please compare all your models using Prettytable library
from prettytable import PrettyTable

x = PrettyTable(["Vectorizer" , "max_depth","min_sample_split" , "AUC",
"F1 Score"])

x.add_row(["BOW", 50 ,500, "82.73%", 0.92])
x.add_row(["TFIDF",50 ,500,"83.06%", 0.92])
x.add_row(["AVG-W2V",10 , 500,"83.11%", 0.92])
x.add_row(["TFIDF-W2V",10 , 500,"79.82%", 0.92])

print(x.get_string(title="DTC"))
```

| Vectorizer | max_depth | min_sample_split | AUC    | F1 Score |
|------------|-----------|------------------|--------|----------|
| BOW        | 50        | 500              | 82.73% | 0.92     |
| TFIDF      | 50        | 500              | 83.06% | 0.92     |
| AVG-W2V    | 10        | 500              | 83.11% | 0.92     |

|   |           |  |    |  |     |  |        |  |      |  |
|---|-----------|--|----|--|-----|--|--------|--|------|--|
|   | TFIDF-W2V |  | 10 |  | 500 |  | 79.82% |  | 0.92 |  |
| + | -----+    |  |    |  |     |  |        |  |      |  |

In [ ]: