

09 Amazon Fine Food Reviews Analysis_RF

February 11, 2019

1 Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454 Number of users: 256,059 Number of products: 74,258 Timespan: Oct 1999 - Oct 2012 Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective: Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative? [Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

2 [1]. Reading Data

2.1 [1.1] Loading the data

The dataset is available in two forms 1. .csv file 2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to “positive”. Otherwise, it will be set to “negative”.

```
In [2]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

In [2]: # using SQLite Table to read data.
con = sqlite3.connect('database.sqlite')

# filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
# SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 500000 data points.
# you can change the number to any other number based on your computing power

# filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000 """, con)
```

```

# for tsne assignment you can take 5k data points

filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 5000

# Give reviews with Score>3 a positive rating(1), and reviews with a score<3 a negative rating(0)
def partition(x):
    if x < 3:
        return 0
    return 1

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)

```

Number of data points in our data (5000, 10)

```

Out[2]:
   Id  ProductId  UserId  ProfileName \
0   1  B001E4KFG0  A3SGXH7AUHU8GW  delmartian
1   2  B00813GRG4  A1D87F6ZCVE5NK  dll pa
2   3  B000LQOCHO  ABXLMWJIXXAIN  Natalia Corres "Natalia Corres"

   HelpfulnessNumerator  HelpfulnessDenominator  Score  Time \
0                      1                      1      1  1303862400
1                      0                      0      0  1346976000
2                      1                      1      1  1219017600

   Summary  Text
0  Good Quality Dog Food  I have bought several of the Vitality canned d...
1  Not as Advertised  Product arrived labeled as Jumbo Salted Peanut...
2  "Delight" says it all  This is a confection that has been around a fe...

```

```

In [3]: display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", con)

```

```

In [4]: print(display.shape)
display.head()

```

(80668, 7)

```

Out[4]:
   UserId  ProductId  ProfileName  Time  Score \
0  #oc-R115TNMSPFT9I7  B007Y59HVM  Breyton  1331510400  2

```

1	#oc-R11D9D7SHXIJB9	B005HG9ETO	Louis E. Emory "hoppy"	1342396800	5
2	#oc-R11DNU2NBKQ23Z	B007Y59HVM	Kim Cieszykowski	1348531200	1
3	#oc-R1105J5ZVQE25C	B005HG9ETO	Penguin Chick	1346889600	5
4	#oc-R12KPB0DL2B5ZD	B0070SBE1U	Christopher P. Presta	1348617600	1

	Text	COUNT(*)
0	Overall its just OK when considering the price...	2
1	My wife has recurring extreme muscle spasms, u...	3
2	This coffee is horrible and unfortunately not ...	2
3	This will be the bottle that you grab from the...	3
4	I didnt like this coffee. Instead of telling y...	2

```
In [5]: display[display['UserId']=='AZY10LLTJ71NX']
```

```
Out [5]:
```

	UserId	ProductId	ProfileName	Time \
80638	AZY10LLTJ71NX	B006P7E5ZI	undertheshrine "undertheshrine"	1334707200

	Score	Text	COUNT(*)
80638	5	I was recommended to try green tea extract to ...	5

```
In [6]: display['COUNT(*)'].sum()
```

```
Out [6]: 393063
```

3 [2] Exploratory Data Analysis

3.1 [2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

```
In [7]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)
display.head()
```

```
Out [7]:
```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator \
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2
3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2
4	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2

	HelpfulnessDenominator	Score	Time \
0	2	5	1199577600

1	2	5	1199577600
2	2	5	1199577600
3	2	5	1199577600
4	2	5	1199577600

	Summary \
0	LOACKER QUADRATINI VANILLA WAFERS
1	LOACKER QUADRATINI VANILLA WAFERS
2	LOACKER QUADRATINI VANILLA WAFERS
3	LOACKER QUADRATINI VANILLA WAFERS
4	LOACKER QUADRATINI VANILLA WAFERS

	Text
0	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
1	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
2	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
3	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
4	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8) ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```
In [8]: #Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False)
```

```
In [9]: #Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first')
final.shape
```

```
Out[9]: (4986, 10)
```

```
In [10]: #Checking to see how much % of data still remains
(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

```
Out[10]: 99.72
```

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

```

In [11]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)

display.head()

Out[11]:
```

	Id	ProductId	UserId	ProfileName	\
0	64422	B000MIDR0Q	A161DK06JJMCYF	J. E. Stephens	"Jeanne"
1	44737	B001EQ55RW	A2V0I904FH7ABY		Ram

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0	3	1	5	1224892800	
1	3	2	4	1212883200	

	Summary	\
0	Bought This for My Son at College	
1	Pure cocoa taste with crunchy almonds inside	

	Text
0	My son loves spaghetti so I didn't hesitate or...
1	It was almost a 'love at first bite' - the per...

```

In [12]: final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]

In [13]: #Before starting the next phase of preprocessing lets see the number of entries left
print(final.shape)

#How many positive and negative reviews are present in our dataset?
final['Score'].value_counts()

(4986, 10)

```

```

Out[13]: 1    4178
         0     808
         Name: Score, dtype: int64

```

4 [3] Preprocessing

4.1 [3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags

2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```
In [14]: # printing some random reviews
        sent_0 = final['Text'].values[0]
        print(sent_0)
        print("="*50)

        sent_1000 = final['Text'].values[1000]
        print(sent_1000)
        print("="*50)

        sent_1500 = final['Text'].values[1500]
        print(sent_1500)
        print("="*50)

        sent_4900 = final['Text'].values[4900]
        print(sent_4900)
        print("="*50)
```

```
Why is this $[...] when the same product is available for $[...] here?<br />http://www.amazon.
=====
I recently tried this flavor/brand and was surprised at how delicious these chips are.  The be
=====
Wow.  So far, two two-star reviews.  One obviously had no idea what they were ordering; the oth
=====
love to order my coffee on amazon.  easy and shows up quickly.<br />This k cup is great coffee
=====
```

```
In [15]: # remove urls from text python: https://stackoverflow.com/a/40823105/4084039
        sent_0 = re.sub(r"http\S+", "", sent_0)
        sent_1000 = re.sub(r"http\S+", "", sent_1000)
        sent_1500 = re.sub(r"http\S+", "", sent_1500)
        sent_4900 = re.sub(r"http\S+", "", sent_4900)

        print(sent_0)
```

```
Why is this $[...] when the same product is available for $[...] here?<br /> /><br />The Victor
```

```
In [16]: # https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-remove-all
        from bs4 import BeautifulSoup
```

```

soup = BeautifulSoup(sent_0, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1000, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1500, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_4900, 'lxml')
text = soup.get_text()
print(text)

```

Why is this \$[...] when the same product is available for \$[...] here? />The Victor M380 and M

=====

I recently tried this flavor/brand and was surprised at how delicious these chips are. The be

=====

Wow. So far, two two-star reviews. One obviously had no idea what they were ordering; the ot

=====

love to order my coffee on amazon. easy and shows up quickly.This k cup is great coffee. dca

```
In [17]: # https://stackoverflow.com/a/47091490/4084039
```

```

import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\'re", " are", phrase)
    phrase = re.sub(r"\'s", " is", phrase)
    phrase = re.sub(r"\'d", " would", phrase)
    phrase = re.sub(r"\'ll", " will", phrase)
    phrase = re.sub(r"\'t", " not", phrase)
    phrase = re.sub(r"\'ve", " have", phrase)
    phrase = re.sub(r"\'m", " am", phrase)
    return phrase

```

```
In [18]: sent_1500 = decontracted(sent_1500)
```



```
print(sent_1500)
print("="*50)
```

Wow. So far, two two-star reviews. One obviously had no idea what they were ordering; the other was
=====

```
In [19]: #remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
sent_0 = re.sub("\S*\d\S*", "", sent_0).strip()
print(sent_0)
```

Why is this \$[...] when the same product is available for \$[...] here?
 />
The Victor

```
In [20]: #remove spacial character: https://stackoverflow.com/a/5843547/4084039
sent_1500 = re.sub('[^A-Za-z0-9]+', ' ', sent_1500)
print(sent_1500)
```

Wow So far two two star reviews One obviously had no idea what they were ordering the other was

```
In [21]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
# instead of <br /> if we have <br/> these tags would have reumoved in the 1st step

stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves',
                'you'll', 'you'd', 'your', 'yours', 'yourself', 'yourselves', 'he', 'him',
                'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself',
                'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', 'they',
                'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had',
                'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as', 'until',
                'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through',
                'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'over',
                'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any',
                'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too', 'very',
                's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'do',
                'do', 've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't",
                "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mightn',
                "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't",
                'won', "won't", 'wouldn', "wouldn't"])
```

```
In [22]: # Combining all the above stundents
from tqdm import tqdm
preprocessed_reviews = []
# tqdm is for printing the status bar
for sentence in tqdm(final['Text'].values):
    sentence = re.sub(r"http\S+", "", sentence)
```

```

sentence = BeautifulSoup(sentence, 'lxml').get_text()
sentence = decontracted(sentence)
sentence = re.sub("\S*\d\S*", "", sentence).strip()
sentence = re.sub('[^A-Za-z]+', ' ', sentence)
# https://gist.github.com/sebleier/554280
sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwords)
preprocessed_reviews.append(sentence.strip())

```

100%|| 4986/4986 [00:01<00:00, 3137.37it/s]

In [23]: preprocessed_reviews[1500]

Out[23]: 'wow far two two star reviews one obviously no idea ordering wants crispy cookies hey'

[3.2] Preprocessing Review Summary

In [6]: *## Similarly you can do preprocessing for review summary also.*

5 [4] Featurization

5.1 [4.1] BAG OF WORDS

In [25]: *#BoW*

```

count_vect = CountVectorizer() #in scikit-learn
count_vect.fit(preprocessed_reviews)
print("some feature names ", count_vect.get_feature_names()[:10])
print('='*50)

final_counts = count_vect.transform(preprocessed_reviews)
print("the type of count vectorizer ",type(final_counts))
print("the shape of out text BOW vectorizer ",final_counts.get_shape())
print("the number of unique words ", final_counts.get_shape()[1])

```

```

some feature names  ['aa', 'aahhhs', 'aback', 'abandon', 'abates', 'abbott', 'abby', 'abdomina']
=====
the type of count vectorizer  <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer  (4986, 12997)
the number of unique words  12997

```

5.2 [4.2] Bi-Grams and n-Grams.

In [26]: *#bi-gram, tri-gram and n-gram*

```

#removing stop words like "not" should be avoided before building n-grams
# count_vect = CountVectorizer(ngram_range=(1,2))
# please do read the CountVectorizer documentation http://scikit-learn.org/stable/mod

```

```

# you can choose these numebrs min_df=10, max_features=5000, of your choice
count_vect = CountVectorizer(ngram_range=(1,2), min_df=10, max_features=5000)
final_bigram_counts = count_vect.fit_transform(preprocessed_reviews)
print("the type of count vectorizer ",type(final_bigram_counts))
print("the shape of out text BOW vectorizer ",final_bigram_counts.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_bigram_counts.get_shape()[0])

```

```

the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (4986, 3144)
the number of unique words including both unigrams and bigrams 3144

```

5.3 [4.3] TF-IDF

```

In [27]: tf_idf_vect = TfidfVectorizer(ngram_range=(1,2), min_df=10)
tf_idf_vect.fit(preprocessed_reviews)
print("some sample features(unique words in the corpus)",tf_idf_vect.get_feature_names())
print('='*50)

```

```

final_tf_idf = tf_idf_vect.transform(preprocessed_reviews)
print("the type of count vectorizer ",type(final_tf_idf))
print("the shape of out text TFIDF vectorizer ",final_tf_idf.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_tf_idf.get_shape()[0])

```

```

some sample features(unique words in the corpus) ['ability', 'able', 'able find', 'able get',
=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (4986, 3144)
the number of unique words including both unigrams and bigrams 3144

```

5.4 [4.4] Word2Vec

```

In [28]: # Train your own Word2Vec model using your own text corpus
i=0
list_of_sentence=[]
for sentence in preprocessed_reviews:
    list_of_sentence.append(sentence.split())

```

```

In [42]: # Using Google News Word2Vectors

```

```

# in this project we are using a pretrained model by google
# its 3.3G file, once you load this into your memory
# it occupies ~9Gb, so please do this step only if you have >12G of ram
# we will provide a pickle file wich contains a dict ,
# and it contains all our courpus words as keys and model[word] as values
# To use this code-snippet, download "GoogleNews-vectors-negative300.bin"
# from https://drive.google.com/file/d/0B7XkCwpI5KDYNlNUTTlSS21pQmM/edit
# it's 1.9GB in size.

```

```

# http://kavita-ganesan.com/gensim-word2vec-tutorial-starter-code/#.W17SRFazZPY
# you can comment this whole cell
# or change these variable according to your need

is_your_ram_gt_16g=False
want_to_use_google_w2v = False
want_to_train_w2v = True

if want_to_train_w2v:
    # min_count = 5 considers only words that occurred atleast 5 times
    w2v_model=Word2Vec(list_of_sentence,min_count=5,size=50, workers=4)
    print(w2v_model.wv.most_similar('great'))
    print('='*50)
    print(w2v_model.wv.most_similar('worst'))

elif want_to_use_google_w2v and is_your_ram_gt_16g:
    if os.path.isfile('GoogleNews-vectors-negative300.bin'):
        w2v_model=KeyedVectors.load_word2vec_format('GoogleNews-vectors-negative300.b
        print(w2v_model.wv.most_similar('great'))
        print(w2v_model.wv.most_similar('worst'))
    else:
        print("you don't have gogole's word2vec file, keep want_to_train_w2v = True,

[('snack', 0.9951335191726685), ('calorie', 0.9946465492248535), ('wonderful', 0.9946032166481
=====
[('varieties', 0.9994194507598877), ('become', 0.9992934465408325), ('popcorn', 0.999275088310

In [36]: w2v_words = list(w2v_model.wv.vocab)
print("number of words that occurred minimum 5 times ",len(w2v_words))
print("sample words ", w2v_words[0:50])

number of words that occurred minimum 5 times 3817
sample words ['product', 'available', 'course', 'total', 'pretty', 'stinky', 'right', 'nearby

```

5.5 [4.4.1] Converting text into vectors using Avg W2V, TFIDF-W2V

[4.4.1.1] Avg W2v

```

In [38]: # average Word2Vec
# compute average word2vec for each review.
sent_vectors = []; # the avg-w2v for each sentence/review is stored in this list
for sent in tqdm(list_of_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need t
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence

```

```

        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors.append(sent_vec)
print(len(sent_vectors))
print(len(sent_vectors[0]))

```

100%| 4986/4986 [00:03<00:00, 1330.47it/s]

4986

50

[4.4.1.2] TFIDF weighted W2v

```

In [39]: # S = ["abc def pqr", "def def def abc", "pqr pqr def"]
model = TfidfVectorizer()
tf_idf_matrix = model.fit_transform(preprocessed_reviews)
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

In [41]: # TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this l
row=0;
for sent in tqdm(list_of_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            # tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors.append(sent_vec)
    row += 1

```

6 [5] Assignment 9: Random Forests

Apply Random Forests & GBDT on these feature sets

SET 1:Review text, preprocessed one converted into vectors using (BOW)

SET 2:Review text, preprocessed one converted into vectors using (TFIDF)

SET 3:Review text, preprocessed one converted into vectors using (AVG W2v)

SET 4:Review text, preprocessed one converted into vectors using (TFIDF W2v)

The hyper paramter tuning (Consider any two hyper parameters)

Find the best hyper parameter which will give the maximum AUC value

Find the best hyper paramter using k-fold cross validation or simple cross validation data

Use gridsearch cv or randomsearch cv or you can also write your own for loops to do this task of hyperparameter tuning

```
</ul>
</li>
<br>
<li><strong>Feature importance</strong>
  <ul>
<li>Get top 20 important features and represent them in a word cloud. Do this for BOW & TFIDF.
  </ul>
</li>
<br>
<li><strong>Feature engineering</strong>
  <ul>
<li>To increase the performance of your model, you can also experiment with with feature engineering
    <ul>
      <li>Taking length of reviews as another feature.</li>
      <li>Considering some features from review summary as well.</li>
    </ul>
  </ul>
</li>
<br>
<li><strong>Representation of results</strong>
  <ul>
<li>You need to plot the performance of model both on train data and cross validation data for
    <img src='3d_plot.JPG' width=500px> with X-axis as <strong>n_estimators</strong>, Y-axis as <strong>auc</strong>
    <p style="text-align:center;font-size:30px;color:red;"><strong>(or)</strong></p> <br>
<li>You need to plot the performance of model both on train data and cross validation data for
    <img src='heat_map.JPG' width=300px> <a href='https://seaborn.pydata.org/generated/seaborn.heatmap.html'>seaborn.heatmap</a>
<li>You choose either of the plotting techniques out of 3d plot or heat map</li>
<li>Once after you found the best hyper parameter, you need to train your model with it, and find the best hyper parameter
    <img src='train_test_auc.JPG' width=300px></li>
<li>Along with plotting ROC curve, you need to print the <a href='https://www.appliedaiculture.com/roc-curve/'>ROC curve</a>
    <img src='confusion_matrix.png' width=300px></li>
```

```

    </ul>
</li>
<br>
<li><strong>Conclusion</strong>
    <ul>
<li>You need to summarize the results at the end of the notebook, summarize it in the table for
    <img src='summary.JPG' width=400px>
</li>
    </ul>

```

Note: Data Leakage

1. There will be an issue of data-leakage if you vectorize the entire data and then split it into train/cv/test.
2. To avoid the issue of data-leakage, make sure to split your data first and then vectorize it.
3. While vectorizing your data, apply the method `fit_transform()` on your train data, and apply the method `transform()` on cv/test data.
4. For more details please go through this link.

6.1 [5.1] Applying RF

```

In [6]: # after preprocessing
df = pd.read_pickle("files/preprocessed.pkl")
df.shape

```

Out[6]: (364171, 12)

```

In [7]: df.head()

```

```

Out[7]:
      Id  ProductId  UserId  ProfileName \
138706  150524  0006641040  ACITT7DI6IDDL  shari zychinski
138688  150506  0006641040  A2IW4PEEK02R0U  Tracy
138689  150507  0006641040  A1S4A3IQ2MU7V4  sally sue "sally sue"
138690  150508  0006641040  AZGXZ2UUK6X  Catherine Hallberg "(Kate)"
138691  150509  0006641040  A3CMRKGEOP909G  Teresa

```

```

      HelpfulnessNumerator  HelpfulnessDenominator  Score  Time \
138706                    0                      0      1  939340800
138688                    1                      1      1  1194739200
138689                    1                      1      1  1191456000
138690                    1                      1      1  1076025600
138691                    3                      4      1  1018396800

```

```

      Summary \
138706  EVERY book is educational
138688  Love the book, miss the hard cover version
138689  chicken soup with rice months
138690  a good swingy rhythm for reading aloud
138691  A great way to learn the months

```

```

Text \
138706 this witty little book makes my son laugh at l...
138688 I grew up reading these Sendak books, and watc...
138689 This is a fun way for children to learn their ...
138690 This is a great little book to read aloud- it ...
138691 This is a book of poetry about the months of t...

```

```

CleanedText \
138706 witty little book makes son laugh loud recite ...
138688 grew reading sendak books watching really rosi...
138689 fun way children learn months year learn poems...
138690 great little book read aloud nice rhythm well ...
138691 book poetry months year goes month cute little...

```

```

CleanedSummary
138706 every book educational
138688 love book miss hard cover version
138689 chicken soup rice months
138690 good swingy rhythm reading aloud
138691 great way learn months

```

```

In [8]: from sklearn.model_selection import train_test_split
        from sklearn.grid_search import GridSearchCV
        from sklearn.datasets import *
        from sklearn.metrics import accuracy_score , f1_score , confusion_matrix
        from collections import Counter
        from sklearn.metrics import accuracy_score, roc_auc_score , roc_curve
        from sklearn.model_selection import train_test_split

```

```

In [9]: # take 50k sample data randomly
        sample_data = df.sample(100000)
        sample_data.shape

```

```

Out[9]: (100000, 12)

```

```

In [12]: # sorted the data using time based
         sorted_data = sample_data.sort_values('Time', axis=0, inplace=False)
         sorted_data.shape

```

```

Out[12]: (100000, 12)

```

```

In [13]: sorted_data['Score'].value_counts()

```

```

Out[13]: 1    84258
         0    15742
         Name: Score, dtype: int64

```

```

In [14]: X = np.array(sorted_data['CleanedText'])
         y = np.array(sorted_data['Score'])

```



```

print(X.shape)
print(y.shape)

(100000,)
(100000,)

In [15]: # Simple cross validation
# split the data sent into train and test
train , test , train_y , test_y = train_test_split(X, y, test_size = 0.3, random_state=42)

# split the train data set into cross validation train and cross validation test
train, cv , train_y, cv_y = train_test_split(train, train_y, test_size=0.3, random_state=42)

print("train data = ", train.shape)
print("cross validation = ", cv.shape)
print("test data = ", test.shape)

train data = (49000,)
cross validation = (21000,)
test data = (30000,)

```

In []:

6.1.1 [5.1.1] Applying Random Forests on BOW, SET 1

```

In [11]: # Please write all the code with proper documentation
#BoW
count_vect = CountVectorizer(min_df=20) #in scikit-learn
count_vect.fit(train)
print("some feature names ", count_vect.get_feature_names()[:10])
print('='*50)

bow_train = count_vect.fit_transform(train)
bow_cv = count_vect.transform(cv)
bow_test = count_vect.transform(test)
print("=====Train Data=====")
print("the type of count vectorizer ",type(bow_train))
print("the shape of out text BOW vectorizer ",bow_train.get_shape())
print("the number of unique words ", bow_train.get_shape()[1])
print("=====Cross validation Data=====")
print("the type of count vectorizer ",type(bow_cv))
print("the shape of out text BOW vectorizer ",bow_cv.get_shape())
print("the number of unique words ", bow_cv.get_shape()[1])
print("=====Test Data=====")
print("the type of count vectorizer ",type(bow_test))
print("the shape of out text BOW vectorizer ",bow_test.get_shape())
print("the number of unique words ", bow_test.get_shape()[1])

```

```

some feature names ['ability', 'able', 'absolute', 'absolutely', 'absorb', 'absorbed', 'absor']
=====
=====Train Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (49000, 5927)
the number of unique words 5927
=====Cross validation Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (21000, 5927)
the number of unique words 5927
=====Test Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (30000, 5927)
the number of unique words 5927

```

```
In [16]: from sklearn.ensemble import RandomForestClassifier
```

```
In [15]: n_estimators = [20,40,60,80,100,120]
max_depth = [1,5,10,100,500,1000]
bow_train_auc = []
bow_cv_auc = []
for i in n_estimators:
    for j in max_depth:
        RFDTTC = RandomForestClassifier(n_estimators=i,criterion='gini', max_depth=j)
        RFDTTC.fit(bow_train, train_y)
        # train data
        y_prob_train = RFDTTC.predict_proba(bow_train)[: ,1]
        y_pred = np.where(y_prob_train > 0.5, 1, 0)
        auc_roc_train = roc_auc_score(train_y , y_prob_train)
        print('\nTrain AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (str(j), str(i), auc_roc_train))
        bow_train_auc.append(auc_roc_train)
        # CV
        y_prob_cv = RFDTTC.predict_proba(bow_cv)[: ,1]
        y_pred = np.where(y_prob_cv > 0.5, 1, 0)
        auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
        print('\nCV AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (str(j), str(i), auc_roc_cv))
        bow_cv_auc.append(auc_roc_cv)
        print("="*50)

```

```
Train AUC for max_depth = 1 and n_estimators = 20 is 75.46%
```

```
CV AUC for max_depth = 1 and n_estimators = 20 is 74.33%
```

```
=====
```

```
Train AUC for max_depth = 5 and n_estimators = 20 is 82.62%
```

CV AUC for max_depth = 5 and n_estimators = 20 is 80.42%

=====

Train AUC for max_depth = 10 and n_estimators = 20 is 86.79%

CV AUC for max_depth = 10 and n_estimators = 20 is 84.33%

=====

Train AUC for max_depth = 100 and n_estimators = 20 is 99.77%

CV AUC for max_depth = 100 and n_estimators = 20 is 87.74%

=====

Train AUC for max_depth = 500 and n_estimators = 20 is 99.97%

CV AUC for max_depth = 500 and n_estimators = 20 is 87.50%

=====

Train AUC for max_depth = 1000 and n_estimators = 20 is 99.97%

CV AUC for max_depth = 1000 and n_estimators = 20 is 87.76%

=====

Train AUC for max_depth = 1 and n_estimators = 40 is 70.86%

CV AUC for max_depth = 1 and n_estimators = 40 is 69.53%

=====

Train AUC for max_depth = 5 and n_estimators = 40 is 86.12%

CV AUC for max_depth = 5 and n_estimators = 40 is 84.57%

=====

Train AUC for max_depth = 10 and n_estimators = 40 is 88.74%

CV AUC for max_depth = 10 and n_estimators = 40 is 85.68%

=====

Train AUC for max_depth = 100 and n_estimators = 40 is 99.87%

CV AUC for max_depth = 100 and n_estimators = 40 is 88.92%

=====

Train AUC for max_depth = 500 and n_estimators = 40 is 99.98%

CV AUC for max_depth = 500 and n_estimators = 40 is 89.35%

=====

Train AUC for max_depth = 1000 and n_estimators = 40 is 99.98%

CV AUC for max_depth = 1000 and n_estimators = 40 is 89.35%

=====

Train AUC for max_depth = 1 and n_estimators = 60 is 80.86%

CV AUC for max_depth = 1 and n_estimators = 60 is 80.08%

=====

Train AUC for max_depth = 5 and n_estimators = 60 is 85.07%

CV AUC for max_depth = 5 and n_estimators = 60 is 83.09%

=====

Train AUC for max_depth = 10 and n_estimators = 60 is 90.12%

CV AUC for max_depth = 10 and n_estimators = 60 is 87.25%

=====

Train AUC for max_depth = 100 and n_estimators = 60 is 99.91%

CV AUC for max_depth = 100 and n_estimators = 60 is 89.70%

=====

Train AUC for max_depth = 500 and n_estimators = 60 is 99.98%

CV AUC for max_depth = 500 and n_estimators = 60 is 89.78%

=====

Train AUC for max_depth = 1000 and n_estimators = 60 is 99.98%

CV AUC for max_depth = 1000 and n_estimators = 60 is 89.72%

=====

Train AUC for max_depth = 1 and n_estimators = 80 is 78.81%

CV AUC for max_depth = 1 and n_estimators = 80 is 77.92%

=====

Train AUC for max_depth = 5 and n_estimators = 80 is 88.71%

CV AUC for max_depth = 5 and n_estimators = 80 is 86.62%

=====

Train AUC for max_depth = 10 and n_estimators = 80 is 90.31%

CV AUC for max_depth = 10 and n_estimators = 80 is 87.44%

```

=====

Train AUC for max_depth = 100 and n_estimators = 80 is 99.91%

CV AUC for max_depth = 100 and n_estimators = 80 is 89.71%
=====

Train AUC for max_depth = 500 and n_estimators = 80 is 99.98%

CV AUC for max_depth = 500 and n_estimators = 80 is 90.03%
=====

Train AUC for max_depth = 1000 and n_estimators = 80 is 99.98%

CV AUC for max_depth = 1000 and n_estimators = 80 is 89.88%
=====

Train AUC for max_depth = 1 and n_estimators = 100 is 81.72%

CV AUC for max_depth = 1 and n_estimators = 100 is 80.26%
=====

Train AUC for max_depth = 5 and n_estimators = 100 is 88.07%

CV AUC for max_depth = 5 and n_estimators = 100 is 86.13%
=====

Train AUC for max_depth = 10 and n_estimators = 100 is 90.19%

CV AUC for max_depth = 10 and n_estimators = 100 is 87.12%
=====

Train AUC for max_depth = 100 and n_estimators = 100 is 99.91%

CV AUC for max_depth = 100 and n_estimators = 100 is 90.21%
=====

Train AUC for max_depth = 500 and n_estimators = 100 is 99.98%

CV AUC for max_depth = 500 and n_estimators = 100 is 90.00%
=====

Train AUC for max_depth = 1000 and n_estimators = 100 is 99.98%

CV AUC for max_depth = 1000 and n_estimators = 100 is 90.00%
=====

Train AUC for max_depth = 1 and n_estimators = 120 is 82.28%

```

CV AUC for max_depth = 1 and n_estimators = 120 is 81.19%

=====

Train AUC for max_depth = 5 and n_estimators = 120 is 87.88%

CV AUC for max_depth = 5 and n_estimators = 120 is 86.15%

=====

Train AUC for max_depth = 10 and n_estimators = 120 is 90.83%

CV AUC for max_depth = 10 and n_estimators = 120 is 87.93%

=====

Train AUC for max_depth = 100 and n_estimators = 120 is 99.92%

CV AUC for max_depth = 100 and n_estimators = 120 is 90.10%

=====

Train AUC for max_depth = 500 and n_estimators = 120 is 99.98%

CV AUC for max_depth = 500 and n_estimators = 120 is 90.23%

=====


Train AUC for max_depth = 1000 and n_estimators = 120 is 99.98%

CV AUC for max_depth = 1000 and n_estimators = 120 is 90.20%


=====

```
In [62]: fig, ax = plt.subplots()
         # auc on cv
         print("====CV Data====")
         cv_scores = np.array(bow_cv_auc).reshape(len(n_estimators),len(max_depth))
         df_cm_cv = pd.DataFrame(cv_scores, n_estimators, max_depth)
         sns.set(font_scale=1.4)
         ax = sns.heatmap(df_cm_cv, annot=True, annot_kws={"size": 10}, fmt='.3g',linewidths=.1)
         ax.set_xlabel("max_depth")
         ax.set_ylabel("n_estimators")
         plt.show()
         print("====Train Data====")
         train_scores = np.array(bow_train_auc).reshape(len(n_estimators),len(max_depth))
         df_cm_train = pd.DataFrame(train_scores, n_estimators, max_depth)
         sns.set(font_scale=1.4)
         ax = sns.heatmap(df_cm_train, annot=True, annot_kws={"size": 10}, fmt='.3g',linewidths=.1)
         ax.set_xlabel("max_depth")
         ax.set_ylabel("n_estimators")
         plt.show()
```

=====CV Data=====



=====Train Data=====



```
In [65]: # depth is 500
# https://www.dataquest.io/blog/learning-curves-machine-learning/

import matplotlib.pyplot as plt
%matplotlib inline


plt.style.use('seaborn')

plt.plot(n_estimators, train_scores[:, -2], label = 'Training AUC')
plt.plot(n_estimators, cv_scores[:, -2], label = 'Validation AUC')

plt.ylabel('AUC', fontsize = 14)
plt.xlabel('depth', fontsize = 14)
plt.title('Learning curves for a Decision trees model', fontsize = 18, y = 1.03)
plt.legend()

Out[65]: <matplotlib.legend.Legend at 0x7f9570918668>
```


Learning curves for a Decision trees model



```
In [66]: i = 120
         j = 500
```

```
RFDTC = RandomForestClassifier(criterion='gini' , max_depth=j, n_estimators=i)
RFDTC.fit(bow_train, train_y)
# train data
y_prob_train = RFDTC.predict_proba(bow_train)[: ,1]
fprt, tprt, thresholdt = roc_curve(train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y , y_prob_train)
print('\nTrain AUC for max depth = %s and n_estimators = %s is %0.2f%%' % (str(j),str(i),auc_roc_train))
# CV
y_prob_cv = RFDTC.predict_proba(bow_cv)[: ,1]
fprc, tprc, thresholdc = roc_curve(cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
print('\nCV AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (str(j),str(i),auc_roc_cv))
# Test
y_prob_test = RFDTC.predict_proba(bow_test)[: ,1]
fprts, tprts, thresholdts = roc_curve(test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
```

```

auc_roc_test = roc_auc_score(test_y , y_prob_test)
print('\nTest AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (str(j),str(

```

Train AUC for max depth = 500 and n_estimators = 120 is 99.98%

CV AUC for max_depth = 500 and n_estimators = 120 is 90.33%


Test AUC for max_depth = 500 and n_estimators = 120 is 90.90%

In [67]: # <https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python>

```

import matplotlib.pyplot as plt
plt.clf()
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b' , label='AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y' , label='AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label='AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

```




```
In [68]: print("F1-Score on test set: %0.2f"%(f1_score(test_y, y_pred_test)))
```

F1-Score on test set: 0.94

```
In [69]: df_cm = pd.DataFrame(confusion_matrix(test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[69]: <matplotlib.axes._subplots.AxesSubplot at 0x7f95705407b8>
```



6.1.2 [5.1.2] Wordcloud of top 20 important features from SET 1

```
In [73]: # Please write all the code with proper documentation
feature_imp = RFDTF.feature_importances_
feature_names = count_vect.get_feature_names()
features = dict(zip(feature_imp, feature_names))
features_df = pd.DataFrame.from_dict(features, orient='index')
features_df.columns = ["Words"]
sorted_features = features_df.sort_index(axis=0, ascending=False)
```

```
In [76]: from wordcloud import WordCloud, STOPWORDS

In [102]: # https://www.kaggle.com/adiljadoon/word-cloud-with-python
# top 20 important features
top_20 = sorted_features.head(20)
stopwords = set(STOPWORDS) - set(top_20["Words"])
wordcloud = WordCloud(background_color='black',stopwords=stopwords,max_words=20,max_
print(wordcloud)
fig = plt.figure()
plt.imshow(wordcloud)
plt.axis('off')
plt.show()
```

<wordcloud.wordcloud.WordCloud object at 0x7f95851a59b0>



6.1.3 [5.1.3] Applying Random Forests on TFIDF, SET 2

```
In [104]: # Please write all the code with proper documentation

model = TfidfVectorizer(min_df=20, ngram_range=(1,2))
#tf_idf_matrix = model.fit_transform(train)

print("=====Train Data=====")
tf_idf_train = model.fit_transform(train)
print("the type of count vectorizer ",type(tf_idf_train))
print("the shape of out text TFIDF vectorizer ",tf_idf_train.get_shape())
print("the number of unique words including both unigrams and bigrams ",tf_idf_train
```

```

print("=====CV Data=====")
tf_idf_cv = model.transform(cv)
print("the type of count vectorizer ",type(tf_idf_cv))
print("the shape of out text TFIDF vectorizer ",tf_idf_cv.get_shape())
print("the number of unique words including both unigrams and bigrams ",tf_idf_cv.get_shape()[0])
print("=====Test Data=====")
tf_idf_test = model.transform(test)
print("the type of count vectorizer ",type(tf_idf_test))
print("the shape of out text TFIDF vectorizer ",tf_idf_test.get_shape())
print("the number of unique words including both unigrams and bigrams ", tf_idf_test.get_shape()[0])

# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

```

```

=====Train Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (49000, 13653)
the number of unique words including both unigrams and bigrams 13653
=====CV Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (21000, 13653)
the number of unique words including both unigrams and bigrams 13653
=====Test Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (30000, 13653)
the number of unique words including both unigrams and bigrams 13653

```

```

In [105]: n_estimators = [20,40,60,80,100,120]
max_depth = [1,5,10,100,500,1000]
tfidf_train_auc = []
tfidf_cv_auc = []
for i in n_estimators:
    for j in max_depth:
        RFDTTC = RandomForestClassifier(n_estimators=i,criterion='gini', max_depth=j)
        RFDTTC.fit(tf_idf_train, train_y)
        # train data
        y_prob_train = RFDTTC.predict_proba(tf_idf_train)[: ,1]
        y_pred = np.where(y_prob_train > 0.5, 1, 0)
        auc_roc_train = roc_auc_score(train_y , y_prob_train)
        print('\nTrain AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (str(max_depth), str(n_estimators), auc_roc_train))
        tfidf_train_auc.append(auc_roc_train)
        # CV
        y_prob_cv = RFDTTC.predict_proba(tf_idf_cv)[: ,1]
        y_pred = np.where(y_prob_cv > 0.5, 1, 0)
        auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
        print('\nCV AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (str(max_depth), str(n_estimators), auc_roc_cv))
        tfidf_cv_auc.append(auc_roc_cv)
    print("="*50)

```

Train AUC for max_depth = 1 and n_estimators = 20 is 66.60%

CV AUC for max_depth = 1 and n_estimators = 20 is 65.91%

=====

Train AUC for max_depth = 5 and n_estimators = 20 is 83.63%

CV AUC for max_depth = 5 and n_estimators = 20 is 81.05%

=====

Train AUC for max_depth = 10 and n_estimators = 20 is 87.89%

CV AUC for max_depth = 10 and n_estimators = 20 is 85.10%

=====

Train AUC for max_depth = 100 and n_estimators = 20 is 99.79%

CV AUC for max_depth = 100 and n_estimators = 20 is 89.63%

=====

Train AUC for max_depth = 500 and n_estimators = 20 is 99.94%

CV AUC for max_depth = 500 and n_estimators = 20 is 89.85%

=====

Train AUC for max_depth = 1000 and n_estimators = 20 is 99.98%

CV AUC for max_depth = 1000 and n_estimators = 20 is 89.95%

=====

Train AUC for max_depth = 1 and n_estimators = 40 is 74.33%

CV AUC for max_depth = 1 and n_estimators = 40 is 73.54%

=====

Train AUC for max_depth = 5 and n_estimators = 40 is 87.63%

CV AUC for max_depth = 5 and n_estimators = 40 is 85.07%

=====

Train AUC for max_depth = 10 and n_estimators = 40 is 90.58%

CV AUC for max_depth = 10 and n_estimators = 40 is 87.25%

=====

Train AUC for max_depth = 100 and n_estimators = 40 is 99.86%

CV AUC for max_depth = 100 and n_estimators = 40 is 91.02%

=====

Train AUC for max_depth = 500 and n_estimators = 40 is 99.95%

CV AUC for max_depth = 500 and n_estimators = 40 is 91.36%

=====

Train AUC for max_depth = 1000 and n_estimators = 40 is 99.98%

CV AUC for max_depth = 1000 and n_estimators = 40 is 91.60%

=====

Train AUC for max_depth = 1 and n_estimators = 60 is 78.62%

CV AUC for max_depth = 1 and n_estimators = 60 is 76.60%

=====

Train AUC for max_depth = 5 and n_estimators = 60 is 88.85%

CV AUC for max_depth = 5 and n_estimators = 60 is 86.23%

=====

Train AUC for max_depth = 10 and n_estimators = 60 is 91.26%

CV AUC for max_depth = 10 and n_estimators = 60 is 87.62%

=====

Train AUC for max_depth = 100 and n_estimators = 60 is 99.90%

CV AUC for max_depth = 100 and n_estimators = 60 is 91.67%

=====

Train AUC for max_depth = 500 and n_estimators = 60 is 99.97%

CV AUC for max_depth = 500 and n_estimators = 60 is 92.21%

=====

Train AUC for max_depth = 1000 and n_estimators = 60 is 99.99%

CV AUC for max_depth = 1000 and n_estimators = 60 is 92.23%

=====

Train AUC for max_depth = 1 and n_estimators = 80 is 83.59%

CV AUC for max_depth = 1 and n_estimators = 80 is 82.11%

=====

Train AUC for max_depth = 5 and n_estimators = 80 is 89.71%

CV AUC for max_depth = 5 and n_estimators = 80 is 87.24%

=====

Train AUC for max_depth = 10 and n_estimators = 80 is 91.42%

CV AUC for max_depth = 10 and n_estimators = 80 is 87.81%

=====

Train AUC for max_depth = 100 and n_estimators = 80 is 99.91%

CV AUC for max_depth = 100 and n_estimators = 80 is 91.89%

=====

Train AUC for max_depth = 500 and n_estimators = 80 is 99.96%

CV AUC for max_depth = 500 and n_estimators = 80 is 92.37%

=====

Train AUC for max_depth = 1000 and n_estimators = 80 is 99.99%

CV AUC for max_depth = 1000 and n_estimators = 80 is 92.70%

=====

Train AUC for max_depth = 1 and n_estimators = 100 is 82.75%

CV AUC for max_depth = 1 and n_estimators = 100 is 81.13%

=====

Train AUC for max_depth = 5 and n_estimators = 100 is 89.91%

CV AUC for max_depth = 5 and n_estimators = 100 is 87.92%

=====

Train AUC for max_depth = 10 and n_estimators = 100 is 91.92%

CV AUC for max_depth = 10 and n_estimators = 100 is 88.64%

=====

Train AUC for max_depth = 100 and n_estimators = 100 is 99.91%

CV AUC for max_depth = 100 and n_estimators = 100 is 91.96%

=====

Train AUC for max_depth = 500 and n_estimators = 100 is 99.97%

CV AUC for max_depth = 500 and n_estimators = 100 is 92.69%


```

=====

Train AUC for max_depth = 1000 and n_estimators = 100 is 99.99%

CV AUC for max_depth = 1000 and n_estimators = 100 is 92.59%
=====

Train AUC for max_depth = 1 and n_estimators = 120 is 83.56%

CV AUC for max_depth = 1 and n_estimators = 120 is 81.83%
=====

Train AUC for max_depth = 5 and n_estimators = 120 is 89.63%

CV AUC for max_depth = 5 and n_estimators = 120 is 87.45%
=====

Train AUC for max_depth = 10 and n_estimators = 120 is 92.08%

CV AUC for max_depth = 10 and n_estimators = 120 is 88.71%
=====

Train AUC for max_depth = 100 and n_estimators = 120 is 99.91%

CV AUC for max_depth = 100 and n_estimators = 120 is 92.17%
=====

Train AUC for max_depth = 500 and n_estimators = 120 is 99.96%

CV AUC for max_depth = 500 and n_estimators = 120 is 92.58%
=====

Train AUC for max_depth = 1000 and n_estimators = 120 is 99.99%

CV AUC for max_depth = 1000 and n_estimators = 120 is 92.73%
=====

```

```

In [106]: fig, ax = plt.subplots()
          # auc on cv
          print("=====CV Data=====")
          cv_scores = np.array(tfidf_cv_auc).reshape(len(n_estimators),len(max_depth))
          df_cm_cv = pd.DataFrame(cv_scores, n_estimators, max_depth)
          sns.set(font_scale=1.4)
          ax = sns.heatmap(df_cm_cv, annot=True, annot_kws={"size": 10}, fmt='.3g',linewidths=
          ax.set_xlabel("max_depth")
          ax.set_ylabel("n_estimators")
          plt.show()

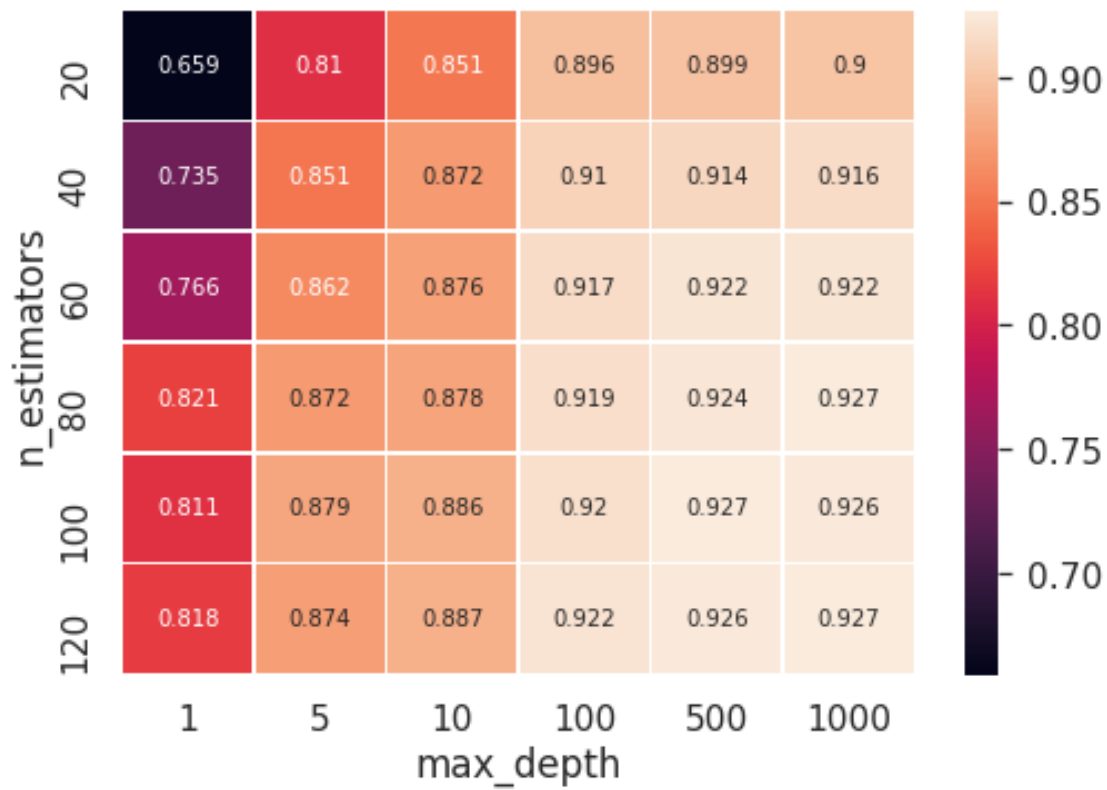
```

```


print("=====Train Data=====")
train_scores = np.array(tfidf_train_auc).reshape(len(n_estimators),len(max_depth))
df_cm_train = pd.DataFrame(train_scores, n_estimators, max_depth)
sns.set(font_scale=1.4)
ax = sns.heatmap(df_cm_train, annot=True, annot_kws={"size": 10}, fmt='%.3g',linewidtht
ax.set_xlabel("max_depth")
ax.set_ylabel("n_estimators")
plt.show()

```

=====**CV Data**=====



=====**Train Data**=====



```
In [107]: # depth is 1000
# https://www.dataquest.io/blog/learning-curves-machine-learning/

import matplotlib.pyplot as plt
%matplotlib inline


plt.style.use('seaborn')

plt.plot(n_estimators, train_scores[:, -1], label = 'Training AUC')
plt.plot(n_estimators, cv_scores[:, -1], label = 'Validation AUC')

plt.ylabel('AUC', fontsize = 14)
plt.xlabel('depth', fontsize = 14)
plt.title('Learning curves for a Decision trees model', fontsize = 18, y = 1.03)
plt.legend()

Out[107]: <matplotlib.legend.Legend at 0x7f956f7525c0>
```

Learning curves for a Decision trees model



```
In [108]: i = 120
          j = 1000
```

```
RFDTTC = RandomForestClassifier(criterion='gini' , max_depth=j, n_estimators=i)
RFDTTC.fit(tf_idf_train, train_y)
# train data
y_prob_train = RFDTTC.predict_proba(tf_idf_train)[: ,1]
fprt, tprt, thresholdt = roc_curve(train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y , y_prob_train)
print('\nTrain AUC for max depth = %s and n_estimators = %s is %0.2f%%' % (str(j),str(i),auc_roc_train))
# CV
y_prob_cv = RFDTTC.predict_proba(tf_idf_cv)[: ,1]
fprc, tprc, thresholdc = roc_curve(cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
print('\nCV AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (str(j),str(i),auc_roc_cv))
# Test
y_prob_test = RFDTTC.predict_proba(tf_idf_test)[: ,1]
fprts, tprts, thresholdts = roc_curve(test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
```

```

auc_roc_test = roc_auc_score(test_y , y_prob_test)
print('\nTest AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (str(j),str

```

Train AUC for max depth = 1000 and n_estimators = 120 is 99.99%

CV AUC for max_depth = 1000 and n_estimators = 120 is 92.75%


Test AUC for max_depth = 1000 and n_estimators = 120 is 92.99%

In [109]: # <https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python>

```

import matplotlib.pyplot as plt
plt.clf()
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b', label='AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y', label='AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label='AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

```




```
In [110]: print("F1-Score on test set: %.2f"%(f1_score(test_y, y_pred_test)))
```

F1-Score on test set: 0.94

```
In [112]: df_cm = pd.DataFrame(confusion_matrix(test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[112]: <matplotlib.axes._subplots.AxesSubplot at 0x7f95708e57b8>
```




6.1.4 [5.1.4] Wordcloud of top 20 important features from SET 2

```
In [113]: # Please write all the code with proper documentation
feature_imp = RFDTF.feature_importances_
feature_names = model.get_feature_names()
features = dict(zip(feature_imp, feature_names))
features_df = pd.DataFrame.from_dict(features, orient='index')
features_df.columns = ["Words"]
sorted_features = features_df.sort_index(axis=0, ascending=False)
```

```
In [114]: # Please write all the code with proper documentation
# https://www.kaggle.com/adiljadoon/word-cloud-with-python
# top 20 important features
top_20 = sorted_features.head(20)
stopwords = set(STOPWORDS) - set(top_20["Words"])
wordcloud = WordCloud(background_color='black',stopwords=stopwords,max_words=20,max_
print(wordcloud)
fig = plt.figure()
plt.imshow(wordcloud)
plt.axis('off')
plt.show()
```

<wordcloud.wordcloud.WordCloud object at 0x7f956a203b38>



6.1.5 [5.1.5] Applying Random Forests on AVG W2V, SET 3

```
In [115]: # Please write all the code with proper documentation
# Train your own Word2Vec model using your own text corpus
##### Train Set #####
i=0
list_of_train_sentence=[]
for sentence in train:
    list_of_train_sentence.append(sentence.split())
##### CV Set #####
i=0
list_of_cv_sentence=[]
for sentence in cv:
```

```

        list_of_cv_sentence.append(sentence.split())
##### Test Set #####
i=0
list_of_test_sentence=[]
for sentence in test:
    list_of_test_sentence.append(sentence.split())
print("Length of Train = ", len(list_of_train_sentence))
print("Length of CV = ", len(list_of_cv_sentence))
print("Length of Test = ", len(list_of_test_sentence))

```

```

Length of Train = 49000
Length of CV = 21000
Length of Test = 30000

```

```

In [116]: w2v_model=Word2Vec(list_of_train_sentence,min_count=15,size=100, workers=4)
print(w2v_model.wv.most_similar('great'))
print('='*50)
print(w2v_model.wv.most_similar('worst'))

```

```

[('fantastic', 0.7984999418258667), ('wonderful', 0.7739488482475281), ('awesome', 0.7482377290
=====
[('greatest', 0.7162206172943115), ('tastiest', 0.7049369812011719), ('best', 0.70081377029418

```

```

In [117]: w2v_words = list(w2v_model.wv.vocab)
print("number of words that occurred minimum 5 times ",len(w2v_words))
print("sample words ", w2v_words[0:50])

```

```

number of words that occurred minimum 5 times 7524
sample words ['coffee', 'smells', 'great', 'first', 'thing', 'morning', 'occasionally', 'enjoy

```

```

In [118]: ##### Train data #####
# average Word2Vec
# compute average word2vec for each review.
sent_vectors_train = []; # the avg-w2v for each sentence/review is stored in this list
for sent in tqdm(list_of_train_sentence): # for each review/sentence
    sent_vec = np.zeros(100) # as word vectors are of zero length 50, you might need
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors_train.append(sent_vec)
print(len(sent_vectors_train))
print(len(sent_vectors_train[0]))

```


100%|| 49000/49000 [02:19<00:00, 351.43it/s]

49000

100

```
In [119]: ##### CV data #####
# average Word2Vec
# compute average word2vec for each review.
sent_vectors_cv = []; # the avg-w2v for each sentence/review is stored in this list
for sent in tqdm(list_of_cv_sentence): # for each review/sentence
    sent_vec = np.zeros(100) # as word vectors are of zero length 50, you might need
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors_cv.append(sent_vec)
print(len(sent_vectors_cv))
print(len(sent_vectors_cv[0]))
```

100%|| 21000/21000 [00:58<00:00, 357.99it/s]

21000

100

```
In [120]: ##### Test data #####
# average Word2Vec
# compute average word2vec for each review.
sent_vectors_test = []; # the avg-w2v for each sentence/review is stored in this list
for sent in tqdm(list_of_test_sentence): # for each review/sentence
    sent_vec = np.zeros(100) # as word vectors are of zero length 50, you might need
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
```

```

        sent_vec /= cnt_words
        sent_vectors_test.append(sent_vec)
    print(len(sent_vectors_test))
    print(len(sent_vectors_test[0]))

```

100%|| 30000/30000 [01:24<00:00, 353.24it/s]

30000

100

In [121]: *# save the datasets as numpy array*

```

w2v_train = np.array(sent_vectors_train)
w2v_cv = np.array(sent_vectors_cv)
w2v_test = np.array(sent_vectors_test)

```

In [122]: n_estimators = [20,40,60,80,100,120]

```

max_depth = [1,5,10,100,500,1000]

```

```

w2v_train_auc = []

```

```

w2v_cv_auc = []

```

```

for i in n_estimators:

```

```

    for j in max_depth:

```

```

        RFDTTC = RandomForestClassifier(n_estimators=i,criterion='gini', max_depth=j)

```

```

        RFDTTC.fit(w2v_train, train_y)

```

```

        # train data

```

```

        y_prob_train = RFDTTC.predict_proba(w2v_train)[:,-1]

```

```

        y_pred = np.where(y_prob_train > 0.5, 1, 0)

```

```

        auc_roc_train = roc_auc_score(train_y , y_prob_train)

```

```

        print('\nTrain AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (s

```

```

        w2v_train_auc.append(auc_roc_train)

```

```

        # CV

```

```

        y_prob_cv = RFDTTC.predict_proba(w2v_cv)[:,-1]

```

```

        y_pred = np.where(y_prob_cv > 0.5, 1, 0)

```

```

        auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)

```

```

        print('\nCV AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (str(

```

```

        w2v_cv_auc.append(auc_roc_cv)

```

```

        print("="*50)

```

Train AUC for max_depth = 1 and n_estimators = 20 is 82.38%

CV AUC for max_depth = 1 and n_estimators = 20 is 81.93%

=====

Train AUC for max_depth = 5 and n_estimators = 20 is 88.58%

CV AUC for max_depth = 5 and n_estimators = 20 is 86.83%

=====

Train AUC for max_depth = 10 and n_estimators = 20 is 96.68%

CV AUC for max_depth = 10 and n_estimators = 20 is 88.37%

=====

Train AUC for max_depth = 100 and n_estimators = 20 is 99.97%

CV AUC for max_depth = 100 and n_estimators = 20 is 86.66%

=====

Train AUC for max_depth = 500 and n_estimators = 20 is 99.97%

CV AUC for max_depth = 500 and n_estimators = 20 is 87.10%

=====

Train AUC for max_depth = 1000 and n_estimators = 20 is 99.97%

CV AUC for max_depth = 1000 and n_estimators = 20 is 87.54%

=====

Train AUC for max_depth = 1 and n_estimators = 40 is 83.84%

CV AUC for max_depth = 1 and n_estimators = 40 is 83.33%

=====

Train AUC for max_depth = 5 and n_estimators = 40 is 88.93%

CV AUC for max_depth = 5 and n_estimators = 40 is 87.36%

=====

Train AUC for max_depth = 10 and n_estimators = 40 is 96.97%

CV AUC for max_depth = 10 and n_estimators = 40 is 89.00%

=====

Train AUC for max_depth = 100 and n_estimators = 40 is 99.98%

CV AUC for max_depth = 100 and n_estimators = 40 is 88.43%

=====

Train AUC for max_depth = 500 and n_estimators = 40 is 99.98%

CV AUC for max_depth = 500 and n_estimators = 40 is 88.37%

=====

Train AUC for max_depth = 1000 and n_estimators = 40 is 99.98%

CV AUC for max_depth = 1000 and n_estimators = 40 is 88.42%

=====

Train AUC for max_depth = 1 and n_estimators = 60 is 81.69%

CV AUC for max_depth = 1 and n_estimators = 60 is 81.05%

=====

Train AUC for max_depth = 5 and n_estimators = 60 is 89.12%

CV AUC for max_depth = 5 and n_estimators = 60 is 87.53%

=====

Train AUC for max_depth = 10 and n_estimators = 60 is 97.15%

CV AUC for max_depth = 10 and n_estimators = 60 is 89.36%

=====

Train AUC for max_depth = 100 and n_estimators = 60 is 99.98%

CV AUC for max_depth = 100 and n_estimators = 60 is 88.79%

=====

Train AUC for max_depth = 500 and n_estimators = 60 is 99.98%

CV AUC for max_depth = 500 and n_estimators = 60 is 89.14%

=====

Train AUC for max_depth = 1000 and n_estimators = 60 is 99.98%

CV AUC for max_depth = 1000 and n_estimators = 60 is 88.91%

=====

Train AUC for max_depth = 1 and n_estimators = 80 is 82.84%

CV AUC for max_depth = 1 and n_estimators = 80 is 82.31%

=====

Train AUC for max_depth = 5 and n_estimators = 80 is 89.38%

CV AUC for max_depth = 5 and n_estimators = 80 is 87.67%

=====

Train AUC for max_depth = 10 and n_estimators = 80 is 97.24%

CV AUC for max_depth = 10 and n_estimators = 80 is 89.54%

```

=====

Train AUC for max_depth = 100 and n_estimators = 80 is 99.98%

CV AUC for max_depth = 100 and n_estimators = 80 is 89.30%
=====

Train AUC for max_depth = 500 and n_estimators = 80 is 99.98%

CV AUC for max_depth = 500 and n_estimators = 80 is 89.28%
=====

Train AUC for max_depth = 1000 and n_estimators = 80 is 99.97%

CV AUC for max_depth = 1000 and n_estimators = 80 is 89.14%
=====

Train AUC for max_depth = 1 and n_estimators = 100 is 84.06%

CV AUC for max_depth = 1 and n_estimators = 100 is 83.45%
=====

Train AUC for max_depth = 5 and n_estimators = 100 is 89.33%

CV AUC for max_depth = 5 and n_estimators = 100 is 87.75%
=====

Train AUC for max_depth = 10 and n_estimators = 100 is 97.23%

CV AUC for max_depth = 10 and n_estimators = 100 is 89.48%
=====

Train AUC for max_depth = 100 and n_estimators = 100 is 99.98%

CV AUC for max_depth = 100 and n_estimators = 100 is 89.43%
=====

Train AUC for max_depth = 500 and n_estimators = 100 is 99.98%

CV AUC for max_depth = 500 and n_estimators = 100 is 89.42%
=====

Train AUC for max_depth = 1000 and n_estimators = 100 is 99.98%

CV AUC for max_depth = 1000 and n_estimators = 100 is 89.42%
=====

Train AUC for max_depth = 1 and n_estimators = 120 is 83.27%

```

CV AUC for max_depth = 1 and n_estimators = 120 is 82.76%

=====

Train AUC for max_depth = 5 and n_estimators = 120 is 89.48%

CV AUC for max_depth = 5 and n_estimators = 120 is 87.88%

=====

Train AUC for max_depth = 10 and n_estimators = 120 is 97.32%

CV AUC for max_depth = 10 and n_estimators = 120 is 89.67%

=====

Train AUC for max_depth = 100 and n_estimators = 120 is 99.98%

CV AUC for max_depth = 100 and n_estimators = 120 is 89.57%

=====

Train AUC for max_depth = 500 and n_estimators = 120 is 99.98%

CV AUC for max_depth = 500 and n_estimators = 120 is 89.62%

=====


Train AUC for max_depth = 1000 and n_estimators = 120 is 99.98%

CV AUC for max_depth = 1000 and n_estimators = 120 is 89.55%


=====

```
In [123]: fig, ax = plt.subplots()
          # auc on cv
          print("====CV Data====")
          cv_scores = np.array(w2v_cv_auc).reshape(len(n_estimators),len(max_depth))
          df_cm_cv = pd.DataFrame(cv_scores, n_estimators, max_depth)
          sns.set(font_scale=1.4)
          ax = sns.heatmap(df_cm_cv, annot=True, annot_kws={"size": 10}, fmt='.3g',linewidths=
          ax.set_xlabel("max_depth")
          ax.set_ylabel("n_estimators")
          plt.show()
          print("====Train Data====")
          train_scores = np.array(w2v_train_auc).reshape(len(n_estimators),len(max_depth))
          df_cm_train = pd.DataFrame(train_scores, n_estimators, max_depth)
          sns.set(font_scale=1.4)
          ax = sns.heatmap(df_cm_train, annot=True, annot_kws={"size": 10}, fmt='.3g',linewidths=
          ax.set_xlabel("max_depth")
          ax.set_ylabel("n_estimators")
          plt.show()
```

=====CV Data=====



=====Train Data=====



```
In [124]: # depth is 500
# https://www.dataquest.io/blog/learning-curves-machine-learning/

import matplotlib.pyplot as plt
%matplotlib inline


plt.style.use('seaborn')

plt.plot(n_estimators,train_scores[:,-2],label = 'Training AUC')
plt.plot(n_estimators,cv_scores[:,-2], label = 'Validation AUC')

plt.ylabel('AUC', fontsize = 14)
plt.xlabel('depth', fontsize = 14)
plt.title('Learning curves for a Decision trees model', fontsize = 18, y = 1.03)
plt.legend()

Out[124]: <matplotlib.legend.Legend at 0x7f955020e588>
```


Learning curves for a Decision trees model



```
In [125]: i = 120
          j = 500
```

```
RFDTc = RandomForestClassifier(criterion='gini' , max_depth=j, n_estimators=i)
RFDTc.fit(w2v_train, train_y)
# train data
y_prob_train = RFDTc.predict_proba(w2v_train)[: ,1]
fpr_t, tpr_t, threshold_t = roc_curve(train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y , y_prob_train)
print('\nTrain AUC for max depth = %s and n_estimators = %s is %0.2f%%' % (str(j),str(i),auc_roc_train))
# CV
y_prob_cv = RFDTc.predict_proba(w2v_cv)[: ,1]
fpr_c, tpr_c, threshold_c = roc_curve(cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
print('\nCV AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (str(j),str(i),auc_roc_cv))
# Test
y_prob_test = RFDTc.predict_proba(w2v_test)[: ,1]
fpr_ts, tpr_ts, threshold_ts = roc_curve(test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
```

```

auc_roc_test = roc_auc_score(test_y , y_prob_test)
print('\nTest AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (str(j),str

```

Train AUC for max depth = 500 and n_estimators = 120 is 99.98%

CV AUC for max_depth = 500 and n_estimators = 120 is 89.34%


Test AUC for max_depth = 500 and n_estimators = 120 is 89.37%

In [126]: # <https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python>

```

import matplotlib.pyplot as plt
plt.clf()
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b' , label = 'AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y' , label = 'AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label= 'AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

```




```
In [127]: print("F1-Score on test set: %.2f"%(f1_score(test_y, y_pred_test)))
```

F1-Score on test set: 0.93

```
In [128]: df_cm = pd.DataFrame(confusion_matrix(test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[128]: <matplotlib.axes._subplots.AxesSubplot at 0x7f95501eff60>
```



6.1.6 [5.1.6] Applying Random Forests on TFIDF W2V, SET 4

```
In [129]: # Please write all the code with proper documentation
```

```
model = TfidfVectorizer()
#tf_idf_matrix = model.fit_transform(train)

print("=====Train Data=====")
final_tf_idf_train = model.fit_transform(train)
```

```

print("the type of count vectorizer ",type(final_tf_idf_train))
print("the shape of out text TFIDF vectorizer ",final_tf_idf_train.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_tf_idf_train.get_shape()[0])
print("=====CV Data=====")
final_tf_idf_cv = model.transform(cv)
print("the type of count vectorizer ",type(final_tf_idf_cv))
print("the shape of out text TFIDF vectorizer ",final_tf_idf_cv.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_tf_idf_cv.get_shape()[0])
print("=====Test Data=====")
final_tf_idf_test = model.transform(test)
print("the type of count vectorizer ",type(final_tf_idf_test))
print("the shape of out text TFIDF vectorizer ",final_tf_idf_test.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_tf_idf_test.get_shape()[0])

# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

```

```

=====Train Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (49000, 42656)
the number of unique words including both unigrams and bigrams 42656
=====CV Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (21000, 42656)
the number of unique words including both unigrams and bigrams 42656
=====Test Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (30000, 42656)
the number of unique words including both unigrams and bigrams 42656

```

```

In [130]: ##### Train #####
# TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

train_tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in
row=0;
for sent in tqdm(list_of_train_sentence): # for each review/sentence
    sent_vec = np.zeros(100) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            #tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf valeus of word in this review

```

```

        tf_idf = dictionary[word]*(sent.count(word)/len(sent))
        sent_vec += (vec * tf_idf)
        weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    train_tfidf_sent_vectors.append(sent_vec)
    row += 1

```

100%|| 49000/49000 [38:08<00:00, 21.41it/s]

In [131]: ##### CV #####

```

# TF-IDF weighted Word2Vec
#tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

cv_tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in th
row=0;
for sent in tqdm(list_of_cv_sentence): # for each review/sentence
    sent_vec = np.zeros(100) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            #tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole courpus
            # sent.count(word) = tf valeus of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    cv_tfidf_sent_vectors.append(sent_vec)
    row += 1

```

100%|| 21000/21000 [2:55:38<00:00, 1.99it/s]

In [132]: ##### Train #####

```

# TF-IDF weighted Word2Vec
#tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

test_tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in
row=0;
for sent in tqdm(list_of_test_sentence): # for each review/sentence
    sent_vec = np.zeros(100) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review

```

```

for word in sent: # for each word in a review/sentence
    if word in w2v_words and word in tfidf_feat:
        vec = w2v_model.wv[word]
        #tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
        # to reduce the computation we are
        # dictionary[word] = idf value of word in whole corpus
        # sent.count(word) = tf value of word in this review
        tf_idf = dictionary[word]*(sent.count(word)/len(sent))
        sent_vec += (vec * tf_idf)
        weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    test_tfidf_sent_vectors.append(sent_vec)
    row += 1

```

100%|| 30000/30000 [23:37<00:00, 21.17it/s]

In [133]: # save the datasets as numpy array

```

tfidf_w2v_train = np.array(train_tfidf_sent_vectors)
tfidf_w2v_cv = np.array(cv_tfidf_sent_vectors)
tfidf_w2v_test = np.array(test_tfidf_sent_vectors)

```

In [134]: n_estimators = [20,40,60,80,100,120]

max_depth = [1,5,10,100,500,1000]

tfidf_w2v_train_auc = []

tfidf_w2v_cv_auc = []

for i in n_estimators:

for j in max_depth:

RFDTTC = RandomForestClassifier(n_estimators=i,criterion='gini', max_depth=j)

RFDTTC.fit(tfidf_w2v_train, train_y)

train data

y_prob_train = RFDTTC.predict_proba(tfidf_w2v_train)[:,-1]

y_pred = np.where(y_prob_train > 0.5, 1, 0)

auc_roc_train = roc_auc_score(train_y, y_prob_train)

print('\nTrain AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (s

tfidf_w2v_train_auc.append(auc_roc_train)

CV

y_prob_cv = RFDTTC.predict_proba(tfidf_w2v_cv)[:,-1]

y_pred = np.where(y_prob_cv > 0.5, 1, 0)

auc_roc_cv = roc_auc_score(cv_y, y_prob_cv)

print('\nCV AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (str(

tfidf_w2v_cv_auc.append(auc_roc_cv)

print("="*50)

Train AUC for max_depth = 1 and n_estimators = 20 is 78.71%

CV AUC for max_depth = 1 and n_estimators = 20 is 78.54%

```

=====

Train AUC for max_depth = 5 and n_estimators = 20 is 85.61%

CV AUC for max_depth = 5 and n_estimators = 20 is 84.12%
=====

Train AUC for max_depth = 10 and n_estimators = 20 is 95.73%

CV AUC for max_depth = 10 and n_estimators = 20 is 86.15%
=====

Train AUC for max_depth = 100 and n_estimators = 20 is 99.96%

CV AUC for max_depth = 100 and n_estimators = 20 is 84.52%
=====

Train AUC for max_depth = 500 and n_estimators = 20 is 99.96%

CV AUC for max_depth = 500 and n_estimators = 20 is 84.66%
=====

Train AUC for max_depth = 1000 and n_estimators = 20 is 99.96%

CV AUC for max_depth = 1000 and n_estimators = 20 is 84.36%
=====

Train AUC for max_depth = 1 and n_estimators = 40 is 80.89%

CV AUC for max_depth = 1 and n_estimators = 40 is 80.80%
=====

Train AUC for max_depth = 5 and n_estimators = 40 is 86.48%

CV AUC for max_depth = 5 and n_estimators = 40 is 85.04%
=====

Train AUC for max_depth = 10 and n_estimators = 40 is 96.23%

CV AUC for max_depth = 10 and n_estimators = 40 is 86.73%
=====

Train AUC for max_depth = 100 and n_estimators = 40 is 99.97%

CV AUC for max_depth = 100 and n_estimators = 40 is 86.10%
=====

Train AUC for max_depth = 500 and n_estimators = 40 is 99.97%

```

CV AUC for max_depth = 500 and n_estimators = 40 is 86.04%

=====

Train AUC for max_depth = 1000 and n_estimators = 40 is 99.97%

CV AUC for max_depth = 1000 and n_estimators = 40 is 85.94%

=====

Train AUC for max_depth = 1 and n_estimators = 60 is 78.85%

CV AUC for max_depth = 1 and n_estimators = 60 is 79.08%

=====

Train AUC for max_depth = 5 and n_estimators = 60 is 86.46%

CV AUC for max_depth = 5 and n_estimators = 60 is 84.97%

=====

Train AUC for max_depth = 10 and n_estimators = 60 is 96.36%

CV AUC for max_depth = 10 and n_estimators = 60 is 87.13%

=====

Train AUC for max_depth = 100 and n_estimators = 60 is 99.98%

CV AUC for max_depth = 100 and n_estimators = 60 is 86.60%

=====

Train AUC for max_depth = 500 and n_estimators = 60 is 99.98%

CV AUC for max_depth = 500 and n_estimators = 60 is 86.63%

=====

Train AUC for max_depth = 1000 and n_estimators = 60 is 99.98%

CV AUC for max_depth = 1000 and n_estimators = 60 is 86.87%

=====

Train AUC for max_depth = 1 and n_estimators = 80 is 79.45%

CV AUC for max_depth = 1 and n_estimators = 80 is 79.33%

=====

Train AUC for max_depth = 5 and n_estimators = 80 is 86.51%

CV AUC for max_depth = 5 and n_estimators = 80 is 85.13%

=====

Train AUC for max_depth = 10 and n_estimators = 80 is 96.42%

CV AUC for max_depth = 10 and n_estimators = 80 is 87.27%

=====

Train AUC for max_depth = 100 and n_estimators = 80 is 99.97%

CV AUC for max_depth = 100 and n_estimators = 80 is 87.07%

=====

Train AUC for max_depth = 500 and n_estimators = 80 is 99.97%

CV AUC for max_depth = 500 and n_estimators = 80 is 86.85%

=====

Train AUC for max_depth = 1000 and n_estimators = 80 is 99.97%

CV AUC for max_depth = 1000 and n_estimators = 80 is 87.11%

=====

Train AUC for max_depth = 1 and n_estimators = 100 is 79.81%

CV AUC for max_depth = 1 and n_estimators = 100 is 79.84%

=====

Train AUC for max_depth = 5 and n_estimators = 100 is 86.57%

CV AUC for max_depth = 5 and n_estimators = 100 is 85.03%

=====

Train AUC for max_depth = 10 and n_estimators = 100 is 96.42%

CV AUC for max_depth = 10 and n_estimators = 100 is 87.44%

=====

Train AUC for max_depth = 100 and n_estimators = 100 is 99.98%

CV AUC for max_depth = 100 and n_estimators = 100 is 87.22%

=====

Train AUC for max_depth = 500 and n_estimators = 100 is 99.98%

CV AUC for max_depth = 500 and n_estimators = 100 is 87.30%

=====

Train AUC for max_depth = 1000 and n_estimators = 100 is 99.98%

CV AUC for max_depth = 1000 and n_estimators = 100 is 87.16%

=====

Train AUC for max_depth = 1 and n_estimators = 120 is 81.13%

CV AUC for max_depth = 1 and n_estimators = 120 is 81.09%

=====

Train AUC for max_depth = 5 and n_estimators = 120 is 86.73%

CV AUC for max_depth = 5 and n_estimators = 120 is 85.14%

=====

Train AUC for max_depth = 10 and n_estimators = 120 is 96.56%

CV AUC for max_depth = 10 and n_estimators = 120 is 87.39%

=====

Train AUC for max_depth = 100 and n_estimators = 120 is 99.98%

CV AUC for max_depth = 100 and n_estimators = 120 is 87.47%

=====

Train AUC for max_depth = 500 and n_estimators = 120 is 99.98%

CV AUC for max_depth = 500 and n_estimators = 120 is 87.34%

=====

Train AUC for max_depth = 1000 and n_estimators = 120 is 99.98%

CV AUC for max_depth = 1000 and n_estimators = 120 is 87.44%

=====

```
In [135]: fig, ax = plt.subplots()
```

```
    # auc on cv
```

```
    print("====CV Data====")
```

```
    cv_scores = np.array(tfidf_w2v_cv_auc).reshape(len(n_estimators),len(max_depth))
```

```
    df_cm_cv = pd.DataFrame(cv_scores, n_estimators, max_depth)
```

```
    sns.set(font_scale=1.4)
```

```
    ax = sns.heatmap(df_cm_cv, annot=True, annot_kws={"size": 10}, fmt='.3g',linewidths=
```

```
    ax.set_xlabel("max_depth")
```

```
    ax.set_ylabel("n_estimators")
```

```
    plt.show()
```

```
    print("====Train Data====")
```


```
    train_scores = np.array(tfidf_w2v_train_auc).reshape(len(n_estimators),len(max_depth))
```

```
    df_cm_train = pd.DataFrame(train_scores, n_estimators, max_depth)
```


```
    sns.set(font_scale=1.4)
```

```
ax = sns.heatmap(df_cm_train, annot=True, annot_kws={"size": 10}, fmt='.3g',linewidtht
ax.set_xlabel("max_depth")
ax.set_ylabel("n_estimators")
plt.show()
```

=====CV Data=====



=====Train Data=====



```
In [136]: # depth is 100
# https://www.dataquest.io/blog/learning-curves-machine-learning/

import matplotlib.pyplot as plt
%matplotlib inline


plt.style.use('seaborn')

plt.plot(n_estimators,train_scores[:,-3],label = 'Training AUC')
plt.plot(n_estimators,cv_scores[:,-3], label = 'Validation AUC')

plt.ylabel('AUC', fontsize = 14)
plt.xlabel('depth', fontsize = 14)
plt.title('Learning curves for a Decision trees model', fontsize = 18, y = 1.03)
plt.legend()

Out[136]: <matplotlib.legend.Legend at 0x7f95506dd080>
```

Learning curves for a Decision trees model



```
In [138]: i = 120
          j = 100
```

```
RFDTc = RandomForestClassifier(criterion='gini' , max_depth=j, n_estimators=i)
RFDTc.fit(tfidf_w2v_train, train_y)
# train data
y_prob_train = RFDTc.predict_proba(tfidf_w2v_train)[: ,1]
fpr_t, tpr_t, threshold_t = roc_curve(train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y , y_prob_train)
print('\nTrain AUC for max depth = %s and n_estimators = %s is %0.2f%%' % (str(j),str(i),auc_roc_train))
# CV
y_prob_cv = RFDTc.predict_proba(tfidf_w2v_cv)[: ,1]
fpr_c, tpr_c, threshold_c = roc_curve(cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
print('\nCV AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (str(j),str(i),auc_roc_cv))
# Test
y_prob_test = RFDTc.predict_proba(tfidf_w2v_test)[: ,1]
fpr_ts, tpr_ts, threshold_ts = roc_curve(test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
```

```

auc_roc_test = roc_auc_score(test_y , y_prob_test)
print('\nTest AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (str(j),str

```

Train AUC for max depth = 100 and n_estimators = 120 is 99.98%

CV AUC for max_depth = 100 and n_estimators = 120 is 87.51%


Test AUC for max_depth = 100 and n_estimators = 120 is 86.89%

In [139]: # <https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python>

```

import matplotlib.pyplot as plt
plt.clf()
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b' , label = 'AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y' , label = 'AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label='AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

```




```
In [140]: print("F1-Score on test set: %0.2f"%(f1_score(test_y, y_pred_test)))
```

F1-Score on test set: 0.93

```
In [141]: df_cm = pd.DataFrame(confusion_matrix(test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[141]: <matplotlib.axes._subplots.AxesSubplot at 0x7f956f14fb00>
```



6.2 [5.2] Applying GBDT using XGBOOST

6.2.1 [5.2.1] Applying XGBOOST on BOW, SET 1

```
In [142]: # Please write all the code with proper documentation
from xgboost import XGBClassifier
```

```
In [143]: n_estimators = [20,40,60,80,100,120]
max_depth = [1,5,10,100,500,1000]
bow_train_auc = []
```

```

bow_cv_auc = []
for i in n_estimators:
    for j in max_depth:
        XGBC = XGBClassifier(n_estimators=i, max_depth=j, learning_rate=0.1)
        XGBC.fit(bow_train, train_y)
        # train data
        y_prob_train = XGBC.predict_proba(bow_train)[: ,1]
        y_pred = np.where(y_prob_train > 0.5, 1, 0)
        auc_roc_train = roc_auc_score(train_y , y_prob_train)
        print('\nTrain AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (s
        bow_train_auc.append(auc_roc_train)
        # CV
        y_prob_cv = XGBC.predict_proba(bow_cv)[: ,1]
        y_pred = np.where(y_prob_cv > 0.5, 1, 0)
        auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
        print('\nCV AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (str(
        bow_cv_auc.append(auc_roc_cv)
        print("="*50)

```

Train AUC for max_depth = 1 and n_estimators = 20 is 73.37%

CV AUC for max_depth = 1 and n_estimators = 20 is 71.95%

=====

Train AUC for max_depth = 5 and n_estimators = 20 is 83.21%

CV AUC for max_depth = 5 and n_estimators = 20 is 80.78%

=====

Train AUC for max_depth = 10 and n_estimators = 20 is 89.53%

CV AUC for max_depth = 10 and n_estimators = 20 is 84.23%

=====

Train AUC for max_depth = 100 and n_estimators = 20 is 99.66%

CV AUC for max_depth = 100 and n_estimators = 20 is 87.45%

=====

Train AUC for max_depth = 500 and n_estimators = 20 is 99.69%

CV AUC for max_depth = 500 and n_estimators = 20 is 87.24%

=====

Train AUC for max_depth = 1000 and n_estimators = 20 is 99.69%

CV AUC for max_depth = 1000 and n_estimators = 20 is 87.24%


```

=====

Train AUC for max_depth = 1 and n_estimators = 40 is 77.93%

CV AUC for max_depth = 1 and n_estimators = 40 is 76.61%
=====

Train AUC for max_depth = 5 and n_estimators = 40 is 87.68%

CV AUC for max_depth = 5 and n_estimators = 40 is 85.03%
=====

Train AUC for max_depth = 10 and n_estimators = 40 is 93.35%

CV AUC for max_depth = 10 and n_estimators = 40 is 87.85%
=====

Train AUC for max_depth = 100 and n_estimators = 40 is 99.94%

CV AUC for max_depth = 100 and n_estimators = 40 is 90.23%
=====

Train AUC for max_depth = 500 and n_estimators = 40 is 99.95%

CV AUC for max_depth = 500 and n_estimators = 40 is 90.22%
=====

Train AUC for max_depth = 1000 and n_estimators = 40 is 99.95%

CV AUC for max_depth = 1000 and n_estimators = 40 is 90.22%
=====

Train AUC for max_depth = 1 and n_estimators = 60 is 79.92%

CV AUC for max_depth = 1 and n_estimators = 60 is 78.48%
=====

Train AUC for max_depth = 5 and n_estimators = 60 is 89.89%

CV AUC for max_depth = 5 and n_estimators = 60 is 87.05%
=====

Train AUC for max_depth = 10 and n_estimators = 60 is 95.20%

CV AUC for max_depth = 10 and n_estimators = 60 is 89.64%
=====

Train AUC for max_depth = 100 and n_estimators = 60 is 99.98%

```

CV AUC for max_depth = 100 and n_estimators = 60 is 91.78%

=====

Train AUC for max_depth = 500 and n_estimators = 60 is 99.99%

CV AUC for max_depth = 500 and n_estimators = 60 is 91.72%

=====

Train AUC for max_depth = 1000 and n_estimators = 60 is 99.99%

CV AUC for max_depth = 1000 and n_estimators = 60 is 91.72%

=====

Train AUC for max_depth = 1 and n_estimators = 80 is 81.63%

CV AUC for max_depth = 1 and n_estimators = 80 is 80.20%

=====

Train AUC for max_depth = 5 and n_estimators = 80 is 91.38%

CV AUC for max_depth = 5 and n_estimators = 80 is 88.39%

=====

Train AUC for max_depth = 10 and n_estimators = 80 is 96.26%

CV AUC for max_depth = 10 and n_estimators = 80 is 90.66%

=====

Train AUC for max_depth = 100 and n_estimators = 80 is 99.99%

CV AUC for max_depth = 100 and n_estimators = 80 is 92.59%

=====

Train AUC for max_depth = 500 and n_estimators = 80 is 99.99%

CV AUC for max_depth = 500 and n_estimators = 80 is 92.50%

=====

Train AUC for max_depth = 1000 and n_estimators = 80 is 99.99%

CV AUC for max_depth = 1000 and n_estimators = 80 is 92.50%

=====

Train AUC for max_depth = 1 and n_estimators = 100 is 83.04%

CV AUC for max_depth = 1 and n_estimators = 100 is 81.55%

=====

Train AUC for max_depth = 5 and n_estimators = 100 is 92.43%

CV AUC for max_depth = 5 and n_estimators = 100 is 89.30%

=====

Train AUC for max_depth = 10 and n_estimators = 100 is 96.96%

CV AUC for max_depth = 10 and n_estimators = 100 is 91.33%

=====

Train AUC for max_depth = 100 and n_estimators = 100 is 100.00%

CV AUC for max_depth = 100 and n_estimators = 100 is 93.02%

=====

Train AUC for max_depth = 500 and n_estimators = 100 is 100.00%

CV AUC for max_depth = 500 and n_estimators = 100 is 92.91%

=====

Train AUC for max_depth = 1000 and n_estimators = 100 is 100.00%

CV AUC for max_depth = 1000 and n_estimators = 100 is 92.91%

=====

Train AUC for max_depth = 1 and n_estimators = 120 is 83.97%

CV AUC for max_depth = 1 and n_estimators = 120 is 82.70%

=====

Train AUC for max_depth = 5 and n_estimators = 120 is 93.22%

CV AUC for max_depth = 5 and n_estimators = 120 is 90.02%

=====

Train AUC for max_depth = 10 and n_estimators = 120 is 97.44%

CV AUC for max_depth = 10 and n_estimators = 120 is 91.81%

=====

Train AUC for max_depth = 100 and n_estimators = 120 is 100.00%

CV AUC for max_depth = 100 and n_estimators = 120 is 93.29%

=====

Train AUC for max_depth = 500 and n_estimators = 120 is 100.00%

CV AUC for max_depth = 500 and n_estimators = 120 is 93.17%

=====


Train AUC for max_depth = 1000 and n_estimators = 120 is 100.00%

CV AUC for max_depth = 1000 and n_estimators = 120 is 93.17%


=====

```
In [144]: fig, ax = plt.subplots()
          # auc on cv
          print("=====CV Data=====")
          cv_scores = np.array(bow_cv_auc).reshape(len(n_estimators),len(max_depth))
          df_cm_cv = pd.DataFrame(cv_scores, n_estimators, max_depth)
          sns.set(font_scale=1.4)
          ax = sns.heatmap(df_cm_cv, annot=True, annot_kws={"size": 10}, fmt='.3g',linewidths=
          ax.set_xlabel("max_depth")
          ax.set_ylabel("n_estimators")
          plt.show()
          print("=====Train Data=====")
          train_scores = np.array(bow_train_auc).reshape(len(n_estimators),len(max_depth))
          df_cm_train = pd.DataFrame(train_scores, n_estimators, max_depth)
          sns.set(font_scale=1.4)
          ax = sns.heatmap(df_cm_train, annot=True, annot_kws={"size": 10}, fmt='.3g',linewidths=
          ax.set_xlabel("max_depth")
          ax.set_ylabel("n_estimators")
          plt.show()
```

=====CV Data=====



=====Train Data=====



```
In [145]: # depth is 100
# https://www.dataquest.io/blog/learning-curves-machine-learning/

import matplotlib.pyplot as plt
%matplotlib inline


plt.style.use('seaborn')

plt.plot(n_estimators,train_scores[:,-3],label = 'Training AUC')
plt.plot(n_estimators,cv_scores[:,-3], label = 'Validation AUC')

plt.ylabel('AUC', fontsize = 14)
plt.xlabel('depth', fontsize = 14)
plt.title('Learning curves for a Desision trees model', fontsize = 18, y = 1.03)
plt.legend()

Out[145]: <matplotlib.legend.Legend at 0x7f9567a308d0>
```

Learning curves for a Decision trees model



```
In [146]: i = 120
          j = 100
```

```
XGBC = XGBClassifier(n_estimators=i, max_depth=j, learning_rate=0.1)
XGBC.fit(bow_train, train_y)
# train data
y_prob_train = XGBC.predict_proba(bow_train)[: ,1]
fprt, tprt, thresholdt = roc_curve(train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y, y_prob_train)
print('\nTrain AUC for max depth = %s and n_estimators = %s is %0.2f%%' % (str(j),str(i),auc_roc_train))
# CV
y_prob_cv = XGBC.predict_proba(bow_cv)[: ,1]
fprc, tprc, thresholdc = roc_curve(cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y, y_prob_cv)
print('\nCV AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (str(j),str(i),auc_roc_cv))
# Test
y_prob_test = XGBC.predict_proba(bow_test)[: ,1]
fprts, tprts, thresholdts = roc_curve(test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
```

```

auc_roc_test = roc_auc_score(test_y , y_prob_test)
print('\nTest AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (str(j),str

```

Train AUC for max depth = 100 and n_estimators = 120 is 100.00%

CV AUC for max_depth = 100 and n_estimators = 120 is 93.29%


Test AUC for max_depth = 100 and n_estimators = 120 is 93.19%

In [147]: # <https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python>

```

import matplotlib.pyplot as plt
plt.clf()
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b', label='AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y', label='AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label='AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

```





```
In [148]: print("F1-Score on test set: %.2f"%(f1_score(test_y, y_pred_test)))
```

F1-Score on test set: 0.95

```
In [149]: df_cm = pd.DataFrame(confusion_matrix(test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[149]: <matplotlib.axes._subplots.AxesSubplot at 0x7f9567a35eb8>
```



6.2.2 [5.2.2] Applying XGBOOST on TFIDF, SET 2

```
In [150]: n_estimators = [20,40,60,80,100,120]
max_depth = [1,5,10,100,500,1000]
tfidf_train_auc = []
tfidf_cv_auc = []
for i in n_estimators:
    for j in max_depth:
        XGBC = XGBClassifier(n_estimators=i, max_depth=j, learning_rate=0.1, colsamp
```

```

XGBC.fit(tf_idf_train, train_y)
# train data
y_prob_train = XGBC.predict_proba(tf_idf_train)[: ,1]
y_pred = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y , y_prob_train)
print('\nTrain AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (s
tfidf_train_auc.append(auc_roc_train)
# CV
y_prob_cv = XGBC.predict_proba(tf_idf_cv)[: ,1]
y_pred = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
print('\nCV AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (str(
tfidf_cv_auc.append(auc_roc_cv)
print("="*50)

```

Train AUC for max_depth = 1 and n_estimators = 20 is 76.56%

CV AUC for max_depth = 1 and n_estimators = 20 is 74.98%

=====

Train AUC for max_depth = 5 and n_estimators = 20 is 84.57%

CV AUC for max_depth = 5 and n_estimators = 20 is 82.51%

=====

Train AUC for max_depth = 10 and n_estimators = 20 is 90.83%

CV AUC for max_depth = 10 and n_estimators = 20 is 86.70%

=====

Train AUC for max_depth = 100 and n_estimators = 20 is 99.86%

CV AUC for max_depth = 100 and n_estimators = 20 is 91.55%

=====

Train AUC for max_depth = 500 and n_estimators = 20 is 99.86%

CV AUC for max_depth = 500 and n_estimators = 20 is 91.13%

=====

Train AUC for max_depth = 1000 and n_estimators = 20 is 99.86%

CV AUC for max_depth = 1000 and n_estimators = 20 is 91.35%

=====

Train AUC for max_depth = 1 and n_estimators = 40 is 78.74%

CV AUC for max_depth = 1 and n_estimators = 40 is 77.29%

=====

Train AUC for max_depth = 5 and n_estimators = 40 is 88.74%

CV AUC for max_depth = 5 and n_estimators = 40 is 86.12%

=====

Train AUC for max_depth = 10 and n_estimators = 40 is 94.33%

CV AUC for max_depth = 10 and n_estimators = 40 is 89.30%

=====

Train AUC for max_depth = 100 and n_estimators = 40 is 99.98%

CV AUC for max_depth = 100 and n_estimators = 40 is 93.23%

=====

Train AUC for max_depth = 500 and n_estimators = 40 is 99.96%

CV AUC for max_depth = 500 and n_estimators = 40 is 92.90%

=====

Train AUC for max_depth = 1000 and n_estimators = 40 is 99.96%

CV AUC for max_depth = 1000 and n_estimators = 40 is 93.12%

=====

Train AUC for max_depth = 1 and n_estimators = 60 is 80.78%

CV AUC for max_depth = 1 and n_estimators = 60 is 79.19%

=====

Train AUC for max_depth = 5 and n_estimators = 60 is 90.97%

CV AUC for max_depth = 5 and n_estimators = 60 is 88.14%

=====

Train AUC for max_depth = 10 and n_estimators = 60 is 95.81%

CV AUC for max_depth = 10 and n_estimators = 60 is 90.80%

=====

Train AUC for max_depth = 100 and n_estimators = 60 is 99.99%

CV AUC for max_depth = 100 and n_estimators = 60 is 93.76%

=====

Train AUC for max_depth = 500 and n_estimators = 60 is 99.99%

CV AUC for max_depth = 500 and n_estimators = 60 is 93.62%

=====

Train AUC for max_depth = 1000 and n_estimators = 60 is 99.99%

CV AUC for max_depth = 1000 and n_estimators = 60 is 93.79%

=====

Train AUC for max_depth = 1 and n_estimators = 80 is 82.39%

CV AUC for max_depth = 1 and n_estimators = 80 is 80.85%

=====

Train AUC for max_depth = 5 and n_estimators = 80 is 92.24%

CV AUC for max_depth = 5 and n_estimators = 80 is 89.31%

=====

Train AUC for max_depth = 10 and n_estimators = 80 is 96.81%

CV AUC for max_depth = 10 and n_estimators = 80 is 91.77%

=====

Train AUC for max_depth = 100 and n_estimators = 80 is 99.99%

CV AUC for max_depth = 100 and n_estimators = 80 is 94.12%

=====

Train AUC for max_depth = 500 and n_estimators = 80 is 99.99%

CV AUC for max_depth = 500 and n_estimators = 80 is 94.06%

=====

Train AUC for max_depth = 1000 and n_estimators = 80 is 99.99%

CV AUC for max_depth = 1000 and n_estimators = 80 is 94.16%

=====

Train AUC for max_depth = 1 and n_estimators = 100 is 83.54%

CV AUC for max_depth = 1 and n_estimators = 100 is 81.92%

=====

Train AUC for max_depth = 5 and n_estimators = 100 is 93.29%

CV AUC for max_depth = 5 and n_estimators = 100 is 90.27%

```

=====

Train AUC for max_depth = 10 and n_estimators = 100 is 97.41%

CV AUC for max_depth = 10 and n_estimators = 100 is 92.36%
=====

Train AUC for max_depth = 100 and n_estimators = 100 is 100.00%

CV AUC for max_depth = 100 and n_estimators = 100 is 94.42%
=====

Train AUC for max_depth = 500 and n_estimators = 100 is 100.00%

CV AUC for max_depth = 500 and n_estimators = 100 is 94.36%
=====

Train AUC for max_depth = 1000 and n_estimators = 100 is 100.00%

CV AUC for max_depth = 1000 and n_estimators = 100 is 94.47%
=====

Train AUC for max_depth = 1 and n_estimators = 120 is 84.78%

CV AUC for max_depth = 1 and n_estimators = 120 is 83.22%
=====

Train AUC for max_depth = 5 and n_estimators = 120 is 94.03%

CV AUC for max_depth = 5 and n_estimators = 120 is 90.91%
=====

Train AUC for max_depth = 10 and n_estimators = 120 is 97.88%

CV AUC for max_depth = 10 and n_estimators = 120 is 92.92%
=====

Train AUC for max_depth = 100 and n_estimators = 120 is 100.00%

CV AUC for max_depth = 100 and n_estimators = 120 is 94.59%
=====

Train AUC for max_depth = 500 and n_estimators = 120 is 100.00%

CV AUC for max_depth = 500 and n_estimators = 120 is 94.52%
=====

Train AUC for max_depth = 1000 and n_estimators = 120 is 100.00%


```

CV AUC for max_depth = 1000 and n_estimators = 120 is 94.65%


=====

```
In [151]: fig, ax = plt.subplots()
          # auc on cv
          print("=====CV Data=====")
          cv_scores = np.array(tfidf_cv_auc).reshape(len(n_estimators),len(max_depth))
          df_cm_cv = pd.DataFrame(cv_scores, n_estimators, max_depth)
          sns.set(font_scale=1.4)
          ax = sns.heatmap(df_cm_cv, annot=True, annot_kws={"size": 10}, fmt='.3g',linewidths=
          ax.set_xlabel("max_depth")
          ax.set_ylabel("n_estimators")
          plt.show()
          print("=====Train Data=====")
          train_scores = np.array(tfidf_train_auc).reshape(len(n_estimators),len(max_depth))
          df_cm_train = pd.DataFrame(train_scores, n_estimators, max_depth)
          sns.set(font_scale=1.4)
          ax = sns.heatmap(df_cm_train, annot=True, annot_kws={"size": 10}, fmt='.3g',linewidths=
          ax.set_xlabel("max_depth")
          ax.set_ylabel("n_estimators")
          plt.show()
```

=====CV Data=====



=====Train Data=====



```
In [152]: # depth is 1000
# https://www.dataquest.io/blog/learning-curves-machine-learning/

import matplotlib.pyplot as plt
%matplotlib inline


plt.style.use('seaborn')

plt.plot(n_estimators,train_scores[:,-1],label = 'Training AUC')
plt.plot(n_estimators,cv_scores[:,-1], label = 'Validation AUC')

plt.ylabel('AUC', fontsize = 14)
plt.xlabel('depth', fontsize = 14)
plt.title('Learning curves for a Desision trees model', fontsize = 18, y = 1.03)
plt.legend()

Out[152]: <matplotlib.legend.Legend at 0x7f956c070ac8>
```


Learning curves for a Decision trees model



```
In [153]: i = 120
          j = 100
```

```
XGBC = XGBClassifier(n_estimators=i, max_depth=j, learning_rate=0.1)
XGBC.fit(tf_idf_train, train_y)
# train data
y_prob_train = XGBC.predict_proba(tf_idf_train)[: ,1]
fpr_t, tpr_t, threshold_t = roc_curve(train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y, y_prob_train)
print('\nTrain AUC for max depth = %s and n_estimators = %s is %0.2f%%' % (str(j),str(i),auc_roc_train))
# CV
y_prob_cv = XGBC.predict_proba(tf_idf_cv)[: ,1]
fpr_c, tpr_c, threshold_c = roc_curve(cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y, y_prob_cv)
print('\nCV AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (str(j),str(i),auc_roc_cv))
# Test
y_prob_test = XGBC.predict_proba(tf_idf_test)[: ,1]
fpr_ts, tpr_ts, threshold_ts = roc_curve(test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
```

```

auc_roc_test = roc_auc_score(test_y , y_prob_test)
print('\nTest AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (str(j),str

```

Train AUC for max depth = 100 and n_estimators = 120 is 100.00%

CV AUC for max_depth = 100 and n_estimators = 120 is 93.91%


Test AUC for max_depth = 100 and n_estimators = 120 is 93.82%

In [154]: # <https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python>

```

import matplotlib.pyplot as plt
plt.clf()
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b' , label = 'AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y' , label = 'AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label= 'AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

```




```
In [155]: print("F1-Score on test set: %.2f"%(f1_score(test_y, y_pred_test)))
```

F1-Score on test set: 0.95

```
In [156]: df_cm = pd.DataFrame(confusion_matrix(test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[156]: <matplotlib.axes._subplots.AxesSubplot at 0x7f956a73fcf8>
```



6.2.3 [5.2.3] Applying XGBOOST on AVG W2V, SET 3

```
In [159]: n_estimators = [20,40,60,80,100,120]
max_depth = [1,5,10,100,500,1000]
w2v_train_auc = []
w2v_cv_auc = []
for i in n_estimators:
    for j in max_depth:
        XGBC = XGBClassifier(n_estimators=i, max_depth=j, learning_rate=0.1, colsamp
```

```

XGBC.fit(w2v_train, train_y)
# train data
y_prob_train = XGBC.predict_proba(w2v_train)[: ,1]
y_pred = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y , y_prob_train)
print('\nTrain AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (s
w2v_train_auc.append(auc_roc_train)
# CV
y_prob_cv = XGBC.predict_proba(w2v_cv)[: ,1]
y_pred = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
print('\nCV AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (str(
w2v_cv_auc.append(auc_roc_cv)
print("="*50)

```

Train AUC for max_depth = 1 and n_estimators = 20 is 82.14%

CV AUC for max_depth = 1 and n_estimators = 20 is 81.69%

=====

Train AUC for max_depth = 5 and n_estimators = 20 is 90.69%

CV AUC for max_depth = 5 and n_estimators = 20 is 88.42%

=====

Train AUC for max_depth = 10 and n_estimators = 20 is 98.83%

CV AUC for max_depth = 10 and n_estimators = 20 is 89.55%

=====

Train AUC for max_depth = 100 and n_estimators = 20 is 99.97%

CV AUC for max_depth = 100 and n_estimators = 20 is 89.05%

=====

Train AUC for max_depth = 500 and n_estimators = 20 is 99.97%

CV AUC for max_depth = 500 and n_estimators = 20 is 89.05%

=====

Train AUC for max_depth = 1000 and n_estimators = 20 is 99.97%

CV AUC for max_depth = 1000 and n_estimators = 20 is 89.05%

=====

Train AUC for max_depth = 1 and n_estimators = 40 is 85.36%

CV AUC for max_depth = 1 and n_estimators = 40 is 84.63%

=====

Train AUC for max_depth = 5 and n_estimators = 40 is 92.40%

CV AUC for max_depth = 5 and n_estimators = 40 is 89.47%

=====

Train AUC for max_depth = 10 and n_estimators = 40 is 99.62%

CV AUC for max_depth = 10 and n_estimators = 40 is 90.45%

=====

Train AUC for max_depth = 100 and n_estimators = 40 is 99.99%

CV AUC for max_depth = 100 and n_estimators = 40 is 90.19%

=====

Train AUC for max_depth = 500 and n_estimators = 40 is 99.99%

CV AUC for max_depth = 500 and n_estimators = 40 is 90.19%

=====

Train AUC for max_depth = 1000 and n_estimators = 40 is 99.99%

CV AUC for max_depth = 1000 and n_estimators = 40 is 90.19%

=====

Train AUC for max_depth = 1 and n_estimators = 60 is 86.57%

CV AUC for max_depth = 1 and n_estimators = 60 is 85.77%

=====

Train AUC for max_depth = 5 and n_estimators = 60 is 93.46%

CV AUC for max_depth = 5 and n_estimators = 60 is 90.12%

=====

Train AUC for max_depth = 10 and n_estimators = 60 is 99.85%

CV AUC for max_depth = 10 and n_estimators = 60 is 90.77%

=====

Train AUC for max_depth = 100 and n_estimators = 60 is 100.00%

CV AUC for max_depth = 100 and n_estimators = 60 is 90.72%

=====

Train AUC for max_depth = 500 and n_estimators = 60 is 100.00%

CV AUC for max_depth = 500 and n_estimators = 60 is 90.72%

=====

Train AUC for max_depth = 1000 and n_estimators = 60 is 100.00%

CV AUC for max_depth = 1000 and n_estimators = 60 is 90.72%

=====

Train AUC for max_depth = 1 and n_estimators = 80 is 87.49%

CV AUC for max_depth = 1 and n_estimators = 80 is 86.66%

=====

Train AUC for max_depth = 5 and n_estimators = 80 is 94.21%

CV AUC for max_depth = 5 and n_estimators = 80 is 90.48%

=====

Train AUC for max_depth = 10 and n_estimators = 80 is 99.95%

CV AUC for max_depth = 10 and n_estimators = 80 is 91.00%

=====

Train AUC for max_depth = 100 and n_estimators = 80 is 100.00%

CV AUC for max_depth = 100 and n_estimators = 80 is 91.01%

=====

Train AUC for max_depth = 500 and n_estimators = 80 is 100.00%

CV AUC for max_depth = 500 and n_estimators = 80 is 91.01%

=====

Train AUC for max_depth = 1000 and n_estimators = 80 is 100.00%

CV AUC for max_depth = 1000 and n_estimators = 80 is 91.01%

=====

Train AUC for max_depth = 1 and n_estimators = 100 is 88.16%

CV AUC for max_depth = 1 and n_estimators = 100 is 87.24%

=====

Train AUC for max_depth = 5 and n_estimators = 100 is 94.80%

CV AUC for max_depth = 5 and n_estimators = 100 is 90.72%

```

=====

Train AUC for max_depth = 10 and n_estimators = 100 is 99.98%

CV AUC for max_depth = 10 and n_estimators = 100 is 91.13%
=====

Train AUC for max_depth = 100 and n_estimators = 100 is 100.00%

CV AUC for max_depth = 100 and n_estimators = 100 is 91.21%
=====

Train AUC for max_depth = 500 and n_estimators = 100 is 100.00%

CV AUC for max_depth = 500 and n_estimators = 100 is 91.21%
=====

Train AUC for max_depth = 1000 and n_estimators = 100 is 100.00%

CV AUC for max_depth = 1000 and n_estimators = 100 is 91.21%
=====

Train AUC for max_depth = 1 and n_estimators = 120 is 88.62%

CV AUC for max_depth = 1 and n_estimators = 120 is 87.64%
=====

Train AUC for max_depth = 5 and n_estimators = 120 is 95.26%

CV AUC for max_depth = 5 and n_estimators = 120 is 90.85%
=====

Train AUC for max_depth = 10 and n_estimators = 120 is 99.99%

CV AUC for max_depth = 10 and n_estimators = 120 is 91.22%
=====

Train AUC for max_depth = 100 and n_estimators = 120 is 100.00%

CV AUC for max_depth = 100 and n_estimators = 120 is 91.32%
=====

Train AUC for max_depth = 500 and n_estimators = 120 is 100.00%

CV AUC for max_depth = 500 and n_estimators = 120 is 91.32%
=====

Train AUC for max_depth = 1000 and n_estimators = 120 is 100.00%


```

CV AUC for max_depth = 1000 and n_estimators = 120 is 91.32%


=====

```
In [160]: fig, ax = plt.subplots()
          # auc on cv
          print("=====CV Data=====")
          cv_scores = np.array(w2v_cv_auc).reshape(len(n_estimators),len(max_depth))
          df_cm_cv = pd.DataFrame(cv_scores, n_estimators, max_depth)
          sns.set(font_scale=1.4)
          ax = sns.heatmap(df_cm_cv, annot=True, annot_kws={"size": 10}, fmt='.3g',linewidths=
          ax.set_xlabel("max_depth")
          ax.set_ylabel("n_estimators")
          plt.show()
          print("=====Train Data=====")
          train_scores = np.array(w2v_train_auc).reshape(len(n_estimators),len(max_depth))
          df_cm_train = pd.DataFrame(train_scores, n_estimators, max_depth)
          sns.set(font_scale=1.4)
          ax = sns.heatmap(df_cm_train, annot=True, annot_kws={"size": 10}, fmt='.3g',linewidths=
          ax.set_xlabel("max_depth")
          ax.set_ylabel("n_estimators")
          plt.show()
```

=====CV Data=====



=====Train Data=====



```
In [161]: # depth is 10
           # https://www.dataquest.io/blog/learning-curves-machine-learning/

import matplotlib.pyplot as plt
%matplotlib inline


plt.style.use('seaborn')

plt.plot(n_estimators,train_scores[:,-4],label = 'Training AUC')
plt.plot(n_estimators,cv_scores[:,-4], label = 'Validation AUC')

plt.ylabel('AUC', fontsize = 14)
plt.xlabel('depth', fontsize = 14)
plt.title('Learning curves for a Decision trees model', fontsize = 18, y = 1.03)
plt.legend()

Out[161]: <matplotlib.legend.Legend at 0x7f956b083668>
```

Learning curves for a Decision trees model



```
In [163]: i = 120
          j = 10
```

```
XGBC = XGBClassifier(n_estimators=i, max_depth=j, learning_rate=0.1)
XGBC.fit(w2v_train, train_y)
# train data
y_prob_train = XGBC.predict_proba(w2v_train)[: ,1]
fpr_t, tpr_t, threshold_t = roc_curve(train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y, y_prob_train)
print('\nTrain AUC for max depth = %s and n_estimators = %s is %0.2f%%' % (str(j), str(i), auc_roc_train))
# CV
y_prob_cv = XGBC.predict_proba(w2v_cv)[: ,1]
fpr_c, tpr_c, threshold_c = roc_curve(cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y, y_prob_cv)
print('\nCV AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (str(j), str(i), auc_roc_cv))
# Test
y_prob_test = XGBC.predict_proba(w2v_test)[: ,1]
fpr_ts, tpr_ts, threshold_ts = roc_curve(test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
```

```

auc_roc_test = roc_auc_score(test_y , y_prob_test)
print('\nTest AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (str(j),str

```

Train AUC for max depth = 10 and n_estimators = 120 is 99.99%

CV AUC for max_depth = 10 and n_estimators = 120 is 91.04%


Test AUC for max_depth = 10 and n_estimators = 120 is 91.12%

In [164]: # <https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python>

```

import matplotlib.pyplot as plt
plt.clf()
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b' , label = 'AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y' , label = 'AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label='AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

```




```
In [165]: print("F1-Score on test set: %.2f"%(f1_score(test_y, y_pred_test)))
```

F1-Score on test set: 0.94

```
In [166]: df_cm = pd.DataFrame(confusion_matrix(test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[166]: <matplotlib.axes._subplots.AxesSubplot at 0x7f956b0a4dd8>
```



6.2.4 [5.2.4] Applying XGBOOST on TFIDF W2V, SET 4

```
In [167]: # Please write all the code with proper documentation
n_estimators = [20,40,60,80,100,120]
max_depth = [1,5,10,100,500,1000]
tfidf_w2v_train_auc = []
tfidf_w2v_cv_auc = []
for i in n_estimators:
    for j in max_depth:
```

```

XGBC = XGBClassifier(n_estimators=i, max_depth=j, learning_rate=0.1, colsamp
XGBC.fit(tfidf_w2v_train, train_y)
# train data
y_prob_train = XGBC.predict_proba(tfidf_w2v_train)[: ,1]
y_pred = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y , y_prob_train)
print('\nTrain AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (s
tfidf_w2v_train_auc.append(auc_roc_train)
# CV
y_prob_cv = XGBC.predict_proba(tfidf_w2v_cv)[: ,1]
y_pred = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
print('\nCV AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (str(
tfidf_w2v_cv_auc.append(auc_roc_cv)
print("="*50)

```

Train AUC for max_depth = 1 and n_estimators = 20 is 78.69%

CV AUC for max_depth = 1 and n_estimators = 20 is 78.72%

=====

Train AUC for max_depth = 5 and n_estimators = 20 is 88.37%

CV AUC for max_depth = 5 and n_estimators = 20 is 86.07%

=====

Train AUC for max_depth = 10 and n_estimators = 20 is 98.53%

CV AUC for max_depth = 10 and n_estimators = 20 is 87.20%

=====

Train AUC for max_depth = 100 and n_estimators = 20 is 99.96%

CV AUC for max_depth = 100 and n_estimators = 20 is 86.70%

=====

Train AUC for max_depth = 500 and n_estimators = 20 is 99.96%

CV AUC for max_depth = 500 and n_estimators = 20 is 86.70%

=====

Train AUC for max_depth = 1000 and n_estimators = 20 is 99.96%

CV AUC for max_depth = 1000 and n_estimators = 20 is 86.70%

=====

Train AUC for max_depth = 1 and n_estimators = 40 is 81.84%

CV AUC for max_depth = 1 and n_estimators = 40 is 81.60%

=====

Train AUC for max_depth = 5 and n_estimators = 40 is 90.60%

CV AUC for max_depth = 5 and n_estimators = 40 is 87.46%

=====

Train AUC for max_depth = 10 and n_estimators = 40 is 99.58%

CV AUC for max_depth = 10 and n_estimators = 40 is 88.48%

=====

Train AUC for max_depth = 100 and n_estimators = 40 is 99.99%

CV AUC for max_depth = 100 and n_estimators = 40 is 88.18%

=====

Train AUC for max_depth = 500 and n_estimators = 40 is 99.99%

CV AUC for max_depth = 500 and n_estimators = 40 is 88.18%

=====

Train AUC for max_depth = 1000 and n_estimators = 40 is 99.99%

CV AUC for max_depth = 1000 and n_estimators = 40 is 88.18%

=====

Train AUC for max_depth = 1 and n_estimators = 60 is 83.49%

CV AUC for max_depth = 1 and n_estimators = 60 is 83.14%

=====

Train AUC for max_depth = 5 and n_estimators = 60 is 91.95%

CV AUC for max_depth = 5 and n_estimators = 60 is 88.25%

=====

Train AUC for max_depth = 10 and n_estimators = 60 is 99.84%

CV AUC for max_depth = 10 and n_estimators = 60 is 88.86%

=====

Train AUC for max_depth = 100 and n_estimators = 60 is 100.00%

CV AUC for max_depth = 100 and n_estimators = 60 is 88.81%

=====

Train AUC for max_depth = 500 and n_estimators = 60 is 100.00%

CV AUC for max_depth = 500 and n_estimators = 60 is 88.81%

=====

Train AUC for max_depth = 1000 and n_estimators = 60 is 100.00%

CV AUC for max_depth = 1000 and n_estimators = 60 is 88.81%

=====

Train AUC for max_depth = 1 and n_estimators = 80 is 84.47%

CV AUC for max_depth = 1 and n_estimators = 80 is 84.03%

=====

Train AUC for max_depth = 5 and n_estimators = 80 is 92.83%

CV AUC for max_depth = 5 and n_estimators = 80 is 88.71%

=====

Train AUC for max_depth = 10 and n_estimators = 80 is 99.94%

CV AUC for max_depth = 10 and n_estimators = 80 is 89.06%

=====

Train AUC for max_depth = 100 and n_estimators = 80 is 100.00%

CV AUC for max_depth = 100 and n_estimators = 80 is 89.15%

=====

Train AUC for max_depth = 500 and n_estimators = 80 is 100.00%

CV AUC for max_depth = 500 and n_estimators = 80 is 89.15%

=====

Train AUC for max_depth = 1000 and n_estimators = 80 is 100.00%

CV AUC for max_depth = 1000 and n_estimators = 80 is 89.15%

=====

Train AUC for max_depth = 1 and n_estimators = 100 is 85.33%

CV AUC for max_depth = 1 and n_estimators = 100 is 84.75%

=====

Train AUC for max_depth = 5 and n_estimators = 100 is 93.53%

CV AUC for max_depth = 5 and n_estimators = 100 is 88.94%
=====

Train AUC for max_depth = 10 and n_estimators = 100 is 99.98%

CV AUC for max_depth = 10 and n_estimators = 100 is 89.24%
=====

Train AUC for max_depth = 100 and n_estimators = 100 is 100.00%

CV AUC for max_depth = 100 and n_estimators = 100 is 89.32%
=====

Train AUC for max_depth = 500 and n_estimators = 100 is 100.00%

CV AUC for max_depth = 500 and n_estimators = 100 is 89.32%
=====

Train AUC for max_depth = 1000 and n_estimators = 100 is 100.00%

CV AUC for max_depth = 1000 and n_estimators = 100 is 89.32%
=====

Train AUC for max_depth = 1 and n_estimators = 120 is 85.87%

CV AUC for max_depth = 1 and n_estimators = 120 is 85.24%
=====

Train AUC for max_depth = 5 and n_estimators = 120 is 94.15%

CV AUC for max_depth = 5 and n_estimators = 120 is 89.10%
=====

Train AUC for max_depth = 10 and n_estimators = 120 is 99.99%

CV AUC for max_depth = 10 and n_estimators = 120 is 89.33%
=====

Train AUC for max_depth = 100 and n_estimators = 120 is 100.00%

CV AUC for max_depth = 100 and n_estimators = 120 is 89.46%
=====

Train AUC for max_depth = 500 and n_estimators = 120 is 100.00%

CV AUC for max_depth = 500 and n_estimators = 120 is 89.46%
=====


Train AUC for max_depth = 1000 and n_estimators = 120 is 100.00%

CV AUC for max_depth = 1000 and n_estimators = 120 is 89.46%


=====

```
In [168]: fig, ax = plt.subplots()
          # auc on cv
          print("=====CV Data=====")
          cv_scores = np.array(tfidf_w2v_cv_auc).reshape(len(n_estimators),len(max_depth))
          df_cm_cv = pd.DataFrame(cv_scores, n_estimators, max_depth)
          sns.set(font_scale=1.4)
          ax = sns.heatmap(df_cm_cv, annot=True, annot_kws={"size": 10}, fmt='.3g',linewidths=
          ax.set_xlabel("max_depth")
          ax.set_ylabel("n_estimators")
          plt.show()
          print("=====Train Data=====")
          train_scores = np.array(tfidf_w2v_train_auc).reshape(len(n_estimators),len(max_depth))
          df_cm_train = pd.DataFrame(train_scores, n_estimators, max_depth)
          sns.set(font_scale=1.4)
          ax = sns.heatmap(df_cm_train, annot=True, annot_kws={"size": 10}, fmt='.3g',linewidths=
          ax.set_xlabel("max_depth")
          ax.set_ylabel("n_estimators")
          plt.show()
```

=====CV Data=====



=====Train Data=====



```
In [169]: # depth is 10
           # https://www.dataquest.io/blog/learning-curves-machine-learning/

import matplotlib.pyplot as plt
%matplotlib inline


plt.style.use('seaborn')

plt.plot(n_estimators,train_scores[:,-4],label = 'Training AUC')
plt.plot(n_estimators,cv_scores[:,-4], label = 'Validation AUC')

plt.ylabel('AUC', fontsize = 14)
plt.xlabel('depth', fontsize = 14)
plt.title('Learning curves for a Decision trees model', fontsize = 18, y = 1.03)
plt.legend()

Out[169]: <matplotlib.legend.Legend at 0x7f956a2614a8>
```

Learning curves for a Decision trees model



```
In [170]: i = 120
          j = 10
```

```
XGBC = XGBClassifier(n_estimators=i, max_depth=j, learning_rate=0.1)
XGBC.fit(tfidf_w2v_train, train_y)
# train data
y_prob_train = XGBC.predict_proba(tfidf_w2v_train)[: ,1]
fprt, tprt, thresholdt = roc_curve(train_y, y_prob_train)
y_pred_train = np.where(y_prob_train > 0.5, 1, 0)
auc_roc_train = roc_auc_score(train_y , y_prob_train)
print('\nTrain AUC for max depth = %s and n_estimators = %s is %0.2f%%' % (str(j),str(i),auc_roc_train))
# CV
y_prob_cv = XGBC.predict_proba(tfidf_w2v_cv)[: ,1]
fprc, tprc, thresholdc = roc_curve(cv_y, y_prob_cv)
y_pred_cv = np.where(y_prob_cv > 0.5, 1, 0)
auc_roc_cv = roc_auc_score(cv_y , y_prob_cv)
print('\nCV AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (str(j),str(i),auc_roc_cv))
# Test
y_prob_test = XGBC.predict_proba(tfidf_w2v_test)[: ,1]
fprts, tprts, thresholdts = roc_curve(test_y, y_prob_test)
y_pred_test = np.where(y_prob_test > 0.5, 1, 0)
```

```

auc_roc_test = roc_auc_score(test_y , y_prob_test)
print('\nTest AUC for max_depth = %s and n_estimators = %s is %0.2f%%' % (str(j),str

```

Train AUC for max depth = 10 and n_estimators = 120 is 99.99%

CV AUC for max_depth = 10 and n_estimators = 120 is 89.49%


Test AUC for max_depth = 10 and n_estimators = 120 is 89.06%

In [171]: # <https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python>

```

import matplotlib.pyplot as plt
plt.clf()
plt.title('Receiver Operating Characteristic')
plt.plot(fprc, tprc, 'b' , label = 'AUC CV = %0.2f' % (auc_roc_cv * float(100)))
plt.plot(fprts, tprts, 'y' , label = 'AUC Test = %0.2f' % (auc_roc_test * float(100)))
plt.plot(fprt, tprt, 'g', label = 'AUC Train = %0.2f' % (auc_roc_train * float(100)))
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

```




```
In [172]: print("F1-Score on test set: %.2f"%(f1_score(test_y, y_pred_test)))
```

F1-Score on test set: 0.93

```
In [173]: df_cm = pd.DataFrame(confusion_matrix(test_y, y_pred_test), range(2), range(2))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
```

```
Out[173]: <matplotlib.axes._subplots.AxesSubplot at 0x7f956d3d05f8>
```



7 Feature engineering

Taking length of reviews as another feature

```
In [18]: # Please compare all your models using Prettytable library
X_review = np.array(sorted_data['CleanedText'])
X_summary = np.array(sorted_data['Summary'])
y = np.array(sorted_data['Score'])
print(X_review.shape)
print(X_summary.shape)
print(y.shape)
```

```
(100000,)  
(100000,)  
(100000,)
```

```
In [19]: # Simple cross validation  
# split the data sent into train and test  
train_review , test_review, train_summary, test_summary , train_y , test_y = train_test_split(train_review, train_summary, train_y, test_review, test_summary, test_y)  
  
# split the train data set into cross validation train and cross validation test  
train_review, cv_review, train_summary, cv_summary , train_y, cv_y = train_test_split(train_review, train_summary, train_y, cv_review, cv_summary, cv_y)  
  
print("train review", train_review.shape)  
print("cv review", cv_review.shape)  
print("test review", test_review.shape)  
print("train summary", train_summary.shape)  
print("cv summary", cv_summary.shape)  
print("test summary", test_summary.shape)  
  
train_review (49000,)  
cv_review (21000,)  
test_review (30000,)  
train_summary (49000,)  
cv_summary (21000,)  
test_summary (30000,)
```

```
In [20]: # bow for summary  
# Please write all the code with proper documentation  
#BoW  
count_vect = CountVectorizer(min_df=15, ngram_range=(1,2)) #in scikit-learn  
count_vect.fit(train_summary)  
print("some feature names ", count_vect.get_feature_names()[:10])  
print('='*50)  
  
bow_train_summary = count_vect.fit_transform(train_summary)  
bow_cv_summary = count_vect.transform(cv_summary)  
bow_test_summary = count_vect.transform(test_summary)  
print("=====Train Data=====")  
print("the type of count vectorizer ",type(bow_train_summary))  
print("the shape of out text BOW vectorizer ",bow_train_summary.get_shape())  
print("the number of unique words ", bow_train_summary.get_shape()[1])  
print("=====Cross validation Data=====")  
print("the type of count vectorizer ",type(bow_cv_summary))  
print("the shape of out text BOW vectorizer ",bow_cv_summary.get_shape())  
print("the number of unique words ", bow_cv_summary.get_shape()[1])  
print("=====Test Data=====")  
print("the type of count vectorizer ",type(bow_test_summary))
```



```

print("the shape of out text BOW vectorizer ",bow_test_summary.get_shape())
print("the number of unique words ", bow_test_summary.get_shape()[1])

some feature names  ['10', '100', '11', '12', '16', '20', '24', '40', '50', 'about']
=====
=====Train Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (49000, 2274)
the number of unique words 2274
=====Cross validation Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (21000, 2274)
the number of unique words 2274
=====Test Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (30000, 2274)
the number of unique words 2274

```

```

In [21]: # bow for review
         # Please write all the code with proper documentation
         #BoW
         count_vect = CountVectorizer(min_df=15, ngram_range=(1,2)) #in scikit-learn
         count_vect.fit(train_review)
         print("some feature names ", count_vect.get_feature_names()[:10])
         print('='*50)

         bow_train_review = count_vect.fit_transform(train_review)
         bow_cv_review = count_vect.transform(cv_review)
         bow_test_review = count_vect.transform(test_review)
         print("=====Train Data=====")
         print("the type of count vectorizer ",type(bow_train_review))
         print("the shape of out text BOW vectorizer ",bow_train_review.get_shape())
         print("the number of unique words ", bow_train_review.get_shape()[1])
         print("=====Cross validation Data=====")
         print("the type of count vectorizer ",type(bow_cv_review))
         print("the shape of out text BOW vectorizer ",bow_cv_review.get_shape())
         print("the number of unique words ", bow_cv_review.get_shape()[1])
         print("=====Test Data=====")
         print("the type of count vectorizer ",type(bow_test_review))
         print("the shape of out text BOW vectorizer ",bow_test_review.get_shape())
         print("the number of unique words ", bow_test_review.get_shape()[1])

```

```

some feature names  ['ability', 'able', 'able buy', 'able drink', 'able eat', 'able enjoy', 'al
=====
=====Train Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (49000, 18450)

```

```

the number of unique words 18450
=====Cross validation Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (21000, 18450)
the number of unique words 18450
=====Test Data=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (30000, 18450)
the number of unique words 18450

```

7.0.1 StackingClassifier

```

In [22]: from sklearn import model_selection
         from sklearn.linear_model import LogisticRegression
         from sklearn.naive_bayes import MultinomialNB
         from sklearn.svm import SVC
         from sklearn.ensemble import RandomForestClassifier
         from mlxtend.classifier import StackingClassifier

In [24]: clf1 = LogisticRegression(penalty='l2', C=0.1)
         clf2 = MultinomialNB(alpha=1)
         clf3 = RandomForestClassifier(n_estimators=120, criterion='gini', max_depth=50)
         lr = LogisticRegression()
         sclf = StackingClassifier(classifiers=[clf1, clf2, clf3], meta_classifier=lr, use_pro

print("3-fold cross validation:\n")

for clf, label in zip([clf1, clf2, clf3, sclf],
                      ['Logistic rgression',
                       'MultinomialNB',
                       'RF Classifire',
                       'Staking Classifier']):
    scores_summary = model_selection.cross_val_score(clf, bow_train_summary, train_y,
                                                    cv=3, scoring='roc_auc')

    scores_review = model_selection.cross_val_score(clf, bow_train_review, train_y,
                                                    cv=3, scoring='roc_auc')

    new_scores = (scores_summary + scores_review) / 2
    print("AUC: %0.2f (+/- %0.2f) [%s]"
          % (new_scores.mean(), new_scores.std(), label))

3-fold cross validation:

AUC: 0.94 (+/- 0.00) [Logistic rgression]
AUC: 0.93 (+/- 0.00) [MultinomialNB]
AUC: 0.91 (+/- 0.00) [RF Classifire]

```

AUC: 0.94 (+/- 0.00) [Staking Classifier]

8 [6] Conclusions

```
In [27]: # Please compare all your models using Prettytable library
         from prettytable import PrettyTable
```

```
x = PrettyTable(["Vectorizer" , "max_depth" , "n_estimators","AUC" , "F1 Score"])

x.add_row(["BOW" , 500 , 120,"90.90%" , 0.94])
x.add_row(["TFIDF",1000 , 120,"92.99%" , 0.94])
x.add_row(["AVG-W2V" , 500 , 120,"89.37%" , 0.93])
x.add_row(["TFIDF-w2v",100 , 120,"86.89%" , 0.93])
print(x.get_string(title="RF Model"))
```

Vectorizer	max_depth	n_estimators	AUC	F1 Score
BOW	500	120	90.90%	0.94
TFIDF	1000	120	92.99%	0.94
AVG-W2V	500	120	89.37%	0.93
TFIDF-w2v	100	120	86.89%	0.93

```
In [28]: x = PrettyTable(["Vectorizer" , "max_depth" , "n_estimators","AUC" , "F1 Score"])
```

```
x.add_row(["BOW" , 100 , 120,"93.19%" , 0.95])
x.add_row(["TFIDF",100 , 120,"93.83%" , 0.95])
x.add_row(["AVG-W2V" , 10 , 120,"91.12%" , 0.94])
x.add_row(["TFIDF-w2v",10 , 120,"89.06%" , 0.93])
print(x.get_string(title="GDBT Model"))
```

Vectorizer	max_depth	n_estimators	AUC	F1 Score
BOW	100	120	93.19%	0.95
TFIDF	100	120	93.83%	0.95
AVG-W2V	10	120	91.12%	0.94
TFIDF-w2v	10	120	89.06%	0.93

```
In [ ]:
```